

Exact Inference Is Not Enough: Opponent Models and Endgames in Six-Player Literature

KRAKEN v1.1

Kausthubh Veldanda

30 August 2026

Abstract

Literature, also called Fish or Canadian Fish, is a six-player, two-team, imperfect-information card game in which players interrogate opponents for named cards and score by declaring the exact distribution of a six-card half-suit among their own team. Every action is public, so the entire hidden state is the deal; there is no legal communication channel, which places the game outside both the two-player zero-sum machinery that solved poker and the fully cooperative machinery developed for Hanabi. We present KRAKEN and the four development cycles that produced it.

Hidden-state inference here reduces to counting assignments in a degree-constrained bipartite system. The general problem is $\#P$ -complete, but real positions carry only 4.1 distinct candidate sets across 25.1 unlocated cards, so a dynamic program returns the exact partition function, marginals, uniform samples and joint probabilities, validated against exhaustive enumeration.

The central result is negative. Making the posterior exactly correct changed nothing in play, -0.008 sets per duplicate deal-pair over 250 pairs, and made it a *worse* predictor of where cards actually sit than the biased sampler it replaced (1.3618 against 1.3408 in mean negative log-likelihood). A uniform posterior is a correct theorem about a false hypothesis: it assumes choices carry no information beyond legality. The old sampler’s bias had accidentally encoded an opponent model, and removing the bias removed the model. Replacing that accident with an explicit one-parameter choice model is worth **+1.92** [+1.48, +2.36] sets per pair, the largest single gain in the study.

Three results follow. The residual errors are not failures to *find* cards: **95.3%** are *allocation* errors, where our team held all six of a half-suit and split them wrongly. Handing a seat the true deal one side at a time — a ceiling obtained by cheating, never a strength figure — prices teammates’ cards at +3.41 against opponents’ +1.31, and we had predicted the reverse. Exact best responses over information sets measure the deployed policy’s endgame exploitability at +0.4405 over 305 positions with one half-suit live, against a deterministic realisation of a policy that mixes. Under the corrected misdeclaration rule (the engine scored it wrongly for five versions), the deployed configuration beats the v0.3 champion by **+2.412** [+2.088, +2.736] sets per pair.

An earlier version of this abstract also reported +2.347 sets per game against an independently developed C++ engine (Nguyen, 2026b) over 10,000 games. That figure is withdrawn. It was an artifact of our own bridge, which is stateless: it replayed the whole public log into a freshly reset agent at every decision, where their arbiter keeps one agent per seat alive for the deal. The two agree exactly on their scripted baselines (0 of 661 decisions) and diverge on every release from v0.4, where their fitted belief arrives, by up to 40.8% of decisions. Priced under a paired design that changes nothing but the transport, the bridge was worth **+3.18** [+3.01, +3.34] sets a game *to us* — more than the entire margin it produced — and their declaration errors fall tenfold, from 0.880 a game to 0.092, once their engine is allowed to keep the state its own arbiter lets it keep. Corrected, KRAKEN **loses** to that engine by -0.70 ; measured independently inside their arbiter through their own bot-package protocol, by -0.62 . Two routes sharing no host, rules or code agree to within 0.08 where the published one is out by three. An absolute measured through a bridge is a statement about the bridge as well as about the engines, which this paper had written down in an appendix and not applied to its own headline.

Version 1.1 took seven pre-registered directions at the residue and shipped none: four were stopped by a futility screen fixed in advance, one is closed by a proof that the belief-space search cannot move 99.96% of the allocation uncertainty, and an audit of our own clustering restates every position-level interval and withdraws one conclusion. Across the whole project, forty-six pre-registrations produced exactly one change to the deployed engine (Appendix A). Two mechanisms measured as real and positive do not ship, because the bar was fixed before the runs and was not amended afterwards.

What survives the correction is a single channel — we ask 1.95 points worse — and locating it **withdrew** a candidate rather than producing one. Under the no-bluff rule a half-suit admits no landable ask exactly when our own team already holds all six of it, so an ask that cannot land is an ask into a half-suit we have already won and not yet declared. We make **13.79%** of our asks there against their 9.81%, and the difference is a declaration *latency* — 16.80 plies from completion to declaration against their 7.84, concentrated in a third of completions — and not a failure of ask

selection: run at their seats, the same measurement says our ask objective avoids those half-suits *better* than theirs does (-0.0101 against $+0.0198$ relative to a random draw from the same menu). Priced at this project’s own measured value of a turn, the channel is worth $+0.458$ [$+0.157, +0.793$] sets a game against a deficit of -0.5250 . The $p(\text{dead})$ term this analysis was written to justify is refuted by its own kill-check. Registered and run anyway in its strongest form (P44), it is a significant **negative** in both populations, -0.3233 and -0.2300 , *while raising the ask hit rate it targets* (0.5245 against 0.5184 on the same deals): under the no-bluff rule a failed ask publicly proves the asker holds another card of that half-suit, so an ask into a half-suit we already own is the engine’s own way of resolving ownership, not waste. The finding replicated on an independent channel (P45): a free-message convention gate raised the hit rate again, 0.5287 against 0.5136 , and lost -0.3067 and -0.8567 . The ownership inference the latency was then blamed on is measured and is not the constraint; the split is. **A measured deficit in a metric is not necessarily a deficit to close.** Nothing ships; there is no v1.2.

Keywords: imperfect-information games; exact inference; opponent modelling; endgame solving.

Contents		10 The shape of the game	58
1 Introduction	2	11 The endgame solved under imperfect information	73
2 Related work	5	12 Discussion	102
3 The game	6	13 Conclusions	113
4 Exact combinatorial inference	7	A The pre-registered thresholds, verbatim	117
5 Exact solving, and a closed form	12	B Corrections this revision makes to its own record	131
6 KRAKEN: the decision procedure	14	C Interoperating with a foreign engine	133
7 Evaluation methodology	19	D Reproducing this work	135
8 Results: what the inference, the objective and the ladder are worth	23	E The champion, and the deployed configuration	136
9 What did not work, what search buys, and what a learner recovers	41	F Information-boundary tests	137

1 Introduction

Literature, also called Fish or Canadian Fish, is a six-player, two-team, imperfect-information card game. Players interrogate opponents for specific cards and score by *declaring* the exact distribution of a six-card half-suit among the members of their own team. It sits in a corner of the game-AI landscape that is unusually poorly served: it is a team game with hidden information and *no legal communication channel*, so neither the two-player zero-sum machinery that solved poker nor the fully cooperative machinery developed for Hanabi applies directly.

This paper covers four iterations of an engine for the game, which is called **KRAKEN**. The bulk of it is v0.4, which built the inference and the evaluation discipline; §11.9 then reports v0.6, whose first act was to discover that this engine had been playing a house variant of the misdeclaration rule since v0.1 and whose remaining work was re-measuring, and in two cases withdrawing, everything that had depended on it. Version 1.0 is v0.6 plus one shipped change (§8.7) and a body of work that did not change the policy at all: it asked where the remaining errors are, and the answers redirected the project away from card-reading entirely.

A note on names, because the identifiers in this paper do not match the title. The engine was called *KV’s FishBot v0.6* until August 2026 and is called **KRAKEN** now, but the shipped configuration is still named `V06_DEPLOYED` in the code and `fishbot4` in the registry. Those identifiers are frozen deliberately: they are what the pre-registration documents in `prereg/` name as their arms, and a registration that fixes

its arms before the data is worth nothing if the arms can be renamed afterwards. The display name moved; the record did not. There is exactly one configuration either way.

The previous version, v0.3, established three results that frame everything here. First, exact belief tracking dominates: an agent that merely chains every logical implication of the public record beats a public-information heuristic by 11.86 sets per duplicate deal-pair. Second, and more surprisingly, search does not help, and the reason is measurable: the standard deviation of a position’s value across possible hidden layouts is about $2.4\times$ the gap between the best and the worst candidate move, so any search that evaluates different candidates against different sampled layouts is ranking luck. Third, acting on that diagnosis, by improving the *objective* rather than the depth of search, produced the study’s largest single gain.

v0.3 closed by naming its own bottlenecks, and this paper attacks them directly. Its belief sampler was documented as *not uniform* over the set of deals consistent with the public record, so every probability the engine reported inherited an unquantified bias. Its ask objective was a hand-weighted sum of three terms. Its claim policy was a fixed confidence threshold. The obvious research programme is therefore: make the inference exact, widen and then learn the objective, and make claiming decision-theoretic.

We carried that programme out, and the headline result is not the one we expected.

1.1 Three senses of “strongest”, kept separate throughout

This paper reports margins against several opponents, and they support different claims. Stating which before any number, so that none of them is read as another:

Lineage strength

The deployed configuration against this project’s own earlier champions, on duplicate deals under one rule dialect. This is established: $+2.412 [+2.088, +2.736]$ sets per pair against the v0.3 champion.

Cross-engine strength

The same configuration against a frozen build of an independently developed engine (Nguyen, 2026b), through a process bridge, under *this* paper’s rule dialect rather than that engine’s. This is measured but narrower than it looks, and narrower than earlier revisions of this paper said. It is one opponent rather than a population; the two dialects differ on out-of-turn declaration, which that project’s own sweep prices at about half a percentage point of its edge; and, newly, §11.10 reports the same pairing measured in *their* arbiter, where the sign reverses. Both numbers are real and neither is a ranking. What separates them is not settled, and the component it turns on is the majority of the margin.

Global standing

How good this engine is at Fish. **Not claimed.** No equilibrium is computed, no upper bound on exploitability exists for any class, and the phrase “near-optimal” does not appear in this paper about any configuration of it. The exact best responses of §11 are lower bounds against a deterministic realisation of a policy that mixes.

1.2 Research questions

RQ1 What is exact hidden-state inference worth in play, and what must a posterior model beyond legality to be worth anything?

RQ2 Where does the residual loss sit — in finding the opponents’ cards, or in allocating our own team’s?

RQ3 Can search over the belief, or a learned ask objective, recover that residue?

RQ4 How exploitable is the deployed policy where the game can be solved exactly, and does a correction fitted against exact values transfer to play?

RQ1 is answered in §8.1 and §8.2, RQ2 in §8.7, RQ3 in §9.3.1 and §9.4.1, and RQ4 in §11 and §11.9. RQ1 and RQ3 return negative answers, RQ2 returns the finding that reframed the project, and RQ4 returns a correction that measures well and does not survive the rule change.

1.3 Contributions

1. **Exact combinatorial inference.** We show that hidden-state inference in Literature reduces to counting assignments in a degree-constrained bipartite system, and that although the general problem is $\#P$ -complete, the specific instances arising in this game are tiny in the dimension that matters: the number of *distinct candidate sets* is on average 4.1 even when 25.1 cards are unlocated. We give an exact dynamic program over that structure that returns the partition function, exact marginals, and exact uniform sampling, and validate it against exhaustive enumeration (Section 4).
2. **A measured account of the disjunctive constraints.** The constraint set contains clauses (“the asker held at least one card of that half-suit”) that the dynamic program does not absorb. We quantify three ways of handling them and show why the obvious two fail, then give an importance sampler whose proposal density is computable in closed form and whose self-normalised weights are consistent for the target (Section 4.4).
3. **Posterior quality measured against ground truth.** Every belief metric in the previous version was indirect. Because the hidden hands exist in the simulator, we can score a posterior directly on how much probability it places on where the cards actually are. We report negative log-likelihood, Brier score, top-1 accuracy and calibration for every inference configuration (Section 4.6).
4. **The central negative-then-positive result.** Making the posterior *exactly* correct under the standard assumption, that players’ choices carry no information beyond legality, made it measurably *worse* at predicting reality than the biased sampler it replaced, and produced no improvement in play. The reason is identifiable: the old sampler’s bias was accidentally informative, encoding a crude opponent model. Replacing the accident with an explicit, single-parameter model of how players choose which half-suit to ask in recovers the loss and then beats both (Sections 4.5 and 8).
5. **A wider, ablatable ask objective,** including two terms that the literature suggests are unstudied: the *exposure* of one’s own already- public cards to the opponent who would receive the turn on a failed ask, and the *signalling* value of an ask, since the choice of half-suit is the only legal communication channel in the game and is simultaneously read by the opposition (Section 6.1).
6. **Methodological corrections** to the previous evaluation harness, including a bias in how timed-out games were handled that we demonstrate can put the true value outside a reported 95% confidence interval (Section 7); a measurement showing the previous version’s experiment cells could not resolve effects below about 0.44 sets per pair, so several of its reported nulls were uninformative rather than negative; and a worked example of a screening study manufacturing a false positive, caught by replication (Section 9.2).
7. **A widened exact solver,** which is the source of two structural facts about the game: the number of live cards is always exactly six times the number of unresolved half-suits, and the perfect-information value has the closed form $\text{sign}(\text{mover}) \times (2f - m)$. The second means the previous version’s headline absolute result, 100% agreement with optimal play in information-resolved positions, was a shallower certificate than it read, and we reproduce it on a corpus that is 79% two-half-suit (Section 5).
8. **The search the variance diagnosis leaves open, and why its obvious form cannot work.** Every search this project has tried samples layouts, and the measured spread across them is why they all failed. A lookahead over the posterior instead is deterministic given the belief, so that variance is structurally absent. We show (Proposition 3) that the natural version of it, one whose transition edits only the asked card’s row, is *exactly* greedy at every depth by an exchange argument, so it is a no-op by construction rather than by measurement, and only the quota coupling between candidates can make such a search non-trivial. Whether the coupled version helps is still open, and the attempt to resolve it produced a methodological result we consider more valuable than the search would have been: a direct measurement, over 24 blocks of an A/A, that this harness’s duplicate-deal intervals cover at their nominal rate and carry no between-run variance component, which required first making the harness able to observe its own null at all (Sections 6.3 and 9.3.1).

2 Related work

2.1 Fish and Literature

We could find no peer-reviewed AI work on Literature, Canadian Fish or Russian Fish. Several open-source implementations with bots exist; the best documented uses Q-learning on a four-player variant with a first-order knowledge encoding. The game family did recently enter a standardised testbed: [Goadrich et al. \(2026\)](#) include Go Fish, and report that “Go Fish and Crazy Eights show little separation between MCTS and random play, likely because CardStock’s determinizations account only for visibility and not for information inferred from action history.” That is a published negative result about exactly the failure mode this paper attacks from the other side — they identify that action history carries information their determinizations ignore; we quantify how much, and put it in.

2.2 The concurrent engine

The one serious exception is [Nguyen \(2026a\)](#): an independently developed C++ engine for six-player, 54-card Literature, with its own seven-version paper series — the engine this paper’s exhibition and foreign-opponent checks play against (§11.9). Read after those checks ran, its documentation turns out to contain this paper’s central finding, arrived at independently on a different architecture: its *exact* rules-posterior is the worst predictor of where cards sit (NLL 1.422 against 1.382 for its behaviour-aware approximation), and substituting the exact object into the same policy loses in play — their phrasing, “exactness and behaviour-awareness are not available in the same object here”, is our Section 8.2 in different words. The convergences run further: their first-order behaviour prior (an ask-tally over every seat, teammates included) is this paper’s choice model with two parameters in place of one; their ungarded determinized search lost 35.7 points to its own blueprint, our sampled searches all lost to their prior (§6.3); and our own deception ladder had to be withdrawn wholesale. An earlier draft added that their “feint” deception feature cost -2.2 on replication. That is withdrawn: in [Nguyen \(2026b\)](#), feint is a near-class member of their opponent *panel*, which their engine beats by $+6.07$ percentage points, and no such figure appears. Where the projects genuinely differ — their deployed belief is an approximate Sinkhorn object seeded by the prior where ours is exact counting reweighted by the model; their declarations price by a learned state value where ours declare at a fixed certainty bar — the disagreements are now measurable across a real engine boundary, which neither project could do alone.

2.3 Determinization, and why double-dummy analysis fails here

Evaluating an imperfect-information position by sampling consistent worlds and solving each one perfectly is the standard technique, and its two failure modes are named: *strategy fusion*, where the searcher implicitly assumes it can act differently in states it cannot distinguish, and *non-locality* ([Frank and Basin, 1998](#)). [Long et al. \(2010\)](#) give the conditions under which it nonetheless works, parameterised by leaf correlation, bias, and a disambiguation factor. Two consequences shape this work. First, our disambiguation factor is unusually high (every ask is a hard constraint on the deal), which is favourable. Second, and decisively, the CLAIM action is a strategy-fusion trap: inside a determinization the searcher knows the true layout, so every claim evaluates as a certain gain and the declared assignment differs between sampled worlds. Claims in v0.4 are therefore never selected by search. [Cowling et al. \(2012\)](#) introduce Information Set MCTS and state plainly that they do not consider belief distributions; v0.3’s own experiments found ISMCTS losing badly to its own prior, which is consistent with both that and with the variance diagnosis it measured.

2.4 Policy-aware inference

The closest published line is [Solinas et al. \(2019\)](#) and [Rebstock et al. \(2019\)](#), who infer hidden cards in trick-taking games from the actions players took, using a learned policy as the likelihood. Their factorised belief does not enforce hand-size constraints; they note the constraint could be added and do not add it. Our setting lets us do both at once: the quota constraints are handled exactly by the dynamic program of Section 4, and the choice likelihood enters as an importance weight on top. [Solinas et al. \(2023\)](#) prove

that exhaustive history filtering is polynomial only when the public tree is sparse; Literature’s is dense, so their route does not port, but the count abstraction makes the constraint side tractable anyway.

2.5 Counting and sampling with fixed margins

The combinatorial problem underneath is counting bipartite degree-constrained subgraphs, equivalently a restricted permanent, $\#P$ -complete in general (Valiant, 1979), with a celebrated approximation scheme (Jerrum et al., 2004) and a literature on sequential importance sampling for contingency tables (Chen et al., 2005) — along with cautionary results showing SIS can be exponentially bad on some instances (Bezáková et al., 2012). We need none of that machinery, because the instances here are small in the dimension that matters. The constraint-programming counterpart is Régis’s global cardinality constraint (Régis, 1996), which achieves arc consistency by a flow argument; our dynamic program is stronger in that it also *counts*, which is what claim probabilities require.

2.6 Team games and common information

Literature is a two-team zero-sum game with pre-play agreement and no in-play communication, which is exactly Celli and Gatti’s regime; the solution concept is a team maximin equilibrium with correlation, not a Nash equilibrium (von Stengel and Koller, 1997; Celli and Gatti, 2018). The convergence guarantees that underwrite two-player zero-sum solvers do not transfer. Hanabi (Bard et al., 2020) is the closest studied game in spirit, and one of its methodological contributions transfers cleanly: cross-play between independently trained partners as the instrument for detecting conventions that work only against copies of oneself (Hu et al., 2020). What does not transfer is the remedy, because Hanabi is fully cooperative while Literature’s only communication channel (the choice of which half-suit to ask in) is *costly* and read by the opposition as well as by one’s partner. We found no published treatment of that combination, and it is the motivation for the *signal* term of Section 6.1.

3 The game

3.1 Rules

Six players sit in a circle in alternating team order: seats $\{0, 2, 4\}$ against $\{1, 3, 5\}$. The deck is partitioned into *half-suits* of six cards each. Our primary variant uses 54 cards and nine half-suits: the eight natural half-suits (2–7 and 9–A of each suit) plus a ninth composed of the four 8s and two individually identified jokers, a Red Joker and a Black Joker. Each player is dealt nine cards. The classic 48-card, eight-half-suit variant is also supported.

On turn, a player either **asks** or **declares**.

An **ask** names an opponent and a specific card. It is legal only if the asker holds at least one card of that card’s half-suit and does not hold the card itself. If the opponent has it, the card transfers *face up* and the asker keeps the turn; otherwise the turn passes to the opponent who was asked.

A **declaration** names a half-suit and states which teammate holds each of its six cards. (The codebase, and earlier versions of this paper, call this move a *claim*; the two words name the same move throughout.) If the declaration is exactly right, the declaring team scores the half-suit. *Any* wrong declaration — a card with an opponent, or all six within the declaring team but the split misstated — awards the half-suit to the opposing team. The requirement to name the exact distribution, not merely to assert possession, is what makes Literature a game about inference rather than collection.

3.2 The rule dialect studied

The rule this engine got wrong for five versions. From v0.1 through the v0.4 measurement campaign, the engine’s default scored the second kind of miss — right team, wrong split — as a *void*: the half-suit was discarded and nobody scored it (`wrong_distribution_outcome="null"`). That is a house variant, not the game: the reference rules, and the one foreign engine we test against, award every misdeclared set to the opponents. The default is now the award rule; the void variant remains implemented and pinned wherever a void-era measurement is reproduced. Every result in this paper is therefore labelled

with its rule era: analyses proved rule-invariant (the closed form, the perfect-information solver and tablebases, whose declarations are certain by construction; posterior quality, which scores predictions rather than declarations) carry over unchanged, in-play measurements from the void era are marked as such, and the numbers the abstract and strategy sections assert are re-measured under the award rule per the pre-registration `prereg/rules_award_baseline.md`. The correction is not cosmetic in play: in the exhibition against Dylan’s FishBot v0.7 the rule touches 466 sets in 1,000 games — one set every other game whose fate the two rules disagree about.

A player with no cards on move passes the turn to a teammate who has cards. If an entire team is out of cards, the other team must claim out the remaining half-suits without asking.

3.3 Termination, and the solution concept

Remark 1 (Termination is a property of the players, not the rules). *Nothing in the rules forces progress: two opponents can pass a card back and forth and return to an identical position, so the state graph contains cycles. Consequently any claim that “every game ends” describes an agent’s stalling rule rather than Literature itself. Our agents carry an explicit progress rule and our harness imposes an action cap; Section 7 explains why the cap must be scored rather than discarded.*

Remark 2 (Solution concept). *Literature is a two-team zero-sum game with pre-play agreement and no in-play communication. The appropriate solution concept is therefore a team-maximin equilibrium with correlation (TMECor) in the sense of Celli and Gatti (2018), not a Nash equilibrium, and the convergence guarantees available for two-player zero-sum games do not transfer. We compute no equilibrium and make no equilibrium claim; every strength statement in this paper is either relative (against a named opponent) or absolute against an exactly solved subgame.*

3.4 Information integrity

An imperfect-information study is worthless if the agents can see the hidden state, and the failure mode is usually subtle rather than flagrant. The boundary is enforced four ways.

Structural. The engine’s GameState object never reaches a policy at all. Policies consume only an Observation: their own hand, the public event log, public hand counts, resolved half-suits, whose turn it is, and the rules.

Proved by reconstruction. Observation.reconstruct rebuilds a seat’s entire view from nothing but (rules, that seat’s initial hand, the public event log). Tests assert the engine-provided observation is *identical* to that reconstruction at every step of full games, for all six seats. In an audit conducted for this paper the property was re-verified over 183,750 seat-plies across eight rule configurations with zero mismatches.

Proved by invariance. Every quantity v0.4 computes — posterior marginals, all eleven ask-objective terms, and the end-to-end policy action — is asserted to be bit-identical when the hidden cards of the other five seats are permuted into a different layout that is equally consistent with the public record. Candidate permutations are drawn from the acting seat’s own belief sampler and then validated by an *independent* replay of the entire public history, so the test does not merely assume the sampler is correct.

Separation of training from acting. Ground truth may be used to *score* (Section 4.6 scores posteriors against the true hands) and to *train*; no acting policy ever sees it. That line is what makes the distinction meaningful rather than decorative.

4 Exact combinatorial inference

4.1 The inference problem, stated exactly

The observation that shapes the whole engine is that *every card movement in Literature is public*. A successful ask names the card that moved; a claim reveals all six locations; nothing changes hands unobserved. It follows that the current location of every card is a deterministic function of the *initial*

deal and the public log, so all hidden-state inference reduces to constraints on the initial deal. Those constraints are of exactly three kinds:

1. **Candidate sets.** Per card c , a set $S_c \subseteq \{0, \dots, 5\}$ of players who may have been dealt it. A failed ask excludes the target; a successful ask pins the card; asking for a card excludes the asker (under standard no-bluff rules); a claim pins all six.
2. **Exact counts.** Every player was dealt exactly q cards (9 or 8); after propagation, player p has a residual quota q_p of the cards that remain unlocated.
3. **Disjunctions.** A legal ask certifies that the asker held *at least one* card of that half-suit at that moment: an OR-constraint over the cards of the half-suit not yet publicly located.

Write F for the set of *free* cards (those not publicly located and not pinned by propagation). A *world* is a map $f : F \rightarrow \{0, \dots, 5\}$ with

$$f(c) \in S_c \quad \forall c \in F, \quad |f^{-1}(p)| = q_p \quad \forall p, \quad (1)$$

that additionally satisfies every OR-constraint. Let Ω denote the set of such worlds and $\Omega_0 \supseteq \Omega$ the set satisfying only (1).

Proposition 1 (Uniform posterior under uninformative choice). *Suppose the deal is uniform over all $\binom{54}{9,9,9,9,9,9}$ partitions, and suppose that at every decision point each player’s choice of action depends on their private hand only through which actions are legal. Then the posterior over worlds given the public history is uniform on Ω .*

Proof sketch. Let h denote the public history and w a world. The prior is uniform, so $\mathbb{P}(w \mid h) \propto \mathbb{P}(h \mid w)$. Factor $\mathbb{P}(h \mid w)$ over the actions in h . Under the stated assumption each factor depends on w only through the legality of the taken action, and card transfers are deterministic given w and the action. Hence $\mathbb{P}(h \mid w)$ takes the same value for every w that makes all of h legal and reproduces every observed outcome, which is precisely Ω , and is 0 otherwise. $\square$

Proposition 1 is the target that v0.3’s sampler was *intended* to hit and, by its own documentation, missed. It is also, as Section 4.6 shows, not the target one should want, because its hypothesis is false, and interestingly so.

4.2 Why the counting is tractable

Counting the elements of Ω_0 is an instance of counting perfect matchings in a degree-constrained bipartite graph, equivalently a restricted permanent, which is $\#P$ -complete in general (Valiant, 1979). The instances arising here, however, are small in the only dimension that matters. Measured over the 360 decision points of the stored corpus (`results/infer_position_stats.json`):

quantity	mean	median	p90	max
free cards $ F $	25.1	25	42	45
distinct candidate sets	4.1	4	6	9
active OR-constraints	4.7	5	9	14
cards per OR-constraint	3.17	3	5	5

Cards sharing a candidate set are *exchangeable*. Grouping them collapses the problem from an assignment over $|F| \leq 45$ labelled cards onto a *group-count matrix* $k \in \mathbb{N}^{G \times 6}$, where k_{gp} is the number of cards of group g dealt to player p . G is small: at most 9 across those 360 positions, with a 99th percentile of 8. That is an *observation about this game*, not a bound — nothing here proves G cannot be larger, and the dynamic program does not assume it is not.

4.3 The dynamic program

Let group g have candidate set S_g and size n_g , so $\sum_g n_g = |F|$. The number of worlds realising a given k is a product of multinomials, hence

$$Z = |\Omega_0| = \sum_{\substack{k: \sum_p k_{gp} = n_g \forall g \\ \sum_g k_{gp} = q_p \forall p, \quad k_{gp} = 0 \text{ if } p \notin S_g}} \prod_{g=1}^G \frac{n_g!}{\prod_p k_{gp}!}. \quad (2)$$

We evaluate (2) by dynamic programming over *players*, carrying as state the vector $r \in \prod_g \{0, \dots, n_g\}$ of cards still unassigned in each group. Define $\Phi_0(r) = \mathbf{1}[r = (n_1, \dots, n_G)]$ and

$$\Phi_{p+1}(r') = \sum_{\substack{\kappa: \sum_{g \in A_p} \kappa_g = q_p \\ r' = r - \kappa, \quad \kappa_g \leq r_g}} \Phi_p(r) \prod_{g \in A_p} \binom{r_g}{\kappa_g}, \quad A_p = \{g : p \in S_g\}, \quad (3)$$

so that $Z = \Phi_6(0)$. The product of binomials telescopes across the six player steps to exactly the multinomial in (2), so the recursion is exact rather than approximate.

Two implementation notes make (3) fast. First, the constraint $\sum_g \kappa_g = q_p$ is enforced with an auxiliary axis counting cards taken so far, so the compositions κ are never enumerated: player p 's step is $O(\sum_{g \in A_p} \min(n_g, q_p))$ array operations. Second, the state space $\prod_g (n_g + 1)$ is small in practice, because G is small and not because $|F|$ is. An earlier draft put it at “median 272, mean 577 over real positions”; those figures back nothing in `results/infer_position_stats.json`, which stored the *quota* lattice $\prod_p (q_p + 1)$ under a similar name and never the group product at all. The group product is now measured alongside it, and the two are different quantities rather than two estimates of one.

What the dynamic program yields.

- **Partition function** $Z = |\Omega_0|$, i.e. the exact size of the information set, computed in closed form. This makes quantities that are normally estimated (such as the disambiguation factor of Long et al., 2010) directly computable.
- **Exact marginals.** A backward pass B_p mirroring (3) gives $\mathbb{E}[k_{gp}]$ by a forward-backward sweep that carries one accumulator per group, so all $6G$ expectations come out of a single pass. Since cards within a group are exchangeable, $\mathbb{P}(f(c) = p) = \mathbb{E}[k_{g(c)p}] / n_{g(c)}$.
- **Exact uniform sampling.** With B available, each player's whole take is drawn from its exact conditional, so a draw is exactly uniform on Ω_0 .
- **Exact joint probabilities.** Pinning named cards reduces n_g and q_p by one each, so the probability that a declared claim distribution is exactly right is a ratio of two evaluations of (2). Because cards are labelled throughout, no extra combinatorial factor is needed.

Validation. On 400 randomly generated systems the dynamic program reproduced the brute-force count *and* every per-card marginal to within 10^{-9} , 400/400; the pinned-count routine matched brute force on 200/200. An exact uniform sampler drawn 40,000 times matched the exact marginals to 0.0000 at four decimal places.

Remark 3 (Completeness). *Because a marginal of zero is a proof of impossibility, the dynamic program is a complete propagator for the candidate-set-plus-quota subsystem: any placement it assigns probability zero is genuinely impossible, and any placement it assigns probability one is genuinely forced. v0.3's propagator was explicitly documented as sound but not complete for this subsystem. The disjunctive constraints remain outside this guarantee.*

4.4 The disjunctive constraints

The OR-constraints are the part of the system (1) that the dynamic program does not absorb. Three routes are available, and we measured all three rather than choosing on intuition.

(a) **Ignore them.** On 40 real positions, exact marginals computed on Ω_0 differed from the truth on Ω by a mean L_1 error of 0.127 per card — *worse* than v0.3’s biased sampler at 0.074. So the clauses carry a great deal of information and cannot be dropped.

(b) **Rejection.** Draws that are exactly uniform on Ω_0 can be filtered to be exactly uniform on Ω . A typical clause covers 3.35 live cards, each landing with the relevant player about one time in five, so a clause holds with probability roughly 0.53 and five independent clauses hold about 4% of the time. Measured acceptance on real positions: mean 0.28, median 0.17, and below 0.05 on 16% of positions. Rejection is therefore exact but unaffordable.

(c) **Fold them into the dynamic program.** Refining the groups by OR-membership keeps the cards within a block exchangeable and makes each clause a local condition on one player’s step, which can be enforced by inclusion–exclusion over that player’s clauses. The refined state space, however, has a bad tail: median 5,760 but p90 230,000 and p99 2.7×10^7 . Exact in the median case, hopeless in the tail.

4.4.1 Sequential importance sampling with closed-form densities

We therefore sample. Cards are assigned in a fixed order, so each world has exactly one generating path and its proposal density $q(w)$ is the product of the per-card choice probabilities. The proposal is quota-proportional (the natural approximation to uniform) multiplied by a boost for choices that would satisfy an as-yet-unsatisfied clause, and forced when a clause has a single live card left. Self-normalised importance weights

$$\tilde{w}_i \propto \frac{L(w_i)}{q(w_i)}, \quad \hat{\mathbb{E}}[\varphi] = \frac{\sum_i \tilde{w}_i \varphi(w_i)}{\sum_i \tilde{w}_i} \quad (4)$$

are then consistent for any target $\propto L$; with $L \equiv 1$ the target is the uniform posterior of Proposition 1.

Two properties are required and both are verified against exhaustive enumeration rather than argued:

- **Support completeness.** Every world in Ω must have $q(w) > 0$, since a weight cannot repair mass the sampler never reaches. We check this *analytically* for every world of every small system (random hit-testing cannot distinguish “unreachable” from “rare”) over 55 systems and 1,797 worlds, all reachable.
- **Density correctness.** The density used to build the weights must be the density the sampler realises. Analytic and empirical path densities agree to a worst gap of 0.0097 over 20 systems.

Weighted marginals then matched exhaustive enumeration to a worst error of 0.029 at 4,000 draws, while the *unweighted* estimator over the same draws was measurably biased — a control without which the test would not be known to discriminate. The Kish effective sample size is $\text{ESS}/N = 0.52$ on real positions (Section 8.3), so the cost of steering the proposal is substantial and, more importantly, is reported rather than assumed. An earlier draft put it at 0.81 here while Section 8.3 put it at 0.52; the 0.52 is what `results/ess_probe.json` measures over 641 real decisions, and the two sentences had disagreed with each other for as long as both existed. Drawing the batch one card at a time across all draws in lockstep, rather than one draw at a time, makes this $6.8\times$ faster at the median for identical output ($11.5\times$ on average, over 60 positions); before that change the scalar loop was 67% of the policy’s total runtime.

4.4.2 A correction: option (c) is in fact tractable

Our estimate above that folding the clauses into the dynamic program is unaffordable was measured on the wrong decomposition, and a separate study overturned it. Refining by *half-suit* is coarser than necessary: what actually needs to stay together is a connected component of clauses that share cards, and OR-free cards can be processed one at a time. With that decomposition and a dynamic program over the residual-quota lattice, an exact backend refused none of 349 real positions at a work budget of 10^7 , costing a median of 7.5 ms and a p90 of 17.5 ms.

That result is worth stating in its own right, because it inverts the usual shape of an accuracy-versus-compute study: *the frontier here is not a curve, it is a point*. The exact backend sits below and to the left

of every sampled one, and there is no compute budget above roughly 8 ms at which a sampler is the right answer. Below that budget the best available option is a hybrid — exact where it fits, sampled where it does not, which still beats every pure sampler at matched cost (L_1 error 0.133 at 5.5 ms against 0.258 for v0.3’s sampler at 3.2 ms).

A mode that promised exactness it could not deliver. The shipped posterior takes the counting DP where no OR clause is active and the importance sampler everywhere else, and reports which it used. It also accepted `mode="exact"`, which forced the DP *whatever the clause set looked like* and then set `Posterior.exact = True` for the result. The DP does not represent OR clauses (that is the entire reason this module contains a sampler), so its draws come from a strict superset of the feasible worlds. At the first position of one seeded game where a clause bites, 30% of those draws are worlds the belief state has already *proved* impossible. The mode a caller would reach for when they wanted a guaranteed-correct answer was the one mode that could quietly return a wrong one under a correct label. It now refuses. Nothing in the shipped agent used it, which is why it survived: a trap with no current caller is invisible to every test that exercises the paths that are taken.

We nevertheless ship the importance sampler, and the reason is the subject of Section 4.5: the exact backend computes the *uniform* posterior, and the uniform posterior is not the target we want. What the measurements in Section 9.1 add is that the difference between an exact posterior and a well-sampled one is worth nothing measurable in play, while the difference between the uniform target and an opponent-model target is worth 1.9 sets per deal-pair. Exactness was the cheaper thing to buy and the less valuable one.

The natural synthesis is available and untried: the choice model (6) is a product of per-(player, half-suit) depth factors, and a depth is a function of one half-suit’s count vector. So if the items of the dynamic program are taken to be whole half-suits rather than clause components, the likelihood factorises over exactly the same structure as the constraints do, and exact inference *under the opponent model* costs the same dynamic program with reweighted item tables. We did not build it, because the evidence above says it would buy approximately nothing in strength; we state it because it is the right thing to build if that evidence ever changes.

4.5 Opponent modelling: relaxing the false hypothesis

Proposition 1 rests on a hypothesis that is plainly false: that a player’s choice of action depends on their hand only through legality. It does not. A player who asks in a half-suit is more likely to be deep in it.

Under the simplest defensible choice model — a player picks which half-suit to ask in in proportion to how many cards of it they hold —

$$\mathbb{P}(\text{ask in } H \mid \text{hand } h) = \frac{d_H(h)}{|h|}, \quad (5)$$

and the denominator is public, hence identical across worlds. The likelihood of the whole observed ask sequence is therefore proportional to a product of depths, giving the importance weight

$$\log L(w) = \gamma \sum_{(a,H)} n_{aH} \log(d_{aH}(w)), \quad (6)$$

where n_{aH} counts asks by player a in half-suit H and $d_{aH}(w)$ is a ’s depth in H under world w . The single parameter γ controls how sharply the model is believed; $\gamma = 0$ recovers Proposition 1 exactly, so the entire idea is one ablatable number. Because (6) decomposes over cards, it costs one integer increment per card inside a draw.

Remark 4 (The obvious refinement does not help). *Depth in (6) is evaluated on the initial deal, whereas the choice model is really about what the player held at the moment they asked. Cards move, so the two differ. We implemented the at-ask-time version, which needs only one replay of the public log per decision, not per candidate world, because a card the propagator has not pinned was never located at any time and so is exactly the card the sampler draws, and it is not better. Posterior negative log-likelihood is*

1.3536 against 1.3522 for the initial-deal version at matched settings, and 1.3391 against 1.3373 at the best settings of each. In duplicate-deal play at matched $\gamma = 0.35$ the two are likewise indistinguishable (-0.050 , $[-0.377, +0.277]$ over 400 pairs).

That is not the end of the story, and Section 8.5 is where it goes: at $\gamma = 1.0$ the at-ask covariate is demonstrated, $+0.102$ $[+0.006, +0.197]$ over 6000 pre-registered pairs. The two statements are compatible and the difference between them is the finding. What is demonstrated is the pair (at-ask depth together with a sharper γ), and neither component is claimed on its own, because the pre-registration bound them together on the argument that a covariate measuring the right thing deserves to be believed more sharply. Fixing γ at the incumbent 0.35 and swapping only the covariate is the comparison this remark makes, and it buys nothing. The theoretically-correct quantity buys nothing over the simpler one, so the simpler one ships.

We flag one methodological point here because it nearly became a false finding. Our first at-ask-time implementation counted only cards that had visibly changed hands, silently omitting every card the propagator had deduced to be someone’s — the large majority. Depths came out far too small, and the variant measured as catastrophic on both instruments (posterior NLL 2.22 against 1.35; -1.54 sets per deal-pair in play). Two independent instruments agreeing on a large effect is normally reassuring; here they were agreeing about a bug. What caught it was that the magnitude was implausible for a refinement, not that either measurement looked wrong.

4.6 Scoring a posterior against the truth

Every belief metric in v0.3 was indirect: sampler cost, agreement with an exact endgame solver, or downstream win rate. In simulation the hidden hands are available, so a posterior can be scored on the question it is actually trying to answer: how much probability does it place on where the cards really are?

We score over cards that are still genuinely uncertain for the acting seat — a card the propagator has already pinned is scored perfectly by every configuration and would only dilute the comparison — using mean negative log-likelihood of the true holder, the six-way Brier score, top-1 accuracy, and a reliability curve in deciles. Reliability must bin the probability assigned to a *fixed* outcome; binning by the probability assigned to the true holder would make every entry a hit by construction and report perfect calibration for any model whatsoever.

5 Exact solving, and a closed form

The exact solver is the project’s only source of *absolute* truth: in a position small enough to solve, there is a right answer, so an engine that disagrees has a defect rather than a preference. v0.3’s solver enumerated card placements and stopped at roughly seven live cards. Widening it produced three results that change how the previous version’s absolute benchmark should be read.

5.1 There is no such thing as a seven-card position

Claims strip all six cards of a half-suit from every hand, so

$$\text{live cards} = 6 \times (\text{unresolved half-suits}) \quad \text{exactly, always.} \quad (7)$$

Every “maximum live cards” threshold in the previous codebase was therefore a binary switch: any value from 6 to 11 selects the same policy. v0.3 solved the one-half-suit layer and refused everything else, and the two-half-suit layer is $6^{12} \times 6 \approx 1.3 \times 10^{10}$ card placements, which is why.

5.2 The count abstraction

Proposition 2 (Card identity is irrelevant under perfect information). *Let σ be any permutation of the deck that maps each half-suit onto itself. Then σ is an automorphism of the perfect-information game, and consequently the value of a position depends only on how many cards of each live half-suit each player holds.*

Proof sketch. σ acts on states by relabelling hands. $\text{ASK}(t, c)$ is legal in s iff $\text{ASK}(t, \sigma(c))$ is legal in $\sigma(s)$: the target and all hand counts are unchanged, “the asker holds a card of this half-suit” is preserved because σ fixes half-suits setwise, and “the asker does not hold c ” is preserved because σ is a bijection of each hand. The ask succeeds iff the target holds c iff $\sigma(\text{target})$ holds $\sigma(c)$, and the transfer commutes with σ . $\text{CLAIM}(h, a)$ maps to $\text{CLAIM}(h, a \circ \pi^{-1})$ with π the restriction of σ to h ; exactness, the “some card sits with an opponent” case and the wrong-distribution case are all preserved, and both states lose exactly the six cards of h . PASS and turn normalisation depend only on which hands are empty. Hence $V(\sigma(s)) = V(s)$, and two states have equal per-(half-suit, player) counts iff some such σ relates them. $\square$

A half-suit’s counts are a composition of 6 into 6 ordered parts, of which there are 462, so a layer with m unresolved half-suits has at most $462^m \times 6$ abstract states: 2,772 for $m = 1$ and 1,280,664 for $m = 2$, against 279,936 and 1.3×10^{10} card placements. Both are solved once, in full, as dense tables — a genuine Fish tablebase, after which a query is an array lookup. Measured on the same 300 one-half-suit positions, and this is one cross-check rather than two: the old solver agreed with the new one on both the value and the complete optimal-action set in 300/300 cases, with no refusals, at 2,240 CPU-seconds against 0.052. That is **43,000** $\times$ by total CPU time and **10,900** $\times$ by median per position (990 ms against 91 μ s). An earlier draft reported “131 further positions” at 4,205 $\times$; there is no second corpus, and 4,205 is not a ratio anything in `results/exact2_cross_check.json` produces.

5.3 A closed form for the perfect-information value

Both full tables match, exactly, the formula

$$V = \text{sign}(\text{mover's team}) \times (2f - m), \quad (8)$$

where m is the number of live half-suits and f the number of them in which the team on move holds at least one card.

This is not only an empirical match at the layers we can enumerate. It holds at every m , and the reason is that it turns on three properties of the *rules* rather than of any position, none of which needs a table:

- A. Footholds never grow.** `check_legal` refuses an ask unless the asker already holds a card of that half-suit, so a team holding none of h can never receive one. Its foothold set only shrinks, and the mover’s team can therefore claim at most the f it has *now*.
- B. A team with no foothold in h cannot deny h .** All six cards of h sit with the opponents, so every assignment it could submit reveals an opponent holder and `_apply_claim` awards the half-suit to the opponents — under both misdeclaration rules. Under the legacy null variant the further point was that no spite-null is available (a certain -1 cannot be converted into a 0); under the award baseline nulls do not exist at all, which only strengthens the premise. Hence the opponents take all $m - f$ whatever the mover does.
- C. The mover takes all f without surrendering the turn.** A hit retains the turn, a claim never moves it, and a claim that empties the claimant passes within the team. With perfect information every ask can be made to hit, so the mover drains and claims each of its f half-suits, and only when its team is cardless does the turn normalise to an opponent — who by symmetry takes the rest.

A and B bound V above by $2f - m$; C bounds it below. Each is checked against the engine rather than against a reading of it, and the optimal perfect-information policy that falls out is one line: *take a hitting ask if one exists; otherwise claim a half-suit your team wholly owns; otherwise pass to a teammate with cards*. Played from 360 random positions spanning $m = 1$ to $m = 9$ it realises $2f - m$ every time, and at $m = 1$ and $m = 2$ it agrees with the exact solver on every position tested — the one step in the argument anchored to a solver rather than to internal consistency.

5.4 The perfect-information game is trivial

The closed form is worth stating as a value because the value is degenerate. A team lacks a foothold in a half-suit only if all six of its cards landed in the opponents’ twenty-seven, with probability $\binom{27}{6}/\binom{54}{6} = 1.15\%$, so at a fresh deal

$$\mathbb{E}[V] = 2 \left(9 - 9 \binom{27}{6} / \binom{54}{6} \right) - 9 = +8.794 \quad (9)$$

of the nine available: **whoever moves first takes everything**.

Premise C is the constructive step and the one an adversary could attack, so it is worth testing against opponents that are trying to win rather than against the greedy rule itself. The premise is stronger than the proof needs: the mover never surrenders the turn, so the opponents do not act at all until its team is cardless, and only the *mover’s* team needs to know the deal. The oracle duel reported in §12.4 is that experiment. Across 80 games against the champion, in every one of the 40 where an oracle seat moves first its team’s realised set count equals its foothold count at the deal — 40/40 exactly, not on average. Mean footholds were 8.94 of nine against the 8.897 that (9) implies.

The other forty games are the sharper result. There a *champion* moves first, holding the advantage that (9) prices at +8.79 of nine, and converts essentially none of it: the oracle team still finishes with 8.95 of its 8.97 footholds. A first move that is decisive under perfect information is worth nothing at all to a player who cannot see the deal.

Two things follow for the project. First, extending the tablebase to larger m would buy no ground truth about Fish, because the closed form already answers every layer and the answer is a different game’s answer. Second, the inference ceiling reported in §12.4 should be read accordingly: its +17.483 per deal-pair was framed as what perfect card-reading is worth, and it is closer to (8) being played out — seven of its ten per-deal differentials are exactly $18 = 2 \times 9$.

This has an uncomfortable consequence for the previous version’s headline absolute result. v0.3 reported that its belief agents choose a provably optimal move in 100% of information-resolved positions, but every one of those positions had a single live half-suit, where (8) reduces to “the team on move wins it”. The score was real, and it is reproduced here, but it is a shallower certificate than it reads. The two-half-suit slice is the one with teeth, and the corpus used in Section 8.6.2 is 79% two-half-suit for that reason.

5.5 Bellman-optimal is not optimal in a loopy game

Fish’s state graph is cyclic, so the Bellman operator has many fixpoints and value-preservation is not optimality: if a side can force +1, every action keeping the value at +1 looks optimal, including one that walks a cycle forever and banks nothing. Measured on the full tables, in **28.6%** of decided two-half-suit states some value-preserving action makes no progress. At one half-suit the figure is 0%, which is why v0.3 never noticed. The solver therefore carries an attractor rank and the agent asks for a *progress-optimal* action, and Section 8.6.2 reports agreement against that stricter criterion.

Two further corrections to v0.3. First, although the state graph is cyclic, *optimal play always terminates* under perfect information: cycles exist, but no optimal line needs one. Second, the fixpoint reached by value iteration was checked to be unique in real Fish layers — eight different initialisations (zeros, $\pm m$, ± 5 , three random) and three sweep orders at $m = 1$ and $m = 2$ all land on the attractor value, so v0.3’s zero-initialisation and dictionary ordering were never load-bearing, even though in general they could have been.

6 KRAKEN: the decision procedure

6.1 The ask objective

6.1.1 A wider basis

v0.3 scored an ask as $\mathbb{P}(\text{success}) + 0.06 \cdot \text{depth}$, later adding a turn-risk and a scarcity term. Those terms are in incommensurable units — a probability, a card count, a normalised hand-size difference, a share, so the weights combining them have no meaning beyond “what won a sweep”. That is consistent with

term	origin	feature value
suit	v0.3	own depth in H (raw card count)
turn	v0.3	$(1 - \pi) \cdot$ negated normalised hand-size excess of the target
scarce	v0.3	$2(\bar{s}_H - \frac{1}{2})$, with $\bar{s}_H$ our team's expected share of H
reveal	v0.3	-1 if H has not previously been shown by us
deplete	v0.3	$\pi \cdot (1 - \text{target hand}/q)$
expose	new	$-(1 - \pi) \cdot$ our publicly located cards the target could take
claim	new	$\pi \cdot (\prod_{c \in H} \mathbb{P}(c \text{ with our team}))^{1/6}$
info	new	$4\pi(1 - \pi) \cdot$ holder entropy of H
certain	new	1 if $\pi = 1$
concent	new	largest single teammate share of our team's holding in H
signal	new	$2(\bar{s}_H - \frac{1}{2})$ if H has not previously been shown by us

Table 1: The ask objective basis. Every term is zero by default, so setting a weight to zero ablates exactly that idea and nothing else — a property enforced by test, for the reason given in Section 7.

what v0.3 measured: a clean inverted-U in each weight, which is the signature of a *tie-breaker* rather than an objective.

v0.4 keeps the additive form, because it is ablatable and because the previous version's evidence says the success probability really is dominant, but widens the basis to eleven terms and computes all of them from the posterior in a single vectorised pass. Writing π for $\mathbb{P}(\text{success})$ and H for the asked card's half-suit, the score is

$$\text{score}(a) = \pi + \sum_j w_j f_j(a), \quad (10)$$

with the terms in Table 1. Two deserve comment because they are, as far as the surveyed literature shows, unstudied.

Exposure. v0.3's turn-risk term penalised handing the turn to an opponent with many cards. Hand size is a proxy; the quantity that actually matters is how much that opponent can *immediately take from us*, which we can compute: cards of ours whose location is already *public* and that sit in a half-suit the opponent can legally ask in. Both parts are public knowledge, so using them is inference, not leakage.

Signalling. An ask publicly certifies that the asker holds a card of that half-suit. In a game with no legal communication channel, that is the only channel there is, and it is simultaneously read by the opposition. v0.3 modelled only the cost (its `reveal` term, worth a marginal +0.29). The benefit is that a teammate learns where a card of a half-suit our team is winning must be, which is exactly the localisation that lets a set be declared. So the sign should depend on who gains more, and the natural proxy is our team's share of the half-suit. Cost and benefit are carried as independent weights so that the data, not the story, decides.

6.1.2 An alternative in the units that are actually paid out

The additive form has a principled competitor. What the game pays is half-suits, so define, for each unresolved half-suit,

$$V(H) = \mathbb{P}(\text{our team ends up scoring } H) - \mathbb{P}(\text{they do}), \quad (11)$$

which is measured in sets, is bounded in $[-1, 1]$, and whose sum over half-suits is the final differential. An ask is then worth the expected change it produces,

$$\text{score}(a) = \pi [V(H \mid \text{success}) - V(H)] + (1 - \pi) [V(H \mid \text{failure}) - V(H)], \quad (12)$$

with every quantity in the same unit. The substantive difference is that this is a *marginal* value rather than a level: “our team is winning this half-suit” is a level, and what should drive a decision is how much one more card moves the outcome.

V is obtained by supervised learning rather than by guessing a functional form: self-play games supply, for every unresolved half-suit at every decision, a feature vector computable from the acting seat’s posterior alone, labelled with how that half-suit eventually resolved. That is roughly nine labels per decision, which is abundant and cheap, in contrast to the rollout regression that v0.3’s variance analysis made expensive. The counterfactual posteriors in (12) are local updates on the half-suit’s marginal block: a success pins the asked card to us, a failure removes the target from that card’s candidate set.

6.2 Declaring

v0.3 swept the claim-confidence threshold directly and found two things. First, claiming eagerly is a real, measurable mistake: dropping the bar from 0.97 to 0.70 costs 0.45 sets per deal-pair. Second, and more usefully, thresholds from 0.85 to near-certainty play *identically*: over 150 duplicate deals the 0.97 and 0.999 policies never once diverged.

The mechanism is worth stating, because it justifies the design rather than leaving 0.97 as a magic number. Suppose our team genuinely holds all six cards of a half-suit but we cannot place them within the team. What does waiting cost? An opponent cannot take the set by claiming it — a claim by a player whose team does not hold all six *gives* the set to us. So the only cost of waiting is running out of opportunity, and the benefit is that continued play localises teammate holdings. Waiting is close to free, which is exactly what the sweep measured and what v0.3’s expected-value model missed by treating the resolution of a split as a coin flip.

The corollary is that the leverage in declaring is not the threshold. It is (i) getting the declared distribution right — misdeclarations cost roughly 0.6 sets per deal-pair at v0.3’s strength *under the void-era rule these games were played with*, and under the award baseline each one is a full-set swing, so the stakes have doubled since that measurement; either way a misdeclaration is a bookkeeping failure rather than an unlucky gamble, and (ii) playing the forced declarations well, where no threshold is involved at all. v0.4 therefore keeps a high threshold and changes what happens underneath it:

- The declared distribution is the posterior MAP, found by shortlisting on the per-card marginals and scoring the shortlist with the *joint* posterior. The product of marginals is not the joint — cards compete for the same quota slots, so per-card modes can be jointly impossible.
- Where the clause set is empty the joint probability is computed exactly by the dynamic program of Section 4; where it is not, deduction still settles the extremes exactly, because a zero under the OR-free system implies a zero under the smaller OR-constrained one, and likewise for one.
- Forced declarations maximise expected sets rather than confidence, and the rule sets the objective: under the award baseline every wrong outcome costs the same set, so the ranking is $2\mathbb{P}(\text{exact}) - 1$ — maximise $\mathbb{P}(\text{exact})$ alone, while under the void-era variant it was $\mathbb{P}(\text{exact}) - \mathbb{P}(\text{an opponent holds one})$. The evaluator reads the rule off the observation and prices both.

That second probability was, until recently, an independence product, and the first bullet is the sentence that condemns it. The claim evaluator returned the *joint* for “this split is right” and $\prod_c \sum_{p \in \text{team}} \mathbb{P}(p \text{ holds } c)$ for “the half-suit is ours at all”, in the same tuple, from the same method, three lines below its own docstring saying the product of marginals is not the joint. Under the void-era rule in force when this bug lived, the forced ranking was $\mathbb{P}(\text{exact}) + \mathbb{P}(\text{ours}) - 1$, so the two carried equal weight; and their difference was read as “ours but wrongly split”, a quantity subtracted across two different distributions and therefore able to come out negative, which it did, and a clamp hid it. (Under the award baseline $\mathbb{P}(\text{ours})$ cancels out of the forced ranking entirely, which shrinks this bug’s surface but does not excuse it: the voluntary path and the void variant still consume the joint.)

Measured over 29406 half-suit queries in 60 games of real play, the two disagree on 27902, and on 86% of those the product *overstates* the joint — by a median of 0.008 and by as much as 0.313. Overstating “our team holds it all” is the direction that makes a declaration look safer than it is. Over

the 5885 positions in the same games where a forced claim could arise, the product produced a negative “wrongly split” 7 times and the correction changes which half-suit is declared at 31 of them. The decision counts are small, and neither is what makes the case: an expected value assembled from two different distributions is wrong at any magnitude, and the magnitude is reported so that the correction is not mistaken for a result.

The screening tier is unaffected and deliberately so. It returns a product for *both* elements, so the pair is internally consistent; what could not stand was one of each.

The screen itself carries a correctness claim, so we measured that too. The middle tier skips the joint query when the product of the per-card MAP marginals falls below 0.35, and returns the product as the half-suit’s probability. Its justification is the comment “most half-suits are nowhere near claimable” — an assertion about a distribution, inside a tier whose entire premise is that the product and the joint differ. If the product understates by enough, a half-suit whose true probability clears the 0.97 threshold is returned below 0.35, no claim is made, and nothing records that a certain set was left on the table.

Over 11687 screened half-suits in 25 games, the number that turn out to be claimable is **0**: the largest true joint among them is 0.581, against a threshold of 0.97. The shortcut is sound. What makes it sound is the margin and not agreement between the two quantities (the gap reaches +0.255 on individual half-suits), so the tier is safe at these settings and would not survive a much lower threshold. That is now a stored measurement (`results/claim_screen_check.json`) with a test that fails if a screened half-suit ever turns out to be claimable, rather than a comment.

“Waiting is close to free” is true about sets and false about information. The mechanism above is right that an opponent cannot take a half-suit our team wholly owns: a claim by a player whose team does not hold all six *gives* the set to us. But the same fact that makes waiting safe makes it sterile. `GameState.legal_asks` requires the asker to hold a card of the half-suit they are asking in. Once our team holds all six, no opponent holds one, so *no opponent can ever ask in it again*. Every public event that could still localise the split inside our team has been removed from the game by our own success at collecting it.

One channel remains and it is not free: an ask of our own that we know will fail. Under the no-bluff rule it certifies that the asker does not hold the named card, which is exactly the fact a partner needs, and it costs the turn (`fish4/perpetual.py`). So a half-suit that becomes ours before we can place it is a half-suit whose split we will know at the deadline exactly as well as we know it now, unless we spend a turn.

This changes what “waiting” means at the two ends of the game. It remains free in sets at risk. It is not free in information, and the bill arrives at the forced declaration, where no legal ask exists and the split must be named from whatever was known when the last opponent ran out of cards.

Finally, a hard architectural constraint taken from the literature: a claim must never be selected by determinized search. Inside any determinization the searcher knows the true layout, so every claim evaluates as a certain gain and the declared assignment differs between sampled worlds — textbook strategy fusion ([Frank and Basin, 1998](#)). Claims in v0.4 are always decided from the posterior.

6.3 Search over the belief

Every search this project has tried — PIMC, Information-Set MCTS, and two value-network variants — evaluates a candidate move by sampling hidden layouts and solving each one, and every one of them lost to its own prior. v0.3 did not leave that as a mystery: it measured the standard deviation of a position’s value across possible layouts at about $2.4\times$ the gap between the best and the worst candidate move, so a searcher scoring different candidates against different sampled worlds is ranking sampling noise. Common random numbers removed the deficit and never produced a surplus.

That diagnosis rules out a family of designs, not search itself. What it rules out is sampling. So the design it leaves open is a lookahead whose state is the *posterior* rather than a sample from it: the marginal matrix $M[c, p] = \mathbb{P}(\text{player } p \text{ holds card } c)$, with transitions that are local edits of that matrix. Given a

belief such a search is a deterministic function (run it twice on one position and it returns the same vector), so the variance that killed the previous attempts is structurally absent rather than merely reduced.

6.3.1 What it computes

A Literature turn is a run of asks that ends the instant one misses. Greedy scoring maximises the immediate success probability; what a possession is actually worth is how many cards it banks before the turn flips. Define

$$W(B, d) = \max_{a \in \mathcal{A}(B)} p_a [1 + W(B \mid a \text{ succeeded}, d - 1)], \quad W(B, 0) = 0, \quad (13)$$

where $\mathcal{A}(B)$ is the legal-ask set implied by the belief and $p_a = M[\text{card}(a), \text{target}(a)]$. The failure branch contributes nothing to W because a miss ends the possession; what a miss costs positionally is already priced by the incumbent objective's turn-risk and exposure terms, and duplicating it here would double-count.

W is in units of cards expected to be banked in the rest of this possession. It enters the policy as an additive bonus, not as a replacement:

$$\text{score}(a) = \underbrace{\pi_a + \sum_j w_j f_j(a)}_{\text{the incumbent objective, unchanged}} + w_{\text{look}} \cdot p_a W(B \mid a \text{ succeeded}, d - 1). \quad (14)$$

The reason it is additive is Section 9.1: replacing $\mathbb{P}(\text{success})$ with a learned half-suit value — correct units, a calibrated model — costs 7.4 sets per deal-pair. Whatever else this game rewards, the probability that the ask lands is the objective and not a term in one. At $d \leq 1$ the bonus is identically zero, so the policy reproduces the incumbent decision for decision; that is enforced by test, for the reason Section 7 gives.

6.3.2 Why the obvious version is provably a no-op

The natural transition is to collapse the asked card's row to a point mass on us. That is exact, and on its own it is *inert*, which is worth stating because it determines whether the design can do anything at all.

Proposition 3 (Greedy optimality under an inert transition). *Suppose the transition edits only the asked card's row, so that $p_{a'}$ is unchanged for every $a' \neq a$, and suppose every candidate remains available after any other is taken. Then (13) is maximised by taking candidates in descending order of p , and its value equals that of the greedy policy at every depth.*

Proof sketch. Under the hypotheses, a depth- d chain is a choice of an ordered d -subset of a *fixed* multiset of probabilities, with value $p_1(1 + p_2(1 + \dots)) = \sum_{k=1}^d \prod_{i \leq k} p_i$. Exchanging an adjacent pair (p_i, p_{i+1}) with $p_i < p_{i+1}$ changes the value by $(\prod_{j < i} p_j)(p_{i+1} - p_i) \geq 0$, so descending order is optimal; and the greedy policy selects exactly that order. $\square$

Proposition 3 is the null this section is tested against, and it is why a lookahead built only on the point-mass update would be worth measuring only as a bug-check. Its hypotheses fail in three places, of which the third is the substantive one:

1. a card taken cannot be asked for again, and one card is askable from several targets, so candidates do disappear in correlated groups;
2. a target can run out of cards, which removes every candidate aimed at them;
3. **the quota system couples the candidates.** Learning that the target held this card makes it less likely they hold the others — they are now known to hold one card fewer, and that probability mass has to go somewhere.

The third is handled by scaling the target’s column down to the count they now have and renormalising each affected row across the remaining players: one sweep of the same proportional fitting the exact posterior does globally. A row whose whole mass sits on the shrunk column renormalises straight back to where it started, which is correct — a publicly located card cannot be made less certain by a counting argument. The sweep costs one column scale and one row normalisation per node, and it is separately ablatable, so Proposition 3 is not merely argued but measured: the no-coupling cell in Table 10 is the experiment.

6.3.3 Cost

The branching factor of a real mid-game position is large: measured over 615 decisions from six self-play games at the champion configuration, the number of legal asks has mean 45.0, median 45, p90 81 and maximum 102, and exceeds 8 in 94.6% of decisions. (An earlier draft of this section said the branching factor “averages 8”. That figure is the mean legal-action count of the *exact-solver corpus* in `results/exact_positions_v2.json`, which consists of endgame positions with one or two live half-suits — an endgame statistic read as a general one.)

Continuations are beamed to the best 4 candidates by immediate p . The root ply is *not* beamed, because the policy needs a bonus for every candidate it is ranking, so a naive depth-3 tree costs $n(1 + b + b^2) = 21n$ expansions, about 945 at the median $n = 45$, rather than the $b + b^2 + b^3 = 84$ a uniformly-beamed tree would give.

Three quarters of those expansions were doing no work. At the last ply the continuation is $W(\cdot, 0) = 0$ by definition, so $q = p(1 + 0) = p$ and the maximum over the beam is simply the largest probability — a number the sort key has already computed. The implementation nonetheless performed up to b full transition/undo cycles there, each copying the marginal matrix and running a column scale and a 54-row renormalisation, to rediscover it. Short-circuiting that ply, and fusing the row renormalisation to avoid a boolean-index round trip, leaves the returned vector bit-identical (maximum absolute difference 0.0 over 24 real positions and every candidate) and costs $n(1 + b) = 225$ expansions instead of 945. Measured median cost falls from about 14 ms per decision to 6.7 ms (mean 8.3, p90 18.1, over 40 positions, minimum of three repetitions each) on the contended host of Section 4.

Against the configuration the lookahead runs actually used — 160 draws, the `fish4/agent4.py` default, since no lookahead cell in `results/v04_duels.jsonl` overrides `n_draws` — the relevant posterior cost is the 31.9 ms row of Table 3, not the ~ 100 ms rows, which are all 512-draw. So the bonus now adds roughly a fifth to a decision. An earlier draft of this section claimed 15% by comparing the pre-optimisation cost against a 512-draw budget these runs never used, which was wrong twice in opposite directions; the duelling results in Table 10 were all produced at the slower 945-expansion cost, since the optimisation is exact and postdates them.

7 Evaluation methodology

Notation, fixed here and used throughout. A bracketed pair $[\cdot, \cdot]$ is always a 95% confidence interval. A figure written $x \pm h$ is always a 95% *half-width*, so $x \pm h$ and $[x - h, x + h]$ say the same thing; where a standard error is meant it is named “standard error” in words. An earlier draft used $\pm$ for both, in the same document and occasionally on the same page, which meant a reader could not tell a 68% figure from a 95% one without opening the source. Set differentials are per duplicate deal-pair unless the text says *per game*, which is used for cross-engine results where there is no pairing.

7.1 Duplicate deals

Every result here uses **duplicate (paired) deals**: each deal is played twice with the teams swapped, on identical cards, an identical rotated starting seat, and identical agent randomness. Only the policy assignment differs, so per-pair set differentials are i.i.d. across deals and confidence intervals mean what they say.

run	arm	identical games	ρ	efficiency
γ policy specificity (G1)	B_none	19.9%	0.089	1.1×
γ policy specificity (G1)	C_measured	21.2%	0.088	1.1×
signalling gate	C_measured	79.8%	0.853	6.8×
signalling gate	B_incumbent	83.4%	0.877	8.1×
doomed-ask gate	B2_mid	86.1%	0.889	9.0×
doomed-ask gate	B_defer	88.1%	0.899	9.9×
forced search, self-play	B_full	98.4%	0.996	268×
forced search, vs v0.7	B_full	99.1%	0.998	414×

Table 2: How many times the games an unpaired design would need to match the precision each paired run achieved. “Identical games” is the share of games the two arms finished on the same margin, and it orders the table: a knob that fires on a rare branch leaves nearly every game bit-identical and pairing removes almost all the variance, while γ re-weights every posterior sample on every decision, the games diverge immediately, and pairing removes nothing.

Earlier versions of this section justified that design by asserting that raw win rates in a card game are dominated by the deal. That is the standard argument for duplicate play and it is borrowed from bridge. It is false here, and the head-to-head corpus was large enough to say so without playing a game.

7.2 The deal decides nothing

The 10,000-game match against v0.7 (§11.9) played every deal from both seat parities, so in the second game our engine holds exactly the cards the opponent held in the first. Writing D_d for whatever the deal is worth to whoever sits on the even seats and E for play noise,

$$m_{\text{even}}(d) = \mu + D_d + E_e, \quad m_{\text{odd}}(d) = \mu - D_d + E_o, \quad (15)$$

so $\text{var}(m_{\text{even}} + m_{\text{odd}}) = 2\sigma^2$ and $\text{var}(m_{\text{even}} - m_{\text{odd}}) = 4\text{var}(D) + 2\sigma^2$. The deal’s share of one game’s variance is then $\text{var}(D)/(\text{var}(D) + \sigma^2)$, which under this model is identically $-\rho$, the correlation of the two parities’ margins.

Measured over 5,000 deals: $\rho = +0.0127$ $[-0.0150, +0.0404]$, so the deal’s share is -1.3% $[-4.0\%, +1.5\%]$. The sample variances are 15.267 against 14.884 for sum and difference, and the implied $\text{var}(D)$ is -0.096 — zero, measured with noise. **Which cards you are dealt does not measurably decide a game of six-handed Fish** (`scripts4/deal_luck.py`, `results/deal_luck.json`).

The mechanism is not mysterious. There are no high cards: every half-suit is worth exactly one set, so a hand cannot be strong, only awkwardly or conveniently distributed. And cards move continuously, so whatever the deal arranged is largely dissolved by the middlegame. The deal does leave a *symmetric* trace — some deals are clumped enough that asks land more often for everybody, and our ask hit rate correlates $+0.087$ $[+0.060, +0.115]$ across the two parities of a deal, but a symmetric texture is not an advantage to either side, which is why it appears in the hit rate and not in the margin. A deal can be textured without being unfair.

Two consequences follow, and the second is the useful one.

Swapping seats buys nothing. Duplicating a deal across seat parities is worth $0.99\times$ the games, because the antisymmetric deal advantage it exists to remove is not there.

Holding the knob fixed buys between $1.1\times$ and $414\times$, and the rate the knob fires sets which. That is a different pairing and it is the one every ship decision in this project rests on: arms A and B play the identical deal from the identical seats against the identical opponents and differ by one parameter. Table 2 prices it on every multi-arm journal the repository holds (`scripts4/pairing_value.py`).

The planning consequence is worth stating plainly, because this project has sized runs by effect size alone and that is the wrong variable. **The sample size a screen needs is governed by how often its knob changes a decision, not by how large the effect is.** The forced-search secondary resolved a $+0.0180$




Figure 1: What actually sizes a duplicate-deal run. *Left:* the variance-reduction factor from pairing against the share of deals on which both arms produced an identical margin — that is, on which the knob under test never changed a decision. The relationship is tight over more than two orders of magnitude. *Right:* the same factor against the size of the effect being measured, which is the quantity everyone reaches for when planning a run and which does not determine the answer: effects below 0.15 sets per pair span the entire range on their own. Same runs as Table 2.

effect to ± 0.012 on 1,000 games; G1, at 1,600 games and the same design, could resolve nothing finer than ± 0.19 . That is a sixteen-fold difference in precision from the same number of deals, and it is entirely attributable to one knob firing on 0.9% of games and the other on 80%.

7.3 How much can an experiment of size n resolve?

The per-pair standard deviation of the set differential between two policies of the strength we compare (v0.3’s champion against its own baseline) is **3.869 sets**, measured over 300 pairs. The distribution is wide: the theoretical bound is ± 18 and observed pairs stayed within ± 11 , but a single pair says almost nothing. The consequences are stark:

effect (sets/pair)	pairs @ 80% power	pairs @ 90%	hours @ 80%
0.10	11 751	15 731	3.9
0.20	2 938	3 933	0.98
0.30	1 306	1 748	0.44
0.50	471	630	0.16
1.00	118	158	0.04

Read the other way: the 600-pair cells of the v0.3 champion search could resolve effects down to about 0.44 sets per pair, so every “no significant difference” it reported below that size was *uninformative rather than negative*. We therefore report, alongside every interval, the smallest effect the run could have detected.

That standard deviation is not a property of the game. The table above, and every power calculation in this paper up to this point, treats 3.8 sets per pair as a constant of the deal population. It is not. It is a property of a *pair of policies* that disagree constantly, and the experiments most in need of careful sizing are the ones whose arms barely disagree at all. We noticed this only when a screening cell measured a per-pair standard deviation of 2.328 — seven standard errors from the figure it had been sized against.

The mechanism is immediate once stated. Under common random numbers, a deal-pair on which the two arms play *identically* contributes a difference of exactly zero. Writing s for the share of pairs on which they diverge at all,

$$\text{Var}(D) \approx s \cdot \text{Var}(D \mid \text{diverge}), \quad \text{so} \quad \text{sd} \approx \sqrt{s} \cdot \text{sd}(D \mid \text{diverge}).$$

That much is arithmetic and follows from the design. The empirical content is what the second factor does. Across the 46 cells of this study that store their per-pair differentials it moves little in absolute terms — 3.86 ± 0.29 , a relative spread of 7.6% against 15.5% for the raw standard deviation, and our first draft stopped there and called that the finding. It does not survive being looked at properly.

Three things weaken it, and the third matters for use. *Most of the residual spread is not structure*: the sampling noise of an sd estimate at these sample sizes is alone 2.8%. *The comparison is carried by a handful of cells*: restricted to the 26 cells with divergence share above 0.83, where most of this study lives, the raw spread is 3.38% and the conditional spread is 3.45%, the same number to within noise, and the correlation between share and sd collapses from +0.95 to −0.03. Where share barely varies the decomposition has nothing to explain, and it explains nothing. And *the conditional term is not constant*: it drifts with share, $\text{corr} = +0.79$, fitted as $2.53 + 1.66s$.

That drift matters precisely where the rule gets used, because the rule is only useful when extrapolated to a low-divergence contrast. At $s = 0.31$ the flat model predicts $\text{sd} = 2.18$ and the fitted line predicts 1.68 — a 30% over-estimate. The direction is the saving grace: over-predicting noise asks for more pairs than are needed, never fewer, so a run sized this way is over-powered rather than under-powered. But a rule invoked to justify a *smaller* experiment than the A/A figure implies should not be quoted as though it were calibrated at the point where it is used, and it is not.

One further caveat on what s means. It is measured as $\mathbb{P}(D \neq 0)$, and D is zero on many pairs where the two arms genuinely diverged and the deal still ended level — zero is not an isolated atom in these histograms. So the cheap substitute the rule advertises, counting disagreements at the level of decisions rather than of finished pairs, is a strictly larger quantity than the one the constant was calibrated against, and using it directly would inflate the predicted noise further.

What survives is narrower than we first wrote and still worth having: the per-pair standard deviation is a property of the comparison rather than of the game, the $\sqrt{s}$ scaling is exact by construction, and sizing every experiment against a figure measured between two maximally-different policies is systematically conservative for the near-identical comparisons that most need care.

The extrapolation was untestable, and then it was tested. When the above was written this study had no cells below $s = 0.44$, so the calibration of the constant at low share (the only place the rule is ever *used*) was not established, and the paragraph said so. The gated re-take run has since landed two blocks at $s = 0.233$ and $s = 0.236$, and the prediction that the flat model over-states low-share noise by roughly 30% was recorded in `jobs/PREREGISTRATION_retake_gate.md` *before any pair of that run was played*. Scoring it out of sample, over the three cells now below $s = 0.5$:

cell	s	measured sd	flat	drift-corrected
re-take gate, block 0	0.233	1.443	1.863	1.409
re-take gate, block 1	0.236	1.440	1.875	1.420
ungated re-take penalty	0.440	2.328	2.560	2.163
mean relative error			+23.1%	−3.6%

The flat constant over-predicts by 23% where it is used; the drift-corrected line lands within 4%. This is the only out-of-sample prediction in the paper that was written down before the data existed and then checked against it, and it is worth more than the several in-sample fits around it. Sizing at low s should use $\text{sd} = (2.53 + 1.66s)\sqrt{s}$, and the flat form remains the conservative choice — it asks for more pairs than are needed, never fewer.

The correlation between s and the raw standard deviation is +0.952, and we report it only to say that it is *not* evidence. If sd really is $c\sqrt{s}$ for roughly constant c , that correlation is the correlation of s with $\sqrt{s}$, which is near 1 over any range — it is implied by the identity rather than a test of it. A reader will compute it anyway, so it is better stated and defused than omitted.

Sizing therefore reduces to predicting one quantity, s , which can be counted from *decisions* rather than from finished pairs and is correspondingly cheap. The consequence is not cosmetic. A candidate that perturbs 10% of pairs has $\text{sd} \approx 1.23$, not 3.80: sized on the constant it would need $9.5\times$ the pairs, which

is the difference between an experiment that looks impossible and one that is routine. We used exactly this to decline a run — a sixth cell in a family of five failures, whose largest possible effect we could show in advance was inside what the cell would resolve, so that it would have returned a null whatever the truth was, and to size its honest replacement at 2000 pairs rather than the 6200 the constant demanded.

The direction of the error is worth naming: sizing on the A/A figure is *conservative*, so nothing already published here is too narrow. What it cost was experiments never run, and nulls recorded from cells that could not have found anything. Those are the expensive kind of mistake precisely because they leave no trace in any interval.

7.4 Anytime-valid stopping

Screening many candidate terms invites peeking. A fixed- n 95% t -interval, monitored after every deal, rejects a *true null* 68.8% of the time on an empirical null and 92.1% on a skewed one, with a median false stop after 11 and 3 deals respectively. We therefore use anytime-valid procedures: an empirical-Bernstein confidence sequence and a betting test supermartingale (Waudby-Smith and Ramdas, 2024). Measured type-I error under continuous monitoring over 4000 null experiments each monitored at every one of 4000 deals: 2.1–2.4% for the betting test and 0.3–0.5% for the confidence sequence, against a nominal 5%.

7.5 Two corrections to the previous harness

Timed-out games must be scored, not dropped. Nothing in Literature forces progress, so a harness needs an action cap. v0.3 dropped any deal-pair containing a capped game. Dropping conditions on the outcome, which is precisely what a paired design exists to avoid, and an audit conducted for this paper found it biases rather than merely blunts: replaying the same 60 deals at reduced caps, the dropped pairs had a true mean differential of -7.14 against -10.13 for the kept ones, and at one cap the reported 95% interval *excluded* the true value. v0.4 scores a capped game on the half-suits actually resolved, with the unresolved ones counting for nobody — the same semantics the exact solver gives an unbroken cycle.

An ablation must reduce to its baseline. v0.3’s most instructive methodological failure was a confounded ablation: the ablated agent silently normalised a second term differently, so every cell compared two changes at once and produced tight, confident, and completely wrong intervals. The confound was caught only by noticing an implausible monotonic pattern in the magnitudes. We therefore enforce by test that a zero weight reproduces the baseline *decision for decision* over real positions, and separately that a non-zero weight demonstrably changes some decision. Every term in Table 1 passes both.

8 Results: what the inference, the objective and the ladder are worth

Figure 2 is this section on one axis: every in-play effect the paper settles, each measured by the paired duplicate-deal design of §7, sorted by size. The subsections that follow take the rows in turn; the orange row is measured in §11.9.

8.1 Exact inference alone changes nothing

The first experiment is the one the research programme predicted would pay. Holding the ask objective at v0.3’s champion weights and changing *only* the inference — exact counting where the clause set is empty, unbiased importance sampling where it is not, in place of a sampler documented as non-uniform — gives a set differential of -0.22 per deal-pair, 95% CI $[-0.96, +0.52]$, over 100 duplicate deals. Statistically indistinguishable, and the interval is wide enough that only effects above about 1.1 sets per pair could have been detected.

This is a genuine negative result about a change that is provably correct. The explanation is in the next subsection.

8.2 Posterior quality against ground truth

Scoring posteriors against the true hidden hands over 517 decisions and 13,640 uncertain-card predictions (Table 3) shows what happened.




Figure 2: Every settled in-play effect, on one axis. Blue rows ship in the deployed configuration; gray rows were evaluated and not shipped — at-ask-time depth a null at $n = 6000$, the learned whole-game weights significantly worse, and the endgame ask correction shipped in the void era and *withdrawn* when the award baseline re-priced it below zero (§11.9). The orange row is that correction against a foreign engine, Dylan’s FishBot v0.7, under the matched award rule: a tie whose interval excludes the sibling-measured void-era $+0.12$ — the paper’s recurring finding that gains measured inside an engine family shrink outside it. Intervals are 95% confidence intervals over paired deals; n counts pairs.

Two things stand out.

First, the exactly-uniform posterior is *worse* at predicting reality than the heuristic sampler it replaced: 1.3618 against 1.3408 at a matched draw budget. That is not a contradiction. Proposition 1 is a correct theorem about a false hypothesis. v0.3’s sampler seeded the OR-clauses by picking a live card of the asked half-suit for the asker, which systematically over-weights worlds in which askers are deep in the suits they asked in — an accidental opponent model. Removing the bias removed the model with it.

Second, making the model explicit and tunable recovers the loss and passes both: 1.3251 at $\gamma = 0.45$. And the improvement is not a large-budget artefact. At 160 draws the opponent model takes the sampler from 1.3731 to 1.3476, which is better than v0.3’s deployed 32-draw configuration (1.4399) by a wide margin.

Third, and separately: ignoring the OR-clauses is by far the worst option tested, confirming the design decision of Section 4.4 from the other direction.

8.3 How much of the sampling budget survives reweighting

Table 3 scores the posterior. This asks what it costs to get there, and finds a diagnostic that has been computed on every decision this engine has ever made and reported on none of them.

The posterior is drawn by importance sampling and self-normalised. Self-normalisation is consistent, not free: if the proposal is a poor match for the target, a few draws carry most of the weight and the batch behaves like far fewer than n independent worlds. Kish’s effective sample size measures that, and the sampler already computes it.

At the champion configuration, 160 draws are worth 83.7 effective, 52% efficiency, with a p10 of 64.8, a median of 87.0 and a p90 of 100.0 over 641 decisions across 36 seat-deals. The exact DP is almost never the path taken: 9 decisions in 641. So the precision of every $\mathbb{P}(\text{success})$ in this engine is the precision of that one sampled batch.

inference	NLL ↓	Brier ↓	top-1 ↑	ms/decision
exact DP, clauses ignored	1.4776	0.7568	0.287	50.4
v0.3 sampler, 32 draws	1.4399	0.7065	0.386	5.9
v0.4 importance sampler, 160	1.3731	0.7006	0.396	29.5
v0.4 importance sampler, 512	1.3618	0.6955	0.403	92.3
v0.3 sampler, 96 draws	1.3552	0.6891	0.406	16.6
v0.3 sampler, 512 draws	1.3408	0.6833	0.416	87.8
v0.4 + opponent model, 160, $\gamma=0.45$	1.3476	0.6856	0.401	31.9
v0.4 + opponent model, 512, $\gamma=0.30$	1.3332	0.6791	0.415	99.8
v0.4 + opponent model, 512, $\gamma=0.45$	1.3251	0.6747	0.413	98.6
v0.4 + opponent model, 512, $\gamma=0.60$	1.3248	0.6744	0.411	100.8

Table 3: How much probability each inference configuration places on where the cards actually are. Lower NLL and Brier are better. A bold NLL marks the three figures the text turns on. The comparison that matters is the first two: the *exactly correct* uniform posterior at 512 draws (1.3618) is worse than v0.3’s *biased* sampler at the same budget (1.3408). The third is the best configuration found. **Ten of the fourteen configurations scored in that campaign are shown;** the four omitted are $\gamma \in \{0.15, 0.30, 0.60, 0.80\}$ at 160 draws, and the last is the point of saying so — at $\gamma = 0.80$ on 160 draws the model scores 1.3955, *worse* than switching it off (1.3731). Two later campaigns scored 21 more configurations, 35 in all; the winning cell is a grid maximum over that set, reported without a multiplicity correction. All three files are named in §D.

A control this paragraph used to claim, and did not have. Earlier drafts read “83.7 effective ... against 99.9 with the opponent model disabled”, and reported 14 exact decisions in 646 for that arm. No such arm was ever run. `results/ess_probe.json` holds four rows, all at $\gamma = 0.35$, differing only in the proposal tilt; the 99.9 is the *p90 of the champion row itself* (99.986), and the 14 in 646 has no source at all. A percentile of one arm had been written up as the mean of another, and the paragraph that followed then explained a gap between them.

That is worth more than a correction, because the argument it supported was sound and the number under it was not. The efficiency loss from steering the proposal is real and is measured, 52%, and the remedy discussed below was tested and did make the marginals worse. What was never measured is how much of the loss the opponent model accounts for, and with $\gamma > 0$ changing the target the exact DP is built to count, the comparison is not a straightforward one to run. It is left open rather than asserted.

The gap has an obvious explanation and an obvious remedy, and the remedy is a trap. The proposal is quota-proportional and knows nothing about the choice model, so the whole tilt of the target lands in the weights; folding one step of the likelihood into the proposal is textbook, and it does exactly what the textbook says — effective sample size rises to 105.8, past even the untilted $\gamma = 0$ figure. Partial strengths are worse than either end, the signature of a proposal distorted but not matched.

Measured against a high-precision reference, the twist makes the marginals $3.4\times$ worse.

Effective sample size measures how *flat* the weights are, and a proposal that covers one corner of the target well and the rest not at all has beautifully flat weights over the corner it covers. The one-step twist overshoots because the likelihood’s depth-zero term is a numerical floor of $\log 10^{-9}$: the astronomical $0 \rightarrow 1$ ratio makes the proposal read “this slot is empty now” as “this slot ends empty” and drives nearly every draw into one configuration. The parameter ships at zero, verified bit-identical to the untwisted sampler, and the tests pin both halves of the result so that the next person to notice the efficiency gap does not have to rediscover why closing it fails.

Is precision still worth buying? v0.3 found belief precision unsaturated, but with a biased sampler and a different estimator, so it does not transfer. Measured here against a 16,000-draw reference, L1 error per card falls as $n^{-0.475}$ from 40 draws to 1280 with no bias floor anywhere:

draws	40	80	160	320	640	1280
L1 per card	0.248	0.174	0.128	0.091	0.066	0.048

Two things follow. The exponent is the pure Monte Carlo law, which independently confirms the sampler is unbiased — a bias floor would show as the curve flattening, and it does not. And precision at the operating point of 160 draws is nowhere near saturated: the prize is real and can be bought directly with draws rather than with a cleverer proposal.

The policy can feel it, and it is the largest single change here apart from the opponent model. The duel that paragraph declined to claim was pre-registered before its screening cell finished — deliberately, because the belief-space search had already cost four rounds partly by being sized against a number a screen had inflated, and the only reliable defence is to fix the design before the number exists. It was sized against a *minimum interesting effect* of +0.15 sets per deal-pair rather than against any estimate, on the ground that an effect smaller than the already-shipped search delivers, for triple the sampling budget, is not one this engine should adopt.

Six blocks of 1000 duplicate deal-pairs, $n_draws = 480$ against 160, everything else identical:

$$+0.340 \text{ sets per deal-pair, } 95\% [+0.243, +0.436].$$

All six blocks are positive, Cochran's $Q = 2.92$ on 5 degrees of freedom ($p = 0.71$), $I^2 = 0\%$, $\hat{\tau} = 0$. The entire interval clears the threshold fixed in advance. For scale, this is more than three times the belief-space search, which took four rounds and 6000 pre-registered pairs to establish at +0.104.

The 400-pair screen of this same cell had read +0.160. The confirmatory run reads +0.340 — *higher*. That is what an unselected screen looks like, and it is worth stating only because every screen this project *selected* on inflated instead.

What it costs, and why that is a separate decision. The pre-registration committed in advance that a demonstrated effect would not automatically move the default, because n_draws trades strength against latency and the public table has a request budget. Measured on 90 positions, a decision costs

$$3.57 \text{ ms} + 6.34 \mu\text{s per draw,}$$

a fit whose residuals reach 0.21 ms across 40 to 1920 draws, and are worst at the low end — −0.21 at 40 draws and +0.20 at 160, against +0.03 and −0.02 at 480 and 1920. So the straight line is a good description of the expensive end and a rough one exactly where the shipped budget sits, which is worth knowing before using it to price a change at 160. The ratio quoted next does not depend on it: it is measured per decision, not read off the fit. At the shipped 160 the sampler is only a fifth of a decision, so tripling it costs +1.9 ms — $1.38\times$, not $3\times$. That headroom is finite and the same fit says so: by 1920 draws a decision is already $3.3\times$ the shipped one. With the effect and the price both measured, the default moved to 480.

That fit is not an explanation, and the difference matters: if the fixed term were setup, it could be hoisted and draws would become nearly free. Timing the sampler alone, on one already-built sampler across 1 to 1440 draws, gives $2.50 \text{ ms} + 8.0 \mu\text{s per draw}$ — about $60 \mu\text{s}$ of fixed cost for each of the 42 free cards it must place. That is per-card numpy dispatch. Sequential importance sampling walks the free cards one at a time, so batching makes each *step* wider without making the number of steps smaller, and the per-card constants are already cached. There is nothing there to hoist.

The consequence is a design rule with a number attached: going from 160 to 480 draws costs 2.2 ms inside the sampler, while going from one decision to two costs 2.5 ms before a single draw is taken. Draws are cheap and decisions are expensive, so under a latency budget a better posterior at one decision beats more decisions with worse ones. The belief-space search of Section 6.3 already obeys this without having been designed to: it never samples, and operates on the single posterior marginal matrix it is handed.

One caution on those two fits. The per-draw slope of the sampler alone ($8.0 \mu\text{s}$) exceeds that of the whole decision containing it ($6.34 \mu\text{s}$), which is impossible until the position mix is accounted for: per-draw work scales with the number of free cards, and the two runs drew positions averaging 42 and 28 of them. Scaling by that ratio predicts $5.4 \mu\text{s}$ against the 6.34 measured, which restores the ordering and

γ	pair score	set diff / pair	95% CI
0.10	0.613	+0.993	[+0.470, +1.517]
0.30	0.667	+1.847	[+1.230, +2.463]
0.45	0.700	+1.860	[+1.336, +2.384]
0.60	0.657	+1.300	[+0.795, +1.805]
0.90	0.723	+1.827	[+1.199, +2.455]

Table 4: Opponent model against the identical policy with the model disabled, 150 duplicate deal-pairs per cell. Every cell excludes zero. The dose-response rises to a plateau by $\gamma = 0.30$ and stays there over a threefold range, which is what one wants from a model with a single free parameter — and unlike the ask-objective terms of v0.3, which showed a sharp inverted-U, there is no weight in this range that is destructive *in play*. Belief quality is a different estimand and does have one: at $\gamma = 0.80$ on a 160-draw budget the model scores a worse NLL than switching it off (§8.2).

accounts for most of the gap but not all of it — the remainder is the rest of a decision that also grows with the draw count. We report it as a partial reconciliation rather than a clean one.

Two of this paper’s four documented winner’s curses were paid for by sizing an experiment against an inflated screen. This is the case where pre-registering against a threshold instead cost nothing and returned the study’s largest result.

8.4 The opponent model in play

Posterior accuracy is not strength. Table 4 measures the model by duplicate-deal play against the otherwise-identical $\gamma = 0$ policy.

The cleanest attribution. Table 4 measures the model against copies of itself. The sharper experiment holds the *opponent* fixed at the v0.3 champion and varies only whether the model is on, at identical deal seeds and identical agent seeds, 250 duplicate deal-pairs each:

v0.4 configuration	pair score	set diff / pair	95% CI
exact inference, $\gamma = 0$	0.502	−0.008	[−0.482, +0.466]
exact inference, $\gamma = 0.30$	0.684	+1.920	[+1.479, +2.361]

The first row is as close to exactly zero as a 250-pair experiment can report. Provably correct inference, replacing a sampler its own authors documented as biased, moved the result by eight thousandths of a set per deal-pair. The second row differs from the first in one parameter.

8.5 The choice model, measured rather than assumed

Section 8.4 establishes that the model is worth +1.9 sets per deal-pair, which is the largest single effect in this engine. It rests on one sentence: *a player asks in a half-suit in proportion to how many cards of it they hold*. Proportional means an exponent of one, and that exponent was never measured. It was chosen because it is the one value whose denominator is public (the depths of a hand sum to its size), which is an excellent reason to try it and no reason whatever to believe it.

There is a specific reason to doubt it. Under an objective dominated by $\mathbb{P}(\text{success})$, what makes a half-suit worth asking in is how many cards it still offers: holding five of six offers exactly one card to ask for, holding one of six offers five. Depth and opportunity pull in opposite directions, so linear-in-depth is not obviously right even in sign.

The design. The policy is available, so the propensity can be measured instead of argued. Over 200 self-play games we record, at every ask, which half-suits were legal, the asker’s depth in each *on the initial deal*, and which was chosen. That is a conditional-logit design in which the alternatives are known from the rules rather than inferred, and it yields 17,005 asks where at least two half-suits were legal — 72,091 alternatives in all. Fitting $\mathbb{P}(\text{ask in } H) \propto \text{depth}_H^\alpha$ by exact conditional likelihood gives α directly, with $\alpha = 1$ the shipped model and $\alpha = 0$ legality carrying all the signal and depth none.

half-suits resolved	0	1	2	3	4	5	6–8
α	2.00	1.28	1.20	0.68	0.30	0.41	−0.02
clustered SE	0.08	0.08	0.12	0.12	0.13	0.18	0.16
decisions	4410	2881	2527	2437	1958	1541	1251

Table 5: The fitted exponent of the choice model against how far the deal has run. The shipped model is a constant 1. It is wrong at both ends, in opposite directions: at the opening the true exponent is about 2 and the model understates its own strongest signal, and by the endgame it is 0 and the model applies a tilt that is not there.

Standard errors are block-bootstrapped over games. This is not a refinement: the likelihood’s own curvature treats every decision as an independent draw, and 17,005 decisions come from 200 deals whose members share a layout, six hands and one set of policies. The correction doubles every interval, and the symptom was visible beforehand — a smooth curve through the per-band estimates returned a χ^2 several times its degrees of freedom.

What it says. Pooled over the whole game the best constant is $\alpha = 1.207 \pm 0.046$ — close enough to 1 that a constant never looked broken, and far enough from either end of Table 5 that it could not be right. Split at the halfway point, $\alpha = 1.331 \pm 0.049$ early against 0.232 ± 0.118 late. The likelihood prefers the fitted constant to $\alpha = 0$ by 1402 nats and to $\alpha = 1$ by 40.

Most of the decay is the covariate, not the player. The obvious reading of Table 5 is behavioural: a player asks where they are deep only if they had a choice, and late, holding two cards with most of the deck resolved, legality binds. That reading is largely wrong, and the same 17,005 decisions say so.

The fit conditions on depth *on the initial deal*, because that is what the sampler assigns. Cards move, so as a game runs on that covariate drifts from the hand the asker actually held — mean absolute disagreement rises from 0.16 cards at the opening to 0.66 by the endgame, with correlation falling from 0.77 to 0.44. Noise in a covariate attenuates its coefficient, and noise that grows with the game manufactures precisely this decay. Refitting on depth at the moment of the ask:

covariate	opening	1–2	3–4	5–8	pooled
initial-deal depth	2.00	1.25	0.52	0.23	1.207
depth at the ask	2.90	2.32	1.72	1.34	2.195

The decay survives and shrinks, and never approaches zero. Late asks do carry depth signal; what decays is how well the initial deal still describes the hand it is about.

The better covariate is free. Depth at the moment of the ask looked expensive (it appeared to need the public history replayed against every candidate world), and it is not, because of a fact about the rules. Cards move only when an ask succeeds, and a successful ask is entirely public: the card, the asker and the target are all in the log. So $\text{depth}_{\text{at ask}}(p, H) = \text{depth}_{\text{initial}}(p, H) + \delta(p, H, t)$ with δ *identical in every hypothesised world*. The sampler goes on assigning initial deals and the likelihood adds a constant known before any drawing. Verified against the engine on 23,268 (player, half-suit, time) triples across six games, with no mismatches.

Measured as a description, the difference is not marginal: at-ask-time depth fits the same choices 4,654 nats better, reaching 6,057 nats above uniform where the initial deal reaches 1,403. The shipped covariate captures under a quarter of what its own one-parameter family can.

This is not the at-time mode of Section 9.1, which computes belief-located depth and adds a present-tense count to free cards drawn over the initial deal; that mode measured -1.544 sets per deal-pair over 250 pairs. Whether the exact version plays better is a duel, and is not claimed here.

Is the power law the right shape? Fitting one functional form well is not evidence that it is the right one. The saturated alternative gives every depth value a free coefficient; it nests the power law, so it can never fit worse, and it is the best that anything conditioning on depth alone can do. On the same decisions the power law gives up 110 nats to it on initial-deal depth and 269 on depth at the ask.

Set against the 4,654 nats that changing the covariate is worth, the shape is not what limits this model — the covariate matters roughly seventeen times more. That is the useful bound: further gains in the choice model need a different conditioning variable or the normaliser below, not a better curve. (The initial-deal figure is a lower bound; that fit had not converged, and its depth-5 coefficient rests on five alternatives in the whole corpus.)

What the engine drops, and where it sits. The likelihood is written as proportional to a product of depths, and the proportionality hides a denominator: a proper choice model divides by the sum of depth^α over the half-suits the player could legally ask in. At $\alpha = 1$ that sum is the player’s hand count — public, identical in every world, and it cancels. That is the original argument for an exponent near 1, and it holds for every covariate here: the depths over live half-suits sum to the hand count whether depth is read at the deal, now, or at the ask.

The cost of being elsewhere had never been measured. Across worlds drawn from the posterior at 30 real positions, the coefficient of variation of the dropped denominator is

α	0.35	0.60	0.80	1.00	1.207	1.60	2.195
mean CV	0.100	0.063	0.032	0.000	0.035	0.107	0.226

a V centred exactly on 1, with the $\alpha = 1$ column serving as the control that the measurement is reading the right quantity. Two readings follow. The shipped $\gamma = 0.35$ pays more for the omission (0.100) than the fitted exponent on its own covariate would (0.035 at 1.207) — the approximation that motivated an exponent near 1 is being violated more at the setting in use than at the one the data prefer. And $\gamma = 1.0$ is not just another point on the sweep of Table 4: it is the unique exponent at which this likelihood is correctly specified. That sweep never tested it; 0.90 and 1.50 bracket it at 150 pairs each, an MDE of 0.87, which separates nothing.

Every number above was measured on one population, and against a foreign policy the exponent changes sign. The corpus behind Table 5 is this engine playing a copy of itself. That is the population the opponent model is used in *when the table is all one engine*, and it is also the one population in which the assumption cannot fail: a policy that asks where it is deep, modelled by a policy that assumes opponents ask where they are deep, will confirm itself.

The obvious control had never been run. Repeating the identical design — same alternatives, same covariate, same conditional likelihood, same block bootstrap over deals, with only *whose asks are recorded* changed — against v0.7 over 6,090 of its asks in 150 cross-engine games gives

depth on the initial deal	1	2	3	4
observed / expected	1.258	0.762	0.548	0.656

and a fitted $\alpha = -1.0041$, 95% $[-1.1434, -0.8648]$, against this engine’s own +1.2071 on the same estimator. The signs are opposite and the exponents lie 2.21 apart. Re-measured through the repaired bridge (P43, R0) the exponent is -1.1055 $[-1.2531, -0.9578]$ over 6,281 asks; the intervals overlap and the sign is unchanged. Applying that exponent to their seats is the worst arm in the study (-0.90 against them, -2.00 in self-play; §11.11). A sampler applying $\gamma = +0.35$ to that seat is not merely mis-scaled; it re-weights the world set in the wrong direction on every one of that seat’s asks, and every marginal the ask objective and the declaration gate read is drawn from the result.

The mechanism was written down in this project before it was checked, and then not checked. Section 8.5’s own premise notes that depth and opportunity pull opposite ways — holding five of six leaves one card to ask for, holding one of six leaves five, so linear-in-depth was never obviously right *in sign*. v0.7’s objective is dominated by a single coefficient on $\mathbb{P}(\text{success})$ roughly three times its next largest, which is precisely the objective that prefers the shallower half-suit because it offers more cards to name. What the self-play corpus measured was not a law of the game but a fixed point of one policy against itself.

Two conclusions, and the second is the load-bearing one. The narrow one is that the limit on this model is not the curve’s shape (worth 110 nats) or even the covariate (worth 4,654): it is that a single global exponent cannot describe a mixed table, and both of those bounds were computed inside the population where the exponent happens to be constant. The broad one is a warning this paper is in a position to give and few others are: *an opponent model validated only in self-play has been validated against the assumption it encodes*. Nothing in the fit diagnostics of §8.5, not the 1402-nat margin over $\alpha = 0$, not the saturated-alternative comparison, not the phase split — could have detected this, because all of them are computed on the same self-consistent corpus. It took a second engine.

We are deliberately not shipping a per-seat exponent on the strength of this. The measurement identifies where our inference is wrong about *one* opponent, and §11.9 records what happened the last time this project turned an opponent-model discrepancy into a knob: the gains grew with the dose and evaporated against a third engine, because what was being measured was deception of a sibling’s inference rather than strength. A per-seat exponent fitted on v0.7 and validated on v0.7 would reproduce that error exactly. It needs its own pre-registration and a foreign opponent that is not the one it was fitted on.

It got one, and the answer is that describing his asks well and beating him are different problems.

The pre-registration is `prereg/gamma_policy_specific.md`, written before a game was played, and it fixed two things in advance: that nothing would ship on a run whose treatment arm was both fitted on v0.7 and measured against v0.7, and that the smallest effect worth adopting was +0.15. Three arms, identical in every respect but one scalar, played on 1,600 games (800 deals in both seat parities), each arm seeing the identical deal from the identical seat:

arm	γ	sets/game	paired difference vs shipped
A, shipped	+0.35	+2.3613	—
B, no model	0.00	+1.9563	−0.4050 [−0.5911, −0.2189]
C, his exponent	−1.00	+1.2588	−1.1025 [−1.2897, −0.9153]

Both directions lose, and the fitted one loses by more than twice as much as switching the model off. The exponent that maximises the likelihood of v0.7’s asks is, in play against v0.7, worse than an exponent with the wrong sign and worse than no opponent model at all.

That is not a paradox once the two objects are named apart. The fit is a *marginal* description of which half-suit a policy names, estimated on depth in the initial deal. The sampler uses it as a *conditional* likelihood, applied to a world set that has already been cut down by the no-bluff certifications, the failed asks and the resolved sets, and it is scored not by how well it predicts the next ask but by where it puts posterior mass over card locations. A re-weighting can be better calibrated about his choices and worse about his cards. Nothing in the estimator we used, which is the standard one, has any purchase on that distinction: it was never asked to.

So the honest position is narrower than either the worry or the reassurance. The self-play validation of §8.5 really was circular, and finding that out took a second engine. But the repair the circularity seemed to imply (adopt the exponent the foreign data want) is measurably worse than the error, and the pre-registration is what stops that from being discovered after the fact. Arm C does not clear +0.15; G2, against the v0.3 champion, therefore does not run, and nothing ships.

In play it is real and not worth buying. Everything above is an argument from likelihood and from a normaliser, not from a duel, so the configuration it points at — at-ask-time depth at $\gamma = 1.0$, the better covariate at the exponent where the model is correctly specified — was pre-registered and run. Six blocks of 1000 duplicate deal-pairs against the champion, sized against a minimum interesting effect of +0.15 fixed in advance, with the run committed to proceed whatever its five screening cells said. They said −0.05, +0.08, +0.32, −0.13 and +0.01, every one of them uninformative at ± 0.53 , which is why they were never load-bearing.

$$+0.1015, \quad 95\% [+0.0063, +0.1967].$$

The interval excludes zero, so by the test fixed in advance the effect is demonstrated. Three things belong next to that, and all three were decided before the data:

It is *below the bar it set itself*. The pre-registration named +0.15 as the smallest effect worth adopting, roughly what the belief-space search delivers, and this is +0.10, with the interval's upper end barely reaching it. Real, and not worth buying. The threshold was chosen before the number, so that judgement is not a reaction to it.

The blocks *disagree*. Cochran's $Q = 12.9$ on 5 dof ($p = 0.024$), $I^2 = 61.4\%$, $\hat{\tau} = 0.150$ — the same size as the pooled effect. Two explanations fit that on its own: the harness understates its uncertainty, or this particular effect varies with the deal population. They are distinguishable, because the identical six-block design has now been run three times and the A/A null twenty-four times:

run	pooled	Q	p	$\hat{\tau}$
lookahead settling	+0.104	4.08	0.538	0.000
precision, 480 vs 160	+0.340	2.92	0.713	0.000
at-ask depth, $\gamma = 1.0$	+0.102	12.94	0.024	0.150
A/A null, 24 blocks	0.000	18.41	0.735	0.000

Every other run gives $\hat{\tau}$ *exactly* zero, including two real effects of comparable and larger size, and so does the null. So the harness is not understating its uncertainty by default, and a deal-population-dependent effect is the better explanation here. We state the multiplicity rather than leave it implicit: four homogeneity tests at the 5% level give roughly a one-in-five chance of at least one significant result under a complete null, so one significant Q out of four points at where to look and does not settle it.

The verdict rests on *two of the six blocks*. Leaving out block 0 gives $[-0.041, +0.168]$ and leaving out block 4 gives $[-0.060, +0.148]$; dropping either one alone puts zero back inside. No block may be dropped for its result and none is (the pre-registered pool is the pool), but under $I^2 = 61\%$, with the interval clearing zero by 6.7% of its own half-width, saying so costs nothing and omitting it would be a choice.

The random-effects interval is $[-0.052, +0.255]$ and *includes zero*. We report it and it decides nothing. The fixed-effect pool was fixed as primary before any of this was known, and adopting a wider estimator once the narrow one is inconvenient is the failure this whole protocol exists to prevent — it is the winner's curse run backwards, choosing the analysis that gives the answer one now prefers. Both numbers are printed; the pre-registered one governs.

So the normaliser argument survives contact with play, and the engine keeps $\gamma = 0.35$ and initial-deal depth anyway. The two parameters were argued for jointly and the screens cannot separate them, so the pre-registration binds them to move together or not at all; they do not move.

The competing story is refuted by the same data. If asks went where the opportunity was rather than where the depth was, propensity would fall as cards held rises. Observed over expected runs 0.46, 1.05, 2.26, 2.64, 2.90 for one through five cards held — monotonically the other way. Whatever else is wrong with the model, its sign is right.

A rule, not a regularity. In 72,091 legal alternatives, not one had an initial-deal depth of zero. That is a theorem rather than a coincidence: asking in a half-suit requires holding a card of it, and cards move only to askers who already hold one, so a player's first card of any half-suit cannot have arrived by play. The consequence matters for the implementation, because the likelihood floors a zero depth at $\log 10^{-9}$ and that floor is therefore excluding worlds the *rules* forbid rather than softening an awkward case. The engine's `allow_bluff_asks` variant does not undo it; that flag widens which cards of a half-suit may be asked for and leaves the half-suit guard alone.

What was changed, and what was not. `gamma_schedule` interpolates between the constant and a weighted quadratic through Table 5 ($\chi^2 = 8.3$ on 4 dof), clamped flat past its vertex because past there the parabola turns up and the measurements do not. Given the paragraph above it is best read as a *dilution correction*: the engine conditions on initial-deal depth, that covariate degrades as the game runs on, and a

v0.4 champion versus	pairs	pair score	set diff / pair	95% CI
v0.3 champion (tuned)	500	0.705	+1.854	[+1.542, +2.166]
— replication	300	0.758	+2.440	[+2.052, +2.828]
— replication, $\gamma = 0.30$	250	0.684	+1.920	[+1.479, +2.361]
v0.3 probabilistic	200	0.792	+2.785	[+2.274, +3.296]

Table 6: The v0.4 champion is `fishbot4` ($\gamma=0.35$): exact/unbiased inference, the opponent choice model, and v0.3’s own ask weights. No new objective term is enabled. Every row was played under the void-era misdeclaration rule; the standard-scoring restatement of this ladder’s headline — the full deployed configuration against the v0.3 champion, both sides pricing misdeclarations correctly — is +2.412 [+2.088, +2.736] over 500 fresh pairs (design R1 of `prereg/rules_award_baseline.md`, `results/award_headline.json`; wholly-held half-suits declared with the wrong split, and so awarded to the opponents: 125 against the deployed side, 146 against v0.3. That counter is the allocation class alone — a declaration naming a card the opponents actually held is a gift rather than a misdeclaration and is counted separately, `fish4/match.py`).

degrading covariate should be believed less. That is a more modest claim than a behavioural one, and a better supported one, and it predicts the term should become unnecessary under the at-ask-time covariate, where nothing drifts. It is normalised by the profile’s mean over the observed distribution of asks, so full strength redistributes the model’s belief across the game without changing how much of it there is in total. That separation is deliberate: γ was already tuned by duels to the broad plateau of Table 4, and a term that moved shape and strength together would leave any positive result unattributable. Zero remains the incumbent, decision for decision.

Whether the reshaped model plays better is a duel, and is not claimed here.

Limitation. This measures the champion against copies of itself, which is exactly the situation the opponent model is used in throughout this paper, and exactly not the situation of a human table. The profile is a fact about this policy’s behaviour, not about Fish.

8.6 Strength, absolute and relative

8.6.1 The strength ladder

Against the two weakest rungs the champion wins every pair: 120/120 against the public-information heuristic (+14.09 sets per pair) and 80/80 against random (+16.30). Those are reported for completeness rather than as evidence; a shutout means the rating gap is a bound, not a measurement.

The headline comparison was run three times at different seeds, and the three estimates span +1.85 to +2.44 — wider than their individual intervals suggest, which is a useful reminder that a 95% interval covers the mean of the deals actually played, and that three such intervals will occasionally fail to overlap. We take the largest experiment as the headline and report the others rather than the most flattering one.

What the strongest configuration is worth, and what nobody has played. Two changes in this version clear a pre-registered bar against the champion: the tripled posterior sampling budget at +0.340 and the belief-space search at +0.104. Each was measured against the champion *alone*, so neither says what their combination is worth, and the combination is not obviously additive, since the search’s whole mechanism is to search a belief the sampler draws.

The stacking run lets the two be chained, because it measured the search *on top of* 480 draws. The two runs use disjoint deal seeds, so they are independent and their standard errors add in quadrature:

link	estimate	pairs
480 draws vs. champion	+0.340 ± 0.049	6000
search on top of 480 draws	+0.072 ± 0.042	6000
chained: search + 480 draws vs. champion	+0.411 ± 0.064	12000

95% CI [+0.285, +0.538] — larger than either change alone and, apart from the opponent model itself, the largest figure in this paper. It is not the largest *single* change: that is the sampler’s +0.340, and this number is two changes at once.

It has now been played directly, which is the only reason the number above can stand unqualified. A chained estimate inherits both links’ assumptions — in particular that each change’s effect does not depend on the deal population the other was measured over. That was plausible, since both pools were drawn the same way, and it was *unmeasured*; this paper’s recurring failure is exactly a plausible unmeasured step carrying a published number. There was also a cheaper reason to run it: **nobody had ever played the shipped configuration against the reference in a single duel**, so every “vs. champion” figure in this document was for a spec that is not the one on the website.

jobs/PREREGISTRATION_combined.md fixes, before any pair: two blocks of 1000 fresh pairs; the A/A 3.796 as the per-pair standard deviation rather than the divergence model, deliberately conservative because this arm differs from the champion by *two* knobs and the model was fitted on arms differing by one; a fixed-effect pool as the only thing that decides; and three outcomes named in advance — agrees, lower, higher, so none could be chosen afterwards. It also states, in advance, that at 2000 pairs against the chain’s 12000 this run *cannot* make the estimate more precise. Its job is to make it honest, and a wider interval is the design working rather than a disappointment.

	estimate	
chained, 12000 pairs	+0.411 [+0.285, +0.538]	indirect
direct , 2000 pairs	+0.357 [+0.191, +0.524]	excludes zero
difference	−0.054 ± 0.107	does not resolve

The chain agrees. The intervals overlap and the difference is half a standard error, so the assumption the chain rests on is now measured and holds, and the indirect caveat can leave this section. Worth noting separately: the design’s sizing assumption held almost exactly (a planned per-pair standard deviation of 3.796 against a realised 3.801), which is unusual enough in this study to be worth one clause, given that the claim-threshold run’s equivalent came in 21% low.

What is *not* settled by this is whether the combination beats the sampler *alone*: that is the stacking run’s own estimate, whose interval contains zero at 68% power. So the defensible ordering remains

$$\text{search} + \text{sampler} \geq \text{sampler} > \text{champion},$$

with the first inequality undemonstrated, and the leftmost quantity now measured directly rather than assembled. The engine names this configuration V04_COMBINED. It is *not* what the site runs, and an earlier draft of this sentence said it was. The deployed configuration is V06_DEPLOYED: the same sampler and depth-3 belief-space lookahead with `claim_forced_exhaustive` on and the endgame correction *off* (`endgame_m` = 0), the correction having been withdrawn under the award rule (§11.9). V04_COMBINED carries that withdrawn arm at `endgame_m` = 2. The run reported here measures the two-knob arm, not today’s deployed spec; both are given verbatim in Appendix E.

It stops paying: precision’s returns diminish. If tripling the sampling budget is worth +0.340, the obvious question is whether tripling it again is worth as much. Pre-registered in jobs/PREREGISTRATION_precision2.md with a *minimum interesting effect* of 0.150 fixed before any pair was played, so that a small positive would be reported as a small positive rather than promoted:

rung	set diff	95% CI
160 → 480 draws, 6000 pairs	+0.340	[+0.244, +0.436]
480 → 1440 draws, 6000 pairs	+0.094	[−0.002, +0.191]
difference	−0.245 ± 0.069	−3.5 SE

Cochran’s $Q = 4.01$ on 5 df ($p = 0.55$), $I^2 = 0\%$, $\tau = 0$: the six blocks of rung 2 agree with each other completely, and with zero.

Rung 2 on its own is **not demonstrated** — its interval includes zero and its point estimate sits below the 0.150 threshold fixed in advance. But the *contrast* is demonstrated, and it is the more interesting quantity. Both rungs multiply the budget by exactly three, so under a log-linear response they would buy the same amount. They do not: the difference between the rungs is -0.245 ± 0.069 , which excludes zero at 3.5 standard errors. Precision’s returns diminish, and they do so somewhere between 480 and 1440 draws.

That is worth putting beside the earlier claim it repairs. The failed-experiment section reported “precision has saturated” on the strength of 200-pair screens around 160 draws, and Section 8.3 refutes that: saturation had not begun at 160. It has begun by 480. The screens reached the right word for the wrong rung, and getting from one to the other took 12,000 pre-registered pairs across two runs.

The default stays at 480. Rung 1 is demonstrated and bought; rung 2 is not demonstrated, costs a further $1.4\times$ on the sampler, and its own pre-registration says a point estimate below 0.150 does not move a default whatever its sign.

A demonstrated effect that is deliberately not shipped. At-ask-time depth at $\gamma = 1.0$ measures $+0.102$ over 6000 pre-registered pairs and its interval excludes zero, so the effect is real. It is also below the 0.15 that `jobs/PREREGISTRATION_at_ask.md` fixed *in advance* as the smallest effect worth adopting, and the upper end of its interval, $+0.197$, barely reaches that. It therefore does not ship, and `depth_mode` stays “initial” in every named configuration. A real effect below a threshold chosen before the data is not a reason to move a default; deciding otherwise after seeing $+0.102$ is exactly what fixing the threshold in advance prevents. A test asserts the default, so the decision cannot quietly reverse itself.

Do v0.3’s objective terms survive better inference? A reasonable worry about the whole result is that v0.3’s hand-tuned tie-breakers were compensating for a poor posterior and would become redundant once it improved. They do not become redundant, but they do become *worth less*. Removing both (setting w_{turn} and w_{scarce} to zero and changing nothing else) costs, over 250 duplicate deals each:

inference	cost of removing both terms	95% CI
without the opponent model	1.660 sets/pair	[+1.240, +2.080]
with the opponent model	1.084 sets/pair	[+0.611, +1.557]

v0.3 measured the same pair at $+1.41$ against its own inference, which sits between the two. So the terms and the opponent model are partially substitutable: about a third of what the tie-breakers were contributing is absorbed once the posterior knows more about opponents’ holdings, and the remaining two-thirds is positional judgement that no amount of card-reading supplies. That is the expected shape if the tie-breakers were doing two jobs — compensating for a blurry posterior, and encoding something genuinely strategic, and it is worth noting that the gains do *not* simply add: 1.85 (the headline) is less than 1.92 (the model alone against the v0.3 champion) plus anything.

8.6.2 Absolute strength: agreement with provably optimal play

Every number above is relative. Table 7 is not. It reports, over 988 exactly solved positions harvested from 260 real games by four different generator policies, how often each agent chooses a *progress-optimal* action — the stricter criterion of Section 5, which excludes value-preserving moves that make no progress.

Four things are worth drawing out.

Every belief-tracking agent is progress-optimal in every information-resolved position, on a corpus four times the size of v0.3’s and 79% of it two-half-suit, and against the stricter criterion. This reproduces v0.3’s strongest correctness claim on harder ground, while Section 5 explains why the one-half-suit version of that claim was shallower than it read.

agent	resolved	uncertain	$m=1$	$m=2$
random	60.8%	32.7%	28.3%	43.9%
heuristic	71.6%	46.8%	34.4%	59.0%
memory	100.0%	61.7%	81.1%	70.1%
probabilistic	100.0%	66.9%	86.8%	73.3%
v0.3 champion	100.0%	67.6%	86.8%	74.0%
v0.4, no opponent model	100.0%	67.5%	85.4%	74.2%
v0.4 champion	100.0%	72.5%	85.8%	78.7%
v0.4 champion, tablebase disabled	100.0%	72.5%	85.8%	78.7%

Table 7: Agreement with progress-optimal play over 988 solved positions, 278 of them information-resolved and 776 with two live half-suits. “resolved” and “uncertain” mean what they did in v0.3; m is the number of live half-suits.

The absolute metric confirms the relative one, independently. Exact inference without the opponent model scores 67.5% under uncertainty against the v0.3 champion’s 67.6%: no change, exactly as the duplicate-deal experiment found and exactly as the ground-truth posterior scores predicted. Adding the opponent model moves it to 72.5%. Three different instruments — paired play, posterior likelihood against the true hands, and agreement with an exact solver — agree on which component does the work.

The tablebase contributes nothing here, and we report that rather than claiming the speedup as strength. Disabling it reproduces the champion’s counts in all six columns of Table 7 exactly — 278 resolved-optimal, 585 uncertain-optimal, 515 uncertain-progress, 182 at $m=1$, 611 at $m=2$. That is consistent with changing no decision and is not a proof of it: the comparison is of aggregates, and two different decision sets can produce the same totals. The per-decision check would be stronger and was not run. In an information-resolved position the marginals are one-hot and the ordinary policy already plays it correctly; in an uncertain position the tablebase refuses by construction. Its value is a $10,900\times$ reduction in the median cost of the endgame path and insurance against a class of error v0.3 demonstrably suffered, not extra strength.

Where the remaining gap is. At $m=2$ under uncertainty the champion agrees with optimal play 78.7% of the time. The missing fifth is not in the endgame in the sense v0.3 meant (these positions *are* the endgame), it is in choosing correctly when the hidden information genuinely has not resolved.

8.7 Where the strength is missing

8.7.1 How much the ask objective leaves on the table

Table 7 is absolute but narrow: it needs a position the exact solver can reach, and it scores agreement with a *perfect-information* optimum rather than value. This measures the objective where it actually operates (positions with hidden information and no solvable endgame) by running one step of policy iteration and reporting the gap.

At a decision with belief b and policy π , the value of ask a is $Q(a) = \mathbb{E}_{w \sim b}[V^\pi(w, a)]$: draw worlds from the posterior, take a , and play the deal to the end with π in all six seats. The one-step regret is $\max_a Q(a) - Q(\pi(\text{root}))$, and it is zero exactly when the objective already picks the best ask available *with respect to its own continuation*. It bounds what any better tie-break could buy.

This is not strategy fusion. Every seat inside a rollout sees only the public history plus the hand the sampled world gives it; nothing consults the world it is being played in. Determinization becomes strategy fusion when the searcher conditions its plan on the sampled world, and here the world only generates observations.

We harvest 40 on-policy positions with at least five half-suits resolved, where the action set is small enough to evaluate *exhaustively* — 18.4 legal asks on average, all of them, rather than a shortlist the objective proposed. Thirty-seven yielded an ask decision rather than a claim.

The estimator cannot be the obvious one. $\max_a \hat{Q}(a)$ over eighteen noisy estimates is biased upward by about two standard errors, so a policy that already played perfectly would measure a regret near $+1.8$.

On synthetic rollouts with twenty-five *truly equal* actions the naive estimator reports +1.83 sets of regret that does not exist. This is the same selection bias Section 9.3.1 is about, arriving at a fourth level: inside the tool built to measure the objective.

Cross-fitting removes it. Choose the challenger on half the sampled worlds, score it on the other half, both ways round; the half that scores it played no part in choosing it, so the estimate is unbiased for the value of the action actually selected, which is at most the best available, making the result a lower bound and never an inflated one. On the same synthetic null it returns -0.004 ± 0.012 over five runs of 4000 trials, and given one action truly worth +2.0 it recovers +0.88.

The result.

cross-fitted one-step regret	-0.097 ± 0.105 , 95% $[-0.302, +0.109]$
naive maximum over actions	+0.579
selection bias, measured	+0.676 sets

The selection bias is the solid finding. Two thirds of a set per decision is the entire difference between reporting that the objective throws away most of a set and reporting that it throws away nothing detectable.

What the interval can and cannot say. Its resolution is set by the noise, and getting that right changed the reading. A first calibration modelled each action’s rollout noise as independent, which is not the design: every action runs through the same worlds with the same seat seeds. Measured properly the paired spread across actions is 1.253 sets against a raw per-world spread of 1.437, so common random numbers barely help — changing the root ask sends the deal down a genuinely different path, and there is little shared variation left to cancel. At that noise the estimator recovers 21% of a true half-set edge and 57% of a full one.

So the honest reading is not that the objective is optimal. It is that its one-step regret is under roughly half a set per decision and this design cannot resolve anything smaller. The same procedure applied to a *random* legal ask returns $+0.128 \pm 0.130$, indistinguishable from zero, which confirms the limit is the design and not the policy: at this budget one-step lookahead cannot reliably beat a coin flip either.

Ranking. Over the same rollouts, the objective’s ordering of asks correlates $+0.164 \pm 0.072$ with what those asks turn out to be worth, against $+0.133 \pm 0.072$ for $\mathbb{P}(\text{success})$ alone. The full objective ranks slightly better than its dominant term, by an amount well inside noise, which is what Section 8.6.1’s tie-break story predicts, and is not on this evidence more than that.

What this says about where to spend effort. Three measurements point the same way, and they are worth stating together because the conclusion is not visible in any of them alone. The objective ranks barely better than $\mathbb{P}(\text{success})$ by itself. Almost every term added *beside* $\mathbb{P}(\text{success})$ (Table 9) is a null. And the largest single effect in this study, +0.340 sets per deal-pair, comes from tripling the sampling budget, which does not touch the objective at all: it makes the $\mathbb{P}(\text{success})$ the objective is already dominated by more accurate.

The reading we take is that in this game the leverage is in the belief rather than in what is computed from it. That is a claim about where the remaining strength is, and it is the reason the next experiments in this line buy posterior quality rather than richer ask features.

We are explicit that the regret bound above contributes *nothing* to this argument. It is per decision, while the play results are per deal-pair, and the two are not commensurable by any factor we can defend — a deal contains dozens of decisions, and one-step improvements do not add. A regret under half a set per decision is compatible with enormous headroom. The argument rests only on the three measurements that are in the same units as each other.

8.7.2 Where the declarations actually go wrong

Section 6.2 predicted where the leverage in declaring would be: not the threshold, but getting the distribution right and playing the forced declarations well. It never checked. A paired margin cannot check it, because a paired margin is blind by construction to anything both arms do equally badly — the difference cancels it. So we built a different instrument.

The path ledger. Every declaration reaches the board through one of four branches of `fish4/agent4.py`, and the ledger (`scripts4/path_ledger.py`) attributes each one to its branch and scores that branch alone. Because every set is declared, sets are conserved and an unattributed declaration is a visible hole rather than a silent omission. Over 1,000 games against v0.7, our seats only:

path	n	per game	wrong	error rate
exact (tablebase)	778	0.778	0	0.000
voluntary (bar 0.97)	3704	3.704	1	0.000
the doomed-ask gate	317	0.317	89	0.281
forced (no legal ask)	178	0.178	103	0.579

The two paths that carry 90% of all declarations have *one* error between them in 4,482 attempts. Every other error, 192 of them, lives in 495 declarations, 9.9% of the total. The prediction was right, and the threshold really is not where the money is.

The engine had a second claim policy, and it did not ask the question. The third row is not a tuned decision. `agent4.py`'s doomed-ask branch fires when the ask it was about to make cannot land, and declares if $\mathbb{P}(\text{exact}) \geq 0.5$ — bypassing the 0.97 bar entirely. `best_candidate` is an argmax over $\mathbb{P}(\text{exact})$ that never reads the second element of the tuple it returns, and that element is $\mathbb{P}(\text{our team holds all six})$ — the number `forced_claim` prices its own version of the same decision with, seventy lines below in the same module. The bar had been 0.5 since v0.3, appeared in no results file, no pre-registration and nowhere in this paper, and the comment above it conceded that “whether a higher bar plays better is an open empirical question, deliberately not settled here by fiat”.

A screen said the gate is *least* accurate exactly where $\mathbb{P}(\text{team}) = 1$, which is precisely where Section 6.2 argues waiting is nearly free. So we pre-registered deferring it there (`prereg/stuck_claim_gate.md`), fixed the ship bar at +0.15 and the withdrawal conditions in advance, and ran 1,000 games per arm on identical deals.

The mechanism moved decisively and the margin did not:

per game, paired	defer at 0.97	middle rung at 0.70
gate declarations	−0.2240 [−0.2558, −0.1922]	−0.2730 [−0.3075, −0.2385]
forced declarations	+0.0840 [+0.0599, +0.1081]	+0.1000 [+0.0728, +0.1272]
wrong declarations	−0.0640 [−0.0847, −0.0433]	−0.0660 [−0.0878, −0.0442]
set margin	+0.0580 [−0.0177, +0.1337]	+0.0400 [−0.0396, +0.1196]

The gate's own error rate falls from 0.281 to 0.086; the voluntary path absorbs the deferrals and stays at one error in 3,795; gate claims fall by 0.224 against a rise of 0.084 in forced, so the deferral is not being paid for at the deadline. A third of our remaining wrong declarations disappear, with an interval nowhere near zero, and neither margin interval clears zero, so under the pre-registration **nothing ships** and the knob stays at its inert default.

At two sets of differential per avoided error the ledger predicts +0.128 and the duel measures +0.058. The prediction sits inside the interval, so this run neither confirms nor refutes the conversion; what it establishes is that whatever the conversion is, it does not clear a bar fixed in advance at +0.15 on a thousand games. The pre-registration's expected-outcome paragraph, written before the run, predicted exactly this and added that a result above +0.15 should have prompted suspicion rather than pleasure.

Two experiments have now produced that shape, and the arithmetic of it is measurable. The gate is the second. The first was the signalling protocol of §10.4. Both were filed at the time as “a real reduction that buys no sets”, which was honest at that power and has since turned out to be wrong about both. The signalling protocol's registered confirmation is +0.1180 [+0.0325, +0.2035] and +0.1220 [+0.0291, +0.2149], and the deferred gate was re-run at power on the current engine under `prereg/defer_gate_at_power.md`: over 2,000 deal clusters it returns +0.0455 [+0.0134, +0.0776],

clear of zero (`results/signal_vs_defer.json`). Both buy sets. Neither clears the +0.15 bar, which is the reason neither ships, and the thousand-game runs that produced the original filings could not have separated “buys nothing” from “buys a twentieth of a set”.

A paired margin is a *net*. An arm that defers declarations avoids some errors, worth $+v$ sets each, and leaves half-suits live, which costs something per deferral; the duel reports the sum. Regressing the paired margin difference on the paired *error* difference separates them, and both halves are quantities this paper has assumed rather than measured — the two-sets arithmetic of §11.9 is exactly the assumption $v = 2$. It costs no games: the journal is already on disk (`scripts4/error_value.py`).

	defer at 0.97	middle rung at 0.70
value of one avoided error	+1.7898 [+1.5927, +1.9870]	+1.8388 [+1.6433, +2.0344]
cost when it avoids nothing	−0.0565 [−0.1237, +0.0106]	−0.0814 [−0.1513, −0.0114]

And without a model, since the regressor takes four values and a line through four levels is close to a two-point comparison anyway — the paired margin by how many errors the arm avoided *on that deal*:

on deals where the arm	n	paired margin
added an error	10	−2.0000 [−3.2396, −0.7604]
avoided none	930	−0.0473 [−0.1019, +0.0072]
avoided one	48	+1.5000 [+0.8216, +2.1784]
avoided two or more	12	+4.1667 [+2.1819, +6.1515]

Three things follow. The two-sets arithmetic is right, and now measured rather than assumed: an *added* misdeclaration costs exactly −2.0000 sets, which is the award rule’s own arithmetic recovered from play rather than asserted from the rulebook. Deferring a declaration costs about 0.36 to 0.43 sets each, and for the heavier arm that cost excludes zero — a price this project did not have, and precisely the quantity the pre-registration called out as unattributable from the margin alone. It is attributable, just not from the margin alone: 0.0640 errors avoided at 1.79 is 0.1146, less 0.0565 of tempo, is the +0.0580 the duel reported.

The third is the weakest and is stated as a direction rather than a figure. Avoiding an error returns less than adding one costs, +1.50 against −2.00, which would follow if a misdeclaration you decline to make does not hand you the set — the half-suit stays live and can still be lost. At $n = 48$ and $n = 10$ the two intervals overlap comfortably and this is a hypothesis the design was not built to settle.

What the intercept is, precisely, because the first draft of this paragraph overclaimed it. Both arms play every deal, so the error difference is a deterministic function of the deal and conditioning on it selects *deals* rather than assignments, which does make the estimate causal on that subpopulation. It is not the cost of deferring in general: the subpopulation is selected by an outcome, and deals where no error was avoided defer about 30% less often than average (0.156 against 0.224 per game). The per-declaration figures above therefore divide by that subpopulation’s own deferral rate. Dividing by the overall rate, which is what a first pass did, understates the price of a deferral by exactly that 30%.

The fourth row is a search problem, and it is the one thing this section produced that shipped. A forced declaration — cards in hand, no legal ask anywhere — is wrong 57.9% of the time. The engine reaches its split there the same way it reaches every other one: `best_for_half_suit` shortlists two candidate holders per card and scores three combinations. At the last live half-suit that is a shortlist over a space small enough to enumerate completely, and enumerating it is not a change of policy but a better search of the same objective.

`claim_forced_exhaustive` does exactly that: at or below one live half-suit it enumerates the whole team assignment space under the joint, fixes any card the propagator has already pinned, and (the property that makes it a search fix rather than a new policy) never returns a split the joint scores below the incumbent’s. Registered at `prereg/forced_exhaustive.md` with the ship bar and the guards fixed first, it cleared in both populations:

the same 16,200 declarations, cut two ways




Figure 3: The lever everyone reaches for, and the one that carries the risk. *Left:* misdeclaration rate against how many of the six the declarer holds. It is not even monotone — holding all six is the safest state there is, so the curve turns over at the right-hand end, and any policy tuned on this covariate is tuned on a confound. *Right:* the same declarations against how many of the six have *never been publicly located*. That one runs from zero errors in 4,980 declarations up to eleven percent. The dip at six is 30 errors in 342 declarations and its interval covers the point before it. Bars are 95% Wilson intervals, which is why the zero cell has a visible upper limit rather than a degenerate one.

population	effect	interval
self-play, 2,400 games	+0.0233	[+0.0133, +0.0334] sets/game
vs. v0.7, last declaration correct	28/87 → 37/87	paired +0.0090 [+0.0031, +0.0149]
vs. v0.7, set margin	+0.0180	[+0.0063, +0.0297]

with accuracy in every other live-half-suit bucket unmoved, which was guard 2 and which it passed exactly. It is in V06_DEPLOYED.

Two things about those intervals are worth noticing, because they look too tight for 1,000 games and they are not. The knob fires only on forced declarations at the last half-suit, so 99.1% of games finish bit-identical in both arms and pairing removes almost all the variance — $414\times$, the top row of Table 2. And the effect is correspondingly small: a rare branch played better is a rare improvement, however cleanly measured. Both facts are the same fact, and §7.2 is where it is priced in general.

What the declaration is actually risking, and a lever that looked free and was not. Of our 0.1759 wrong declarations a game, 0.1676 are allocation class (our own team held all six and we named the wrong split) against 0.0083 ownership errors. That is 95.3%, and it is a different problem from reading the opposition: every card is held by someone who knows they hold it, so the team has the answer and no member of it does.

The obvious lever is that a declaration may be made by any member of the team, on their own turn. Someone holding four of the six guesses two; someone holding one guesses five. Over 1,800 self-play games and 16,156 wholly-held declarations, 30.4% were made by a player a teammate could have out-informed, so the opportunity is real. The lever is not:

The error rate *rises* with the declarer’s own holding, so deferring to the better-placed teammate would move declarations up that curve. Six is trivially perfect — hold every card and there is nothing to name.

The right variable is on the right, and it is selection that connects them. A player holding one card of a half-suit only declares it when the other five are already pinned; holding five leaves exactly one card unaccounted for, and if that card was dealt and never asked for it is a coin flip between two teammates on a set that feels certain. Held at fixed “never publicly located” the k -effect flattens: in the best-populated stratum (4 unlocated, $n = 2,344$) the error rate across $k = 0 \dots 5$ runs 0.050, 0.068, 0.057, 0.057, 0.042, 0.077, with no trend. **What a declaration risks is not how many cards you are missing but how many of them have never moved.**




Figure 4: Every misdeclaration this engine makes is compelled. Split by the path the declaration arrived on: the two paths it takes because it *wants* to be exactly perfect over 14,678 declarations between them, and both paths where it declares because it *must* — the endgame gate, and a forced declaration with no cards left to ask with — carry the whole error rate. This is why the headroom in §8.7 is not a calibration problem. The policy already declines every declaration it is entitled to decline.

declarer's own cards	<i>n</i>	error		<i>n</i>	error
1	2,640	0.017	<i>cards never located:</i> 1	4,980	0.000
2	2,024	0.048	2	4,234	0.006
3	1,282	0.042	3	2,990	0.032
4	1,494	0.063	4	2,344	0.055
5	1,830	0.068	5	1,266	0.111
6	6,468	0.000	6	342	0.088

Table 8: The declarer-holding lever against the mediator, on the same 16,156 wholly-held declarations. *Left:* error rate by how many of the six the declarer holds — the lever, which is not monotone and peaks at five. *Right:* by how many have never been publicly located, which is what the risk actually tracks. Figure 3 plots both with intervals.

The split is not frozen. It is natural to say that once a team holds all six no opponent may legally ask in the half-suit, so nothing further can inform the split. The first clause is right and the conclusion does not follow. Hand counts are public and `fish/beliefs.py::_propagate` uses them: a player whose quota reaches zero is excluded from every unresolved card, and a player whose quota equals what they could still be holding has those cards pinned. So the split keeps resolving after it freezes, not through anything happening in that half-suit, but through the rest of the game draining hands.

Measured, over the same corpus, as cards pinned to a specific player beyond the declarer's own hand and beyond anything a successful ask publicly located: 46.6% of wholly-held declarations derive nothing, 33.2% derive one card, 15.1% two, 5.1% three or more. 8,626 of 16,156 wholly-held declarations, 53.4%, carry at least one derived card, and the error rate falls from 0.046 with none to 0.011 with one and 0.000 with three. **Deriving cards after the freeze is the common case, not the exception.** It is what makes a declaration safe. The doomed-ask gate had already implied as much and we read it the other way: it deferred 0.224 declarations a game while adding only 0.084 forced ones, so about 0.14 moved into a *confident* declaration, which is only possible if waiting supplied something.

Is that where the headroom is, though? The paragraphs above say where our *errors* are. Those are not the same question, and the difference is settleable by handing a seat information it should not have and seeing what it does with it.

Every figure in this paragraph is a bound obtained by cheating. The arms are `fish4/oracle.py` handed the true deal; the opposition never cheats. They are not strength measurements, they do not belong in a ladder, and they must not be read beside an honest margin. Over 600 games per arm on identical deals (`prereg/information_ceiling_split.md` → `results/ceiling_split.json`):




Figure 5: What perfect card-reading is worth, one side of the table at a time. **Every value here is obtained by cheating:** the arm is handed the true deal and the opposition never is, so these bound what information could buy and are not strength measurements. They never appear in a ladder. The two partial arms are told almost exactly the same *number* of cards — 40.7 against 40.4 a game — and the teammate side is worth 2.6× the opponent side anyway. We predicted the reverse. The two do not sum to the whole (4.72 against 6.61), so the sides are complements rather than separable contributions.

told the true location of	cards pinned/game	ceiling over honest
its teammates' cards	40.7	+3.4100 [+3.1625, +3.6575]
its opponents' cards	40.4	+1.3067 [+1.0070, +1.6063]
everything	59.8	+6.6067 [+6.4004, +6.8129]

The middle column is the point. The two partial arms are told *almost exactly the same number of cards* and one is worth 2.6 times the other, so this controls for how much a seat knows and isolates what it knows about. The allocation/ownership split in the error ledger does translate into a split in information value, and the answer to the question in this paragraph's title is yes.

Two cautions, both of which the pre-registration owes the reader. The teammate arm's zero allocation errors are *definitional* — a seat that knows every teammate's cards and its own reads the split off rather than inferring it, so that cell could not have come out any other way and confirms nothing. And the arms are sub-additive, +4.7167 against +6.6067 for omniscience, where the registration braced for the opposite: knowing where a card is *not* pays only if you can also act on it.

9 What did not work, what search buys, and what a learner recovers

9.1 Everything else we tried

Table 9 is the rest of the study. It is mostly nulls, and it is reported in full because a table of only the things that worked would misrepresent how much was tried. The ten arms of the v1.2 programme (P43–P45) are not in this table: each was registered against a dual-population bar (vs SESTINA through BRIDGE_REV 3 and in self-play against the v1.1 champion) rather than as a self-play cell, and each is accounted for in §11.11 and Appendix A, the six that reached a duel with their intervals; none cleared either half of the +0.15 bar. Each cell here is a duplicate-deal duel against the champion with exactly one thing changed, and the last column is the smallest effect that cell could have resolved at 80% power given a per-pair standard deviation of 3.869 sets.

That column uses one standard deviation for every row, and we now know it is a property of the comparison rather than of the game. Since the correction is conservative (the constant is the noisiest case), it can only have made these cells look *less* informative than they were, which raises the question of whether any of them was in fact resolvable. Rescoring each cell that stored its per-pair differentials against its own measured noise moves the resolvable effect by a median factor of 0.92, and by 0.60 in the quietest, and changes no verdict at all.

That last clause was where we stopped, and it was the wrong check. “ $|\text{est}| > \text{MDE}$ ” is a test at $\alpha \approx 0.005$, because the MDE is $(z_{0.975} + z_{0.80}) \text{se} = 2.80 \text{se}$; everywhere else in this paper a cell is said to

what changed	setting	pairs	set diff	95% CI
<i>inference precision</i>				
draws per decision	96 (from 160)	160	−0.300	[−0.852, +0.252]
draws per decision	320 (from 160)	200	−0.035	[−0.578, +0.508]
count mode	$\sqrt{\cdot}$, $\gamma=0.70$, 160	200	+0.480	[−0.042, +1.002]
count mode	$\sqrt{\cdot}$, $\gamma=0.70$, 320	200	−0.165	[−0.701, +0.371]
no-declaration signal [†]	$\lambda = 0.9$	200	+0.205	[−0.301, +0.711]
no-declaration signal [†]	$\lambda = 1.8$	200	+0.190	[−0.313, +0.693]
<i>new ask-objective terms</i>				
exposure	0.40	160	−0.113	[−0.660, +0.435]
claim progress	0.30	160	−0.206	[−0.807, +0.395]
signalling	+0.25	160	−0.300	[−0.890, +0.290]
signalling	−0.25 (conceal)	160	+0.056	[−0.530, +0.642]
information gain	0.10	160	−0.394	[−0.932, +0.144]
concentration [‡]	0.15	160	−0.037	[−0.653, +0.578]
turn-risk retune	0.35 (from 0.60)	160	−0.075	[−0.670, +0.520]
scarcity retune	0.35 (from 0.20)	160	+0.125	[−0.381, +0.631]
<i>claiming and the endgame</i>				
claim threshold	0.99 (from 0.97)	200	+0.010	[−0.010, +0.030]
informed pass choice	on	160	+0.013	[−0.012, +0.037]
tablebase	one half-suit only	200	+0.005	[−0.005, +0.015]
tablebase	disabled entirely	200	+0.005	[−0.005, +0.015]
<i>the learned half-suit value objective</i>				
blended	$w_{\text{value}} = 0.5$	200	+0.410	[−0.086, +0.906]
blended	$w_{\text{value}} = 1.5$	200	+0.040	[−0.502, +0.582]
replacing the objective	pure ΔV	200	−7.355	[−7.875, −6.835]

Table 9: Everything else. Only the last row resolves, and it resolves against us. Every other interval contains zero at a minimum detectable effect of roughly 0.8 sets per pair, so these are “not shown to matter at this budget”, not “shown not to matter”.

have resolved something when its 95% *interval excludes zero*, which is 1.96 se. We had been auditing the table against a bar 43% higher than our own definition, and since the gap between the two MDEs is a median 5% of the MDE, finding no change was close to arithmetically guaranteed.

Under the definition the rest of the paper uses, **five cells change**: the retake penalty at −0.340 [−0.663, −0.017], a decisive lookahead cell at +0.307, and three blocks of pre-registered runs. All five sit below the MDE their row quotes and have intervals excluding zero. The MDE column says what a cell could *detect* at 80% power, which is a stricter and different thing from what it did *resolve*, and reading the column as “this cell found nothing” understates them. The retake cell is the sharpest case: another pre-registration in this study treats that −0.340 as an established effect and designs a follow-up against it, while this table files it as a null. Both readings cannot be right, and the interval is the one that governs.

None of this widens a published interval — each cell’s interval was always computed from its own sample variance, and the constant only ever fed the MDE column. What it corrects is how that column should be read. The cells that recorded only a mean and an interval cannot be re-examined at all, which is itself the argument for storing per-pair data, and why v0.4 does.

Four of these deserve comment.

[‡] **This row is retained as history and does not describe the code.** Both halves of it were incapable of finding an effect, and the reason is worth more than the row.

The *formula* could not express the preference it was written for. `concent` at v1 was `ctx.team_concentration[hs]`: one number per half-suit, identical for every candidate ask in it and independent of the target and of who would end up holding the card. A term taking the same value on every ask in a half-suit can tilt the choice of half-suit and nothing else. Worse, its sign is backwards in the case it exists

for — when the concentration sits with a *teammate*, taking a card breaks it up, and v1 scored that ask highest precisely because the half-suit was concentrated. This is the defect claim carried at v1, and it takes the same remedy: corrected in place to the expected *change* the ask would cause, with `TERM_VERSIONS["concent"]` bumped to 2 so that `stale_terms()` flags every stored harvest fitted against the old column.

And the *weight* was inert. Measured over 596 decisions with a real choice and 30,065 candidate asks (`scripts4/concent_scale.py` → `results/concent_scale.json`), the corrected feature has median magnitude 0.0299 and a median within-decision spread of 0.0677. At the screened weight of 0.15 it changes which ask is taken on 1.7% of decisions; it takes a weight near 1.0 to reach 8%. A 160-pair run of a knob acting on one decision in sixty, reporting an interval four times the ship bar in each direction, is a statement about the harness rather than about the idea.

The retest mattered because this is the only term in the basis that points at the error we actually make. Of our 0.1759 wrong declarations a game, 0.1676 are allocation class (our own team held all six and we named the wrong split) against 0.0083 ownership errors, and a half-suit held entirely in one hand has no split to name. It was registered at `prereg/concentration_v2.md` at 4,000 games, with arms at 0.60 and 1.50 and with the mechanism as a *withdrawal condition* rather than a secondary: allocation errors had to fall and concentration at declaration time had to rise.

It ran and it is withdrawn, on condition 1. At $w = 0.60$ the effect is $-0.0430 [-0.1369, +0.0509]$ and at 1.50 it is $-0.1445 [-0.2490, -0.0400]$, worse than shipped. Allocation errors were required to fall and rose monotonically instead: 0.1557, 0.1772, 0.1827. Condition 2 passed — concentration really did rise, $0.8106 \rightarrow 0.8168$, so the corrected feature computes what it claims and the *strategy* is what is wrong, which is a more useful outcome than a null.

The reason is the resource the argument for the term ignored. `GameState.legal_asks` requires the asker to hold a card of the half-suit, so concentrating a team's holding into one hand narrows the set of half-suits that hand can ask in, and being unable to ask is the definition of being forced. Forced declarations rise monotonically with the weight, $0.167 \rightarrow 0.197 \rightarrow 0.218$ a game, and the forced path is 47–49% wrong. The term buys a marginally better split and pays for it in the one currency that keeps a seat able to act. (Stated as consistent-with rather than proven: the run does not instrument askable half-suits per seat, and the monotone dose response in the harm is the strongest evidence here.)

[†] **These two rows measured nothing, and we found out afterwards.** The no-declaration term was assembled into the importance weights and then overwritten: the batch sampler built it first and the opponent-model branch that follows did `log1 =` rather than `log1 +=`, so whenever any non-self player had already asked, 632 of 641 decisions in `results/ess_probe.json`, the term was discarded before use. In the handful of decisions where it survived, it addressed the half-suit by column indices built against the caller's free-card order and applied to a matrix whose columns are in the sampler's own sort; six of eight columns differ on a typical position, so the “did the opponents take this whole half-suit” test read a mixture of cards from other half-suits. And the scalar reference sampler stored the term and never applied it at all, so the batch/scalar agreement test could not see any of this.

Both cells are therefore reported as *invalid rather than null*: they are what a disabled feature scores, which is the champion plus sampling noise, and the intervals above are consistent with exactly that. All three bugs are fixed and pinned by tests that fail against the previous code; the feature has not been re-measured, and until it is there is no result here in either direction. Nothing else in this paper is affected, because `opp_lambda` defaults to zero and no other cell sets it.

Precision appeared to have saturated, and it had not. v0.3 found that raising the sampled-world count from 32 to 96 was worth +0.54 sets per pair and concluded that belief precision had not saturated. Under v0.4's inference these screens said it plainly had: halving or doubling the draw budget around 160 changed nothing measurable, and the natural reading was that v0.3 had been buying accuracy rather than precision (more draws partially averaging out a *bias*), so that once the bias was gone, sampling noise was no longer binding.

That reading is refuted, by a pre-registered 6000-pair run reported in Section 8.3: going from 160 draws to 480 is worth +0.340 sets per deal-pair, 95% CI $[+0.243, +0.436]$, all six blocks positive,

$I^2 = 0\%$. Precision had not saturated; the screens that said so were 200-pair cells resolving about ± 0.53 , which cannot distinguish $+0.340$ from zero and returned exactly the null they were bound to return. Their agreement with each other was not corroboration — underpowered cells agree on nothing in particular, and four of them agreeing on a null is what four underpowered cells do. The paragraph is kept in its original form above because the inference it drew was reasonable on the evidence available and wrong anyway, which is this paper’s most frequently repeated lesson and is more useful stated against a concrete case than in the abstract.

The word was right and the rung was wrong, which is the more interesting correction. A second pre-registered run at $480 \rightarrow 1440$ draws returns $+0.094 [-0.002, +0.191]$, and the contrast against rung 1’s $+0.340$ is -0.245 ± 0.069 , so precision’s returns *do* diminish, just not where the screens said. Saturation has not begun at 160 and has begun by 480. See Section 8.3.

One cell technically resolves, and should still be ignored. Lowering the claim bar from 0.97 to 0.90 gives $+0.035 [+0.007, +0.063]$: an interval excluding zero, and an effect of one fortieth of a set per deal-pair. After forty-odd cells at a nominal 95% level, an interval that narrowly excludes zero at a magnitude this small is not evidence of anything, and we report it rather than promoting it — particularly having already watched a larger apparent effect fail to replicate twice (Section 9.2). The neighbouring cell in the other direction (0.99) gives $+0.010 [-0.010, +0.030]$, which is the shape of a policy that barely changes its decisions at all. This one has since been run properly, and Section 9.4.3 reports what an unselected 2000 pairs did to it.

The claim threshold and the tablebase are exact no-ops. The intervals $[-0.010, +0.030]$ and $[-0.005, +0.015]$ are not small effects; they are almost no divergence at all. Raising the claim bar from 0.97 to 0.99 changes essentially no decisions, reproducing v0.3’s bimodality finding on a much better posterior. Disabling the exact endgame solver entirely changes essentially nothing, which agrees exactly with the absolute benchmark (Section 8.6.2), where it changed not one of 988 decisions.

Replacing the objective with the learned value is catastrophic. The half-suit value model is good — held-out log-loss 0.727 against 0.821 for the class prior and 0.743 for the best single feature, and a mean calibration bias of -0.006 . That last figure flatters it, and the flattery is worth one sentence because “well-calibrated” is doing work in the argument below. Five of its ten deciles are miscalibrated at the 95% level: it is over-confident at both ends (decile 0 says 7.6% and delivers 0.5%; decile 9 says 92.7% and delivers 90.0%) and under-confident in the middle (deciles 4 and 5 both run about 3 points high). The mean bias is near zero because those errors cancel, not because each decile is right. Scoring asks by the expected change in it, in units of sets, loses 7.4 sets per deal-pair. Blending it in at low weight is harmless and does nothing. This is the sharpest confirmation of v0.3’s central claim about this game that we have: the probability that an ask lands is not one term among several, it is the objective, and a decently-calibrated model of what a half-suit is worth is no substitute for it.

The eleven-term basis bought nothing. Neither the two terms motivated by the literature survey (exposure, signalling) nor the four motivated by inspection improved on v0.3’s three. The signalling term is worth singling out because its sign was the interesting question: $+0.25$ (advertise a holding in a half-suit our team is winning) scored -0.300 and -0.25 (conceal it) scored $+0.056$, so if anything the concealment direction is favoured, but neither interval comes close to excluding zero.

9.2 The near-misses, and the one that resolved

Two cells in Table 9 had positive point estimates whose intervals only just included zero, which is exactly the situation in which a screening study manufactures a false positive if the best cell is promoted without re-measurement. We reran both at 500 pairs on fresh seeds (minimum detectable effect 0.48 rather than 0.77) together with their combination:

candidate	screen (200 pairs)	retest (500 pairs)	retest 95% CI
blended value objective, 0.5	+0.410	−0.208	[−0.534, +0.118]
$\sqrt{\cdot}$ counting, $\gamma = 0.70$	+0.480	+0.212	[−0.092, +0.516]
both together	—	+0.490	[+0.184, +0.796]
replication, fresh seeds	—	+0.068	[−0.254, +0.390]
replication again	—	−0.066	[−0.395, +0.263]

When this was written the combination was the only cell in the study (out of more than twenty-five) whose interval excluded zero in the improvement direction, which is exactly why it looked like something. On a fresh set of 500 deals it returned +0.068, and on a second fresh set −0.066. It was noise.

That framing no longer holds, and the reason it stopped holding is the point. Three cells in this paper now exclude zero upward — the pre-registered lookahead at +0.104, posterior precision at +0.340 and at-ask depth at +0.102, and every one of them was *designed* to be able to, at thirty times this cell’s sample size with the analysis fixed in advance. The distinction that matters is not how many intervals clear zero, it is whether clearing zero was a possible outcome of a design or the best of twenty-five looks.

An earlier draft of that sentence said *four*, and counted the claim threshold’s +0.035 among them. That was the error the sentence is about, committed inside the sentence about it: the claim-threshold figure is a screening cell selected out of 103 for excluding zero, which is the *other* category entirely. It is the eighth appearance of the pattern this paper keeps documenting, the fourth of them one level above where the error was last caught, and the only one caught by a *run* rather than by rereading — the confirmatory 2000 pairs of Section 9.4.3 put the effect at +0.002.

We report this at length because it is the single most useful methodological observation the study produced. Twenty-five screening cells at a nominal 95% level will produce roughly one apparent winner under a complete null, and this one arrived with a respectable-looking interval, a plausible story (two independently-motivated changes that stack, exactly as v0.3’s two winning terms had) and a point estimate twice the sum of its parts, which should have been the tell. Promoting it would have shipped a change worth nothing and, worse, would have anchored the next version’s baseline on it.

9.3 Search over the belief, measured

9.3.1 Belief-space search: four rounds to one answer

Section 6.3 describes the one search design the variance diagnosis does not rule out, and predicts, in Proposition 3, that a version of it without the quota coupling would collapse to greedy. Both halves were measured. Table 10 is the whole record, and it took 9,700 deal-pairs across four rounds — 800 in screens, 1500 in replications, 1400 decisive, 6000 in the pre-registered settling run — to answer a question a 200-pair cell appeared to answer at the start.

The answer is that the search helps, by about a tenth of a set per deal-pair. The section is long because the route to that number is the more useful part: the first cell overstated it by a factor of five, the next three rounds disagreed with each other, and the disagreement turned out to be a fact about selection rather than about the harness or the game.

Three runs of one configuration disagree. The screening cell, at 200 pairs, returned +0.570 and excluded zero. Following the replication policy of Section 9.2 we reran the identical configuration twice more at 500 pairs, changing nothing but the seed streams. The two retests returned −0.044, 95% CI [−0.321, +0.233], and +0.386, 95% CI [+0.086, +0.686]. The second of those excludes zero. The first does not, and its interval excludes the value the second reports.

That is not a result about the lookahead. It is a result about the experiment. Three estimates of one quantity, differing only in seeds, are mutually inconsistent: Cochran’s $Q = 7.03$ on 2 degrees of freedom, $p = 0.030$, with $I^2 = 71.5\%$, so roughly seven-tenths of the spread between these runs is *not* the sampling variation their own intervals model.

The within-run intervals are not the problem. The per-pair standard deviation recovered from each cell separately is 3.21, 3.17 and 3.42 sets. Those agree with each other and with the 3.869 figure used

cell	pairs	set diff	95% CI	MDE
depth 3, $w = 0.25$	200	+0.570	[+0.125, +1.015]	0.75
depth 2, $w = 0.25$	200	+0.345	[−0.083, +0.773]	0.75
depth 3, $w = 0.60$	200	−0.080	[−0.583, +0.423]	0.75
depth 3, <i>no quota coupling</i>	200	+0.070	[−0.369, +0.509]	0.75
depth 3, $w = 0.25$ — retest	500	−0.044	[−0.321, +0.233]	0.48
— retest again	500	+0.386	[+0.086, +0.686]	0.48
<i>no coupling</i> — retest	500	+0.112	[−0.164, +0.388]	0.48
depth 3, $w = 0.25$ — decisive A	700	+0.307	[+0.063, +0.551]	0.40
— decisive B	700	+0.001	[−0.240, +0.243]	0.40
depth 3, $w = 0.25$ — settle, block 1	1000	+0.179	[−0.034, +0.392]	0.34
— block 2	1000	−0.009	[−0.217, +0.199]	0.34
— block 3	1000	+0.118	[−0.089, +0.325]	0.34
— block 4	1000	+0.241	[+0.040, +0.442]	0.34
— block 5	1000	+0.064	[−0.147, +0.275]	0.34
— block 6	1000	+0.028	[−0.175, +0.231]	0.34

Table 10: Belief-space lookahead against the v0.4 champion, which is the identical policy with the bonus disabled. MDE is the smallest effect that cell could have resolved at 80% power, so a null row says the run did not detect an effect of that size, not that there is none. This table uses the per-pair standard deviation of 3.796 measured over the 4800 A/A pairs reported later in this section, rather than the 3.869 used elsewhere in this paper: the settling run was powered against 3.796, and a table reporting a run against a different resolution from the one it was designed for would misstate what it could see. The two differ by under 2%.

throughout this paper, so each interval correctly describes the deals that run actually played. The natural inference is that the variance the intervals miss sits *between* runs; the DerSimonian–Laird estimate of such a component from these three cells is $\tau = 0.268$ sets per deal-pair. *We measured that inference directly and it is wrong* (see below), but it is worth following through first, because the two pooled answers it produces are the reason the question mattered enough to measure.

Pooling all 1200 pairs on that assumption gives two answers, and only one of them is usable:

pooling	estimate	
fixed effect	+0.226 [+0.041, +0.411]	excludes zero
random effects	+0.278 [−0.083, +0.638]	includes zero

A fixed-effect pool assumes the runs are three samples of one true value, which is precisely what Q appears to reject; the random-effects interval, which adds τ^2 to every cell, would then be the one reflecting where a fourth run might land, and it includes zero. Read that way the effect is unresolved and could be reported neither as working nor as a null, which is where we left it, and which is what made the size of τ worth establishing rather than assuming.

The three cells prompted a question about the harness, and the answer is that the harness is fine.

The reading above, that a duplicate-deal cell has a variance component its own interval does not model, would, if true, make the minimum-detectable-effect table of Section 7 optimistic for every single-cell result in this paper. That is too consequential to leave inferred from three points, so we measured it.

The instrument is an A/A: one policy against a copy of itself, where the true differential is exactly 0. That required a change to the harness, and the reason is worth recording, because it means this question could not previously have been asked. Under the default seat-based seeding an A/A plays a bit-identical game in both halves of a duplicate pair, so its differential is $(a - b) + (b - a) = 0$ for every deal — a point mass, not a small effect correctly measured as zero. The harness could not observe its own null. Giving each side an independent randomness stream (`independent_seeds`, off by default so the archive stays comparable) makes the null a real random variable.

Twenty-four blocks of 200 pairs, each with its own deal seeds and its own agent seeds — differing in exactly the way two separate runs differ:

nominal 95% intervals covering the truth	23/24 = 95.8% [79.8, 99.3]
Cochran's Q	18.41 on 23 df, $p = 0.735$
I^2	0.0%
τ , between-run standard deviation	0.0000
grand mean of the 24 estimates	-0.0215 (truth: 0)
per-pair standard deviation, 4800 pairs	3.796 (this paper uses 3.869)

Coverage is nominal. The observed spread of the block estimates, 0.2454, is *smaller* than their own standard errors predict, 0.2701; there is no excess variance to attribute. A duplicate-deal interval means what it says.

We retract the between-run variance. An earlier version of this section reported $\tau = 0.268$ sets per pair from the three cells above and argued that it made every small result in this paper more fragile than it looks. That was a Type I error. The $Q = 7.03$, $p = 0.030$ which produced it is exactly the one-in-thirty event, and we promoted it to a structural claim about the harness on the strength of three points, which is the error this paper devotes Section 9.2 and two of its four lessons to warning against, committed inside the section doing the warning. The minimum-detectable-effect table stands as written.

The lookahead is left genuinely open, on better terms. If the cells are homogeneous (and a harness that produces no excess heterogeneity under the null is strong evidence that they are), then fixed-effect pooling is the valid analysis, and over all 1200 pairs the estimate is +0.226, 95% CI [+0.041, +0.411].

We did not report that as a result. It excludes zero by 0.04; this configuration had already produced one screening cell that failed to replicate; and reaching for the pooling that happens to resolve, immediately after the pooling that happened not to, is the same error in the opposite direction. So we fixed a size in advance and ran it: two further cells of 700 pairs each, on fresh seeds, pre-registered as decisive.

They returned +0.307, 95% CI [+0.063, +0.551], and +0.001, 95% CI [-0.240, +0.243].

The verdict, on the test we committed to. Pooled, those two cells give +0.153, 95% CI [-0.019, +0.324], which *includes zero*. The experiment designed in advance to settle this does not settle it, and we report that rather than the alternative below, because a post-hoc pool does not get to overrule a pre-registered one.

What the fuller picture says, and why we still do not claim it. Dropping the original screening cell — it was *selected* for resolving, so including it in any pooled estimate is winner's curse, and pooling the four unselected cells over 2400 pairs gives +0.153, 95% CI [+0.022, +0.285]. The point estimate is identical to three decimal places across both analyses, which is fairly strong evidence that a small real effect exists, of order a seventh of a set per deal-pair. It is not evidence that we have demonstrated one.

Dropping that cell also disposes of the heterogeneity this section opened with: Q falls to $p = 0.063$ over the four unselected cells and $p = 0.081$ over the two decisive ones, neither significant. Combined with the A/A measurement above, the story is simply that one cell was selected for being large and the rest is sampling noise — no between-run variance, and nothing wrong with the harness.

The sizing was wrong, by the error this section is about. We chose $n = 2600$ on the grounds that its minimum detectable effect of 0.21 would separate +0.23 from zero. But +0.23 was the pool *including* the selected cell — the inflated estimate. Against the unbiased +0.153, a correctly powered test needs roughly 5000 pairs. We committed the winner's curse in the power calculation for the experiment whose purpose was to correct for the winner's curse, which is the third time in this study the same error has appeared one level up from where it was last caught. It appears a fourth time in Section 8.7.1, inside the estimator built to measure the ask objective, where taking a maximum over eighteen noisy action values would have reported two thirds of a set of regret that does not exist. The pattern is worth naming: each occurrence is not a repetition of the last but a level above it, and none was caught by being more careful — each was caught by measuring the thing that had been reasoned about.

The pre-registered run, and the answer. So we wrote the analysis down before running it. `jobs/PREREGISTRATION_lookahead.md` fixes, in advance: six blocks of 1000 fresh pairs; sizing against the *unbiased* +0.153 rather than any figure containing the selected cell; a per-pair standard deviation of 3.796 from the A/A rather than the older 3.869; a fixed-effect pool of the six new blocks as the only thing that decides; no cell droppable for its result; and no further run added to chase significance. The last six rows of Table 10 are that run.

	estimate	
primary — the six new blocks, 6000 pairs	+0.104 [+0.020, +0.189]	excludes zero
secondary — all ten unselected cells, 8400 pairs	+0.119 [+0.048, +0.190]	excludes zero

The effect is demonstrated. Three things belong beside that sentence.

The estimate has fallen at every stage — +0.570 at the screening cell, +0.153 pooled over the unselected cells, +0.104 now. That is the winner’s-curse decay in its textbook shape, and the pre-registered figure is the one with no selection in it.

The run was sized against +0.153, whose 80% power at 6000 pairs needs a true effect of 0.137. Against a true +0.104 it was under-powered, and the lower bound of +0.020 says how narrowly it cleared. A replication would be worth more here than another decimal place.

And the heterogeneity that opened this section is gone: $Q = 4.08$ on 5 degrees of freedom, $p = 0.54$, $I^2 = 0\%$, $\tau = 0$. Six blocks of one configuration agree within their own intervals, which is what the A/A predicted and what the original three cells did not. The disagreement really was one selected outlier and sampling noise.

The implementation shipped disabled while the effect was undemonstrated. On the strength of the primary test that default can now change — for a gain of about a tenth of a set per deal-pair, against the cost measured in Section 6.3.3, which is a trade rather than a free win.

What survives regardless. Proposition 3 does not depend on any of this. A possession-chain search over a belief whose transition edits only the asked card’s row is exactly greedy at every depth, so the obvious version of this idea is a no-op by construction rather than by measurement. That is a structural statement about why search keeps failing in this game, it removes a whole branch of the design space rather than a point in it, and it is the part of Section 6.3 we would keep. Its empirical shadow is the no-coupling cell, and the shadow is worth stating carefully, because the proposition entails greedy equivalence *at a node* and the step from there to a null set differential is a modelling inference rather than a theorem. Setting `couple=False` also does not restore the proposition’s hypotheses outright — the transition still zeroes the asked row and decrements the target’s count. What is demonstrated is narrower and still to the point. Instrumenting the search itself over 60,277 multi-branch nodes on champion trajectories (`scripts4/greedy_shadow.py`), the coupled arm departs from the highest-probability branch at 452 of them, 0.750%; the uncoupled arm departs at 1, 0.0017% — a factor of 452. Its set differential is a null at both 200 and 500 pairs (+0.070 and +0.112, both intervals comfortably containing zero), which is consistent with that.

The single uncoupled departure is worth a sentence, because it is real — relative gap in q of 1.9×10^{-3} , three orders of magnitude above floating point, and Proposition 3 predicts exactly zero. It is not a contradiction. `couple=False` does not restore the proposition’s hypotheses, as the module says: the transition still decrements the target’s count, so a player reduced to zero cards drops out of `legal_asks` and the available branch *set* changes downstream. That is a channel the exchange argument does not cover. The empirical shadow is therefore a near-null with a mechanism rather than a null, and an earlier draft reporting it as “53,273 of 53,273” overstated it — those counts were in the paper and in no results file, which is how a number stays unchallenged. We would also withdraw a remark from an earlier draft that the coupling “changed no decision at all”: over 1,110 champion-trajectory positions the coupled and uncoupled arms disagree on 0.54% of decisions, so observing no disagreement in a sample of 59 is the modal outcome rather than evidence of equality.

opponent	agent	pairs	set diff / pair	95% CI
memory (off-distribution)	$\gamma = 0.30$	150	+4.727	[+4.119, +5.335]
	$\gamma = 0$	150	+3.793	[+3.247, +4.340]
anti-model exploiter	$\gamma = 0.30$	200	+6.570	[+6.113, +7.027]
	$\gamma = 0$	200	+5.735	[+5.237, +6.233]

Table 11: The opponent model against opponents it was not shaped by. `memory` acts only on facts it can prove and has no probabilistic component at all. The anti-model exploiter is the v0.4 policy with the depth term *inverted* ($w_{\text{suit}} = -0.35$), so it deliberately asks in half-suits it is shallow in — the exact behaviour the model assumes does not happen.

Does it still pay once precision is bought? The two demonstrated features of this version (the belief-space search and the larger posterior sampling budget) were each measured against the champion *alone*, so neither says anything about the other. They could add, or the sharper belief could already be earning what the search was earning. Pre-registered in `jobs/PREREGISTRATION_stack.md` and written at one of six blocks, which is the only moment at which writing such a thing proves anything:

	set diff	95% CI
lookahead on top of 480 draws, 6000 pairs	+0.072	[−0.010, +0.154]

Cochran’s $Q = 1.94$ on 5 df ($p = 0.86$), $I^2 = 0\%$, $\tau = 0$ — the six blocks agree completely.

The interval contains zero, and the pre-registration fixed in advance how that must be read. At 6000 pairs this run has 68% power against the +0.104 it was sized against, so failing to exclude zero is a *failure to resolve an effect of the size assumed* and is not evidence that the search stops paying. The script that prints this verdict refuses to print the second sentence. It is a genuine limit of the design rather than a finding, and the honest summary is that we do not know.

What the run does bear on is additivity, and only weakly. The search on its own measured +0.104; on top of the sampler it measures +0.072, a difference of −0.032 with the two intervals overlapping. As far as this study can tell the two features add rather than overlap — with one point on each side of the comparison, which is exactly as weak a statement as it sounds. Neither `WEB_DRAWS` nor `V04_STRONGEST` moves on this result in either direction: both features are independently demonstrated, and this run is about what may be *claimed*, not about what ships.

`w_lookahead` still defaults to 0, and that is now a choice about cost rather than about evidence. The effect is demonstrated; it is worth about a tenth of a set per deal-pair for roughly 6.7 ms per decision (Section 6.3.3), so a caller under a latency budget should decide deliberately instead of inheriting. The engine exposes the enabled configuration under its own name rather than by moving the default, because every “vs champion” cell in this paper names the default spec and moving it would redefine what those runs measured.

9.3.2 Is the opponent model just fitting its own kind?

The model is a hypothesis about how opponents choose, and the opponents it was measured against in Table 4 are copies of the same policy. Two tests separate a real effect from self-play overfitting.

Against memory, a policy with no probabilistic component whatsoever, the model is still worth about +0.93 sets per pair. Against an opponent built specifically to violate its assumption it is still worth about +0.84. The second number deserves care: the exploiter is a weak policy in absolute terms (inverting the depth term costs it a great deal), so what the comparison shows is that the model is not *harmed* by an opponent designed to mislead it, not that it cannot be exploited by a stronger adversary. A genuinely coordinated best-responder is future work.

rule	comparison	pairs	set diff	95% CI
declare at any time	v0.4 vs v0.3 champion	200	+2.175	[+1.647, +2.703]
declare at any time	γ vs no γ	200	+1.935	[+1.450, +2.420]
48-card deck	v0.4 vs v0.3 champion	200	+1.080	[+0.557, +1.603]
wrong split $\rightarrow$ opponents	v0.4 vs v0.3 champion	200	+1.690	[+1.191, +2.189]

Table 12: The result holds under every rule variant tested, including the declare-at-any-time rule as the game is usually written down. Note the 48-card row: with eight half-suits instead of nine, the margin roughly halves, which is what one would expect if the advantage is inferential — a smaller deck leaves less to infer.

9.3.3 Does it survive the rule variants?

The engine supports several rule options, and a result that only holds under one configuration is a weaker result. Table 12 repeats the headline comparison under three of them. One row has since changed status: the games of this section were played under the void-era misdeclaration default, so “wrong split $\rightarrow$ opponents” was a variant when it was measured — it is now the baseline (§3), and the row is the measurement the rule correction’s continuity expectation was registered against. Its +1.690 carries one caveat worth stating: both agents’ declaration policies at the time still priced misses with void-era expected values, so the row shows the margin survives the rule, not that either side played it optimally — the restated headline under the award baseline, with both sides pricing it correctly, is design R1 of `prereg/rules_award_baseline.md`.

The declare-at-any-time rows required an engine change. v0.3’s game loop only ever asked the player on move, so although the rule was accepted by the configuration it could not be exercised by any policy, and it matters most precisely where it is unreachable, for a seat that has run out of cards and therefore never gets a turn of its own while a teammate still has some. v0.4 polls off-turn seats for a claim they can make *by deduction alone*, which costs a belief update rather than a posterior per seat, and offers it identically to both generations so that the comparison stays a comparison of one thing.

9.4 Learning the objective from data

9.4.1 Learning the ask objective

v0.3’s stated top priority for this version was to learn the ask objective rather than hand-weight it. We did, by common-random-numbers paired rollout regression: 874 positions, 83,168 games played to the end, all pairs of candidate asks within a position regressed on the difference of their feature vectors, with standard errors clustered by position (the naive standard errors are on average $2.00\times$ too small, since pairs from one position share candidates, worlds and rollout seeds).

The machinery works: common random numbers cut the variance of a pairwise comparison to 0.64 of independent worlds, bringing the noise-to-signal ratio from v0.3’s 2.4 to 1.19; a permutation test over position-level sign flips gives $p = 0.003$ against a null R^2 of 0.005; and a torch MLP on the same paired target improves validation R^2 from 0.234 to 0.249, so the objective’s *shape* is not the limitation either.

The learned weights lose. Played against the incumbent over 120 duplicate deal-pairs, the best of three variants scored -2.183 sets per pair $[-2.746, -1.621]$.

The diagnosis is specific, and it is a calibration failure of the target rather than of the fit. Three of the eleven coefficients have answers already established by *play*: v0.3’s 600-pair sweeps put turn at 0.60 and scarce at 0.20 and measured the pair as worth +1.41 sets per deal-pair. The regression returns $\text{turn} = 0.157$ with a clustered standard error of 0.058, so an effect the size play has already measured would sit ten standard errors from zero, and the dataset does not see it. More starkly, position-centred rollout value rises by only +0.101 sets across the entire range of $\mathbb{P}(\text{success})$, and an unpinned fit puts that coefficient at 0.063.

The likeliest mechanism is the continuation policy. A rollout has to be finished by a policy that can attach to a determinized mid-game position, and the belief tracker cannot: it is anchored on the initial deal and refuses. That leaves a public-information heuristic, which throws away most of the value of a

marginal card, so a card won by a good ask is largely squandered before the game ends and the ask stops mattering to the final differential. The binding constraint is not precision (pairing already fixed that), it is that the continuation is too weak for the difference between two asks to survive to the end of the game.

This is the same shape as v0.3’s search failure, one level up. There, the problem was that the evaluation target could not rank asks; here, a much more careful attempt to build a better target hits the same wall, and locates it in the rollout policy rather than in the statistics.

The mechanism was right and the wall was not there. Everything above stands except one clause, and that clause carried the whole line. “The belief tracker cannot” is false. `BeliefState` is anchored on the initial deal, but it does not have to be *handed* one: given the current hand and the real public history, `initial_hand()` back-computes a deal consistent with both, and the tracker attaches to that. The stored positions did not carry the public history (they recorded its *length*), so the strong policy had nothing to attach to, and a property of the dataset was read as a property of the engine.

Repeating the paper’s own measurement with the full v0.4 policy finishing every rollout, over 110 positions and 3056 scored candidate asks: position-centred rollout value rises by $+0.681$ with a deal-clustered standard error of 0.207 across the range of $\mathbb{P}(\text{success})$, against the $+0.101$ above.¹

Those two numbers must not be quoted against each other, and we nearly did. The $+0.101$ was measured over the `LEARNING` harvest, whose positions span the whole deal — median two half-suits resolved, only 29% with four or more, while this measurement harvests late positions, four or more by construction. Attributing the gap to the continuation, when the position distribution also differs, is a two-factor change read as one: the error this section is about, committed inside the correction to it.

The control is cheap and settles it. Running the public-information heuristic on the `SAME` 110 positions, the same determinized worlds, the same seeds and the same 3056 root asks gives a slope of $+0.040 \pm 0.045$ — consistent with the published $+0.101$ and indistinguishable from zero. Because the two arms score identical asks at identical positions, $\mathbb{P}(\text{success})$ is bit-identical between them, and the difference of the two within-slopes is exactly the within-slope of the differenced outcome. That paired contrast is

$$+0.641 \pm 0.145 \quad (4.4 \text{ standard errors, clustered by position}).$$

So the continuation is the cause, and the earlier comparison was right for a reason rather than by a coincidence of position mix.

One thing that contrast changes twice, and it had to be checked rather than argued. The engine arm finishes the game *and* is handed the real public log; the heuristic arm finishes the game *and* starts blind to every card the table has already watched change hands, because an empty log is the knowledge set `PublicInfoHeuristic` was defined and audited against and is what the original learning rollouts gave it. So the two arms differ in information as well as in policy, and reading the whole $+0.641$ as “the continuation policy” would be this section’s own two-factor error committed one level further in — inside the correction to the correction.

A third arm settles it: the same weak heuristic, on the same 110 positions and 3056 root asks, with the public log seeded. It scores $+0.084 \pm 0.071$, beside the unseeded heuristic’s $+0.040$ and nowhere near the engine’s $+0.681$. Because all three arms score identical asks at identical positions, the contrasts are exact within-slopes of differences on the same x and therefore add:

$$\underbrace{+0.597 \pm 0.145}_{\text{policy alone}} + \underbrace{+0.044 \pm 0.069}_{\text{the log alone}} = +0.641.$$

Handing the weak policy the whole public history moves the target by an amount this design cannot distinguish from zero. The asymmetry was real; it is not what the contrast measures.

This arm exists because of an audit note about something else entirely — that seeding the history also shortens the stall rule’s effective window, since stalled scans the last 80 actions and the seeded

¹That standard error is clustered on the *deal*, not the position. The harvest emits consecutive plies, so these 110 positions come from **four** deals; clustering on the position gave ± 0.142 , and with a t critical value at three degrees of freedom the interval is $[+0.027, +1.335]$ rather than $[+0.401, +0.962]$ — $2.33\times$ wider, and still excluding zero. Every standard error in this section carries the same correction (`scripts4/cluster_audit.py`).

prefix eats a mean of 10.2 of them. That particular asymmetry turns out to be inert: over 13,290 engine-continuation decisions the rule fired *zero* times, because the prefix eats ten actions and the engine finishes in twenty-six more, so nothing approaches eighty without a claim. It fires at 1.02% of the heuristic arm's decisions instead, which is the same fact as the 181 plies, but chasing it is what surfaced the information asymmetry, which was the one that mattered.

One caution on that pairing, because it briefly said the opposite. Averaging the per-position slope *differences* without weights gives $+0.020$ with a standard error of 0.246, apparently a null. That is a different estimand, not a more careful version of the same one: an unweighted average is dominated by positions carrying almost no contrast in $\mathbb{P}(\text{success})$, whereas the reported slopes weight each position by its own sum of squared centred $\mathbb{P}(\text{success})$. One position with four scored asks contributes a slope difference of -21.7 , because a slope fitted to four points is not an estimate of anything. The robust summaries agree with the weighted one: the *median* per-position difference is $+0.299$ and the engine's slope is the larger at 63% of positions. A pairing that changes the weighting is not a pairing.

We first computed that unweighted figure as $+0.194 \pm 0.160$, which is what a silent bug produced. Positions were being kept only when their clustered standard error came out positive — ordinary-looking defensive code, and wrong, because within a single cluster the CR0 numerator is $(x'e)^2$ and $x'e$ is the OLS normal equation, hence exactly zero in exact arithmetic. The filter kept the 97 of 110 positions where floating-point roundoff happened to leave a nonzero residual dot product, and among the 13 it discarded was the -21.7 outlier. A defensive filter on a quantity that is analytically zero selects on roundoff, and it flattered the comparison it was printed to stress-test. No single position carries it (the largest leave-one-out move is 0.06), and the run's first 40 positions, which had been measured on their own first, agree with the remaining 80 within 1.7 standard errors, so it is one slope rather than an average of two regimes. Per-position slopes nonetheless run from -6.4 to $+3.7$: it is that between-position spread, and not the number of asks scored, that sets the interval, which is why tripling the run barely narrowed it.

The mechanism is measurable too, and turns out to be almost embarrassingly concrete. Handed the same position, the same world and the same root ask, the heuristic needs 181 plies on average to reach a terminal state where the engine needs 26 — a factor of seven. The extra plies are exchanges that undo each other: nothing stops the heuristic trading a card back and forth until its stall rule fires, and every one of those plies is a chance for the root ask to stop mattering to the final differential. Neither policy comes near the 400-action cap (324 and 56 at worst), so no part of this is capped games being scored on partial boards.

So the diagnosis was correct about the mechanism (a weak continuation does flatten the target), and wrong about it being a wall. The target the regression could not learn from was flat because of what finished the game, and finishing the game properly un-flattens it by a factor of nearly seven. The learning line is re-opened rather than abandoned, at a cost: with six belief trackers in the loop a rollout takes 240–500 ms against 17–23 ms for the heuristic, roughly $20\times$.

Two further things had to be true and only one of them was obvious. The stored worlds must satisfy everything the public log implies, not merely preserve hand counts, which is all the public heuristic ever needed. A successful ask locates a card; a *failed* ask asserts, under the no-bluff rule, that the asker holds at least one card of that half-suit, and a count-preserving shuffle violates that constraint routinely. Worlds must therefore come from the belief sampler. And the failure had to stay loud: seeding a rollout with an empty history does not error, it anchors six trackers on a determinized mid-game hand as though it were the deal and returns a number computed from a game that could not have happened. That configuration is refused rather than run.

A second obstacle, which the continuation fix does not touch. With the target no longer flat, the obvious next step is to ask which of the eleven objective terms it responds to, and the answer is that this design cannot say. Fitted one at a time over the 3076 rows of that run — twenty more than the 3056 the slope above uses, because a slope needs two scored asks at a position and twenty rows sit alone at theirs — several features look more informative than $\mathbb{P}(\text{success})$ itself: *expose* $+1.807$, *deplete* $+1.387$, *certain* $+0.919$, against $\mathbb{P}(\text{success})$'s $+0.681$ (standard errors 0.536, 0.249, 0.223 and 0.207) — all four deal-clustered, and all four still excluding zero. Fitted together, $\mathbb{P}(\text{success})$ turns *negative*, at

-0.329 ± 0.264 (that one is quoted as first computed, position-clustered, because the sentence after it says the column should not be read at all). Neither column should be read, and the diagnostic that says so belongs beside them: $\mathbb{P}(\text{success})$ carries a variance inflation factor of **13.5** here, meaning 92.6% of its within-position variation is reproduced by the other ten terms. It is not an independent regressor; it is computed from the same posterior the features are, and correlates with `deplete` at 0.854 and `certain` at 0.738. A coefficient that is not identified does not become identified by being reported, and the negative sign is an artefact of the basis rather than a finding about play.

This matters for what happens next. The diagnosis above was right — a weak continuation does flatten the target, and finishing rollouts properly un-flattens it sevenfold, but it was not the only thing standing in the way. Near collinearity of the feature basis is a *second* obstacle, independent of the first, and fixing the continuation does nothing about it: the design matrix is nearly singular whatever finishes the game, so the fitted weights are unstable however much signal the target carries. The whole eleven-term fit explains 11.7% of within-position variance, so there is signal to work with; what there is not is a way to attribute it to individual terms as they currently stand.

The two most-quoted figures in this paper were the two nothing could check. `check_paper_numbers.py` watches over a hundred figures against the results files that produce them. It could not watch the abstract’s headline (the +1.85 margin over the previous champion), because that lives in the JSONL duel pool while every other watched figure lives in a JSON file, and the loader only read JSON. It could not watch Table 7 either, because that table’s *rates* were computed in the L^AT_EX from counts the pipeline stored, so the number the paper prints was never a number the pipeline held. Both are correct; both sat outside every drift check for reasons as dull as a file format and a division done in the wrong place. This is the third time the same shape has appeared — after `settle_verdict.py` printing the lookahead headline and persisting nothing, and the deadlock quartet existing in no file at all, and each time the unwatched figure was among the most load-bearing in the document, for the same reason: the numbers a paper repeats most are the ones nobody re-derives.

The manifest cannot find its own gaps, so the check is now inverted as well. `scripts4/unwatched_claims.py` extracts every number the paper asserts in **bold** — its strongest claims, the ones a reader takes as measured, and requires each to be either pinned by the manifest or exempted with a *stated reason*; an exemption with no reason is indistinguishable from an oversight, which is the failure being guarded against. Running it the first time returned nineteen unpinned figures. Six were version numbers, row labels or a saturated column; five belonged to the two figures above; and the remaining *eight* were ordinary measurements nobody had thought to watch — among them the A/A per-pair standard deviation that every power statement in this paper divides by, the stalling fraction that motivates the progress-optimal criterion the absolute-strength table is scored against, and the champion posterior’s own NLL and Brier scores.

A supervisor a bystander could switch off. The rollout pass is restarted by a shell supervisor whose liveness check was `pgrep -f "learn_ask_objective.py rollout --runv2"`. That matches the *full command line of every process*, so a one-line loop written to *watch* the pass, one whose own command line contained that string, was itself a match. The supervisor concluded the job was alive and stopped restarting it; the pass stayed dead for as long as the watcher ran, and was found only because a separate progress check reported zero workers. The widening script used the same pattern to choose a PID to kill.

This is the mirror image of a failure this paper already records, in which the same function returning `False` on any exception made a guard fail *open*. Both come from reading “this string appears somewhere in the process table” as “this job is running”. The check now walks `/proc` and counts a process only when it is a Python interpreter whose own `argv` carries every token, excluding the caller and its ancestors; its test asserts both directions, including that `pgrep` really does match the bystander it is given.

One term of the eleven will not be fitted, and the reason is worth one sentence. The `claim` feature’s formula was corrected after this harvest was taken — it multiplied by the product over all *six* cards of the half-suit including the one being asked for, so it scored exactly zero on provably certain steals, the asks

basis	CV R^2	vs. the full basis
full, eleven terms and $\mathbb{P}(\text{success})$	0.0959	—
strongest three, selected <i>within</i> each fold	0.1030	$+0.0070 \pm 0.0038$
drop terms until every VIF < 5	0.0964	$+0.0005 \pm 0.0017$
ridge, $\lambda = 1$	0.0968	$+0.0009 \pm 0.0015$
$\mathbb{P}(\text{success})$ alone	0.0430	-0.0530 ± 0.0234
principal components, $k = 4$	0.0150	-0.0809 ± 0.0306

Table 13: Eighteen candidate bases for the ask objective, scored by leave-one-position-out cross-validated within-position R^2 . Folds are at the position level because asks inside one position share worlds, seeds and a value level, so an ask-level fold would leak. Six representative rows; selecting the strongest three terms *within* each fold is the only basis that beats the full one, and by less than twice its own standard error.

it exists to reward. Its stored column therefore describes a feature the engine no longer computes. The loader refuses such a column by default, and the fit stage names it, zeroes it, and records that it was not fitted, because a ridge coefficient of exactly 0.0 for `claim` reads identically to “fitted and came out at zero”, and those are different statements. The other ten terms are unaffected; re-harvesting to recover the eleventh is a day of rollouts for a term the champion already weights at zero, and the paragraph below is the reason that is not the bottleneck.

“So use a different basis” was reasoning, and the measurement says no. The natural conclusion from the paragraph above (and the one we wrote first) is that the re-run needs a different basis, or a penalty, or terms dropped, rather than more rollouts. It is cheap to check, because the completed run already holds every scored ask, so we checked it before spending anything. Eighteen candidate bases, each scored by *leave-one-position-out* cross-validated within-position R^2 (position-level folds, because asks inside a position share worlds, seeds and a value level): The best of them beats the basis in use by $+0.0070 \pm 0.0038$, an interval that *contains zero*. Ranked by point estimate a three-term basis wins; measured against the spread across positions, nothing here is distinguishable from what the objective already uses, and picking one off this table would be selecting on noise.

That is the textbook consequence of collinearity and it is worth stating plainly, because it is easy to get backwards: near-collinearity makes individual *coefficients* unidentified without making *predictions* worse. A reduced basis is therefore not the fix, and neither is a penalty — ridge does nothing at any λ that does not also destroy the fit. Two smaller things fall out of the same table, and the first has to be stated carefully because it reads, at a glance, like the opposite of this paper’s central advice.

Dropping terms by variance inflation removes $\mathbb{P}(\text{success})$ *first*, and costs nothing measurable ($+0.0005 \pm 0.0017$). That is **redundancy, not unimportance**, and the two are easy to confuse. The variance inflation factor of 13.5 says precisely that 92.6% of $\mathbb{P}(\text{success})$ ’s within-position variation is already reproduced by the other ten terms, which is unsurprising, since all eleven are computed from the same posterior over the same worlds. A regression cannot miss a quantity its other regressors have already encoded, and it does not follow that the quantity does not matter. Section 12.1’s finding that the probability the ask lands *is* the objective rests on *play* (scoring asks by anything else loses by -7.4 sets per deal-pair), and nothing in this table bears on it. What the table says is narrower: given ten features already carrying that information, an eleventh copy of it adds nothing.

And every principal-component basis is far worse, so the predictive directions in this design are not its high-variance directions — rotating to the biggest components discards exactly the signal.

We record this as the seventh appearance of the pattern, and the shortest-lived: the sentence it refutes was written earlier the same day, in the paragraph immediately above, by the same argument the paper keeps warning about. A plausible mechanism that generates no experiment is not a finding, even when the mechanism is real, and here the mechanism *is* real. The variance inflation factor of 13.5 stands. What does not stand is the inference we drew from it about what to do next.

The first version of this comparison also got its own selection wrong, which is worth one sentence because it is the same error at a third level. It ranked terms by $|\beta \cdot \text{sd}|$ on a fit to *all* the data and then cross-validated the winner, so the selection had already seen every held-out position. Selecting inside each fold instead is what the table above reports.

Collinearity is not a licence to explain away every disagreement, and it does not cover the one that started this section. *turn* is the term play has measured most directly, and it is precisely the term collinearity cannot be blamed for: its variance inflation factor is 1.4, and it varies between candidate asks at 111 of the 113 positions the feature fit retains (the slope's 110, plus three the within-position fit keeps and the slope does not). It comes back at $+0.010 \pm 0.154$ all the same, where duplicate-deal play puts it at 0.60. That gap is real but not yet a finding in either direction: these positions are late by construction (four or more half-suits resolved), and holding the turn may simply be worth less when most of the board is already decided, so a coefficient near zero here is not in itself evidence against a weight of 0.60 across a whole game. What can be said is that the explanation is not the one immediately to hand. Reaching for the nearest available mechanism is how this paper's recurring error starts, and it would have started again here.

We record this as the fifth appearance of the pattern this paper keeps documenting, and the first in which reasoning about a quantity did not merely inflate it but retired an entire research direction. The diagnosis was well-argued, specific, mechanistic, and supported by a number. It was also never tested, because testing it required exactly the thing it asserted was impossible.

The weights were played, which settles the line. Everything above is estimation. `jobs/PREREGISTRATION_learned_weights.md` fixes, before the fit was run and before any weight vector had been looked at, that the *validation duel is the test*: two blocks of 1000 fresh pairs, the pinned-*p* vector only, the shipped champion as the baseline rather than the weaker `fishbot4()` defaults `v0.4` validated against, and a fixed-effect pool as the only thing that decides.

	estimate	
v0.4's attempt, 120 pairs (<i>different design</i>)	−2.183	best of three variants
this run , 2000 pairs	−0.745 [−0.914, −0.576]	entirely below zero

The learned weights lose, and this is the stronger negative. The two blocks agree exactly ($I^2 = 0\%$, Q at $p = 0.33$), the realised per-pair standard deviation of 3.857 matches the 3.8 the design assumed, and the arms diverge on 87.6% of pairs — an ask-weight change alters decisions constantly, which is why the A/A figure is the right one here and why this run resolves what it was built to resolve.

Read against the diagnosis it was built to test, the result is unambiguous. `v0.4` blamed the weak rollout continuation; that diagnosis was measured, found right about the mechanism, and the continuation was replaced — the target's slope on $\mathbb{P}(\text{success})$ went from $+0.101$ to $+0.681$ on late positions and $+0.355$ across the whole harvest. *And the weights still lose.* So the continuation was a real obstacle and not the binding one. The binding one is the basis, which this section measured before the duel was played and which the table above says a smaller basis does not fix.

What the vector is not. The fit moves several terms with permutation $p \leq 0.023$ — `expose` to $+0.825$ from an incumbent of zero, `certain` to -0.502 , `signal` to -0.249 , and the individual signs are not findings. `certain` correlates 0.738 with $\mathbb{P}(\text{success})$, which carries weight 1.0 by convention, and the within-position feature fit reported earlier in this section gives the same term the *opposite* sign on a different population (`results/target_feature_fit.json`, $+0.649$ against this -0.502). That is the collinearity this section already documented for $\mathbb{P}(\text{success})$'s own coefficient, arriving a second time.

The single most informative row is *turn*: incumbent $+0.600$, learned $+0.108$. The regression cuts by more than 80% the weight on the term duplicate-deal play values most highly and the term collinearity cannot be blamed for, and then loses by three quarters of a set per pair. Whatever the regression is fitting, it is not what wins games.

One term of the eleven carries 0.0 by construction rather than by measurement: `claim`'s stored column predates the correction to its formula, so it was named, zeroed and excluded from the fit at every

stage. The incumbent weights it at zero too, so the duel is unaffected, but the fit is a ten-term fit and is reported as one.

A cut-off resolved before it could be chosen. The pre-registration fixed the fit at “every position the pass has completed at the moment the duel queue drains”. The queue drained at 743 of 1023 and the pass did not stop there, because the supervisor that widens it exists precisely to make draining the queue speed the pass up. Worse, that timing had been moved earlier the same session by a repair to the supervisor’s liveness check, so the trigger named a number we had influenced. The amendment resolving it to the full 1023 was written at 919, before any fit had been run, and committed to computing the 743 fit anyway as a robustness check. That check came back with *no sign flips*, a largest single move of +0.135 and total L_1 movement of 0.394 across eleven terms: the vector is stable, and the ambiguity turned out not to matter. Learning that afterwards is the only order in which it is worth anything.

9.4.2 The split posterior is badly under-confident

Section 6.2 sets the claim threshold at 0.97 and argues for it from a specific mechanism: while our team genuinely holds all six cards of a half-suit an opponent cannot take it (a claim by a team that does not hold all six *gives* the set away), so the only cost of waiting is running out of opportunity. Waiting is close to free. v0.3 measured the consequence and could not explain it: thresholds from 0.85 to 0.999 play identically, and over 150 duplicate deals 0.97 and 0.999 never once diverged.

Scoring the posterior against ground truth explains it. Over 40 games we recorded, at every position where a half-suit was in fact entirely within one team, the posterior probability the engine assigned to its own MAP split and whether that split was the true one — 2719 decisions on 359 distinct half-suits, ground truth used for analysis only:

posterior says	decisions	actually the true split
[0.000, 0.500)	2062	66.5%
[0.500, 0.700)	203	93.6%
[0.700, 0.850)	75	97.3%
[0.850, 0.970)	24	100.0%
[0.970, 1.000]	355	100.0%

The interior cut points used to be 0.700/0.825/0.970, where 0.825 was 1 minus the null rate on stuck half-suits — a decision bar imported from Section 10.1. Re-measuring that null rate therefore moved two rows of a table about the *sampler*, for a reason having nothing to do with the sampler. How well a stated probability is calibrated is a property of the posterior alone, so the bins are now fixed, the bar is reported separately where it belongs, and the whole table is stored in `results/stuck_claim_value.json` rather than printed, which is the failure `scripts4/check_verdicts.py` exists to catch, found here in a table rather than a verdict.

Two things follow. The distribution is **bimodal** — 76% of decisions sit below 0.5 and 13% above 0.97, with almost nothing between, which is why the threshold is insensitive across the whole range v0.3 swept. There is nothing there to be sensitive to. That is the explanation the sweep lacked, and it is a fact about the posterior rather than about the claim rule.

And the posterior is **badly under-confident**. It says “below one half” and is right two thirds of the time; it says “between a half and 0.7” and is right 93.6% of the time. The claim threshold is the one place in this engine where a split probability is consumed as a *number* rather than as a ranking, and a threshold set on a miscalibrated number does not mean what it says. This is a finding about the sampler, not about the claim policy, and it is where we would look next.

What we did not compute, and why. The obvious next step is to price the change: declaring on splits the engine currently waits on, against waiting and accepting the measured 23.4% misdeclaration rate on stuck half-suits, would beat waiting wherever the MAP is right more than 76.6% of the time under the void-era rule this section’s games were played with (a wrong declaration cost 0 there; under the award

baseline it costs -1 , the eventual stuck declaration risks the same -1 , and the break-even lands elsewhere — one more reason the refusal below stands), which the table says it is, everywhere above 0.5. Carried through, that arithmetic returns a number of the size this paper calls demonstrated.

It is not reported, because the two numbers are measured over different populations. The 23.4% is the null rate for half-suits a team is *stuck* on in Section 10.1’s sense: it can *prove* it holds all six and still cannot place the split. The population in the table above is every half-suit that *is* entirely within one team, provable or not — 8.97 per game out of nine, so very nearly all of them. Subtracting one from the other is a difference between two different things, and it produces a publishable-looking number by doing so. This is the same error as quoting the learning harvest’s $+0.101$ against a late-position $+0.681$ (Section 9.4.1), and it is recorded here because we made it, saw the size of the answer, and went looking for why it was too good. Pricing this properly needs the null rate re-measured over this population, and until then there is no estimate and no queued run.

9.4.3 The declaration threshold: the one screen that resolved, re-run

Section 9.1 reports a cell that lowers the claim bar from 0.97 to 0.90 and returns $+0.035$ [$+0.007$, $+0.063$] over 400 pairs. It is the only small, cheap, unexplained positive in a table the paper describes as mostly nulls, and it contradicts a documented conclusion: `fish4/claim4.py` argues that waiting to claim is close to free, and reports that over 150 duplicate deals the 0.97 and 0.999 policies never once diverged.

That cell was selected out of 103 for excluding zero, so it is inadmissible as evidence and `jobs/SCREENING_DISCIPLINE.md` forbids it from sizing its own confirmation. `jobs/PREREGISTRATION_claim_threshold.md` fixes, in advance: two blocks of 1000 fresh pairs; a minimum interesting effect of $+0.02$, deliberately an order of magnitude below the $+0.15$ used elsewhere, because this change is one constant and costs nothing to run; a fixed-effect pool of the two blocks as the only thing that decides; and no block droppable for its result.

	estimate	
screening cell, 400 pairs (<i>selected</i>)	$+0.035$ [$+0.007$, $+0.063$]	excludes zero
primary — two new blocks, 2000 pairs	$+0.002$ [-0.013 , $+0.017$]	includes zero

The effect does not survive. The confirmatory estimate is a twentieth of the screen’s, and the decay of -0.033 , against a two-sample standard error of 0.016, is itself 2.1 standard errors — selection inflation measured rather than inferred.

The two blocks of that primary did not run the same program. `scripts4/check_engine_provenance.py` fingerprints the engine files each duel block executed and compares them within a pool. Two pools in this paper come back mixed, and this is one of them: the two 1000-pair blocks pooled above differ in `fish4/askfeat.py` and `fish4/oppmodel.py`, and block one’s `askfeat.py` digest is not in this repository’s history at all, so those blocks cannot be re-run even in principle. The other mixed pool is the re-take bonus of §12.2, pooled at $+0.013$ [-0.052 , $+0.079$]. A third, the combined arm, was mixed, was caught by the same check, and was re-blocked on one engine under `jobs/PREREGISTRATION_combined_reblock.md`; the replacement landed at $+0.357$ [$+0.191$, $+0.524$], near what the registration predicted, so that mixture had no numerical consequence.

The blocks here agree within their noise — $Q = 1.35$ on 1 df, $p = 0.25$, and that Q is computed by the provenance checker for exactly this reason. It is worth being plain about what agreement does and does not buy: it bounds the damage, and it does not make the pool one measurement of one thing. The interval above describes an average over two programs, which is not what a confidence interval is usually read as. The checker exits non-zero on this repository today, and is deliberately left out of continuous integration, because a check that is permanently red is one people learn to scroll past, but leaving it out of CI is not a reason to leave it out of the paper.

This is a null with a mechanism. A behavioural measurement taken before the run, over 998 decisions carrying a claim candidate, found that raising the bar from 0.97 to 0.90 admits *one extra claim in 998*. The run then diverged on 2.6% of pairs. Two arms that make one different decision in a thousand and finish 97% of deals identically are not a null from thin data; they are the same policy, and `claim4.py`'s argument survives its one piece of counter-evidence.

The design missed its own power target, and says so. The pre-registration derived a per-pair standard deviation of 0.286 from the screening cell's recorded interval, giving a minimum detectable effect of 0.018 — just under the +0.02 it was built to detect. The run's realised standard deviation is 0.346, 21% higher, so the realised MDE is 0.022 and sits *above* the threshold. An effect of exactly +0.02 would have been missed more often than the design intended, and the null is reported against 0.022 rather than against the number that was planned. The interval-derived standard deviation, which reproduces a cell's measured value to within 1% everywhere else in this study, is the one input that failed here, and it failed in the direction that flatters the design.

A prediction recorded in advance, and scored. The divergence model of Section 9.1 was deliberately *not* used to size this run, because every cell it was fitted on diverges on 23–88% of pairs and a rarely-binding threshold does not. The pre-registration records what the model would have implied and predicts it will be wrong in a stated direction: at this per-pair standard deviation the model implies a divergence share of 0.8%, about sixteen pairs in 2000, and the prediction is that the true share is substantially higher *because* a threshold change alters one decision rather than the whole line of play, so its conditional differences must be smaller than the model's constant 3.88. Both halves hold: the measured share is 2.6%, and the measured conditional standard deviation is 2.17.

That divergence share is a ninth of the low end of the range the model was fitted on, which makes this the most extreme out-of-sample point available for the drift correction of Section 9.1, and the correction was fitted before this run existed. The flat conditional term of 3.859 overstates the measured 2.166 by 78%. The drift-corrected form $2.53 + 1.66s$ predicts 2.576, an overstatement of 19%: still wrong, four times less wrong, and wrong in the same direction the correction was built to fix. A model extrapolated an order of magnitude past its data has no business being right, and it is not; what is worth recording is that the term added to it keeps earning its place where it cannot have been tuned.

10 The shape of the game

10.1 Perpetual states, measured

This is the measurement; §10.4 is the question it answers, in three layers of which only the third turns out to be interesting.

Two hundred games per column, with ground truth used for analysis only.

Repetition is universal, not exotic. A position repeated in *every one* of 200 games, and some position recurred eight times in a single game. v0.3 established cyclicity by search, as a possibility; it is in fact a routine feature of ordinary play.

Every game dies before it ends. All 200 games reached a position in which no ask anywhere on the board could succeed. It happens late — about 4.4 actions from the end, roughly 4.2% of plies, but it is universal. The last phase of a game of Fish is not a contest at all; it is a bookkeeping exercise in which the sets are already decided and only the declarations remain.

Termination is not the progress rule's doing. With the agents' stalling rule fully disabled, not one game in 200 hit the action cap, and the mean game length was identical to four significant figures. So although the rules genuinely permit non-termination, and although positions genuinely repeat, in practice a strong policy walks out of every cycle on its own. The convention is insurance, not the mechanism. Consistent with this, widening the stall window from 80 to 200 actions produced 367 exact ties in 400 duplicate deals: the rule almost never fires.

	normal	no progress rule	signalling on
actions per game	104.4	104.4	104.8
games hitting the 4,000-action cap	0	0	0
games in which a position repeated	200	200	200
most repeats of a single position	8	8	8
games reaching a <i>dead</i> position	200	200	200
share of plies fully dead	4.24%	4.24%	4.56%
actions played after going dead	4.4	4.4	4.8
games where a team held an unplaceable set	110	110	116
dead half-suits at peak	3.13	3.13	3.15
nulls per game	0.275	0.275	0.220

Table 14: Repetition and dead positions in 200 games of champion self-play. The last row counts the event “a team held all six and named the wrong split”, which the void rule turns into a null and the award baseline into a gift to the opponents. Earlier drafts reported it as “nulls per game” from a study pinned to the void rule; the same 200 seeds replayed under the award baseline agree on *every* field of this table, the last row included (results/perpetual_study_void.json against results/perpetual_study_award.json). The rule reaches only claim4.forced_claim, where it changes the split ranking from p_{exact} to $p_{\text{exact}} + p_{\text{team}}$; over 284 forced declarations in 800 games of champion self-play it changes the ranking 0 times and the claim returned 0 times, with 166 of the 284 offering only one candidate split at all (scripts4/rule_bite.py).

The deadlock is real, and it is expensive. In 110 of 200 games (55%) a team reached a position where it could prove it held an entire half-suit and still could not place the split. Over the 1800 half-suits of those 200 games, 171 (9.5%) reach that state, and they are nulled **23.4%** of the time [17.7, 30.3] against **0.92%** [0.56, 1.51] for half-suits that never get stuck (a factor of 25), so they account for **73% of all nulls** while being under a tenth of half-suits.

Four figures in this paragraph were wrong in every draft before this one, and the way they were wrong is the reason the paragraph now names its denominators. They read 17.5%, 2.8%, a factor of 6.3, and 27% of all nulls. Those numbers appear in no results file: perpetual_study.py tracked which *teams* were stuck and recorded nothing per half-suit, so nothing in the pipeline could have checked them and nothing did. They were carried from draft to draft as though measured. The direction survives re-measurement — it survives it more strongly, since the true ratio is four times the one claimed, but 27% against a measured 73% is not a rounding difference, and one of these numbers had already been used downstream as a threshold (Section 9.4.2). Every figure here now comes from results/perpetual_study.json and is checked by scripts4/check_paper_numbers.py against it.

The free-signalling escape works, and does not pay. Turning on the protocol of Section 10.4 cuts nulls from 0.275 to 0.220 per game, a 20% reduction, which confirms the mechanism does what the theory says. It buys no wins: against the identical policy with signalling off, over 500 duplicate deals, +0.002 sets per pair [−0.086, +0.090]; restricted to strictly dead positions, +0.012 [−0.000, +0.024]. Those intervals are far tighter than the usual ± 0.48 because the change is a near-no-op — it alters so few decisions that most pairs come out identical.

Two things are going on, and both are worth stating. The strictly-free version of the protocol applies only when the *whole board* is dead, which the seat can prove on about 0.3% of plies — too rare to matter. The relaxed version fires more often but is no longer free: it spends a turn to send one bit, and the tempo it concedes cancels the sets it saves. So the escape is real theory with a null in practice, and the honest summary is that Fish’s stalemate is a genuine structural feature of the game that a strong engine is already handling adequately by other means. One post-script from the rule correction: this ledger was priced in the void era, where each avoided misdeclaration was worth one point of differential; under the award baseline it is worth two, so the duel was re-run rather than re-asserted (design R5 of prereg/rules_award_baseline.md, outcomes fixed before any pair). Over 500 fresh pairs of V06_DEPLOYED with the protocol on against itself with it off, the protocol again cuts misdeclarations, 109 against 138, and the effect moves to +0.0680 [−0.0326, +0.1686]: the sign the re-pricing predicts, at

about the size the arithmetic suggests, and an interval that still straddles zero. Per the registered outcome it stays off; what changed is the honest summary, from “buys no wins” to “plausibly worth about a set every fifteen games, not yet demonstrated at this power” — a properly sized extension is registered future work rather than a post-hoc one.

There is a specific reason to expect that extension to land differently, which is not a reason we had at the time. Both measurements of this protocol fired it at the same threshold. The relaxed mode signals only when the ask the seat would otherwise make has $p \leq 0.15$ (`signal_max_p`), a number chosen when there was no scale to choose it against. §10.5 then supplied one, bucketed by exactly that probability, and it puts a turn at -0.043 below $p = 0.25$ and $+0.004$ in $[0.25, 0.50)$, with standard errors of 0.169 and 0.143 , so 95% intervals of $[-0.374, +0.288]$ and $[-0.276, +0.284]$, which do not establish that a turn is free there so much as fail to resolve anything smaller than about a third of a set, and about $+0.45$ only above 0.50 . The gate sits at less than a third of the way to where this paper’s own measured price of a turn crosses zero, and the two sections have never been read against each other. The sweep in `jobs/j14_signalling.json` tried 0.35 and stopped.

Two other things are now known that were not. An avoided misdeclaration is worth measurably less than the two points the rulebook implies. For the signalling arms themselves it is $+1.2784$ $[+1.0642, +1.4926]$ sets (`results/error_value_signal_gate.json`), an interval that excludes the assumed 2 ; the deferral gate’s own figure is higher at $+1.7898$ $[+1.5927, +1.9870]$ (§8.7.2) and an earlier draft used that one here, which prices one mechanism with another’s exchange rate. On the signalling figure, a signal pays whenever its chance of placing a split we would otherwise get wrong exceeds the turn’s price divided by 1.28 , and below $p_{\text{best}} = 0.50$ that numerator is not resolved by the run that measured it. And the population it would act on is reachable: of the 130 half-suits we declared wrongly over 800 cross-engine games, the median number of *our own* turns between assembling the set and the deadline is 25.5 , with only 16 having none left.

`prereg/deadline_signalling.md` registered exactly that, with the incumbent 0.15 carried as a replication arm — if it did not reproduce the $+0.068$ above, the new rung would mean nothing.

It ran, and the argument was right about the price of a turn and wrong about what was limiting the mechanism. $1,000$ games per arm against `v0.7` on identical deals, both parities, bridge revision 2 , zero fallbacks (`results/signal_gate_confirm.json`):

arm	margin	vs. A
A, signalling off	$+2.4760$	—
B, <code>signal_max_p = 0.15</code>	$+2.5940$	$+0.1180$ $[+0.0325, +0.2035]$
C, <code>signal_max_p = 0.50</code>	$+2.5980$	$+0.1220$ $[+0.0291, +0.2149]$

B replicates: its interval contains the $+0.068$ it was carried to reproduce, so C means something. Wrong declarations fell rather than rose, $0.170 \rightarrow 0.152$ a game. No withdrawal condition fired. And C did not ship: the bar was a point estimate at or above $+0.15$ with the interval clear of zero, and $+0.1220$ clears zero without reaching the bar.

The prediction was that C would beat B, because a threshold three times higher should fire three times more often. It did not, and the path ledger says why. Signalling moves about 210 declarations out of the gated path and into the forced path, and *B and C move almost the same ones*: 75 gated declarations against 78 , a difference of three in a thousand games. Widening the gate by a factor of 3.3 changed the engine’s behaviour on three declarations.

So the gate was never the binding constraint, and the reason is visible once stated: the protocol fires when a seat is stuck, and a stuck seat is one whose best ask is already bad, so $p_{\text{best}} \leq 0.15$ is very nearly implied by the situation it fires in. §10.5 was right about the price of a turn and wrong about which quantity was limiting this mechanism. That retires the hypothesis, which is worth more than the $+0.004$ it bought.

What the effect is made of, conditioned rather than fitted (`scripts4/error_value.py` on this run’s own journal, arm C):

games	paired margin
avoided no error, $n = 876$	+0.0502 [−0.0248, +0.1253]
avoided one, $n = 66$	+1.6061 [+0.9819, +2.2302]
avoided two or more, $n = 6$	+3.6667 [+0.9861, +6.3472]
added an error, $n = 52$	−0.9615 [−1.6399, −0.2832]

The whole effect is avoided errors: the 876 games where nothing was avoided contribute +0.044 with an interval containing zero, the 72 where something was contribute +0.128. And the ceiling on the mechanism is the last row — it *adds* an error almost as often as it avoids one, 52 games against 72. That, not the gate, is what a future attempt would have to attack.

What that row is made of. The same journal carries a per-game declaration path ledger, so the two groups can be separated without replaying anything (`scripts4/signal_error_paths.py`). Signalling drains the *gated* path in both, and by nearly the same amount — −0.635 declarations a game where it adds an error, −0.847 where it avoids one, so that is not what distinguishes them. What distinguishes them is entirely where the drained declaration lands instead:

group (arm C minus A, per game)	voluntary	gated	forced	exact
added an error, $n = 52$	−0.788	−0.635	+1.327	−0.250
avoided an error, $n = 72$	+0.139	−0.847	−0.139	+0.111

Against path error rates of 0.05% voluntary, 0.0% exact, 10.3% gated and 46.3% forced, that is the whole story: the protocol spends a turn on information which arrives before the deadline in 72 games and too late in 52, and the price of “too late” is a 0.05%-wrong declaration becoming a 46.3%-wrong one. So the mechanism is not mis-aimed and not mis-gated — it is running a race it cannot see the clock of, and a future attempt would need a condition predicting whether the information will arrive in time, not a better target or a wider threshold.

Those four rates are the pre-registration’s own, and one of them has since moved. A fresh 1,600 games of that arm, versus the same opponent, put the forced path at 36.6% rather than 46.3%, and the cause is measured rather than guessed: the journal was written twenty minutes before `claim_forced_exhaustive` shipped into the champion, and re-running the new instrument with that field switched back off returns the forced path to 44.6% [37.7, 51.8], against 36.6% with it on (`results/signal_deadline_noexhaustive.json`). The journal records the bridge revision and no engine digest, which is exactly how a figure outlives the program that produced it. Nothing above changes: the decomposition and its 52-against-72 split were taken on one engine and are internally consistent, and on today’s champion, measured on that one run so that both rates come from the same program, a forced declaration is still wrong 273 times as often as a voluntary one (36.6% against 0.13%). The price of “too late” fell; it did not stop being the thing that separates the two groups.

A new instrument answers the question this section leaves open, and its first answer was wrong in a way worth recording. Measuring the state at each signalling decision over 1,600 fresh games (`results/signal_deadline.json`) shows that the stall clock — 80 actions with no half-suit resolved, after which the engine forces a declaration — does *not* separate the two groups at the moment the signal fires: 13.0 actions against 12.1, an interval covering zero. What separates them there is how much game is left, and chiefly how many legal asks the seat still has, 31.9 against 20.5. Read from those observables alone the conclusion is that the mechanism runs out of asks rather than time. It is the wrong conclusion. Counting which of the engine’s two forced routes actually fires puts 185 of the 244 bad episodes on the stall route and only 59 on the no-legal-ask route.

Both readings are of the same data, and the reconciliation is the result. The gate is a per-turn predicate, so a seat with no good ask to make keeps satisfying it and signals again: episodes that end badly carry a mean of 42.5 signals against 5.6 for those that end well, and the stall clock they started at 12 stands at 59 by the last of them. The mechanism is not racing a clock it cannot see — it is running that clock down itself, and low askability predicts which episodes will. What the row needs is therefore not a deadline

predictor but a limit on repeats: a signalling ask proves once that a seat does not hold one named card, and proving it forty more times spends forty actions of the seat’s own window to say nothing new.

That last sentence is right about the information and wrong about the value, which was worth the cost of finding out. Registered in advance and run over 4,000 games, suppressing the repeats makes the engine *worse* by -0.0715 $[-0.1068, -0.0362]$ sets a game (prereg/signal_no_repeat.md, results/signal_no_repeat_9900000.json). A second registered run at a fresh seed base replicated it at -0.1075 $[-0.1429, -0.0721]$ (results/signal_no_repeat.json), and measured the mechanism’s own value on the current engine at $+0.1435$ $[+0.0971, +0.1899]$ — consistent with the $+0.1220$ above, and the figure to quote for signalling after the exhaustive-search change. The two arms are level on declaration errors, 0.1383 a game against 0.1388 , so nothing is paid in mistakes. What the repeats buy is postponement. The engine reaches the signalling branch before the gated declaration, so a seat that can signal signals *instead of* declaring at a bar that is wrong about a quarter of the time, and signalling again next turn defers it again: 0.069 gated declarations a game against 0.299 with signalling off, at a tenth the error rate. Remove the repetition and the deferral goes with it — the suppressed arm returns to 0.226 gated declarations at 23.4% wrong and lands back on the no-signalling baseline, $+2.4395$ against $+2.4450$. Across the three runs of this arm the repeats are 92.7% of all fires — $33,248$ fires over $2,473$ distinct cards in the largest, so the asks that carry no new information are almost all of them, and they are what buys the delay.

What that does *not* establish is where the value goes, and an earlier draft closed here by calling the mechanism “a stall that happens to communicate”. That reads a ledger of one team’s declarations as though it were the whole scoreboard. This instrument counts what our three seats declare and discards what the opposing three do, and the answer was in the half it discarded.

The margin is an identity, and this line was reading a third of it. Under the award rule the NULL_TEAM branch of the engine’s claim resolution is unreachable, so each of the nine half-suits is awarded to exactly one team by exactly one declaration and by nothing else. Writing D for declarations a side makes and W for the ones it loses,

$$\text{margin} = 2(D_{\text{us}} - W_{\text{us}} + W_{\text{them}}) - 9.$$

Three channels, each worth two sets a declaration: how many of the nine we get to declare at all, how many of ours we get wrong, and how many of theirs *they* get wrong. Every instrument in this section reported the middle one and neither of the others. scripts4/margin_identity.py computes all three; the residual is zero to fifteen places on every run this project has stored, which it must be, because the relation is arithmetic and not a model.

Re-reading the arms through it, on a run that counted both sides over 2,000 paired deals (results/where_the_margin_lives.json):

arm	effect	race	ours	theirs
signalling	+0.1190	−0.1655	+0.0445	+0.2400
deferred gate	+0.0540	−0.0645	+0.1040	+0.0145

Signalling’s own-error saving is a fifth of its effect. It buys a quarter of a set a game of the *opponents’* mistakes and hands two thirds of that back in half-suits it never gets to declare, because each signal is a deliberately doomed ask and a doomed ask loses the turn. The deferred gate is the opposite: its whole margin is our own errors, and it moves the opponents’ by $+0.0145$. They are two mechanisms, not two doses of one, which is why they do not stack.

Some of that opponent gain is arithmetic — signalling hands them 0.083 declarations a game and they are wrong on about a fifth of anything they declare. Splitting $w_{\text{them}} = d_{\text{them}} e_{\text{them}}$ at the mean assigns $+0.0362$ to that handover and **+0.1997** to a change in the rate itself, which rises from 21.87% to 24.37% . Deferral moves the same rate by $+0.0001$. So 83% of what signalling buys is that the opponents declare *worse*, not merely more often.

The repeats are the mechanism rather than its waste: suppressing them keeps the own-error saving and returns the opponents' rate to baseline. Read that way the 92.7% of asks carrying no new information are not a cost the protocol pays, they are what it is for.

A budget tests the interior, and the dose-response is convex. Suppressing the repeats is one end of that curve. The other question is whether the first few signals buy most of the confusion while the rest only burn turns, in which case a cap would keep the opponent channel and return the race. That was registered before the parameter existed, with three withdrawal conditions and a stated power limit (`prereg/signal_budget.md`), and run over 4,000 fresh games at a seed base used by nothing else (`results/signal_budget_11700000.json`). All three conditions pass. The mechanism replicates a third time at $+0.1605$ [$+0.1134, +0.2076$] against the registered $+0.1435$, a two-sample z of $+0.50$ carrying both uncertainties; the caps bind in the registered order at 8.94, 1.48 and 0.69 signals a game; and the identity closes on all four arms. The primary is refuted. A per-game budget of six is -0.0780 [$-0.1065, -0.0495$] against the uncapped arm, and the channels say why.

arm	effect	race	ours	theirs
uncapped	$+0.1605$	-0.1585	$+0.0465$	$+0.2725$
budget 6	$+0.0825$	-0.0115	$+0.0795$	$+0.0145$
budget 2	$+0.0645$	-0.0075	$+0.0555$	$+0.0165$

A budget does two of the three things it was meant to do. It returns almost all of the race loss, and it keeps the own-error saving; it improves it, to the best of the four arms, at 2.46% of our declarations wrong against the shipped 3.25%. It then deletes the channel that was most of the gross gain. At 1.48 signals a game the opponents' extra wrong declarations are $+0.0073$ [$-0.0063, +0.0208$], which covers zero, and their error rate is 21.24% against a baseline of 21.08% and an uncapped 24.02%. So the dose-response is convex rather than concave: volume is most of what the mechanism buys, and there is no cheap version of this protocol that keeps the channel intact. Two registrations say so from opposite directions — suppressing repeats deletes the channel, and so does rationing them. A later run at 2.5 times this power sharpens the claim and takes one word out of it: at 1.38 signals a game the channel is $+0.0088$ [$+0.0007, +0.0169$], small but real, so it is not that nothing survives a cap, only that a quarter of proportional survives it. That is the subject of its own registration below.

One figure from that run is worth recording rather than acting on. The replication arrived at $+0.1605$, above the $+0.15$ bar the roster was given before any arm was seen, where the registered run arrived at $+0.1435$, below it. Taking the larger of two agreeing replications is exactly the selection this project is built to avoid. Pooled on both uncertainties they are $+0.1519$ [$+0.1188, +0.1849$], a point estimate sitting on the bar with an interval spanning it. Nothing ships on that, and this registration fixed no ship criterion; the deployed configuration still has `signal_mode="off"`.

Two limits, stated with it. The three channels are an *accounting* and they co-move — a half-suit we stop declaring leaves RACE and arrives in THEIRS carrying their error rate with it, so no channel is what the arm would gain if only that counter moved. And the rate this rests on is one opponent's: a screen over every honest policy in the registry (`results/opponent_error_screen.json`) puts `dylan_v07` at 21.8%, the next at 10.8%, and the rest between 3% and 7%. Sixty deals is enough to rank the pool and not enough to pin a rate; a later section records what happened when one of these figures was used as though it were. The effect is real and replicates across five arms and two engine revisions, every one of them against that same opponent. Whether it is a property of the convention or an exploit of one policy is not something this opponent pool can answer, and a registered attempt to make it answer established why not.

The generality run, and the confound that was in its registration. `prereg/signal_generality.md` put the arm against the only other honest opponent with error volume to move, and against the champion itself as a control that could refute the identity's account without being able to confirm generality. Both ran over 800 paired deals at a fresh seed (`results/signal_generality_ev_claim_12100000`).

json, results/signal_generality_self_12100000.json). The control holds: against itself the arm is +0.0200 [−0.0232, +0.0632], covering zero, so the margin is not arriving from somewhere the identity fails to account for. Against ev_claim nothing moves either, in the margin at −0.0525 [−0.1415, +0.0365] or in the channel the mechanism works through, at −0.0088 [−0.0296, +0.0121], whose upper bound excludes the +0.0458 a full-strength generalisation predicts.

That reads as an opponent-specific effect until the dose is put beside it.

opponent	arm	signals/game	their extra wrong declarations
dylan_v07	uncapped	8.940	+0.1363 [+0.1135, +0.1590]
ev_claim	uncapped	2.171	−0.0088 [−0.0296, +0.0121]
dylan_v07	budget 6	1.477	+0.0073 [−0.0063, +0.0208]
dylan_v07	budget 2	0.686	+0.0083 [−0.0043, +0.0208]
self	uncapped	0.487	+0.0069 [−0.0010, +0.0147]

Every row consistent with zero is a row below three signals a game, and the single row with a channel is the only one that rises above it. That pattern is real and the inference it invites is not: a later registered run at 2.5 times the power finds a channel at 1.38 signals a game, so these rows cover zero because they are underpowered and not because the mechanism has a threshold. The budget result already established that dylan_v07’s own channel disappears at that dose, so the dose alone accounts for the ev_claim null and opponent robustness is not needed to explain it.

The defect is in the registration rather than in the run. It predicted each opponent’s effect by applying the rate multiplier measured against dylan_v07 to that opponent’s own baseline error rate, which assumes the mechanism fires equally often against everyone. It does not. How often a seat is stuck with nothing worth asking is a property of the opponent one is playing, not a parameter of the experiment, and it varies by a factor of eighteen across this grid. The registration fixed the sample size, the verdicts, and two power limits in advance, and did not fix the one quantity that decides whether the contrast is between opponents or between doses.

So the question stands open with a sharper statement of what would close it. Only dylan_v07 supplies the volume at which the mechanism works, and a budget can only lower a dose. Until an opponent other than dylan_v07 can be made to draw eight signals a game, “the convention needs volume” and “dylan_v07 is exploitable at volume” are the same measurement.

Why the dose differs, and why that breaks the design. A screen over five opponents answers it (scripts4/signal_dose_screen.py, results/signal_dose_arms.json): 400 deals each, both arms on the identical deal, with every quantity about the opponent rather than the mechanism read off the arm that does not signal. The dose factors exactly, into three terms that are each opponent-dependent:

$$\text{fires per game} = \underbrace{s_A}_{\text{opportunity}} \times \underbrace{s_B/s_A}_{\text{amplification}} \times \underbrace{f/s_B}_{\text{second gate}}$$

where s_A and s_B are stuck turns a game with the protocol off and on, and f is fires a game.

That equation is an *identity*: the two inner terms cancel and leave f . Writing a “product” column beside a “measured” one, as an earlier revision of this table did, therefore checks nothing at all — the two can differ only by display rounding. What the decomposition buys is which *term* differs by opponent, so only the terms are shown, each with the interval it went two revisions without.

opponent	s_A , stuck turns/game	s_B/s_A	f/s_B
dylan_v07	4.150 ± 1.238	3.02 [2.32, 4.16]	0.896 [0.868, 0.916]
ev_claim	3.005 ± 0.452	1.65 [1.46, 1.88]	0.462 [0.428, 0.497]
search	3.112 ± 0.443	1.64 [1.45, 1.87]	0.390 [0.359, 0.420]
memory	2.770 ± 0.421	1.95 [1.70, 2.25]	0.421 [0.388, 0.455]
self	1.030 ± 0.214	1.44 [1.16, 1.76]	0.423 [0.313, 0.525]

The two ratio columns are cluster bootstraps over deals, since both arms are played on the identical deal and a ratio of aggregates is not a mean. Against `ev_claim` the dose ratio is 4.93 and the three terms contribute 1.38, 1.84 and 1.94. No term dominates, and the first is not what it looks like: we enter the stuck state *less* often against `dylan_v07` than against `ev_claim`, `search` or `memory` (0.517 ± 0.065 episodes a game against 0.703 ± 0.068 , 0.733 ± 0.074 and 0.812 ± 0.082), and the whole of the opportunity difference is that each episode lasts 1.87 times longer.

An earlier revision explained that away by calling board-wide ambiguity and opponent ask-hits “flat” across the five. Measured against their own intervals neither is: ambiguity runs 24.49 ± 0.30 to 26.02 ± 0.21 cards and their asks land on our seats 22.51 ± 0.70 to 27.94 ± 0.74 times a game, both separated several times over. The conclusion survives, and for a better reason than the one it was given. `dylan_v07` leaves our hands the *most* located of the five — the lowest ambiguity in the column — and is the opponent we get stuck against most; `memory` leaves them least located and we get stuck against it least. Whatever drives the opportunity term, it is not that this opponent tells us less about our own cards, because it tells us more.

The second gate is the largest single term and the most interpretable. The protocol needs the stuck state *and* a best ask unlikely to land, and against `dylan_v07` 89.6% [86.8, 91.6] of stuck turns clear that bar against 39 to 46% everywhere else, with no interval on either side coming near the other. When we cannot place a half-suit against this opponent we usually have nothing worth asking anywhere either: the position is more thoroughly exhausted, not merely more ambiguous.

The middle term is the one with consequences beyond this section. A signal is a doomed ask, so it throws away the turn and hands back the same stuck state, and the protocol therefore *extends the state that triggers it* — by $3.02\times$ against `dylan_v07` and 1.44 to $1.95\times$ against the rest. The dose is thus partly endogenous, and it is not a property of the opponent that a design could measure in advance and match across arms. That is precisely what `prereg/signal_generality.md` assumed when it predicted each opponent’s effect by scaling one rate multiplier onto that opponent’s baseline. The registration was not merely unlucky in its power; its independent variable was not independent. Amplification tracks the second gate across these five opponents, which is what a mechanism that spends turns to stay where it is should do, but five points do not fix the functional form and none is claimed.

A matched-dose design therefore cannot be built by choosing opponents. It has to hold fires a game fixed by construction, and the budget parameter of §8.7 caps a dose but cannot raise one.

The matched-dose comparison, and the first generality result in this line. `prereg/signal_matched_dose.md` registered that design before either of its banks existed. The lever that raises a dose is the cheapness gate, `signal_max_p`: the protocol fires only where our best ordinary ask is unlikely to land, and opening it fires on stuck turns where a decent ask remained. It does not change what a signal proves, because `signalling_ask` picks a card in a half-suit our own team provably holds and is doomed whatever the gate says; only the opportunity cost of the turn moves. The price is that the arms are then different policies, and the comparison is of the protocol *tuned to fire N times a game* across opponents rather than of the shipped protocol.

Calibration (`results/matched_dose_calibration.json`, a bank that scores nothing) settled two guesses. The gate is saturated against `dylan_v07`, where opening it from 0.50 to 1.00 moves the dose from 9.242 to 9.040; and against `ev_claim` opening it fully buys only 2.348 to 3.112, because the amplification term does not rise with it. So the highest dose both reach is $D = 3.1$, and the reference opponent is brought down to it with a budget of 16.

That is far below the 8.940 at which the effect was known to exist, so the registration’s feasibility gate applied: the reference opponent’s own channel had to clear zero at D or the study was to be abandoned unrun. It cleared, at $+0.0600$ [$+0.0118$, $+0.1082$], and in clearing it filled the interior of the dose-response that had only three measured points. **This falsifies a sentence in the previous paragraph of an earlier revision, which said no other opponent reaches a volume at which the effect exists: the effect exists at 3.1, and `ev_claim` reaches it.**

The scored run, 2,500 paired deals at a fresh bank (`results/matched_dose_scored.json`), lands both opponents inside the registered 15% dose tolerance at 2.766 and 2.872, with no unfinished games and no bridge fallbacks:

opponent	dose	their extra wrong declarations a game
dylan_v07	2.766	+0.0454 [+0.0326, +0.0582]
ev_claim	2.872	+0.0172 [+0.0048, +0.0296]

Both clear zero, so at equal dose the mechanism raises a second engine's declaration error rate and the effect is *not* specific to dylan_v07. The registration named two ways it could generalise and the interval separates them. Transferring this run's own reference rise of 1.136 points to ev_claim's 7.96% baseline, the same rate *factor* predicts +0.0160, which is inside the interval, and the same rate *rise* in points predicts +0.0423, which is outside it. So the effect is not a fixed number of percentage points: it scales with how often the opponent was already wrong, and an opponent who misdeclares half as often loses about half as much to it.

What that does not establish, having named only two candidates. A registration can only separate the hypotheses it writes down, and the two this one wrote down are not the only ones. A rate bounded by one is more naturally shifted in log-odds than multiplied, and a constant log-odds shift fitted to the same reference point predicts +0.0188 at ev_claim, which is *also* inside the observed interval. That is not a coincidence to be argued away: for small p the odds $p/(1-p)$ are approximately p , so multiplicative and logistic scaling agree wherever the baseline is low, and every opponent in this registry except one has a baseline below 22%. Over that range the two laws differ by less than 20%, which is far less than any interval this study can measure.

They separate in one place. heuristic misdeclares far more often than anything else in the registry — the screen that found it said 72.5%, and a proper measurement below says 62.0%, which is the subject of its own paragraph — and at the rate the registration used, the multiplicative law predicts +0.0426 against the logistic law's +0.0143, a factor of 2.97. prereg/signal_dose_law.md registered that test before the bank existed, naming both predictions and naming in advance the outcome that would hurt most: an interval excluding both, which would mean the transfer is not a one-parameter function of the baseline at all.

It is logistic, and the multiplicative reading is refuted. Run at the same D on 2,500 paired deals (`results/signal_dose_law.json`), heuristic lands at +0.0178 [+0.0048, +0.0308]: covering the logistic prediction and *excluding* the multiplicative one. All four withdrawal conditions hold — the identity closes on the recorded ledger with a residual of zero on every game, the dose lands 0.2% from D , no game was unfinished, and the realised half-width of 0.0130 is inside the registered power limit of 0.0142, so the run separates the laws rather than failing to.

That refutes the reading this section carried in an earlier revision, which called the effect a constant *fraction* of the opponent's baseline errors. It is not; the fraction shrinks as the baseline rises, which is what a bounded rate must do and what the multiplicative law cannot represent.

One of the two registered inputs was a screen figure, and it was wrong. Both predictions were computed from heuristic's baseline of 72.52%, and that number came from `results/opponent_error_screen.json`: a 60-deal descriptive screen, 131 declarations, 95 of them wrong. A screen is good to about four points and this registration quoted it to four decimals. Measured properly on its own bank of 1,000 games and 1,166 declarations (`results/heuristic_baseline.json`), the rate is 62.01% — a miss of 10.5 points, about 2.4 of the screen's own standard errors, so sampling noise rather than a defect in either measurement. The defect is upstream of both: nothing in the pipeline objected when a figure good to ± 4 points was used as though it were exact, and the check that caught it was written after the run rather than before.

Recomputing both predictions at the measured baseline over its measured 1.166 declarations a game gives +0.0389 multiplicative and +0.0182 logistic. The interval still excludes the first and covers the

second, so the registered verdict is also what the properly measured input gives, and the separation the test relied on narrows from a factor of 2.97 to 2.14 without closing. That is a post-hoc robustness check and is worth exactly what post-hoc robustness checks are worth: it could not have rescued a verdict that flipped, and it is reported because the input it corrects is one the registration got wrong.

One parameter, three baselines. The table below reads every baseline off the bank that measured that opponent’s control arm at length, and `scripts4/dose_low_table.py` assembles it from those files (`results/dose_low_table.json`) rather than from arithmetic anybody did by hand — which is how the four figures that stood here in the first draft of this paragraph came to be wrong in their last digit.

opponent	baseline	declarations/game	observed	predicted
ev_claim	7.96%	3.728	+0.0172	+0.0188
dylan_v07	21.08%	3.998	+0.0454	+0.0454
heuristic	62.01%	1.166	+0.0178	+0.0182

The column on the right is one number applied three times, and the number was fitted to the middle row alone: a log-odds shift of +0.0669, taken from `dylan_v07` and carried unchanged to the other two. It predicts `heuristic` to within 0.0004 at a baseline nearly three times as high, and `ev_claim` to within 0.0016 at a bit over a third of it. Refitting by least squares on all three moves the shift to +0.0660, a change of 1.4% over an eightfold range of baseline error rate. So the protocol does not merely transfer between opponents: it shifts the *log-odds* that an opponent misdeclares by a constant, and one parameter measured against one engine predicts the other two.

What that is not is a mechanism. A one-parameter law fitted to three points and chosen from two candidates is a description that survived a test designed to break it, and the test was a good one — the two candidates differ by a factor of two even at the corrected baseline — but nothing here explains *why* the shift should be constant in log-odds, and a third functional form nobody wrote down could fit these three points too. The claim is that the opponent’s baseline error rate is enough to predict the effect at these three baselines, one dose and one rule set, not that it is the only thing the transfer depends on.

Three things travel with that. The dose is 2.8, not the 8.9 the shipped protocol produces against `dylan_v07`, so this is generality at a volume neither engine plays at by default. The test opponent’s arm has its gate fully open, which costs it in tempo: its margin effect is -0.1372 ± 0.0529 against the reference arm’s $+0.0760 \pm 0.0393$, exactly the instrument cost the registration said in advance was not the outcome being measured. And this is two engines under one rule set, not a statement about human declarers.

And along the dose axis it is neither a line nor a threshold. The transfer law was measured with the dose held fixed by construction, so it says nothing about volume — and volume is the one lever this project can turn. Converting every scored arm into the log-odds shift that would produce its measured effect (`scripts4/dose_linearity.py`, descriptive) put four points between +0.0211 and +0.0242 of shift per signal, across a threefold dose range and three engines. One number appeared to describe the whole surface. `prereg/signal_dose_linearity.md` registered the test: a direct replication at power of the arm that produced the 1.48 point, 5,000 paired deals against that arm’s 2,000, with the constant fitted to a single point six times the test dose away, and four verdicts named in advance.

It came back **NEITHER**, which the registration had named as the outcome that would hurt most:

realised dose		1.381	
their extra wrong declarations	+0.0088	[+0.0007, +0.0169]	
linear predicts	+0.0201	excluded	
threshold predicts	0	excluded	

Both halves of that matter, and they point opposite ways.

There is no dead zone. The interval excludes zero, so the mechanism works at 1.4 signals a game — less than a sixth of the shipped dose, and well inside the range this paper had read as inert. The earlier

	v0.3 champion	v0.4 champion
cards gained per possession	1.419	1.295
possessions gaining nothing	44.5%	48.8%
failed-ask share	40.9%	43.2%
asks per game	100.3	96.5
ask accuracy, thirds of the game	54.1 / 59.7 / 64.2%	52.2 / 58.0 / 60.8%
misdeclaration rate, natural half-suits (void era)	4.4%	3.1%
claim delay, median / p90 (actions)	0 / 58	0 / 49
first-move advantage	+0.23 [−0.27, +0.73]	+0.32 [−0.22, +0.86]

Table 15: Homogeneous tables of each champion, 100 games each, on the measures v0.3 used to characterise skill. The v0.4 champion is stronger and looks worse on almost every one of them: it gains fewer cards per possession and fails more asks. These statistics stop tracking strength once the engine is willing to spend a turn buying information.

arms covered zero because they were underpowered, exactly as the linear reading said they would be, and the sentence above about rows below three signals a game describes the power of those rows rather than a property of the mechanism.

And it is not proportional. The effect is a quarter of what a constant shift per signal predicts, and that prediction is excluded. The per-signal shift is +0.0095 here against +0.021 at every dose from 2.8 up, so the surface turns on softly between one and three signals a game and is approximately linear above that. A single constant now survives only on the interval [+0.01823, +0.01826], a window 0.00003 wide where the highest-dose point’s lower bound grazes this one’s upper bound. That is a knife edge, not a fit.

Three things are worth saying about the procedure rather than the result. The constant was fixed in advance from one named point, and it was recorded before the run — in `results/dose_linearity_points.json` — that it already sat *above* what every existing point jointly permitted. Refitting it against the low-dose data the run existed to test would have been the failure registrations are for. All four withdrawal conditions passed: the identity closes with a residual of zero on every one of 20,000 games, the realised dose is 6.5% from the arm being replicated, the half-width of 0.0081 is inside the registered limit of 0.0095, and there were no unfinished games or bridge fallbacks. And two doses do not make a curve: a smooth saturating function through these points would be indistinguishable here and was not tested. What is dead is the pair of readings this project actually wrote down.

One condition was not checked. The registration requires the margin identity to close on every arm, and the scoring instrument as run did not record our own declaration ledger, so that could not be verified from its payload. The primary here is a directly counted quantity and does not pass through the identity, the rule the identity depends on is pinned in the harness, and the instrument now records the ledger and reports a residual — zero on every game of a re-run smoke. But on the 14,300,000 bank the condition is recorded as unverified rather than as passed.

One reading from that run survives its confound, because it does not depend on the opponent channel at all. Against `ev_claim` the protocol is almost purely a deferral device: its split is race −0.2725, ours +0.2375, theirs −0.0175, and our own wrong declarations halve from 0.2350 a game to 0.1163, the largest own-error saving measured anywhere in this line, as gated declarations fall from 0.441 to 0.098. It still does not pay, because it costs more in race than it returns.

One methodological note, recorded rather than buried. The regression form of the same decomposition puts the cost-when-nothing-is-avoided at +0.0990 [+0.0117, +0.1863], excluding zero, where the directly conditioned figure is +0.0502 [−0.0248, +0.1253], which contains it. The line is misspecified: avoiding an error is worth +1.61 and adding one costs only −0.96, so the errors this protocol adds are cheaper than the ones it avoids, and a straight line through both cannot represent that. The conditional mean is the one to believe.

10.2 Game statistics, and a statistic that stops working

Homogeneous tables of each champion, 100 games each, on the measures v0.3 used to characterise skill:

Three of these reproduce v0.3 and one contradicts it.

Reproduced: there is still no measurable first-move or seat advantage, both intervals comfortably containing zero with the starting seat rotated across deals; the synthetic “eights and jokers” half-suit still behaves like any other (mean resolution rank 4.08 against 3.99, null rate 3.0% against 3.1%); and ask accuracy still rises from the opening into the later game, so the opening remains the least-informed phase.

Contradicted, and worth stating carefully: v0.3 named cards-per-possession “the single best summary statistic of skill”, on the evidence that it doubled from the memory tier to the probabilistic tier. The stronger engine here scores *lower* on it. The mechanism is visible in the other rows: v0.4 misdeclares fewer sets (3.1% against 4.4%) and declares sooner (p90 49 against 58), so half-suits resolve faster and the game contains fewer, harder asks — 96.5 per game against 100.3. Cards-per-possession compares two tables on how easy their asks were, and how easy the asks are depends on how quickly the claiming policy is removing the resolved material. Within a tier that shares a claim policy it tracks skill; across engines whose claiming differs, it does not. A statistic that inverts when the thing it is measuring improves is not a measure of skill, and we would retract the v0.3 claim rather than qualify it.

10.3 The engine as an instrument

A strong policy that can only be measured in aggregate is a weak research instrument, and an unusable one. The three things a chess engine gives a user all exist here; they just mean something different in a game whose hidden state is the entire problem.

Evaluation, in sets. The banked differential plus, for every unresolved half-suit, the learned probability that our team scores it minus the probability that theirs does (11). So +1.4 reads as “expect to finish 1.4 sets ahead of the current score”. The underlying model is fitted to outcomes and scores well in aggregate: held-out log-loss 0.727 against 0.821 for the class prior, with a mean bias of -0.006 . It is not, however, calibrated. That mean bias is near zero because errors cancel, not because each decile is right, and five of the ten deciles are individually miscalibrated at the 95% level (§9.1), so the number is a usable ranking and should not be read as a probability. This is the one component that is a good *predictor* and a disastrous *objective* (Section 9.1), which is itself worth knowing: an evaluation function and an objective function are not the same object.

A ranked move list with attribution. Every legal ask, with the probability it lands, the evaluation on each branch, and a per-term breakdown of the policy’s score. The breakdown matters for a research instrument rather than a toy: when a human disagrees with the engine, the disagreement can be attributed to a named term instead of left as an oracle’s verdict.

A principal variation. In chess the PV is the expected line. The Fish analogue is the *possession*: the chain of asks the engine would make if each landed, since a turn in Fish is a run of successful asks. It is computed by pinning each card as it is won and re-ranking, so it is the engine’s actual plan rather than a static list.

Two things a chess engine never has to show, and this one must. The **card table** is the posterior over who holds every unresolved card — where the strength of this version actually lives, so it is displayed rather than hidden. The **declaration table** gives, per live half-suit, the probability our team holds all six, the most likely split, and the exact probability that split is right, which is the quantity the claim decision turns on.

The information boundary survives the user interface. Every panel is built from the seat’s own Observation; the posterior shown is an inference from the public record, not a look behind it. The only exception is the reveal after the last card is public anyway.

Mixed tables. The same server seats any mixture of people and engines at one table, which is what makes the engine usable as an instrument on human play rather than only on self-play. The design point worth recording is that this is not a cosmetic choice: teams in Fish are the alternating seats $\{0, 2, 4\}$ against $\{1, 3, 5\}$, so a request as ordinary as “three of us against three bots” is a constraint on *which* seats the people take, not on how many. A table therefore takes a seat *arrangement* rather than a bot count — one that fills a whole side before crossing, and one that interleaves, and the 3-versus-3 configuration that a group of friends actually wants is the former at $n = 3$. Engine seats move on a clock rather than

instantly, because a decision costs about 4 ms and an unpaced possession appears to a human as a single discontinuous jump. Pacing alone is not enough: a longer gap is only a longer gap unless there is something to read in it, so the table states each move in words together with what it *proved*. That second part is free, and is the one place where the no-bluff rule pays the observer directly: an ask is a truthful statement about the asker’s hand as well as a question about the target’s, so a failed ask establishes that the target does not hold the card, that the asker does not either, and that the asker holds some other card of that half-suit. The default gap is 12 s and is adjustable in play. Freezing the table turns the same control into a step-through — one engine move per request, still frozen afterwards, which is what makes an engine’s whole possession readable, since a possession is a run of successful asks and it is the run, not any single ask, that carries the plan. Each seat is served only its own *Observation*, so the information boundary of Appendix F is enforced by the transport, not merely respected by the client.

Analysis costs 130–270 ms, which is interactive, and the same objects drive the terminal client, the browser client and the batch experiments, so what a user sees is what the experiments measured.

10.4 Does Fish have a stalemate? Three layers of answer

Chess has stalemate and threefold repetition. The natural question is whether Literature has an analogue (a position the game cannot get out of), and the answer comes in three layers, of which only the third is interesting.

10.4.1 Layer 1: the state graph is cyclic

Nothing in the rules forces progress. Two opponents can trade a card back and forth and return to a bit-identical position, so positions can repeat. This is a property of the rules, not of any implementation. It follows immediately that *termination is a property of the players*: any claim that “every game of Fish ends” describes a convention rather than the game. Our policies carry an explicit progress rule and our harness a cap, and Section 10.1 measures how much of “games end” those conventions are actually doing.

10.4.2 Layer 2: under perfect information it does not matter

Cycles exist, but the exact solver of Section 5 shows that no *optimal* line needs one: with all cards visible, optimal play always terminates. Cyclicity on its own is therefore not a strategic phenomenon. The reason is the closed form (8) — footholds only shrink, so the mover simply collects everything reachable and the position marches forward.

10.4.3 Layer 3: dead positions

This layer has no chess counterpart, and it is the one worth having.

Proposition 4 (Dead half-suits). *Call a live half-suit dead when all six of its cards sit with one team. Then no ask in that half-suit can ever succeed.*

Proof. A legal ask in half-suit H requires the asker to hold a card of H . If one team holds all six, no member of the other team holds any card of H , so no opponent of the holding team may ask in it at all. And a member of the holding team may only ask an *opponent*, who by hypothesis holds no card of H , so such an ask fails. No card of H therefore ever changes hands again. $\square$

Call a position *dead* when every live half-suit is dead. By Proposition 4 the material is then frozen for the remainder of the game: every legal ask is guaranteed to fail, the turn merely circulates, and the only state changes left are declarations. This is as close to a stalemate as Fish gets, and, unlike a chess stalemate, it is not a draw. The sets are already decided; they have simply not been claimed.

10.4.4 The apparent deadlock, and why it is not one

A dead position looks like a genuine trap for a team that holds a half-suit but cannot place the split among its own members. No card will move again, so (the argument runs) no new information can arrive, so the team must either gamble on a declaration or stall forever.

That argument is wrong, and the way it is wrong is a strategy.

Information does not only arrive when cards move. Under the no-bluff rule a player may not ask for a card they hold, so **asking for a card proves publicly that you do not hold it**. In a dead position that ask is guaranteed to fail, which means it costs no material whatsoever: there is no card to lose, and the turn it concedes is worthless because no opponent can win a card either. The ask is a *free one-bit broadcast to your partners*.

A team stuck in a dead position can therefore talk its way out. Each member asks for cards of the contested half-suit that they do not hold; after a few such asks the split is common knowledge within the team and the declaration is safe. Literature's only communication channel is normally costly and read by the opposition; in a dead position it becomes free, and the opposition can do nothing with what it hears, because the half-suit is already beyond their reach.

Remark 5 (A definition that made the phenomenon invisible). *Our first implementation defined “our team provably holds this half-suit” as “every card’s holder is pinned and all six are teammates”. That makes “we hold it but cannot place the split” a contradiction in terms, since pinning the holders is placing the split, and the condition never fired once in 1,047 measured plies. The correct condition is weaker and is the whole point: every possible holder of every card is a teammate, which does not require knowing which teammate. Under that definition the situation occurs on 2.6% of plies.*

10.5 The price of a turn, and why it is conditional

A turn and a card are different currencies, and this section prices the first. Half the ask objective prices *tempo* — turn_risk, depth, the retained turn a hit buys, and none of those weights was ever fitted against a measured scale, because none existed. This section supplies one, and the answer is not a constant.

10.5.1 Measuring it

The design is a duplicate deal with one intervention. Both arms play the identical deck with identical agent seeds; at one pre-chosen decision of one pre-chosen seat the treatment substitutes an ask that provably cannot land, taken from the true state so the donation is certain rather than a modelled mistake, and both arms then play on honestly. The difference in the donating team's differential is the price of exactly one turn. Substituting the agent's *own* action instead gives exactly 0.000 on every deal, which is the control that makes the number mean anything.

Two runs of 1500 pairs each put the price at +0.271, 95% CI [+0.092, +0.450] when the donation falls at decisions 0–4, and +0.497, 95% CI [+0.324, +0.671] at decisions 14–21. Tempo value rises through a game.

10.5.2 It depends entirely on whether you have an ask worth making

Bucketing 1464 donations by p_{best} , the success probability of the ask the seat would otherwise have made:

p_{best}	pairs	price	standard error
[0.00, 0.25)	374	−0.043	0.169
[0.25, 0.50)	511	+0.004	0.143
[0.50, 0.75)	179	+0.508	0.258
[0.75, 1.01)	400	+0.415	0.164

Above $p_{\text{best}} = 0.5$ a turn prices at about +0.45 sets. Below it this design resolves nothing: the two low bands carry 95% intervals of [−0.374, +0.288] and [−0.276, +0.284], so anything up to about a

third of a set remains possible there. **What the run establishes is the conditionality, not a zero** — a turn is worth much less to a seat with no good ask than to one with a good ask, and how much less is not measured here. Both halves are champion self-play (`scripts4/turn_price.py` seats six copies of the engine), so this is a statement about how this policy values a turn, not about Fish.

A provenance note this table needs. The four prices and standard errors above are not recoverable from any results file. `scripts4/turn_price.py` persists its summary but not the per-decision p_{best} values the bands are cut on, so the table cannot be rebuilt from the run that produced it; and `scripts4/tempo_regime.py`, which writes `results/tempo_regime.json`, carries these four prices as a hardcoded constant copied from this paper and measures only the bucket *counts* beside them, which it puts at 340/347/155/356 over 1,198 decisions against the 374/511/179/400 over 1,464 printed here. So the one file that appears to back this table in fact contains the table’s own numbers. This is the defect §9.4.1 names in another context (a figure that lives in the paper and in no results file is one that stays unchallenged), and it is recorded here rather than repaired, because repairing it means re-running the measurement and this paper does not have that run.

This is not the timing effect in disguise, and the direction rules that out: a late turn prices *higher* than an early one, so if low- p_{best} positions were merely late positions they would be worth more rather than nothing.

10.5.3 What it explains, and what it cost to find out

The measurement began as a critique of the engine. On 1.5% of decisions (229 of 15,542) the champion plays an ask that cannot land while one with median success probability 0.385 was available — it surrenders the turn for certain. That looked like a defect, and `avoid_doomed_asks`, which restricts the choice to asks that can land and ranks them by the same objective, was pre-registered against it with a predicted +0.05 to +0.15.

Over 8000 duplicate deal-pairs it scored +0.017, 95% CI $[-0.024, +0.059]$: at most +0.06 with 95% confidence, under a fifth of the largest gain this project has demonstrated, and plausibly zero.

Reconciling that with the turn price took four rounds, and three of the four found errors in the bridge rather than in the game. Treating the measured price as the value of a turn understated tempo by a factor of 1.7, since the donation displaces the agent’s normal action, which itself retains the turn only with probability 0.583. Counting firings across all six seats double-counted, since the arm changes only the three seats it plays. Neither helped. The band table did: the branch fires when $p_{\text{best}} = 0$ *by definition*, at the floor of the worthless range, and the substitute’s median $p = 0.385$ sits inside it too. The arm retains turns worth approximately nothing, 38.5% of the time, and its predicted effect is approximately zero.

So the champion’s objective was right and the critique was wrong. Those 1.5% of decisions throw away a *worthless* turn, which is exactly the circumstance in which the objective’s other terms should dominate the probability term, and they do. The discrepancy was a selection effect in the instrument: turns priced at positions sampled without regard to option quality and spent at positions where $p_{\text{best}} = 0$ by construction.

10.5.4 A note on public information

One candidate explanation, tested and set aside, produced a fact about the game worth keeping. Forking each firing and playing the champion’s doomed ask in one branch and the substitute in the other, the substitute sheds 0.344 more card-equivalents of a teammate’s uncertainty, and 0.269 more of each *opponent’s*. An ask is public, and it reaches two partners and three opponents, so the team-weighted balance is

$$2 \times (-0.344) - 3 \times (-0.269) = +0.120 \text{ card-equivalents per firing, favouring the quieter ask.}$$

In six-player Fish a public disclosure is net-negative for the discloser’s team whatever it says, because the opposition outnumbers the partnership three to two. This is the arithmetic behind the no-bluff structure

noted in Section 6.1: signal is welded to move, and what you tell your partner you tell your opponents more of.

11 The endgame solved under imperfect information

These figures were recomputed. The solver’s transposition table keyed nodes on hands, turn, half-suit winners and belief weights, and *not* on the history. The opponents in it are the champion, whose action is a function of its whole observation, so two nodes identical in every key field but reached by different move orders have different continuations and different values. Merging them returned one branch’s value for the other, and the maximisation read those values — on the position that exposed it, the search reported +0.2500 for a move a rollout of the same strategy scores at +0.7500, and visited 154 nodes where the corrected search visits 12,086. The fault was found by an invariant that holds with no ground truth at all: a best response may copy the champion at every information set, so it cannot score below it. Five $m = 2$ positions did.

On the 304 $m = 1$ positions that both the broken and the corrected run solve, the aggregate barely moved — the gain went +0.4403 to +0.4392. That is not because the bug was harmless. It corrupted 43 of those 304 positions individually, and the errors went *both ways*: the fix raised the optimum in 24 and lowered it in 19, so they very nearly cancelled in the mean. The champion figure is identical position for position, as it must be (it is a rollout that never consulted the table), and that identity is what licenses the comparison at all. Everything below is the corrected run.

Every “exact” result above is a statement about the perfect-information game, and Section 10.5’s companion result showed that game is degenerate. The v0.3 headline (belief agents choosing a provably optimal move in 100% of information-resolved positions) was scored against that same perfect-information optimum, and every one of those positions had a single live half-suit. This section measures what the champion does at $m = 1$ when it genuinely cannot see, which is a different question with a different answer.

11.1 An exact best response over information sets

`fish4/exact_ii` enumerates the deals consistent with the public record and optimises over the deviator’s *information sets* (the same action at every deal it cannot tell apart) by backward induction against champion opponents. There is no sampling anywhere in it.

Two assumptions, stated rather than buried. The prior is uniform over consistent deals, which is what the constraint system alone licenses; tilting it by the champion’s own opponent model would make the answer a statement about that model. And the opponents are a *deterministic realisation* of the champion, seeded from a hash of the observation, because the champion samples and a distribution over policies is not a strategy one can best-respond to. That second choice rests on a property that was checked rather than assumed: over 629 real decisions a freshly constructed agent given the same observation computes a bit-identical belief to one attached from the deal.

Two controls run first and gate the rest. Where the belief pins every card there is nothing hidden left, so the exact value must equal the closed form $2f - m$. Over 200 games it does, 344/344. And where the support is small enough to afford it, the champion’s value is computed twice — once by rolling each deal forward, once by walking the same tree the best response walks and playing the champion at the deviator’s nodes. Same strategy, two code paths, so they must agree, and they do: **601/601**. The first control cannot see a fault that needs several deals to show, which is exactly the fault that was there.

A third control tests the *belief* rather than the search, and is the only one here that does. The true deal must be one of the deals the belief admits — if it is not, the belief has excluded the truth and every value computed over that support is a value of a different game. Over 665 positions it is: **665/665**. It is free, because the support is already enumerated to solve the position at all.

pinned positions matching the closed form	34/34
the true deal in the belief's support	375/375
tree and rollout, same champion strategy	116/116
exact optimum in the seat on move	+0.7734
the champion in the same seat	+0.4484
exact gain from deviating	+0.3250
negative entries	0/109

Table 16: The $m = 2$ layer of the exact imperfect-information solve, with the controls that need no ground truth. The three agreement checks are what caught the two solver faults; the zero negative entries are the post-fix state of the defect that produced impossible negative gains.

11.2 What the endgame is worth, and what the champion gets

Across 305 genuinely hidden positions solved exactly (13 skipped for a belief support above 24 deals and 3 over the 300,000-node budget, so 95% coverage):

exact optimum in the seat on move	+0.5084
the champion in the same seat	+0.0679
exact gain from deviating	+0.4405

with median gain +0.3333, maximum +1.5556, and — as the invariant requires, since a best response can always copy the champion — *no* negative entry in 305 positions. That invariant is now a hard failure rather than a printed warning: the run that violated it wrote its results file anyway.

Two readings. First, hidden information is expensive: the closed form values the last half-suit at +1 to the team on move, and played optimally without sight of the cards it is worth +0.51. Some positions are worth -0.5 — won under perfect information, lost under optimal play without it. Second, the v0.3 certificate is not wrong but certifies a much easier thing: where the champion cannot see, it captures +0.068 of the +0.508 available.

The caveat, measured. Best-responding to a realisation one can predict is easier than beating the mixture it came from, so this figure likely overstates exploitability against the champion as it plays. How much it overstates is not settled here, but its relevance is: over 41 hidden $m = 1$ information sets at 12 seeds each, the champion repeats its move every time in only 32% of them, averaging 2.41 distinct actions. It mixes at most of these decisions.

11.3 Two live half-suits, and what it costs to say so

The same solver at $m = 2$. The layer is where the impossible negative gains appeared, and where the second fault lived: a deal the champion has no move in was dropped from the belief and the survivors reweighted, instead of being scored where it stands. At $m = 2$ that case is routine — 2,821 deals over the run, and at $m = 1$ it never happens once, which is why the fault was invisible in the layer above.

This number is not comparable with $m = 1$'s, and the reason is coverage. Only 109 of 313 genuinely hidden positions were solved — 35%. Of the rest, 141 had a belief support above 24 deals and 63 exceeded the 300,000-node budget, and those are the *wide* positions: the ones where the champion knows least. (A further 28 positions over the budget had a support of *one*: those are pinned, so they belong to the control above rather than to the hidden set, and counting them here would understate the coverage. The control therefore ran on 34 of the 62 pinned positions available, not all of them.) Supports of 2 and 3 alone account for 50 of the 109 solved. So +0.3250 is the gain on the easier half of a harder layer, which is exactly where one should expect the champion to look better than it is, and it should not be read beside $m = 1$'s +0.4405 as though the two were measured on comparable samples.

That is the honest limit of exact play in this project as it stands: $m = 1$ is solved at 95% coverage, and $m = 2$ is solved on a third of its positions, selected for being easy.




Figure 6: What an exact deviation is worth, against how much the belief leaves open, over the 305 solved $m = 1$ positions. Gray points are positions (jittered for visibility); blue points are means over support bands with 95% confidence intervals. The gain rises with support — and the unsolved positions lie beyond the right edge of this plot, at a median support of 48 against a solved median of 3. That is why the solved-layer means read as lower bounds, why §11.4 tries to replace that extrapolation with bounds, and why the honest answer is that the plot cannot be read past its own right edge.

Both gains are lower bounds, and the reason is the same coverage that limits $m = 2$. The exact gain rises with how much is hidden. Over the solved positions, splitting each layer at its median belief support:

	narrow support	wide support	Welch t
$m = 1$	+0.3295 ($n = 172$)	+0.5840 ($n = 133$)	+5.37
$m = 2$	+0.2553 ($n = 63$)	+0.4205 ($n = 46$)	+2.06

At $m = 1$ the trend is solid on its own ($r = +0.190$ against support, $t = +3.37$ over 305 positions; Figure 6); at $m = 2$ it is suggestive rather than significant ($r = +0.145$, $t = +1.52$), and the two agree in sign and rough magnitude. The mechanism is the obvious one: the more the belief leaves open, the more an exact answer is worth over a heuristic.

The positions the solver cannot reach are precisely the wide ones — 16 unsolved at $m = 1$ with median support 48, against a solved median of 3; 204 unsolved at $m = 2$ with median support 60 against a solved median of 4. The obvious reading is that both reported means are therefore *lower* bounds on the gain over the whole layer.

That reading is an extrapolation, and §11.4 tries to replace it with bounds that hold on every position rather than a slope read past its data. The attempt does not settle it: the bounds are real, they were validated on 716 positions, and they are too loose to decide. It does establish something smaller and firmer — that the cheap proxy which *looked* like it settled the question cannot, because the proxy loosens as the belief widens. No corrected figure is offered. The honest position is that the direction of this bias is unknown, not that it is known and favourable.

11.4 Bounding the layer instead of solving it

Two thirds of $m = 2$ is unsolved, and the previous section leant on a slope fitted below the cap to say which way that biases the headline. A slope fitted over supports 2–24 and read across a gap running to 60,480 deals is an argument, not a measurement. Two quantities replace it, and both hold by construction on *every* position, reachable or not.

The two bounds. For a position with belief support $\{d\}$, weights $\{w_d\}$ and champion value C :

- **Upper.** $U = \sum_d w_d V_{PI}(d)$, where $V_{PI}(d)$ is the solver run on the single deal d . Solving one deal is the perfect-information relaxation (the deviator is told the deal and may answer differently in each),

and any genuine imperfect-information policy is one the relaxation may also play. Hence $V_{II} \leq U$. No new solver is involved: ExactII on a single state is $V_{PI}(d)$.

- **Lower.** $L = \max_a \sum_d w_d \text{champion}(d | a)$ over root actions legal in every deal: play a , then revert to champion play. One policy is attainable, so $V_{II} \geq L$. The champion's own root action is a candidate, so $L \geq C$ and the bounded gain can never come out negative.

The control. Positions with support at most 12 are also solved exactly, in the same run and on the same state object, and the exact gain must lie inside its own bounds. It does, on 100/100. The control is computed in place rather than joined to the results of §11.3: that file keys its rows by game alone, several to a game and only for the positions it attempted, so any alignment would have been a guess, and a control built on a guessed join checks nothing.

The pre-registered answer is: not settled. Over 248 positions in 59 games, the ones not solved exactly are bounded into $[+0.2976, +1.7455]$. The study's solved mean over the same games is $+0.3214$, which the interval straddles. The bounds are real and they are too loose to decide; `prereg/ii_bounds.md` fixed that reading before the run.

Perfect information is a bad bound for this game. That is the first thing the run establishes, and it is worth stating because the PI relaxation is the standard tool reached for here. Where the truth is known it overshoots by $+0.3278$, against the one-ply policy's $+0.0585$ undershoot — five times looser, on the same positions. On 36 of 248 positions *every* deal exceeded its budget and fell back to the trivial figure, so those are bounded by arithmetic rather than by any search. Above support 150, 82% of all deals fall back.

Exploiting the champion at $m = 2$ is almost always one move. On 88 of the 100 positions solved exactly, playing the single best deviation and then reverting to champion play attains the exact optimum. Over all hundred the gap averages $+0.0585$ (95% CI $[+0.0197, +0.0974]$). The exploitation is a local error and not a sustained plan, which is what makes the cheap bound as tight as it is, and is a fact about the champion worth having on its own.

The same bounds at $m = 1$. Correcting one layer and leaving the other resting on the slope would leave a reader unable to tell whether $m = 1$ was checked and held or simply not checked, so it was checked. Over 200 games the layer has 644 positions, of which 340 are pinned — the belief admits one deal, the gain is zero by construction, and they are counted separately rather than averaged in. Of the 304 with a real belief, 276 are solved exactly here, at a mean gain of $+0.4116$ against the $+0.4405$ of §11 over its 305; the difference is which positions each run reached, not a disagreement. The control holds on 616/616, which with the $m = 2$ layer's 100/100 is 716 positions whose exact value lies inside its own bounds. One deviation attains the optimum on 543 of the 616 — 88%, the same fraction as at $m = 2$, on six times the sample.

Only 13 of the 304 sit above the support cap, so at $m = 1$ there is very little to extrapolate across in the first place: 95% coverage means the question the previous section worries about is nearly moot here. It is the long solvable range, not the unsolved tail, that makes this layer useful — it is what allows the next paragraph to be measured at all.

The cheap instrument loosens as the belief widens, which is why it cannot settle this. Because the one-ply gain is computable everywhere and the exact gain is not, it looked like the only *matched* instrument spanning the support cap. Whether it is depends entirely on whether its looseness is constant in support, and at $m = 2$ that could not be checked: exact values exist only up to support 12 there, over which the gap shows no trend ($t = -0.04$). Running the same bounds on the $m = 1$ layer — cheaper, and solvable exactly out to support 24 — gives the range to check it, and the answer is that the gap **grows**:

support	one-ply gain	<i>n</i>
1–2	+0.1111	54
3–8	+0.2618	63
9–24	+0.3933	52
25–60	+0.3408	33
61–150	+0.1155	19
151 and up	+0.1899	27

Table 17: One-ply gain against belief support at $m = 1$, split at the study’s 24. Reported for completeness and *not* as a fact about Fish: §11.7 shows the cheap instrument’s looseness grows with support ($t = +6.27$), so the fall-off in the last two rows is the instrument losing its grip rather than the game getting easier.

$m = 1$ support	exact minus one-ply	<i>n</i>
2–4	+0.0785	209
5–8	+0.2496	31
9–16	+0.3920	29
17–24	+0.1667	7

Slope +0.0236 per deal, $t = +6.27$ over 276 positions. So a one-ply figure that flattens or falls above the cap is exactly what a loosening instrument produces, whether the true gain is flat or rising. The $m = 2$ null is not evidence that the trend stops; it is evidence that this instrument stops tracking. Reported here as a measurement of the instrument.

For completeness, what the comparison showed before that was known, and which nothing above licenses reading as a fact about Fish. Splitting at the study’s 24:

At or below the cap it averages +0.2541 ($n = 169$); above it, +0.2350 ($n = 79$); Welch $t = -0.42$. Taken at face value that says the rise stops. Taken with the calibration above, it says the instrument stops, and the two readings cannot be told apart from this data.

The comparison was **exploratory** from the start, and the pre-registration records it as such: it was added after the upper bound turned out not to bind, and its bands were chosen after seeing the support distribution. It is left in because the negative result about the instrument is worth having, and because a comparison that was run and then quietly dropped once it stopped supporting a conclusion is the thing this project is trying not to do. The pre-registered verdict is unchanged and stands alone: not settled.

11.5 Where the disagreements are, and one that is provable

Comparing the champion’s move to the exact optimum’s at each hidden decision, 78 of 263 agree (10 more exceeded the 300,000-node budget and are unsolved, not counted either way). Of the 185 disagreements, 78 cost *exactly* zero — ties broken differently, not errors, and the rest cost +0.5609 on average, or +0.2005 spread over every hidden decision. Their shape does not implicate one component: of 93 costly ask-versus-ask disagreements, 47 share the target only, 11 the card only, and 35 neither. No repair is proposed from that.

Thirteen could not be priced at all, because the declared split was in no candidate the solver enumerates, and that set contains every assignment true in at least one consistent deal. Those claims were true in *no* deal the record allows. Checked directly, and the weak version of the check refuted the reading it was built to confirm: testing each card’s own holder mask gives 0/1080. A declaration can satisfy every card’s mask separately and match no complete deal, because the hand *counts* do not add up — the failure `claim4` names as “per-card modes can be jointly impossible”.

	claims	jointly impossible
$m = 1$	60	7 (11.7%)
$m = 2 \dots 9$	480	0

Only at $m = 1$, and structurally: there the hand counts are tight enough that a mis-split is detectably infeasible, while above it there is slack enough that almost any team-only declaration fits some deal. Every such declaration was wrong by construction: $+0.000$ sets each against $+0.968$ for the rest, under the void-era rule these games were played with — under the award baseline each would have scored -1 , a fact the adoption verdict below now carries.

11.6 The repair, and why it is not adopted

A max-flow feasibility test over the live cards decides the question exactly at any layer, and agrees 40/40 with brute-force enumeration where enumeration is available. Supplying the most likely *feasible* declaration instead — filtering alone changes nothing, because `forced_claim` rebuilds a declaration from the masks when no candidate survives — gives, over 6000 pre-registered pairs:

$$+0.028, \quad 95\% \text{ CI } [+0.024, +0.033].$$

Every block positive, every block excluding zero, and 27 wins against 973 ties and 0 losses in a representative block: the arm alters play about 0.09 times a game, so most deals are bit-identical and repairing a claim that could not have won cannot cost anything.

The bar fixed before the arm existed was $+0.05$. **It is not adopted**, and the champion is unchanged. A provable error, a free repair, and an effect smaller than the smallest one this project agreed in advance to care about — all three are the finding. One caveat now attaches: the $+0.028$ was earned under the void rule, where a repaired impossible declaration saved one point of differential; under the award baseline it saves two, so the arm’s value plausibly doubles and clears the bar. Re-running this 6000-pair design under the award rule is the one adoption decision the rule correction can plausibly reverse, and it is queued rather than presumed.

11.7 The all-champion profile is not an equilibrium

Computed with the corrected solver, and therefore not covered by the withdrawal in Section 11.

Section 11’s figures average over *positions*, and a game can contain several $m = 1$ decisions. That is the right unit for asking how well the champion plays the endgame and the wrong one for asking how exploitable it is: summing gains across positions in one game double-counts a continuation the deviator only gets to play once. So this takes exactly one position per game (the first $m = 1$ decision with genuinely hidden cards), and lets the deviator play champion-identically before it and optimally from it.

The value is already a team differential: $+1$ for a half-suit the deviator’s team takes and -1 for one the opponents take, so a half-suit changing hands moves it by 2, exactly as it moves the set differential `scripts4/exploitability.py` reports. Over 60 games:

	per game	95% CI
sampled rollout best response, whole game	-0.4875	$[-1.2442, +0.2692]$
exact, one endgame decision only	$+0.1111$	$[+0.0460, +0.1761]$

The interval excludes zero, so a single seat deviating at a single endgame decision takes a positive amount, and if the all-champion profile were an equilibrium, no single seat could gain by deviating at all. The rollout responder, deviating for a *whole* game, measured -0.4875 : it does not merely fail to find an exploit, it loses ground, which is what `fish4/bestresponse.py` warns strategy fusion can do to a PIMC responder.

Why the per-game figure is the smaller one, and why it is the right one. Of the 60 games, only 22 contain a first $m = 1$ decision with anything hidden. In the other 36 the belief has *pinned every card* by the time the endgame arrives, not a wide support, not a missing position: all 36 are pinned, so the closed form already answers them and the restricted deviator gains exactly zero. Conditional on a hidden decision existing the gain is $+0.2928$; per game it is $+0.1111$. Quoting the conditional figure beside a whole-game bound would compare a mean over a favourable subset against a mean over everything. The

2 games whose first hidden decision exceeded the node budget are excluded from the denominator rather than scored zero: unsolved is not worth nothing.

That split is the better description of the champion. Its inference finishes the endgame outright in 60% of games; where it does not, it leaves about a third of a set on the table.

What this does not bound. The opponents are a *deterministic realisation* of the champion, seeded from a hash of the observation, and best-responding to a realisation one can predict is easier than beating the mixture it came from. The champion mixes at most of these decisions — over 41 hidden $m = 1$ information sets at 12 seeds each it repeats its move every time in only 32% of them. So this bounds the realisation, not the champion as it plays, and the claim is correspondingly narrow: *that* profile is not an equilibrium.

11.8 Playing the endgame from the solution

Section 11.5 measures that the champion picks the exact optimum at under a third of hidden $m = 1$ decisions. `fish4/endgame_ii.py` puts the solver in the agent at exactly those decisions. Duplicate-deal paired against the identical configuration with the flag off, so everything but the policy cancels:

seed block	set differential	95% CI
0 (screen)	+0.0875	[+0.0307, +0.1443]
1	+0.0650	[+0.0033, +0.1267]
2	+0.0750	[+0.0150, +0.1350]
pooled, 1200 pairs	+0.0765	[+0.0422, +0.1108]

Cochran’s $Q = 0.280$ on 2 degrees of freedom: the three blocks are as consistent as independent estimates of one quantity get. That matters here more than the interval does. This project has already had a cell clear the same bar at +0.490, return +0.068 on one fresh block and −0.066 on a second, and turn out to be noise; the replications are what separate this from that.

The firing rate is part of the result. The first version of this measurement returned $+0.000 \pm 0.000$ over 40 paired deals in a single-seat arm, which reads exactly like a clean null and was nothing of the kind — the policy had fired *twice*. Raising the caps from 12 deals and 50,000 nodes to the study’s 24 and 300,000 changed that count from 2 to 2. A seat faces about 0.075 solvable hidden $m = 1$ decisions per game, not because the search is too small but because the champion’s belief has usually pinned the endgame before it arrives — the same fact Section 11.7 reports from the other side, where 36 of 60 games have no hidden first endgame decision at all. Every arm now reports how many decisions the flagged seats faced and how many the solver answered, and a run where it never fires refuses to print a differential.

One mechanism, visible in the raw counts. Over the 1200 void-era games the policy produced 446 misdeclarations against the control’s 473 — a misdeclaration being a declaration with the right team and the wrong split, voided in these games, awarded to the opponents under the current baseline. The solver enumerates only declarations true in some consistent deal, and each avoided misdeclaration is worth twice as much under the award rule as it was here (design R4 of the rule-correction pre-registration re-prices this stack).

Cross-play: the gain is not exploitation of one opponent. The obvious objection is that `fish4/exact_ii` best-responds to the v0.4 champion, so a gain against that champion might be exploitation of it rather than better endgame play. Against V03_CHAMPION — a different agent, which the solver’s opponent model is simply *wrong* about — over 1200 paired deals in three disjoint seed blocks, pooled the same way:

$$+0.0637 \quad 95\% \text{ CI } [+0.0306, +0.0968], \quad Q = 0.729 \text{ on } 2 \text{ df}$$

against +0.0765 for the champion it was built to answer. The first cross-play block alone cleared zero only barely, at +0.0750, [+0.0100, +0.1400], and was used here to weaken a caveat before it had met

the bar the champion arm was held to; it has met it now. The same gain against an opponent it was not modelling is what one would expect of a policy that plays the endgame better, and not what one would expect of an exploit. The null count moves the same way here too: 407 against 443.

The *all* arm cannot settle this and is reported for completeness only: with every seat flagged both teams are identical, so its differential is about zero by construction, and -0.050 ± 0.198 over 40 pairs is what that looks like.

Not promoted, and why. The flag is off, and after the cross-play result only one of the two original reasons still stands.

The first was that a best response is not an equilibrium strategy, so this might measure only that the policy beats *this* champion. Section 11.7 remains a standing reminder that strength against a named opponent and unexploitability come apart, but the cross-play row is direct evidence against that reading here, and the objection is much weaker than when it was written.

The second stands unchanged: turning the flag on would change the champion, and the champion is what `fish4/exact_ii` models. Every exact figure in Sections 11 and 11.7 is a statement about the current champion, and promoting this would invalidate all of them at once. That is a decision with a cascade behind it, to be taken deliberately and re-measured afterwards, not a consequence of a good table.

11.9 The solver as a critic: a defect found, fixed and shipped

Every improvement in this project until now came from a duel. This one came from a proof, and the duel only confirmed it.

The exact solver was built to measure exploitability. It also does something a duel cannot: it says *which move* was wrong, in a position, with a number attached. §11.4 established that beating the champion in the endgame is almost always *one* move — on 88 of 100 exactly solved $m = 2$ positions and 543 of 616 at $m = 1$, a single deviation followed by ordinary champion play attains the exact optimum. A defect that is one move wide is a defect that can be looked at.

The move is an ask, and it is riskier than the champion’s. The improving move is an ask rather than a claim on 131 of 154 improvable $m = 1$ positions and 128 of 138 at $m = 2$, so the defect is in ask selection. Paired within each position — same belief, same hand, so nothing about the position can explain the difference:

	champion’s ask	better ask	paired t	n
$m = 1$	$p = 0.799$	$p = 0.218$	-18.88	122
$m = 2$	$p = 0.611$	$p = 0.219$	-10.02	91

The better ask is riskier on 113 of 122 at $m = 1$ and safer on *none*. On 73 of those 122 the champion asked a card it was **certain** to get, and that was still a mistake.

It is certainty, not probability. The obvious fix (weigh the success probability less) does not work, and the data says so before any duel. The objective is $\text{score}(a) = p(a) + F(a) \cdot w$, and since the argmax is scale-invariant, $p + F \cdot (kw)$ ranks asks exactly as $(1/k)p + F \cdot w$: the existing objective already spans every positive weight on p . Searching that one-parameter family over 388 endgame positions carrying the exact one-ply value of every candidate ask returns $k = 1$. No change.

What helps is the information term, which rewards asks the belief cannot call. Held out by *game*, so positions from one deal cannot sit on both sides:

model	train	held out
champion (0 parameters)	+0.0805	+0.1498
scale k (1 parameter)	+0.0805	+0.1498
info = +2.00	+0.1236	+0.1591
certain = -0.50	+0.1212	+0.1784
full search (11 parameters)	+0.1786	[+0.1281, +0.1884]

The eleven-parameter fit gains most on train and its held-out score is an *interval*, not a number: run at four search seeds with the budget fixed, it lands anywhere from below the champion to above every one-parameter model. That is a property of which random restarts came up, not of the model, and one run of it would have reported whichever number that happened to be. The one-parameter models have no such freedom (their grid is exhaustive), which is a reason to prefer them beyond their score. `info = +2.0` was chosen on train and confirmed by leave-one-game-out cross-validation inside the training games; certain scored better on the held-out set and was *not* chosen, because picking it would be picking by the check.

Three duels, because an effect measured against a weak base does not transfer by assumption. The correction applies only when at most two half-suits are live. Every run was pre-registered with its arm, size, seeds and readings fixed first, and on deals the fit never saw:

	diff (sets)	95% CI	pairs
vs champion, registered	+0.0835	[+0.0338, +0.1332]	2000
vs champion, replication	+0.0980	[+0.0480, +0.1480]	2000
vs champion, pooled	+0.0907	[+0.0555, +0.1260]	4000
on top of what ships	+0.1220	[+0.0711, +0.1729]	2000

The last row is the one that moved a default and the one to quote: the site plays a depth-3 belief-space lookahead, which searches exactly the positions this correction is about, and it might already have been finding these asks. It was not — the effect is if anything *larger* on the stronger base, 15 of 16 blocks positive.

What was predicted and was wrong. The pre-registration said the effect would be small and the test might be underpowered to see it: the correction moves a handful of decisions a game, and offline it was worth +0.0093 in half-suit units. It is worth +0.1220 sets in play, about what running the entire exact solver in the endgame gains (+0.0765). The prediction that the endgame is too thin a slice of the game to matter was wrong, and it is recorded here because it was written down first.

It wins more; it is not harder to exploit. Winning duels and being less exploitable are different claims, and the solver can be asked the second one directly: it takes the opponents’ policy as an argument, so the same positions can be solved twice. Over 285 $m = 1$ positions from champion self-play, each solved against the champion and against the champion plus the correction — same position, same belief, same hand, so nothing about the position can explain a difference:

paired change	estimate	95% CI
the policy’s own value	+0.0189	[−0.0132, +0.0509]
what a best response takes	−0.0002	[−0.0264, +0.0260]
the gap between them	−0.0190	[−0.0587, +0.0206]

Neither clears zero, and the two nulls are not the same null. The own-value change is simply underpowered — a half-width of 0.032 around +0.019, with the sign the in-play result predicts. The best-response change is *pinned*: −0.0002 with a half-width of 0.026, an interval that excludes any reduction the size of the +0.1220 the correction gains in play. A best response still takes the same amount.

So the correction does not buy its winnings by closing the hole. That is a limit on the method rather than on this particular fix, and it is worth stating because the method is the interesting part of this section: an exact best responder used as a *critic* improves average-case play, and nothing about fixing the moves it flags makes the worst case harder to reach. Exploitability and strength are being measured, and they moved independently.




Figure 7: The deception ladder, both measurements of the same knob. Top: each pre-registered rung’s gain over its sibling incumbent, growing with the threshold m — eight of eight blocks positive at every rung. Bottom: the same configurations against the foreign v0.3 champion on 500 shared deals; every rung above the shipped $m \leq 2$ is a significant loss, and always-on loses to v0.3 outright. The panels share an x-axis, so each column is one configuration read against two opponents: the rungs were not measuring strength, they were measuring how much of the game was spent poisoning the sibling’s opponent model.

Pushing the correction up the game measured deception, not strength. What happened when the correction was extended past the endgame is reported in full because it is the sharpest methodological result in this paper, and because half of it is a finding about the game.

The correction applies below a threshold m of live half-suits. A pre-registered ladder raised the threshold one rung at a time, each rung duelling the incumbent, each shipping only on a CI clear of zero. Every rung cleared, and the rungs *grew*: $+0.1220$ to $m \leq 2$, $+0.3025$ to $m \leq 3$, $+0.6995$ to $m \leq 4$, $+1.0125$ to $m \leq 5$, and $+4.2790$ for the amended coarse step to always-on — a bigger single gain than the whole v0.3-to-v0.4 release. Eight of eight blocks positive at every rung, no timeouts, engine digests matched.

All of it was an artefact of *who was on the other side of the table*. Each rung duelled the sibling configuration, and the sibling’s opponent model conditions on which half-suit a player chooses to ask. Information-seeking asks violate that model systematically, so they poison the sibling’s inference; the rungs grew with m because a higher threshold spends more of the game poisoning. Against the v0.3 champion (a foreign opponent whose inference is built differently) with every configuration playing the same 500 deals and per-deal differencing against no correction:

threshold	margin over v0.3	paired Δ vs. no correction
none	+3.006	—
$m \leq 2$	+2.914	−0.092 [−0.448, +0.264]
$m \leq 3$	+2.568	−0.438 [−0.829, −0.047]
$m \leq 5$	+1.526	−1.480 [−1.888, −1.072]
always on	−2.160	−5.166 [−5.643, −4.689]

Monotone, and the mirror image of the sibling ladder (Figure 7): every rung above 2 is a significant *loss* against a foreign opponent, ending at a configuration that loses to v0.3 outright. The always-on rung

was shipped for minutes on the strength of $+4.28$ and reverted when the first cross-play block came back; every rung above 2 is withdrawn, and the shipped threshold is back at $m \leq 2$ — whose foreign-relative effect is a paired tie, and which alone rests on evidence that is not a sibling duel.

Three readings, stated separately. *About the game*: exploiting an opponent model that conditions on action choices is worth several sets per game in six-player Literature, which makes deception a first-class strategic resource here, found by accident and measured to ± 0.24 . *About the method*: a pre-registered ladder with a fixed stopping rule, eight-of-eight block agreement, matched engine digests and growing effects passed every check this paper uses, and measured the wrong quantity at every rung, because no rule asked *against whom* the gain was defined. Robustness of the opponent pool now belongs beside pre-registration in the discipline, not downstream of it. *About the correction*: the $+0.1220$ of the preceding paragraphs is a gain against the engine’s own family, bounded against v0.3 inside $[-0.448, +0.264]$; the exact solver’s diagnosis stands, and what its repair bought is play that is better against the family and no worse outside it.

The same question against an engine nobody here wrote. Everything above still measures the correction against relatives: the sibling ladder, and a v0.3 that shares this project’s rules engine and most of its ancestry. What the withdrawal actually demands is an opponent with *no* shared code, and one exists: **FishBot v0.7** by dylann4500, an independently written C++ engine for the same six-player, 54-card game — different author, different belief representation, different search, its 55-parameter policy frozen upstream. Its published account (Nguyen, 2026b) gives that policy’s fit objective as minimax regret over an opponent panel, not self-play; an earlier draft of this paragraph described it as self-play tuned, which was wrong. It plays under this engine’s arbiter through a process bridge that replays the public history to the frozen policy and legality-checks every move it proposes, substitutions counted rather than hidden; zero were needed across the 2,000 games this section reports. Two caveats travel with the comparison: v0.7 was tuned for a rule variant with out-of-turn declarations and plays here under this paper’s rules, which its policy is told; and it is one foreign opponent, not a population. A third caveat — that the first run scored misdeclarations under this engine’s void-era rule while v0.7’s native game awards them — was removed by re-running the whole design under the award baseline, below.

The check was pre-registered before the first scored game (prereg/foreign_opponent_m2.md), outcomes fixed in advance: the deployed configuration with the correction on, and the same configuration with it off, each played three copies of v0.7 on 250 deals in both rotations, 500 games per arm, the statistic being the on-minus-off difference of set margins on the same deal and rotation. The registered question was reversal, and the answer is no: $-0.0300 [-0.1081, +0.0481]$ against Dylan’s v0.7, at a power sized for deception-scale flips. The sharper reading is what the interval *excludes*: $+0.1220$, the correction’s own sibling-measured size. On a third, unrelated opponent the effect is a tie, so the sibling number does not transfer at anything like face value — the section’s finding again, demonstrated this time on an engine from outside the project entirely. An earlier draft of this sentence attributed to the upstream author a remark that its advantage “does not transfer”; that quotation is withdrawn. It is not in Nguyen (2026b), whose §12.1 argues the opposite — that the advantage is *not* an artefact of self-play coordination, one seat’s upgrade being worth $+1.93$ of a $+4.49$ percentage-point three-seat gain, and cross-play between three independently trained runs moving the result by 0.01 points. Our result here stands on our own measurement and needs no support from theirs.

The context the check ran in: with the correction on, the deployed configuration beats v0.7 by $+2.448$ sets per game ($+2.478$ with it off), taking 64.3% of decided sets across the 1,000 games. That margin is a statement about this game under this paper’s rules against an engine tuned under slightly different ones, and it is quoted as context, not as a ranking of the two projects; the measurement the design paid for is the paired difference, which needs no ranking to be valid.

The same check, rule-matched. When the misdeclaration baseline flipped (§3), the design was re-registered and re-run from scratch under the award rule (design R2 of prereg/rules_award_baseline.md: fresh seeds, its own journal, the rule pinned in the runner rather than inherited), which makes the two engines agree on what every wrong declaration costs — v0.7’s native risk model is now exactly the rule in

force. Over another 500 paired games the answer does not move: -0.0440 $[-0.1344, +0.0464]$ against Dylan’s v0.7 under the award rule — no reversal, and the sibling-measured size is again excluded. The margin context: $+2.544$ sets per game with the correction on, $+2.588$ with it off, 64.3% of decided sets over the 1,000 games — the share to the decimal of the void-era run, on disjoint deals, under a different scoring rule.

What the rule-matched run adds is the accounting of the rule itself: 466 sets in those 1,000 games were misdeclarations — wholly-held half-suits declared with the wrong split, each now a set for the other side, and they divide 160 against this engine to **306** against v0.7. The correction the void rule had been silently granting was worth nearly twice as much to the opponent, so scoring misdeclarations properly *widened* nothing — it is the reason the margin held rather than shrank when the game switched to the rule v0.7 was built for. This paper first attributed that gap to v0.7’s forced declarations, which come from a sampled best guess where this engine declares the exact posterior MAP (§6.2). The next paragraph withdraws that attribution: the split is unchanged by a repair that fixes the forced declarations outright, so it belongs to their voluntary declaration gate instead.

A defect in the bridge, found from the other side, that had been flattering this engine. Every figure above was taken through this project’s bridge into v0.7, and an adversarial read of their C++ — their source, not their paper — found a defect in that bridge rather than in either engine. Forced to declare, their own driver picks the first live half-suit the mover *holds a card in* (engine/src/game.hpp:535); this bridge passed the first live half-suit, full stop. Their bestGuess was therefore periodically asked to name all six owners of a half-suit it held nothing of — a question with no anchor in its own hand, which their own driver never puts them in a position to be asked, and which under the award rule hands the set here.

The defect ran in this engine’s favour, which is the direction a project does not catch by reading its own results and nodding. Priced paired — all 10,000 deals replayed under both bridge revisions, identical deals and identical seating, the same duplicate-deal design every screen in this paper uses (scripts4/bridge_bug_price.py) — it is worth -0.0784 $[-0.0881, -0.0687]$ sets per game: real, and about 3% of the margin. It is also an order of magnitude smaller than the 0.30 sets per game the same read estimated from 40 games, which is what the paired instrument is for. Under the repair their declaration accuracy rises from 76.18% to 78.87%, while this engine’s moves from 96.65% to 96.49% — the control that confirms the repair touches only their side of the bridge.

What the repair does *not* move is the misdeclaration split, and the reason corrects the mechanism the previous paragraph had assumed. The anchorless forced declarations were *ownership*-class errors (half-suits the opposing team still held a card of), whereas the 160-to-306 split counts *allocation*-class errors, wholly-held half-suits declared with the wrong split. Across the same 10,000 paired deals the allocation split is flat: 2,784 against v0.7 before the repair, 2,775 after. The finding that this engine misdeclares roughly half as often survives intact; the explanation for it does not, and the gap is a property of their voluntary declaration gate rather than of a question this bridge was mis-posing. The pre-registered R2 contrast is likewise undisturbed, because the defect is common-mode across its two arms — both the correction-on and correction-off games ran through the same bridge, and a paired difference cannot see an error that appears identically in both halves of the pair. What is restated is the *context* the R2 file reports beside its verdict: the margin and set share, which are absolute rather than paired, and which the corrected head-to-head below supersedes.

The head-to-head, re-measured through the repaired bridge.

WITHDRAWN. The $+2.3466$ below is retracted. §11.10 shows it to be an artifact of this bridge’s statelessness rather than a measurement of the two engines: their agent was reset and its belief rebuilt before every decision it made, which cost them more than the whole margin. Corrected, KRAKEN *loses* this match by about 0.6 to 0.7 sets a game, on two independent routes that agree with each other and not with this table. The section is kept in full, and in place, because the number was published and because how it was found is the useful part; every figure in it should be read as describing their engine under that handicap.

This was the figure the paper stood behind, and the only one taken through bridge revision 2: 10,000 games on the same deals the retracted run used, seats rotated so neither engine opens more often, both engines on their native misdeclaration rule, zero substituted moves and zero unfinished games (scripts4/mega_match.py → results/mega_match.json). Matching the withdrawn run deal for deal was deliberate: it made the replacement a like-for-like correction rather than a differently-powered second opinion. It is worth noticing that every one of those precautions held, and none of them could see this: a defect common to both arms of a paired design is invisible to the design, and an absolute is not paired against anything.

margin (withdrawn, §11.10)	+2.3466 [+2.2928, +2.4004] sets/game
sets	56,733 to 33,267 (63.04% of decided)
games won	80.41% ± 0.78 (8,041W / 0T / 1,959L)
ask hit rate	51.75% against 48.14%
declaration accuracy	96.49% against 78.87%

Every row of that table is an artifact of the bridge, and the paragraph this paper used to print underneath it read the artifact as a property of their engine. It said that the two engines find cards at comparable rates and diverge in knowing when a half-suit is *finished*, and that the 18-point declaration gap was the mechanism of the margin. What the same pairing gives once the bridge sends the hand their reset actually means (§11.10, bridge revision 3):

	retracted (rev 2)	corrected (rev 3)
corrected margin	+2.3466	−0.5250 [−0.6886, −0.3614]
	10,000 games	1,200 games
hit rate, ours / theirs	51.75% / 48.14%	52.98% / 54.93%
declaring right, ours / theirs	96.49% / 78.87%	97.59% / 96.68%

Both readings invert. The 18-point declaration gap is 0.9 points: they declare very nearly as well as we do, and the run that said otherwise was scoring an opponent whose posterior we had corrupted into declaring half-suits its own opponents were holding. And the ask comparison does not merely shrink, it *changes sign* — we were never 3.6 points ahead on asking, we are 1.95 points behind it. The one row this paper treated as the minor term is the only channel in which the two engines actually differ, and it is the channel we lose. (Two registered arms have since closed that gap by more than half and lost by a third of a set each doing it, §11.11; a deficit in the metric is not a deficit to close.)

Nothing in §6.2’s inference-side argument depended on this table, and it is not withdrawn with it; what is withdrawn is the claim that the scoreboard independently corroborated it against this opponent.

How much of the margin is that, exactly. The paragraph above is a qualitative reading of two rows, and it had never been priced. It can be: every declaration in a game is either right or one of two kinds of wrong, and under the award rule each wrong one moves the margin by two sets, so the scoreboard decomposes exactly. Over the same 10,000 cross-engine games the headline margin is measured on (scripts4/margin_decomposition.py):

component	rate	contribution to margin
measured margin	—	+2.3466 [+2.2928, +2.4004]
their ownership-class errors	0.5667/game	+1.1334
their allocation-class errors	0.2775/game	+0.5550
our own wrong declarations	0.1759/game	−0.3518
everything else	—	+1.0100 [+0.9573, +1.0627]

This decomposition is withdrawn in full. It says that fifty-seven per cent of the margin is declaration accounting, and it is arithmetically exact — but it decomposes a margin that does not exist, and every term

in it is a count of wrong declarations this project’s own bridge caused. Their 0.844 wrong declarations a game are 0.09 at rev 3. It is kept here, struck, because a decomposition that was correct arithmetic on corrupted inputs is a better warning than a deleted section: the check that would have caught it is not a tighter interval on any of these rows but a second host for the same pairing. The prose below is the original, verbatim.

Fifty-seven per cent of the margin is declaration accounting. The edge is real and it is nearly fivefold (we declare wrongly 0.176 times a game against their 0.844), but it is a declaration edge, and this project has spent the large majority of its compute on the ask objective. The residual, everything the asking and the inference and the lookahead between them produce once the declarations are netted out, is +1.010 sets per game: a real margin, and a little over two fifths of the headline one.

These figures are computed on the same 10,000 games the headline margin is, which took a correction. They were first taken from a 600-game probe and quoted beside a margin from the large block, which invites a reader to apply a ratio measured on one population to a number measured on another. The two agree — the probe put their error rate at 0.918 a game against 0.844 here, and the share at 61% against 57%, but the version on the headline’s own games is the one to quote.

This is a decomposition of where the sets landed, not a counterfactual. A declaration that did not happen changes every ply after it, and no arithmetic on a finished game recovers those plies; the residual row is what remains after the netting, not the margin a differently-declaring v0.7 would have conceded. The distinction is worth insisting on because the temptation is to read the last row as “our true strength”, and it is not that.

The same account, across his whole ladder. Everything above is one opponent. That was never a limit of the bridge: the shim reads a policy spec and hands it to their factory, which dispatches on the base prefix, so all six of their releases were one string away and only v0.7 had ever been played. Six rungs at 600 duplicate deals each, 7,200 games, zero substituted moves, zero unfinished games, and the declaration identity closing with residual zero on every game (scripts4/dylan_ladder_sweep.py → results/dylan_ladder_sweep.json):

opponent	margin, sets a game	their error rate	their wrong/game
v0.2	+1.5283 [+1.3735, +1.6832]	9.13%	0.365
v0.3	+0.6133 [+0.4597, +0.7670]	4.71%	0.199
v0.4	+1.9867 [+1.8371, +2.1362]	16.36%	0.598
v0.5	+2.1233 [+1.9663, +2.2803]	15.06%	0.574
v0.6	+1.8850 [+1.7333, +2.0367]	15.72%	0.624
v0.7	+2.3600 [+2.2017, +2.5183]	20.71%	0.828

The v0.7 row is an unplanned replication of the headline. It is a different bank, a different harness and a fortieth of the games, and its interval covers the withdrawn +2.3466 above.

Across six opponents the margin tracks their declaration error rate at $r = 0.95$. Declarations a game are similar throughout, 3.66 to 4.22 for them against 4.79 to 5.34 for us, so this is a comparison of rates and not an artifact of how often anyone declares. The single-opponent decomposition above says 57% of the v0.7 margin is declaration accounting; the ladder says the same quantity is what orders the whole pool.

And their engine became easier for ours as it advanced. Their v0.3 is the hardest rung we face and their v0.7 the easiest, a gap of 1.75 sets. Through the identity that gap is 0.629 of their extra wrong declarations, 0.215 of our extra declarations, and 0.028 of our own accuracy: roughly three quarters of it is that v0.7 misdeclares four times as often as v0.3.

That is not a claim that their project regressed, and the distinction matters. Their v0.2 and v0.3 are scripted baselines rather than fitted policies, and their own published head-to-head has v0.6 beating both v0.4 and v0.5 in *their* arena. A policy tuned against its own pool can buy declaration volume that wins there and costs against a different opponent, which is a statement about what each project optimised against rather than about which engine is better.

Three things that make the sweep worth trusting, and one that does not resolve. The instrument check was fixed before the run. Their own head-to-head puts v0.4, v0.5 and v0.6 within ± 0.046 to ± 0.146 sets of each other across five replicates a pairing, so a bridge mistranslating one version would surface as a spurious strength gap on that rung. The measured spread across those three is 0.238 sets against per-rung half-widths near 0.15, which is what their flatness plus our sampling noise predicts.

Their cheat harness is excluded, by base name rather than by keyword. Their tree exposes a base whose three policies their own header calls deliberately-cheating: one inverts its seed to the deal seed and plays clairvoyantly, one shares state across the three team seats, one decodes a teammate's arbitrary tie-break label. The base with no cheat option selected falls through to an honest agent, so a guard matching only the word would admit it and would admit the rest the day the option was renamed. Three further bases are excluded as instruments rather than releases: one carries an online model of its opponent's policy, one requires the opponent's source to invert its transcript, one is an identity control.

The provenance was weaker than we would like, and has since been checked rather than stated. Their repository was available to us only as a single squashed commit, so the parameter vectors compiled into v0.4, v0.5 and v0.6 could not be checked against the ones their published results were generated from, and earlier versions of this paragraph said so. Their full history is now published, which turns the caveat into a measurement.

The instrument is their own reproduction protocol rather than ours: build *their* engine at each release commit and at the commit we measure against, run the identical match on both, and compare every field of the result. Across five rungs at three seeds and 240 games a cell, win rate, mean sets, ask accuracy, declaration counts and events per game agree *exactly* at every cell; the play differences number zero. What moved is their bootstrap interval, in the last place at 5 of 15 cells, and nine JSON fields their older builds did not emit — reporting rather than play. The rungs we measure are the rungs they released.

Two limits survive. A differing field would have localised a change to the policy and the arbiter jointly, since both live in one binary; the pass is clean, so that distinction never had to be drawn. And this verifies neither their published *numbers* for a rung nor the deal banks that produced them — their manifests cite commits absent even from the full history and record a dirty working tree, and nothing on our side can reach those.

What does not resolve is a hole in the table above rather than beside it. *Our* declaration error rate is 8.38% against their v0.4 where every other rung puts it between 2.05% and 4.19%, six standard deviations outside the rest, with our own engine identical in all six cells. It is not a volume effect: our error rate correlates with our declaration count at $+0.71$ across the six, and dropping v0.4 collapses that to $+0.10$, so the correlation is the outlier rather than a trend it sits on. The obvious suspect was the one option their manifest attaches to that rung, and a 250-deal probe kills it — with and without the option the two arms are identical to four decimals, for two independent reasons in their source, and so the anomaly stands with no hypothesis behind it.

What the other 39% is made of. The residual — everything asking, inference and search produce once the declarations net out — has a shape that follows from the rules. Every set is won by declaring a half-suit your team holds, so once misdeclarations are accounted for what remains is the race to *assemble* half-suits, and that race has exactly two factors which multiply: how many asks a team gets to make, and how many of them land. This paper has reported the second since v0.4 and has never reported the first. Over 800 cross-engine games (`scripts4/acquisition.py`):

per game	ours	v0.7	difference
turns	57.250	52.562	+4.688 [+3.962, +5.413]
asks	51.894	48.265	+3.629 [+3.102, +4.155]
asks that landed	26.731	23.066	+3.665 [+3.129, +4.201]
half-suits assembled	5.008	3.416	+1.591 [+1.418, +1.765]
plies holding a set unspoken	78.166	32.440	+45.726 [+41.5, +49.9]

Splitting the 3.665-card gap at the mean assigns $+1.802$ to volume and $+1.863$ to conversion, and an earlier draft of this paragraph read that as two mechanisms of equal size. It is one mechanism

counted twice. A landed ask retains the turn, so asks per possession is $1/(1 - p)$, and the two hit rates alone, 0.5151 against 0.4779, predict an ask-count ratio of $2.062/1.915 = 1.077$ against an observed $51.894/48.265 = 1.075$. The volume term is not a second contribution; it is what the conversion term implies. Our extra 4.7 turns a game are a consequence of asking better, not an independent source of cards, and the two may not be added.

That the product splits evenly at the mean is a property of the arithmetic and not a finding: for any two factors whose ratio is close to one, the symmetric split assigns each about half. The honest statement is the conversion edge: we land 51.5% of asks against 47.8%, and every other row of the table above follows from it.

The last row is the one that costs us. We hold assembled half-suits for 78 plies a game against their 32, and our own declaration accuracy falls monotonically with how long we have held the set:

plies held assembled	declarations	wrong	error rate
0 (never owned it)	6	6	1.000
1–9	2797	46	0.016
10–29	517	27	0.052
30–79	454	23	0.051
80+	238	28	0.118

Holding is not the cause and the direction of the arrow matters: a set you can place, you declare, so “held long” and “could not place” are nearly the same event. What the table establishes is that being unable to place a set is *visible early*. And by §6.2’s channel argument, early is when it can still be acted on — once the half-suit is ours no opponent can ask in it, so nothing they do will place it, and the only remaining channel costs one of our turns.

There is time. Over the 130 half-suits we declared wrongly, the median number of *our own* turns between assembling the set and the deadline is 25.5, and only 16 of 130 were assembled with no turn left to spend. The lever is reachable in 88% of the cases that go wrong.

Whether spending those turns pays is a different question and this paper has already answered a version of it, though not in the negative. The void-era duel of §10.4’s signalling protocol returned +0.002 [−0.086, +0.090] sets per pair; the registered confirmation under the award rule returned +0.1180 [+0.0325, +0.2035] and +0.1220 [+0.0291, +0.2149], both clear of zero (§10.1). The protocol is switched *off* in the deployed configuration because it misses a ship bar fixed in advance, not because it does nothing. What the measurements above add is not a new mechanism but a different *trigger*: that protocol fires when the current turn is cheap, and the errors are concentrated in half-suits identifiable a median of 25 turns before the turn gets cheap. Whether a deadline-priced trigger changes the answer needs its own pre-registration, and the error-value result of §8.7.2 says what it will be measured against: one avoided misdeclaration is worth about 1.8 sets, and a spent turn has to cost less than that times the chance the signal actually places the split.

What their errors are made of, which is not what we guessed. The obvious hypothesis about the ownership-class leak (the 0.650 sets a game they hand over by declaring a half-suit we still hold a card of) was that our *silence* causes it: under the no-bluff rule an ask certifies a holding, so a half-suit we have never asked in is one their posterior has little reason to place us in. That is backwards. Over the same games, the 398 cards of ours sitting inside their 390 ownership errors split as

publicly pinned to us by an ask we had won	0
never publicly moved since the deal	398

Their engine does not misplace a card it has seen move; it only ever loses one with no public trace. And the leak concentrates where we have spoken, not where we have been quiet. Weighting by exposure (plies during which our team held exactly one card of a live half-suit) gives 45.4 ownership errors per 1,000 plies in half-suits we had asked in against 0.675 where we had not, a rate ratio of 67 [42, 107].

Sorted by what had happened in the half-suit before they declared it, 1,927 declarations followed both an ask of ours and a later successful ask that took one of that half-suit off us, and 19.25% of those were ownership-wrong; 411 followed only the taking, and 3.16% were.

That mechanism was ours, and this paragraph had it backwards. What this section originally said is quoted in the next sentence, and it is wrong. “*Our ask certifies that we hold at least one other card of the half-suit. When an opponent later takes one card of it off us, their inference appears to treat that certification as discharged, as though ‘at least one’ had been ‘exactly one’, and every card it then loses is one that never moved in public.*” It was offered as a hypothesis about their engine formed from outside it rather than a reading of their source, and that caution was the right instinct pointed at the wrong party: [Nguyen \(2026c\)](#) read the source, and the certification is not discharged by anything their inference does. It is discharged by *our bridge*, which reset their engine with the hand the seat held at that moment instead of the hand it was dealt (§11.10). A card the seat had since lost was filed as never its own, so when the ask at ply k was replayed that card was excluded from the certificate’s candidate set, and a certificate with one surviving candidate pins a card the asker never held.

Every number in this subsection is therefore a measurement of our own defect, and each is a signature of it rather than a coincidence: the cards are ones that never moved in public because those are the ones a phantom deduction lands on; the rate ratio concentrates in half-suits we had asked in because our ask is what creates the certificate; and the 19.25% against 3.16% split is the defect’s own precondition, an ask of ours followed by a capture of ours, against the capture alone. This project measured the fingerprint of its own bug in four independent ways and read all four as evidence about somebody else’s inference. The subsection is kept, in full, because the forensics were sound and only the attribution was not, and because the moment at which a careful error-analysis of an opponent should have become a suspicion of the instrument is exactly here.

We are not building a feature on this. An exploit of one engine’s inference is exactly the shape of the withdrawn deception ladder of §11.9, and a term that farmed it would grow with its dose against v0.7 and mean nothing anywhere else. It is reported because a margin whose majority component has an identified cause is a different object from one whose does not, and readers of the head-to-head are entitled to know which they have.

What this measurement is, stated with the other project’s own scope in view. Their paper is unusually careful about what it does not claim, and the sentence that matters here is theirs: “Whether FishBot v0.7 is the strongest Canadian Fish agent anyone has built . . . the paper does not establish and does not claim. The comparison class contains only agents this project wrote.” Their limitations section names the missing ingredient exactly — “there are no independently authored six-player Fish agents to compare against, which is why §1.2 confines the paper’s comparative claim to this project’s own lineage”, and lists an independently authored opponent first among the things a comparative claim beyond that lineage would require ([Nguyen, 2026a](#)).

WITHDRAWN. The +2.347 in the paragraph below is the retracted figure of §11.9, and the bridge it credits with one known defect found, priced and repaired had a second, the statelessness that §11.10 prices. Corrected, KRAKEN *loses* this match: -0.6960 through a persistent bridge, -0.6200 inside their arbiter, and -0.5250 [-0.6886 , -0.3614] through the repaired bridge at BRIDGE_REV 3. The paragraph is kept because the scope argument in it survives the change of sign; the number does not.

That ingredient now exists in both directions, and it is worth being precise about what each side may therefore say. We can report a margin against an engine nobody here wrote, which is a stronger kind of evidence than either project’s internal ladders: $+2.347$ [$+2.293$, $+2.400$] sets per game (the figure since withdrawn), 63.0% of decided sets, over 10,000 games under rules both engines implement natively, with zero substituted moves and a bridge whose one known defect had, we then believed, been found, priced and repaired. What that does *not* establish is a global ranking of the two projects, and this paper declines that claim for the same reasons theirs does: one opponent is not a population, their engine was frozen before ours had the exhibition bridge to test against, and neither project has run the other’s configuration search against its own. What the number does establish is narrower and, for both papers,

new: the first six-player Literature result measured across an engine boundary rather than inside one lineage, which is precisely the instrument their limitations section asks for, and the one this paper spent §11.9 demonstrating that a project cannot build for itself.

Under the award rule the correction itself no longer earns its place, and is withdrawn. Designs R3 and R4 of the same pre-registration close the loop. R3, the declaration-threshold sanity: against 0.999 the shipped 0.97 bar produces *bit-identical* games — 0 of 200 pairs diverge, +0.0000 exactly, and against 0.90 the difference is +0.0400 [−0.0278, +0.1078]; the void-era finding that any high bar plays the same survives the rule change, as its rule-invariant mechanism said it should. R4, the correction stack re-priced: the registered expectation was that the sign would hold and the size could move. The expectation was wrong, and it is reported because it was written first: under the award rule the deployed configuration WITH the $m \leq 2$ correction scores −0.1040 [−0.2184, +0.0104] against the same configuration without it, over 500 fresh pairs, and the correction arm misdeclares *more* (123 against 106, again wholly-held splits; both arms are counted the same way, so the comparison is what it looks like). The mechanism is not mysterious: the correction’s $d_{\text{info}} = 2.0$ was fitted against solver targets whose continuations priced a wrong declaration at zero, and the game no longer plays that price. Taken with the foreign tie above (−0.0440 under the matched rule), the correction fails, under the baseline rule, the CI-excludes-zero bar every shipped knob in this project is held to. It is withdrawn from the deployed configuration (V06_DEPLOYED in the registry is V04_COMBINED without it), and the exact-solver diagnosis behind it, which is rule-invariant, awaits a refit against award-rule targets: the solver journals now carry the rule in their resume fingerprints so that refit cannot silently mix eras. The void-era +0.1220 stands as what it always was, a measurement of the void-era game.

Their one load-bearing mechanism, ported and measured here: it does not transfer. The other project’s attribution study puts a single coordinate’s leave-one-out contribution at +1.69 [+0.80, +2.57] of a +4.73 percentage-point whole, and states in terms that the stronger claim, that the gain *lives* in that coordinate, is not supported at its power (Nguyen, 2026b); an earlier draft of this sentence credited the coordinate with essentially the whole gain, which overstates what they claim. The mechanism is a *contestation* term that scores a candidate ask by $(1 - p_{\text{hit}}) \times (\text{opponent-held mass of the half-suit}/6) \times (\text{cards of it nobody can place}/6)$ and carries a strongly positive weight, deliberately spending likely-miss asks in opponent-dominated, still-ambiguous half-suits (Nguyen, 2026a). Because their engine and ours now play the same game under the same rules, that claim is directly testable outside the lineage it was measured in, which is the kind of question neither project could ask alone.

Design R6 of the pre-registration ports it and sweeps the dose in both directions, with two arms of a second idea alongside — a *silence prior* that down-weights sampled worlds in which a live half-suit already sits wholly within one team, on the ground that a team holding all six and able to place them would usually have declared. Seven arms plus an unmodified baseline each played 500 games against three copies of v0.7 on identical deals and rotations, paired against the baseline: 4,000 games, zero substituted moves.

Every arm is negative; four are significantly so, and the largest positive dose is the worst cell in the table. No arm clears the registered bar, so by the outcome fixed in advance neither knob ships, stages 2 and 3 do not run, and the negative is reported rather than swept into a footnote. The strongest reading the data supports is monotone in the wrong direction: the more of their mechanism we add, the more of our margin over them we lose.

This is the third time this paper has watched a gain fail to survive a change of opponent family, and the first time in the direction that costs *us* something to admit. The deception ladder’s rungs were real against siblings and vanished against v0.3; the endgame correction was real under one scoring rule and reversed under another; and their contestation term is, by their own careful measurement, worth several points inside their engine and is worth nothing but damage inside ours. The mechanism is not thereby refuted (it plainly works where it was fitted), and the honest conclusion is narrower and more useful than either “it works” or “it does not”: an ask-scoring coordinate tuned against one belief representation is not a portable strategic insight, even between two engines playing an identical game. Their belief is a

arm	paired Δ vs. baseline	margin over v0.7
baseline (neither knob)	—	+2.732
contestation -1.0	-0.4200 $[-0.7137, -0.1263]$	+2.312
contestation -0.3	-0.2400 $[-0.5085, +0.0285]$	+2.492
contestation $+0.3$	-0.3200 $[-0.5790, -0.0610]$	+2.412
contestation $+1.0$	-0.2920 $[-0.6261, +0.0421]$	+2.440
contestation $+3.0$	-1.1240 $[-1.4735, -0.7745]$	+1.608
silence $\delta = 0.7$	-0.1960 $[-0.4775, +0.0855]$	+2.536
silence $\delta = 0.9$	-0.2880 $[-0.4883, -0.0877]$	+2.444

Table 18: Seven arms against three copies of v0.7 on identical deals and rotations, each paired against an unmodified baseline: 4,000 games, zero substituted moves. Every arm loses, and the two largest weights lose most — a dose-response, which refutes the mechanism rather than merely failing to find it.

Sinkhorn approximation seeded by a behavioural prior; ours is exact counting reweighted by a choice model. A term that pays in the first is apparently already priced in the second.

What did not change. FishBot4’s own defaults. The champion is the reference point for every measurement in this paper, and moving it would silently reprice all of them at once; the correction remains in the archival V04_COMBINED spec exactly as measured (its withdrawal from the *deployed* configuration under the award rule is the previous paragraph’s, and the registry’s, record). A test requires whole games to be *bit-identical* when the knob is off, and requires the first divergence to fall at two live half-suits or fewer when it is on — checked by replaying the shared prefix, not by asking the agent where it was.

11.10 The same pairing in their arbiter, where the sign reverses

Appendix C argues that a bridge is asymmetric: whoever owns the game loop owns the rules, the RNG and every tie-break, so if only one side can host, “our engine is stronger” and “our arbiter favours us” are not separable claims. That argument has been in this paper since the bridge was built, and until now it was only an argument. Both directions existed, but only one of them ever measured anything: every cross-engine number above was taken with our engine hosting, because there was no way to seat our policy in a batch run of theirs.

There is now. That project’s engine has acquired an external-bot protocol, so our champion can be packaged, installed into their engine, and played from their command line as an ordinary seat. It passes their own conformance check. Under their arbiter, their deals, their seat rotation and their dialect, over 1,200 games against the same frozen configuration:

	<i>our arbiter</i>	<i>their arbiter</i>
games	10,000	1,200
our margin (sets/game; the left figure is the withdrawn one)	+2.3466	−0.6200
our games won	80.41%	40.08% [37.33, 42.83]
our declarations/game	5.005	4.027
their declarations/game	3.995	4.832
ours declared right	96.49%	99.13%
theirs declared right	78.87%	97.91%

The interval on the second column excludes 50%, so the reversal is not sampling noise at this size. **Neither column is a ranking of the two engines, and the honest reading of the pair is that this paper cannot supply one.**

Where the reversal lives, and the explanation that does not survive. It is the same component the margin decomposition already identified: their declaration accuracy, which moves by 19.0 points between the two columns while ours moves by 2.6 — and a declaration *count* that reverses with it. In our arbiter

the nine half-suits split 5.005 to 3.995 our way; in theirs, 4.027 to 4.832 theirs. Whoever declares a half-suit is, apart from the errors, the side that scores it, so this row is very nearly the scoreboard restated, and it is the row the reversal is made of. Since 57% of the withdrawn +2.3466 is declaration accounting (§11.9), a component that large moving that far is the whole reversal.

The standing explanation for their high error rate here was the dialect. Their policy routes about 3.4 declarations a game through an out-of-turn channel our rules do not have and our bridge therefore never polls, and a policy leaning on a channel it is denied should play worse. That is testable inside their own engine, which takes a flag for it, with nothing of ours involved — their policy, their arbiter, their deals, their opponent, one flag changed. It returns the opposite:

their arbiter	out-of-turn/game	their wrong declarations/game
their dialect	3.28	0.1250
--no-out-of-turn	0.00	0.0750
their --legacy bundle	0.00	0.0729
<i>our arbiter, our dialect</i>	0.00	0.8442

Taking the channel away makes their declarations *more* accurate, not less; their own dialect sweep agrees in direction, pricing no-out-of-turn at +0.52 points of their edge. So the dialect is not what separates 0.075 from 0.844, and the caveat this paper has attached to the head-to-head since the bridge was built has the sign of this component backwards. Nor is it the opponent: facing our champion in their own arbiter their error rate is 0.101 a game, *lower* than the 0.125 they make against their own v0.6, so meeting this engine does not by itself degrade their declarations.

The control that had to run first, and the shape it found. Before any of that can mean something about strength, one alternative has to be excluded: that the package simply does not carry our policy, and we lose there because we arrive weakened. A bot crippled in transit loses to everything. So our champion was run against their *whole* released ladder in their arbiter, and the shape compared against the same ladder measured in ours.

It is not crippled — it beats their v0.3 by +1.2972 inside their own arbiter. But the two ladders do not merely differ in level. **They run in opposite directions, and so does the quantity underneath them.**

their version	<i>our arbiter</i>		<i>their arbiter</i>	
	our margin	their decl. error	our margin	their decl. error
v0.2	+1.5283	9.13%	—	—
v0.3	+0.6133	4.71%	+1.2972	13.65%
v0.4	+1.9867	16.36%	—	—
v0.5	+2.1233	15.06%	−0.3528	2.00%
v0.6	+1.8850	15.72%	−0.4806	1.53%
v0.7	+2.3600	20.71%	−0.6200	2.09%
<i>Spearman rank correlation against their version number</i>				
our margin		+0.771		−1.000
their decl. error		+0.771		−0.400

In *their* arbiter their ladder behaves as a version ladder should: each release beats us by more than the last, perfectly monotone over the four rungs measured. In *ours* it is inverted — their later releases do *worse* against us — and their declaration error rate climbs monotonically with their own version number, from 4.71% at v0.3 to 20.71% at v0.7.

The individual rungs make the inversion sharper than the correlations do. Their v0.3 is their *best* declarer in our arbiter and their *worst* in their own; their v0.7 is their worst in ours and among their best in theirs. The two arbiters do not disagree about the level of a quantity here. They disagree about its order.

A project’s successive releases getting monotonically worse at the one thing this game scores, and only when played through our bridge, is not a shape that improving strength produces. The

natural place to look is what their v0.4 introduced and every later version built on — their fitted belief — and how it fares on the event stream our bridge replays into it.

What that leaves, and why it is reported as open. Our arbiter, our bridge, or an interaction of the two. This paper is in a poor position to wave that through: §11.9 is the record of a bridge defect destroying a set of results, and the head-to-head has already been re-measured once for another.

This paragraph is superseded and is kept only to date the reasoning. It read: “57% of +2.3466 is declaration accounting; their declarations are four times worse against us at v0.7 than at v0.3 while their own arbiter says the opposite; and we cannot currently say why. We are not withdrawing the number.” The number is withdrawn (§11.10), and “we cannot currently say why” was the correct state of knowledge at the time and the wrong place to stop.

One clause of it needs more than superseding, because it is a claim about method rather than a number. “The paired contrasts that share the bridge are unaffected either way” is the defence this paper offered for every foreign-opponent result it kept, and Nguyen (2026c) is right that it does not hold as stated. A shared component cancels from a paired contrast only when its effect is *additive*; a corrupted posterior in the opponent is not, because it changes which positions the two arms even reach. Their report gives the clean demonstration from this paper’s own numbers: the README’s dialect table has off-turn declaration costing us −0.300 sets against the corrupted engine, while against the native one in their arbiter the same rule is worth +0.111 to us. A gap measured against a corrupted opponent is not the gap against the real one, and it can invert.

What survives is narrower and worth stating exactly. Contrasts in which *both* arms are ours — self-play, duplicate-deal pairs against our own champion, the ablations of §10.1 run against `ev_claim` — never had the bridge in the loop and are untouched. Contrasts run against `dylan_v07` are not rescued by their pairing and are withdrawn with the absolutes: R2, R6, the dialect gaps, and the contestation port of Table 18.

One caveat on the reverse cells survives the control and is worth keeping. Our champion reaches their arbiter through a package that answers off-turn declaration polls only on certainties — our policy has no out-of-turn behaviour to offer, having never been developed against the channel — so our side in the right-hand columns is *a* realisation of our champion rather than the one in the left. That is why the argument above rests on the *trend across their versions*, which holds our side fixed within each arbiter, rather than on comparing levels between them.

The measurement is also in tension with the mechanism proposed earlier in this section, that our ask certifies “at least one other card of this half-suit” and a later take against us discharges it. That is a property of our *asking* rather than of our arbiter, so it predicts the same errors in their arbiter, and they are not there — and it offers no reason why the effect should grow with *their* version number, which is the fact that now most needs explaining.

The mechanism, and it was in our own docstring. The experiment that separates these is a decision-level comparison of their engine against itself, and it identifies the cause. *Our shim is stateless*. It spawns a fresh process for every decision and replays the entire public log into a freshly reset agent — the right shape for a website that reconstructs each session from its log, and the shape every cross-engine number in this paper was measured through. Their arbiter instead builds one agent per seat per deal and feeds it events as they happen. Those two are the same thing only if their agent is a pure function of (reset state, event sequence).

It is not, and the boundary is exactly where the two ladders stop agreeing:

their version	decisions	divergent	rate	DECL→ASK	ASK→DECL
v0.2	323	0	0.00%	0	0
v0.3	338	0	0.00%	0	0
v0.4	415	81	19.52%	5	1
v0.5	336	137	40.77%	4	0
v0.6	292	36	12.33%	2	1
v0.7	296	84	28.38%	9	0

Identical on both scripted baselines; divergent on every version from v0.4, which is the version their fitted belief arrives at.

Which piece of their state, and whether it could be repaired. That boundary names a release, not a mechanism, and the two are not the same question. The mechanism matters because some kinds of lost state are cheap to restore from outside their engine and some are not: an RNG our fresh process re-seeds could be handed a per-decision seed derived from the event count, and the streams would line up. Their factory gates the two RNG consumers behind spec options, so the parity measurement can be re-run at the corners (`scripts4/statefulness_mechanism.py`). Divergences are split into the band where both paths ask but ask differently and the band where one declares and the other asks, because the ablations move one and not the other:

arm	overall	ask ordering	declaration gate
the frozen release	23.52%	20.63%	2.64%
their search off	16.61%	13.53%	2.71%
their tie-breaking off	27.96%	24.59%	3.25%
both off	15.78%	11.60%	4.18%
and w12=0	13.92%	9.74%	4.18%

Their determinized search is the largest identified term and it lives *entirely* in the ask-ordering band — `V06Agent::resetV6` re-seeds `srng` from the agent seed, so a fresh process per decision draws the same determinizations every turn where their arbiter's stream advances. Their tie-breaking contributes nothing, and their own source says why: at `rtie=1` the tie RNG is *constructed* per decision from a rolling hash of the public stream and the event count, so it is replayable by construction, and only `rtie=2` would not be. The third arm neutralises feature 12 of their scored vector, which asks whether this is the half-suit they asked for last time: it is assigned inside their `chooseAsk` and never by their `observe`, so through a stateless bridge it is dead at every decision despite carrying weight 3.12582 in the frozen vector. It is real and small.

Three further candidates are excluded by reading their source rather than by spending an arm on an inert code path. Their belief consumes no randomness at all in the released configuration, which runs `BeliefMode::Fast`; and `Belief::sinkhornDisj` re-initialises its marginals from the constraint set on entry, so it does not warm-start. *An earlier version of this section named a warm-started Sinkhorn/IPF fit as one of two candidates. That was wrong, and is withdrawn.* Their dead-ask memory, which their reset does wipe, is gated behind `deadAsk`, `deadInSearch` and v0.7's `dead7`, all off in the frozen spec.

The finding is the band that does not move. Everything identified above lives in ask ordering. The declaration gate — the band that actually costs sets, since a declaration taken a turn early is a wrong one — is untouched by every mechanism found, and 13.92% of decisions still differ with all of them switched off. So the convenient answer is not available: this bridge cannot be made honest by re-seeding, because the term that matters is not an RNG. Running their engine as their arbiter runs it is the remedy, and the reason the corrected figures in this section were measured that way rather than by patching the instrument that produced the retracted one.

The *direction* is what closes the loop. Across the four affected versions the stateless path declares where the persistent path asks twenty times, and asks where it declares twice. More declarations, taken on a belief rebuilt from scratch, is precisely the shape that produces a 20.71% declaration error rate here against 2.09% at home — and it produces it only for the versions that have a belief to rebuild, which is why their ladder inverts through our bridge and not through theirs.

What it was worth, in sets. A different move is not automatically a worse one, so the divergence above is a fact about decisions and not yet about the scoreboard. The paired contrast prices it: 2,000 pairings — 1,000 deals at each of two seat assignments — with every pairing played twice on identical cards, identical seats and identical agent seeds, our champion unchanged in both arms, the *only* difference being whether their engine keeps its state between decisions (`scripts4/bridge_statefulness_price.py`).

our margin, stateless bridge	+2.4800
our margin, persistent bridge	−0.6960
the bridge was worth, to us	+3.1760 [+3.0142, +3.3378]
their wrong declarations/game, stateless	0.8800
their wrong declarations/game, persistent	0.0925

Zero fallbacks and zero unfinished games in both arms. The stateless arm reproduces the published figure on fresh deals (+2.4800 against the withdrawn +2.3466), which is what makes the other column believable rather than a harness artifact.

The bridge was worth more than the entire published margin. Their declaration errors fall by a factor of ten when their engine is allowed to keep the state their own arbiter lets it keep, and the head-to-head reverses.

Two independent routes now agree, and they agree against us. The persistent-bridge figure was measured in *our* arbiter under *our* dialect over a Python transport. The reverse-arbiter figure was measured in *their* arbiter under *their* dialect through their own bot-package protocol. They share no code path, no rules, no deal generator and no host:

	our margin	games
our arbiter, <i>stateless</i> bridge (published, withdrawn)	+2.3466	10,000
our arbiter, <i>persistent</i> bridge	−0.6960	2,000
their arbiter, our bot package	−0.6200	1,200

The counts are that route’s own games, which for the persistent bridge is one per pairing rather than the paired design’s total across both its arms — an earlier version of this table gave that total, twice over, and so credited the corrected figure with twice the games behind it.

The two routes that do not reset their engine every decision agree to within 0.08 sets a game. The one that does is the outlier, by three. They are also much the smallest two: the retracted figure has the largest sample here and is the one that is wrong, which is the point rather than an embarrassment to it.

This is a retraction of §11.9’s headline. The +2.3466 is not a measurement of the two engines. It is a measurement of their engine *with its belief destroyed and rebuilt before every decision it makes*, which is a policy neither project has ever intended to run and which nobody could reproduce outside this bridge. On the evidence above the honest statement of the cross-engine result is:

Under this paper’s rule dialect, KRAKEN v1.1 loses to SESTINA v1.0 by roughly 0.6–0.7 sets per game.

What survives is narrower and still worth having. The paired contrasts that share the bridge are untouched, because both of their arms were handicapped identically — which is exactly the distinction Appendix C drew, in a sentence written long before there was any reason to think it applied to this paper’s own headline. And the negative results of §9 and §9.3, which are internal contrasts against our own champion, never went through this bridge at all.

11.11 Where the ask deficit lives, and why the ask objective cannot reach it

The corrected table leaves exactly one channel in which the two engines differ, and it is the channel we lose: they hit 54.93% of their asks and we hit 52.98%. After the first registration of the v1.2 programme (P43, Appendix A) put three candidates at that 1.95 points and moved none of them — the worst, the opponent’s own fitted exponent applied to its seats, cost −0.90 against them and −2.00 in self-play; removing the model entirely cost −0.59 and −1.07 — the obvious next step was to find a fourth. The programme ran to three registrations and ten arms, P43–P45, and this section reports the last two. P43’s screen, on the dual-population design every later arm used (300 deals on both parities per population; candidate minus champion, sets a game, vs SESTINA through BRIDGE_REV 3 and in self-play against the v1.1 champion):

arm	vs SESTINA	self-play
opponent_gamma = 0.0 (model off, their seats and ours)	−0.5867 [−0.891, −0.283]	−1.0667 [−1.286, −0.847]
opponent_gamma = $\hat{\alpha}$ (their fitted exponent, their seats)	−0.9033 [−1.207, −0.600]	−2.0033 [−2.212, −1.794]
avoid_doomed_asks	−0.0933 [−0.196, +0.010]	−0.0900 [−0.317, +0.137]
depth_mode = at_ask	−0.0033 [−0.228, +0.221]	−0.0867 [−0.313, +0.140]

No arm clears either half of the bar. Two lessons outrank the nulls. The exponent the second arm applied is R0’s re-measurement of §8.5’s fit through the repaired bridge, $-1.1055 [-1.2531, -0.9578]$ over 6,281 asks, so the opposite-sign finding survives its instrument and acting on it is still the worst arm in the study: a correct measurement of an opponent does not license a change to us. And removing the model entirely costs -0.59 and -1.07 , so the mis-signed model is an asset — though the first arm switched it off on our two teammates as well as on their three seats, because `gamma_team` defaults to `opponent_gamma`, and which side of the table that cost belongs to is what the fourth registration (P46, registered and running as this is written, §13.2) separates. The `depth_mode` row is also the defect that taught this paper to read a zero-width interval as the signature of an experiment that never ran: on its first run the value fell through silently to the default, the arm was the champion, and `fish4/agent4.py` now raises on an unrecognised value. What follows is the attempt at a fourth arm, and it ends by withdrawing the candidate rather than registering it.

None of the three instruments here is pre-registered, and none licenses an arm. They are reported because the second contradicts the first, the third contradicts the second, the reason it does is structural rather than statistical, and the conclusion is that a change this project was about to register would have pushed on the one part of the channel where it is already ahead of the engine that beats it.

11.11.1 A decomposition that says execution, and is blind where it matters

The gap has two possible shapes. We might choose the same asks and land them less often (*execution*), or choose differently (*selection*). The Oaxaca–Blinder identity separates them exactly, given a stratifier s :

$$r_t - r_o = \underbrace{\sum_s f_o(s) [r_t(s) - r_o(s)]}_{\text{execution}} + \underbrace{\sum_s [f_t(s) - f_o(s)] r_t(s)}_{\text{selection}}.$$

Over 38,212 asks in 400 games (`scripts4/ask_deficit_anatomy.py`), our hit rate is 53.06% against their 55.31%, a gap of 2.25 points, and six stratifiers agree:

stratifier	execution	selection
phase (half-suits resolved)	+0.0228	−0.0002
options available	+0.0208	+0.0017
half-suits to choose from	+0.0238	−0.0012
our depth in the named half-suit	+0.0273	−0.0047
our team’s cards in the named half-suit	+0.0077	+0.0149
target’s hand size	+0.0221	+0.0004

Five of the six put essentially the whole gap in execution, with no unmatched mass in any of them. The reading at the time was that we choose the same half-suits and are worse inside them.

The sixth row is the one that was right, and the instrument is the reason it took two more runs to see it. Conditioning on *our own team’s* holding in the named half-suit moves two thirds of the gap into selection — and that stratifier’s top bucket, where the team holds all six, has a hit rate of **exactly zero on both sides**, because an ask into such a half-suit cannot land. A stratum where $r_t(s) = r_o(s) = 0$ contributes exactly zero to *both* terms of the identity, however far apart the two sides’ shares of it are. Ours is 12.61% against their 9.52%, and the decomposition prices that difference at nothing directly: it can only see it in the strata the displaced mass came from. **The decomposition is structurally blind to the one stratum the deficit lives in**, which is not a bug in the arithmetic but a property of using the opponent’s rate as the reference price.

11.11.2 A miss anatomy that says selection

The arbiter sees every hand, so each miss can be classified by what was legally available instead: another opponent held the named card (wrong target); the named opponent held another card of that half-suit we could legally have asked for (wrong card); both; some other legal pair in that half-suit would have hit (wrong pair); or nothing in that half-suit was legally askable from anyone (*unavoidable*). Over 28,513 asks in 300 games (`scripts4/ask_miss_anatomy.py`), 30.64% of our misses are unavoidable against 21.94% of theirs. Per ask rather than per miss that is 14.5% against 10.1%: a 4.4-point gap in a channel whose whole size is 1.3 points in that block.

So the deficit is not in which card or which opponent we name once the half-suit is fixed. It is in asking into half-suits that are already exhausted from our side. That is a hypothesis about the ask objective, and it is worth nothing until one prior question is settled.

11.11.3 The kill-check: did the objective have the information?

A half-suit is dead for us when every legal ask in it must miss. That is a joint event over card locations, so a per-card marginal cannot express it; it is evaluated by drawing 128 worlds from the same `BeliefState` the agent samples from and taking the share in which every legal ask in that half-suit misses. The true deal is read only to score the estimate. Over 1,200 games and 98,326 ask decisions (`scripts4/ask_deadness_signal.py`):

The belief knows. $p(\text{dead})$ separates dead half-suits from live ones at $\text{AUC} = 0.9518$, and a proxy that costs no draws at all — the share of the half-suit already accounted for on our own side — reaches 0.8883. On our 6,519 asks into a truly dead half-suit, some alternative scored strictly lower 96.63% of the time, and switching to the belief’s favourite lands in a live half-suit 93.76% of the time against 84.75% for a uniformly random *other* half-suit on the same menus (+9.00% [+8.33, +9.71]).

That is the answer the programme was built to find, and it is not the finding, because the same measurement run at SESTINA’s seats says the engine that beats us does not use the signal either — and uses it *less*. Their $p(\text{dead})$ is not their belief, which we cannot query; it is what an agent with their information could have known, computed by the machinery we judge ourselves by. In ground truth, per ask decision:

	chosen was dead	<i>menu</i>	<i>selection</i>
KRAKEN v1.1	0.1327	0.1428	− 0.0101 [−0.0139, −0.0063]
SESTINA v1.0	0.0945	0.0746	+ 0.0198 [+0.0171, +0.0226]
ours − theirs	+0.0382	+0.0682	−0.0299

The *menu* column is the dead rate of a half-suit drawn at random from the ones that seat could legally ask into; *selection* is what the objective does relative to that draw, paired on the same decisions. Our dead-ask rate is 3.82 points worse than theirs. Our menu is 6.82 points worse and our selection from it is 2.99 points **better**: we steer away from dead half-suits, they steer into them, and they win. **A $p(\text{dead})$ term in the ask objective would push harder on the one channel where we are already ahead of the engine we are behind.**

On the belief scale the two sides look the other way round (+0.0166 ours against +0.0268 theirs, both positive), and that comparison is the confounded one: $p(\text{dead})$ is an estimate and the estimator is not equally sharp at the two sets of seats (AUC 0.9518 at ours, 0.8819 at theirs), so a cross-side reading on that scale prices the estimator as well as the objective. Within a side it is paired and sound; across sides only ground truth settles it, and adding that column is the whole reason the run was repeated.

11.11.4 What a dead ask actually is

The menu term is not about asking at all, and the reason is structural. Under the no-bluff rule an ask is legal only for a card of a half-suit the asker holds another of, so every legal ask in a half-suit misses

exactly when no opponent holds any card of it — and the six cards of an unclaimed half-suit are all in somebody’s hand. Therefore:

A half-suit is dead for us if and only if our own team holds all six of it.

A dead ask is not a badly chosen ask into a contested half-suit. It is an ask into a half-suit our team has already won and has not yet declared. The equivalence is asserted at every decision by the instrument below rather than argued, so a rules variant that breaks it (bluff asks, a claim rule that leaves cards live) fails loudly instead of quietly changing what the numbers mean.

Completion is also *absorbing*: once a team holds all six, no opponent holds a card of that half-suit, so no opponent can legally ask into it, so no card can leave it. The state persists until the owning team declares. That makes the interval from completion to declaration a latency rather than a window that can close on its own, and makes the asks spent into it meanwhile a clean price for the delay. Over 1,200 games and 10,724 completions (`scripts4/completion_latency.py`):

	KRAKEN v1.1	SESTINA v1.0
half-suits completed per game	4.208	4.728
mean plies to declare it	16.80	7.84
median plies to declare it	3.0	1.0
never declared	0.00%	0.00%
asks spent into it meanwhile	1.520	0.996
those, as a share of all asks	13.79%	9.81%
completions sat on 10+ plies	34.63%	20.92%
share of dead asks from those	87.38%	75.99%

We sit on a half-suit we already own for more than twice as long, and it is a tail rather than a general slowness: a third of our completions account for seven eighths of our dead asks. Every completion is declared eventually on both sides, so this is latency and not abandonment.

Is the channel the right size? The excess is +1.689 dead asks a game. This project has a measured price for exactly one thing of that kind — §10.5’s value of a single donated turn, +0.2713 [+0.0922, +0.4505] sets over 1,500 paired deals — and propagating both uncertainties (a paired bootstrap on the per-game difference, the standard error on the price) puts the excess at +**0.458** [+0.157, +0.793] sets a game against a measured deficit of −0.5250 [−0.6886, −0.3614]. The deficit sits inside the interval.

That is an accounting of a difference and not a counterfactual. Removing a dead ask changes every ply after it; their 4.7 a game are not zero and nothing here says ours could be; and the turn price was measured with our own engine on both sides at a decision chosen without reference to what the agent wanted. It answers “is this channel big enough to be the story”, which is a smaller question than “is it the cause” — and the smaller question is the one these numbers can carry.

11.11.5 P44: the two ends of the latency, and a hit rate that is not worth having

The analysis above is exploratory and licenses nothing. What it licensed is a registration (P44, Appendix A), and the registration found a defect worth attacking: **every doomed-ask mechanism in this engine** — the stuck claim gate, the signalling ask, `avoid_doomed_asks` — **is conditioned on the agent being certain the best ask cannot land**. That is 1.5% of decisions. On the asks that actually are dead the belief’s own estimate averages 0.3377, so none of it opens. The apparatus is aimed at the right idea through a gate that almost never fires.

Two arms, each replacing that certainty test with a graded one, each bit-identical at its default: **D1** restricts the ask when the top-scoring one’s half-suit is ours with probability at least 0.5; **D2** declares a half-suit we are certain we own ($p_{\text{team}} \geq 0.99$) at exactness 0.77 rather than 0.97, that threshold being a break-even from this paper’s own turn price and declaration accuracy rather than the best of a sweep.

D2 never reached a duel, and that is the result. It fired on 0.235 declarations a game against a futility bar of 0.250 fixed before the run. It fires so rarely because it gates on knowing the set is ours, and *we usually do not know*: averaged over every dead half-suit on the menu, the belief puts only 0.3120 on it being entirely ours (and 0.3377 on the ones we go on to ask into). We often do not know we own them (the belief puts 0.3120 on it), and that is real; but supplying the knowledge does not open the gate. Relaxing the ownership test from 0.99 all the way to 0.20 at the same 0.77 exactness adds 0.103 declarations a game, flat below 0.77 and under half the futility bar (§11.11.7). The gate is waiting on the split. That is not what the 95.3% allocation-error figure invites a reader to assume — that figure counts declarations we made and got wrong, not the ones we never made — and it moved the target, briefly, from splitting to ownership inference. The measurement that followed moved it back: ownership is not the binding constraint of the gate, the split is.

D1 reached the duel and failed in both populations. 600 pairings, 1,800 games, zero fallbacks:

	vs SESTINA	self-play
D1 dead_ask_threshold = 0.5	−0.3233 [−0.527, −0.119]	−0.2300 [−0.448, −0.012]

Not a null — a significant negative in both populations. And it *worked*: the candidate’s ask hit rate is 0.5245 against the champion’s 0.5184 on the same deals. D1 did exactly what it was built to do, landed more asks, and cost a third of a set a game doing it.

So the ask hit rate is not a quantity to maximise. Under the no-bluff rule a failed ask publicly proves the asker holds another card of that half-suit, which is precisely the fact a partner needs to place a split. An ask into a half-suit our team wholly owns is not waste: it is this engine’s own mechanism for resolving the ownership it does not yet know it has. Removing it raised the metric and lowered the score.

That closes a reading this section’s own opening invites. The corrected head-to-head puts us 1.95 points behind on ask hit rate and calls it the only channel in which the two engines differ; P44 shows that closing that gap directly makes us weaker. **A measured deficit in a metric is not necessarily a deficit to close**, and the one instrument that could tell the difference was a duel with a bar fixed in advance — not any amount of further anatomy.

11.11.6 What this closes

The ask deficit is real, it is worth about half a set a game, and it is not an ask problem. It is the visible cost of a declaration latency: our team completes half-suits and holds them, and every ask made into one while it waits is certain to fail. That puts it downstream of the residue this paper has reported since v1.0 — 95.3% of our own wrong declarations are allocation errors, where the team held all six and named the wrong split (§8.7.2). The team has the answer and no member of it does, so it waits; and while it waits it asks into its own half-suits.

P44 shot down what the exploratory analysis pointed at, and P45 replicated it on an independent channel (§11.11.8): two arms, two registrations, two mechanisms, both raising the ask hit rate by one to one and a half points and both losing about a third of a set against SESTINA. **The ask objective was not the place.** A $p(\text{dead})$ penalty is the graded form of `avoid_doomed_asks`, it fires at 4.95% of decisions where the incumbent fires at 1.5%, and it is worse than the incumbent rather than better: −0.3233 and −0.2300 against −0.0933 and −0.0900. The channel did not need a bigger lever. It needed not to be pulled.

What survives is the other end, and it is narrower than what this section set out to find. The latency is real and expensive; it is caused by the split: relaxing the gate’s ownership test all the way to 0.20 adds only 0.103 declarations a game, so the positions D2 could not find are missing because the exactness band 0.77–0.97 is empty, not because ownership is unknown; and it is not fixable by asking differently, because the arm built on that premise made things worse while improving the metric it targeted. The ownership inference itself — the 0.3120 — was the natural next target, and the two subsections that follow report what happened when it was measured rather than argued about. Three estimators of it were

put at the same 178,354 half-suit cells against the dealt truth; the sharpest of them scores worst; all three cross 0.99 on the same 4.17% of owned half-suits; and the gate is flat in ownership from 0.77 to 0.20. Neither an estimator swap nor an ownership recalibration survives that, so the ownership inference is not an estimator problem and the gate is not an ownership problem. The one mechanism aimed at the split — the locating convention — was then nominated as P45 and lost in both populations.

11.11.7 The ownership inference, measured — and it was never the constraint

P44’s reading — that we sit on completed half-suits because we do not know we own them — sent the next step at the ownership inference, and the honest first move was to measure it rather than to register an arm against it. Under the no-bluff rule a half-suit admits no landable ask exactly when our own team holds all six, so “dead” and “entirely ours” are one event, and this repository estimates it three different ways: the independence product over per-card team marginals that the ask features compute and D1 gated on (`fish4/askfeat.py`); an unweighted Monte Carlo over the constraint-feasible set, which is what produced the published 0.3120; and the weighted joint `Posterior.prob_all_with`, which the engine has implemented and tested all along and nothing shipped reads, at 0.20 ms a decision. The instrument (`scripts4/ownership_estimators.py`, `results/ownership_estimators.json`) puts all three at the same decisions against the dealt truth: 1,000 games at BRIDGE_REV 3 on an independent seed block, 41,512 decisions, 178,354 live half-suit cells.

estimator	mean when genuinely ours	AUC	Brier
independence product	0.4035	0.9704	0.0525
uniform joint	0.3230	0.9549	0.0639
weighted joint	0.3523	0.9641	0.0592

The uniform joint reproduces the published 0.3120 at 0.3230 on an independent block, which is the check that makes the other two rows comparable rather than merely adjacent. Then the surprise: **the sharpest available estimator is the worst of the three at this job**. Paired and clustered by decision, the weighted joint against the product is **−0.0512** [−0.0517, −0.0506]. That is not a defect in `prob_all_with`. The six cards of a half-suit compete for the same quota slots, so their team-membership events are *negatively* correlated; the independence product is therefore an over-estimate of the true joint, and over-estimating is rewarded by any score that only asks how high the estimator goes on positives. The product is over-confident, and it beats the correct object because the metric pays for over-confidence.

None of it matters for the gate. All three cross 0.99 on **4.17%** of genuinely owned half-suits, so no declaration gate is rescued by reading a different number. And the gate was never starved by ownership in the first place. `claim4`’s voluntary gate fires on $p_{\text{team}} \geq \text{owned_p_team}$ and $p_{\text{exact}} \geq \text{owned_threshold}$; P44 set the exactness knob to 0.77, left ownership at 0.99, and measured 0.235 declarations a game against its 0.250 bar. Sweeping the *other* knob on the quantity the gate really reads, holding the registered 0.77 and excluding declarations the champion already makes:

owned_p_team	flagged	precision	new declarations/game
0.99 (P44’s)	14	1.0000	0.0140
0.95	19	1.0000	0.0190
0.90	37	0.9730	0.0370
0.85	60	0.9667	0.0600
0.77 down to 0.20	103	0.9223	0.1030

Flat below the registered exactness. Relaxing the ownership test all the way to “more likely than not it is *not* ours” buys 0.1030 new declarations a game — under half the futility bar P44 fixed in advance, and the bar was not widened. The binding constraint is not ownership. It is the exactness band: we are almost never 77–97% sure of an exact split. Either we are at 0.97 and already declaring, or we are far below 0.77. So P44’s recorded explanation is half right and points the wrong way: we do often fail to

know we own a completed half-suit, and that is real, but supplying the knowledge does not open the gate, because the split is what the gate is actually waiting on. Two candidate arms — swapping the estimator, recalibrating the ownership cut — were stopped for about twenty-five minutes of compute instead of two registrations and roughly 7,000 pairings.

11.11.8 P45: the free message, and a channel closed from both ends

What that measurement leaves standing is the split, and the engine’s one mechanism aimed at the split is the pre-play convention (row 3 of the closures in §12.3; P5–P10 in Appendix A). Its record is mostly bad, and one entry in it is a specification error: the aimed convention lost +1.750 to the champion (its encoder-only ablation +1.467 [+0.818, +2.116] over 120 pairs) through a gate that compared *probability of success* while the agent ranks asks by the objective’s scores, so a gate of “0.05” was routinely paying a median 0.3596 objective units out of a range of about 1.5. The gate now reads the same currency the pick was made from, and the one specification that document fixed and never executed is the **free-message gate**: `convention_max_cost = 10-9` in the re-priced units, so the agreed card is named only when it *ties* the card the objective already chose and the message costs literally nothing in the objective’s own terms. The decoder half was not re-run, because the free read — listen to the 35.3% of asks that name the agreed card by coincidence, send nothing — had already settled at 3,000 pairs, -0.002 [-0.127, +0.123]. Anything F1 won had to come from the sender.

P45 registered that one arm (`prereg/kraken_v12_free_message.md`), with the prior stated against it and a futility screen on *carry* rather than on decisions changed: a free-message gate swaps only on ties, and on a tie the champion picks uniformly from the tied pool, so a replay comparison counts RNG tie-breaking as gate firing — a ten-game probe read 16.45% by that route against a true carry gain of 0.7 points. The screen passed with a carry gain of +5.80% over the champion on the same deals (39.5% against 33.7%), reproducing within a point the 40.1% the earlier document had predicted for this gate on a harness it was not fitted to. The duel then lost, in both populations: 600 pairings, 1,800 games, zero fallbacks.

	vs SESTINA	self-play
F1 <code>convention_max_cost = 10⁻⁹</code>	-0.3067 [-0.607, -0.006]	-0.8567 [-1.069, -0.645]

Three things in that table were not predicted, and writing the prediction down first is what makes each visible.

Wrong about the size, not the verdict. The registration predicted a null in both populations, centred near the free read’s -0.002, on the reasoning that F1 is that same channel plus 4.8 points of carry. F1 is not a null; in self-play it is nearly seven times the free read’s point estimate. Sending changes the game in a way that receiving does not, and the free read’s null does not bound it.

The dual population fired the other way. P45 added the SESTINA population because every convention duel to date had been self-play only, and said in advance that a clearing arm should be distrusted in self-play, where our seats share a code book and the opponent does not. The opposite happened: **self-play is where the arm loses most**, by a factor of 2.8. The population added as a check on optimism was the optimistic one.

A free message is not free, and the objective is why. F1 costs nothing in the objective’s own units by construction and loses 0.86 sets a game in self-play, so whatever it pays, scores cannot see it. Offered as a hypothesis: the gate replaces a *random* tie-break with a *deterministic* one, and all three of our seats resolve ties the same way. The champion’s uniform draw from the tied pool is not a placeholder — it decorrelates three seats that share an objective — and the free-message gate removes it at every tie. That is the same shape of error as the mis-priced gate this arm was written to avoid: a cost genuinely zero in the quantity being measured and non-zero in the game.

And the ask hit rate rose again, and lost again. F1’s ask hit rate is **0.5287** against the champion’s 0.5136 on the same deals; D1’s was 0.5245 against 0.5184. Two registrations, two channels — one filters asks by ownership, the other renames the card on a tie — and both raise the hit rate by one to one and a half points and lose about a third of a set a game against SESTINA. P44’s “the ask hit rate is not a quantity to maximise” was a single arm and could have been the arm. It has now replicated on an independent channel, seed block and registration, and it is the most durable finding either programme has produced. Side by side (hit rate on the same deals; margin vs SESTINA in sets a game, 600 pairings each):

registration	arm	arm	champion	vs SESTINA
P44	D1 (dead_ask_threshold = 0.5)	0.5245	0.5184	−0.3233
P45	F1 (convention_max_cost = 10 ^{−9})	0.5287	0.5136	−0.3067

What this closes. The convention channel. With the free read’s null at 3,000 pairs, both halves are now measured in play: **receiving is worth nothing and sending costs, at every gate ever tried.** The 1.72 bits an ask carries remains a true fact about the rules and remains unrealisable by this engine. V06_DEPLOYED is byte-for-byte unchanged, and KRAKEN v1.1 still loses to SESTINA v1.0 by −0.5250. Ten arms across three registrations have now been put at that deficit, and the two channels a reader would reach for next — the ownership inference and the convention — are closed by measurement rather than by argument.

12 Discussion

12.1 What this says about playing Fish

The results have readings that do not require an engine, and we state them because a study of a game people play ought to say something to the people playing it.

Watch where opponents ask, not just what they ask for. This is the finding of the paper, and it is actionable at the table. A player who asks in a half-suit is probably deep in it, in rough proportion to how many of its cards they hold, and that single inference is worth more than every other improvement we made combined. Concretely: when an opponent asks in hearts, do not only note the hard fact that they hold at least one heart — shift weight toward their holding several. The engine’s version of this is one parameter, and it is worth about 1.9 sets per duplicate deal-pair against a strong opponent that ignores it.

An exact rule for the perfect-information endgame. Once every remaining card is public, the game is solved in closed form: the team on move scores every live half-suit in which it holds at least one card, and the opponents score the rest. Footholds only shrink and a well-informed asker never misses. The practical corollary is that in a fully-known endgame the only thing that matters is *how many live half-suits your team has a foothold in when it is your turn*, not how many cards, not which cards.

Do not agonise over when to declare, but never guess. Reproducing v0.3’s finding on a much better posterior: raising the declaration bar from 0.97 to 0.99 changes essentially no decisions. The reason is structural rather than statistical. If your team genuinely holds all six cards of a half-suit, no opponent holds any card of it, so no opponent can legally ask in it, and a declaration by a team that does not hold all six gives the set to you. Waiting is close to free, and that mechanism is rule-invariant. What the award rule changes is the price of impatience: a guessed split that used to void the set now hands it to the other team, a two-set swing in the differential. Wait until you know the split; the rule correction doubled the fine for not waiting.

Two tie-breakers really do hold up. When two asks look about equally likely to land, prefer the one that, if it misses, hands the turn to the opponent with fewer cards, and prefer to fight in half-suits your team is already winning. v0.3 measured that pair at +1.41 sets per deal-pair; with an inference layer good enough to make the primary criterion much sharper, they are still worth +1.08. Better card-reading does not make positional judgement redundant.

The endgame is bookkeeping, so do the bookkeeping. Every one of 200 measured games reached a position where no ask anywhere could succeed: the sets were decided and only the declarations were left. Two consequences for a human. First, the last few turns are not a contest, so spend your attention earlier. Second (and this is the one people get wrong) if your team holds a half-suit outright, *no opponent can legally ask in it*, so it cannot be taken from you by any ask or by any opposing declaration. There is never a reason to rush that declaration. In our void-era games, half-suits that reached “ours but we cannot place the split” were misdeclared 23.4% of the time against 0.92% otherwise (a factor of 25), and under the award rule each of those is no longer a voided set but a set *donated*. Almost every one is a self-inflicted wound, and the wound now costs twice as much.

You can talk, and most players do not realise it. You may not ask for a card you hold. So asking for a card *proves to the table* that you do not have it, and when your team already owns the half-suit, the ask cannot lose you anything, because no opponent holds a card of it to give. In that situation an ask is a free, legal message to your partners telling them exactly where a card is not. Use it when you are stuck on a split. (Our engine cut its misdeclarations about 20% doing this in void-era measurement and it did not gain games there; re-measured under the award rule, where each save is worth two points of differential instead of one, the cut persists — 109 misdeclarations against 138. The 500-pair *screen* put the value at $+0.0680$ [$-0.0326, +0.1686$], which does not clear zero; the registered confirmation that followed does, at $+0.1180$ [$+0.0325, +0.2035$] and $+0.1220$ [$+0.0291, +0.2149$] (§10.1). So the mechanism’s value is established and positive. It stays off in the engine because that registration’s ship bar was a point estimate at or above $+0.15$, and the bar was not amended after the fact, not because the effect is unproven. For a human, who is far more likely to misplace a split than an engine is, the trade is almost certainly worth it.)

Spend your effort on card-tracking, not on cleverness about which ask to make. Three measurements in this study point the same way and none of them is visible alone. The full eleven-term objective ranks candidate asks barely better than $\mathbb{P}(\text{success})$ by itself. Nearly every term we added *beside* $\mathbb{P}(\text{success})$ is a null. And the largest single effect in the whole study apart from the opponent model itself, $+0.340$ sets per deal-pair, comes from tripling the sampling budget, which does not touch the objective at all: it only makes the $\mathbb{P}(\text{success})$ the objective is already dominated by more accurate.

For a human that is a claim about where to put attention. Knowing more precisely where the cards are beats reasoning more carefully about what to do with what you know. The engine’s version of “know more precisely” is more Monte Carlo draws; yours is tracking the failed asks, which are the constraints that do the work. It is not that positional judgement is worthless (the two tie-breakers above are real), but that between an hour spent on card-reading and an hour spent on ask selection, the measurements say the first hour is worth several of the second.

Do not hold back the re-take. A card an opponent has just taken from you is a card you know they hold, so asking for it back is a certain ask, and a certain ask keeps your turn. The case against is that the exchange is public and repeating it tells the table this half-suit is contested while neither side nets a card. Strong players describe doing it. We measured it, and it loses: a penalty on re-taking cost -0.340 sets per deal-pair at a small weight and -1.015 at a large one, the loss growing monotonically with how firmly the policy withholds. Two variants of playing to the score — more variance when behind, more caution when ahead — also failed to beat playing the same way regardless.

The situation is not rare, which is what makes the result usable: a re-take is available at 11% of positions, and at 18% of decision points the acting seat is already two or more exchanges deep with some opponent. So this is advice about a common decision, not a corner case. Take the certain ask.

We are careful about one thing here. What was measured is a penalty applied to *every* re-take, including the first, and the argument for withholding is really about a *repeated* exchange. The engine can now gate the penalty on the exchange recurring, and we can say in advance what that version could be worth at most, $+0.159$ sets per deal-pair, because the gate touches only the 5% of positions where a re-take is available and no card has yet come back the other way, which is 47% of what the ungated penalty flags. That is small enough that the 200-pair cells this study used for screening could not have resolved it either way, which is why no such cell was run.

It has now been run at 2000 pairs, pre-registered in `jobs/PREREGISTRATION_retake_gate.md` with the design, the five prior failures and the analysis all fixed before a pair was played. The gated penalty scores **-0.004** sets per deal-pair $[-0.067, +0.059]$ against the champion — indistinguishable from not having it. Against the ungated -0.340 the difference is $+0.336 \pm 0.169$.

Two readings, and the paper takes the cautious one. The gate does remove the harm: the ungated form's loss does not survive restricting the penalty to a repeated exchange, which is what its own argument always said it should. But the recovery is $+0.336$ where the gate's model caps it at $+0.159$, and a gate cannot recover more harm than it removes flags for. One of those numbers is wrong, and the likelier one is the ungated -0.340 : it comes from a 200-pair screen whose interval $[-0.665, -0.015]$ barely clears zero, which is precisely the configuration that inflates an estimate — the sixth appearance in this paper of the same selection effect. So the honest statement is that withholding the re-take, in the only form its argument ever justified, neither helps nor hurts; and that the harm the ungated version appeared to do was partly an artefact of the screen that measured it.

The advice does not change, and it is now better supported. Withholding does not pay. Take the certain ask.

Trading hard does not pay either. Every cell above *withholds*. The opposite policy — take the card straight back, trade hard inside the duel, which is what strong players more often describe doing — had never been measured. It has now been, at 2000 pre-registered pairs (`jobs/PREREGISTRATION_retake_bonus.md`), and it returns **+0.013** $[-0.052, +0.079]$: the seventh null in this family. Against the gated penalty's -0.004 the difference is $+0.017 \pm 0.046$. The two policies act on *disjoint* slices of the same decision in opposite directions, and both are null, which says something neither says alone — the re-take decision does not repay any policy at all, in either direction.

What that null is about was fixed before the run, because the reachable base rate was measured first. A re-take is a *certain* ask, so the objective already scores it at $P(\text{success}) = 1$ (the term this whole paper says dominates), and it is *already the chosen ask* at 69% of the positions where it is legal. A bonus can therefore act at only 3.5% of positions, and where it can, the median score gap to the leader is **0.000**: at the median reachable position the weight is breaking an exact tie between two asks the objective values equally. Breaking a tie has expected value zero unless the objective is systematically wrong in a way correlated with re-taking. So this is a null about tie-breaking, and it is emphatically *not* a measurement of “trade hard in a duel” as a human would mean it. The pre-registration says all of this before the first pair, which is the only reason the sentence can be written now.

The practical advice is therefore unchanged in content and much better supported in kind: the re-take is a certain ask, so take it, and do not spend any thought on whether to make a policy of taking it or of holding back. Neither policy is worth anything. Spend the attention on the asks that are not certain.

Both of those figures moved, and the reason is worth recording rather than quietly restating. The depth statistic originally counted successful asks between this seat and one opponent in *either direction separately*, so an opponent taking two cards off this seat and giving nothing back scored the same depth as a card going out and coming home — opposite situations, since in the first that opponent has plainly netted two cards. Over 1023 harvested positions, 37.8% of the positions where the gate opened had no two-way exchange at all with the opponent it opened on. Requiring cards to have moved both ways takes the “two or more exchanges deep” rate from 57% to 18% and *raises* what the gate could recover from $+0.098$ to $+0.159$, because sparing a first re-take now correctly includes every position where nothing has come back yet. This is the same defect the gate was built to fix, one level down: an implementation that did not match its own stated hypothesis. The unit test that should have caught it was named for exactly this property and its example happened to run both ways, so it passed against code that did not have it.

A warning about the obvious idea. Scoring asks by “how much does this improve my team's chance of winning this half-suit”, which sounds like the right objective, is measured in the right units, and rests on a well-calibrated model — is a disaster: -7.4 sets per deal-pair against simply asking for the card most likely to be there. In this game, the probability that the ask lands is not one consideration among several. It is the objective, and everything else is a tie-break.

γ_{opp} held at the incumbent 0.35; 1,557 decisions, paired by decision




Figure 8: Why the teammate-side model was withdrawn. Both panels are paired against the incumbent cell by *decision* — cards inside one decision share a belief, so the card is not the independent unit — with γ_{opp} held at the shipped 0.35. Sharpening the model on our own side buys a small improvement in the proper score around $\gamma_{\text{team}} = 0.7$ and pays for it by naming the true holder less often at every setting. The allocation decision reads the argmax, so that is a worse belief for the job it has to do.

12.2 Failed and abandoned experiments

Recorded deliberately, because in this project the negative results have repeatedly been the informative ones. The v1.2 programme’s ten registered arms (P43–P45), none of which cleared its bar, are catalogued in §11.11 and Appendix A rather than here.

One γ per side of the table, refuted before it was played. Two v1.0 results point hard at the teammate side of the belief. 95.3% of our residual errors are allocation errors, and a seat handed its teammates’ cards gains $2.6\times$ what a seat handed its opponents’ cards gains (Figure 5). The choice model weights every seat’s asks by one exponent $\gamma = 0.35$, and that one number does two different jobs: on the opponents’ side it sharpens the read that picks the next ask, on our own side it sharpens the read that places a split. The obvious move is to price them separately, with $\gamma_{\text{team}} > \gamma_{\text{opp}}$.

It does not work, and the reason it does not is worth more than the knob would have been. Registered at `prereg/gamma_split.md` with two conditions fixed before the numbers were read (the paired NLL interval must clear zero *and* the paired top-1 interval must not fall below it), and run over 1,557 decisions at positions the incumbent reaches, scoring every cell on the same positions at 720 sampler draws.

That last clause is a measurement condition and not a detail, which we learned after this section was written. A paired belief difference *scales with the draw budget it is scored at*, and does not converge in the range anyone would run: the same intervention on the same transcripts measures -0.0064 $[-0.0217, +0.0090]$ (a null) at 180 draws and -0.0493 $[-0.0649, -0.0337]$ at 2,880, with every doubling above 720 still significant on a paired contrast (`results/channel_precision_plateau.json`). So a figure of this kind is a property of an intervention *and a budget*, and both belong beside it. At the 480–720 this project uses, 75–78% of the 2,880-draw value is present, so no result here is materially attenuated; the steep part of the curve is below 360, which nothing uses. Why the scaling happens is open (two explanations were registered and refuted), and it is recorded in `RESEARCH_FRONTIER.md` rather than resolved here.

The predicted cell (0.35, 0.70) improves NLL by -0.0119 $[-0.0150, -0.0088]$ and worsens top-1 by -0.0093 $[-0.0128, -0.0057]$. Condition 2 fails. Figure 8 is the whole result: the proper score and the hit rate move in opposite directions, so the NLL gain is probability mass being spread rather than the read improving.

Two cells do pass both conditions, and the best of them is $(\gamma_{\text{opp}}, \gamma_{\text{team}}) = (0.70, 0.70)$ — on the diagonal, and so not a split at all. It beats the one off-diagonal cell that passes by more than $3\times$. What the run found is that the incumbent γ is too low on this pool, which Table 3 already reported for the posterior

in general and which the engine declined to act on because the belief improving did not carry into play. The single off-diagonal survivor has $\gamma_{\text{team}} < \gamma_{\text{opp}}$: believe our own side *less*, the reverse of the hypothesis.

Withdrawn, with no duel played. The teammate value the ceiling split measured is real and remains unclaimed; what this says is that reweighting the existing depth heuristic is not the route to it. The route that remains is a genuinely different teammate model — conditioning on our own policy’s likelihood rather than on a depth heuristic, which is the only formulation that uses the fact that our teammates run *our* policy and we possess it.

An exact posterior, by enumerating the disjunctions. Since the largest measured effect in this study comes from making the posterior more accurate, an *exact* posterior is the obvious thing to want, and there is a standard route to one. The counting DP is exact over candidate masks and quotas; what forces sampling is the OR constraints the no-bluff rule leaves behind. The complement of an OR is a conjunction of exclusions, which the DP absorbs natively, so inclusion–exclusion over the k OR constraints gives both the exact partition function and exact marginals in 2^k DP passes and no sampling at all.

We measured the two quantities that decide it before writing any of it. On 100 real decisions the median is $k = 6$ OR constraints (max 14; *none* has zero), and one DP pass returning marginals costs 0.86 ms. So the median decision would cost 55 ms against the sampler’s 3.8, and only 15% of decisions would cost no more than sampling does.

The reason it fails is structural rather than a constant factor, which is why no amount of optimisation rescues it: the cost of exactness grows as 2^k in the same k that makes sampling hard. Where the sampler struggles — many OR constraints, a tightly coupled support — enumeration is exponentially out of reach, and where enumeration is cheap there is barely a constraint to handle and the sampler is already close to exact. The two methods are easy in the same places and hard in the same places, so neither can rescue the other. Recorded because this is an idea that will occur to the next person too, and it took ten minutes to kill with two measurements rather than a week to kill with an implementation.

Exact inference, on its own. The headline effort of this version (replacing a documented-biased sampler with provably correct inference) bought nothing in play (Section 8.1) and was measurably worse at predicting reality (Section 8.2). We report it as a failure of the *hypothesis*, not of the implementation: the implementation is validated against exhaustive enumeration in four separate ways. What failed was the assumption that the uniform posterior is the right target.

Rejection sampling for the disjunctive clauses. Exactly correct, and unaffordable: measured acceptance mean 0.28, median 0.17, below 0.05 on 16% of positions. Abandoned in favour of importance sampling with closed-form densities.

Folding the clauses into the exact dynamic program. Sound, and tractable in the median case (refined state space 5,760), but with a tail that makes it unusable as a default: p90 230,000 and p99 2.7×10^7 . Retained as an option, not as the path.

A naïve batch rewrite of the sampler. Vectorising the draw across the batch gave only $1.4\times$ until the dictionary materialisation of the draws was removed as well; the correct fix was to carry the whole batch as an integer matrix end to end, after which the same change gave $5.1\times$. Worth recording because the first measurement would have justified abandoning a change that was in fact the right one.

An under-converged fit that looked plausible. The first half-suit value model was fitted by plain gradient descent on unstandardised features and reported a held-out log-loss *worse than the class prior* while producing entirely reasonable-looking coefficients. Feature scales spanning two orders of magnitude were the cause. This is the reason the fitting routine now reports against two baselines (the class prior and a single-feature model) rather than against nothing.

A false positive, caught. Of more than twenty-five screening cells, exactly one produced an interval excluding zero in the improvement direction ($+0.490$, $[+0.184, +0.796]$, 500 pairs). It failed to replicate twice on fresh seeds ($+0.068$ and -0.066). We record it as a failed experiment rather than burying it in a methodology note, because it is the shape of the mistake this project is most likely to make next.

change	set diff	95% CI
break the duel, $w_{\text{retake}} = 0.30$	-0.340	[-0.665, -0.015]
break the duel, $w_{\text{retake}} = 1.00$	-1.015	[-1.562, -0.468]
gated on a real duel, 2000 pairs	-0.004	[-0.067, +0.059]
<i>reward</i> the re-take, $w_{\text{retake}} = -0.30$, 2000 pairs	+0.013	[-0.052, +0.079]
tie-breakers up when behind, $w_{\text{behind}} = 0.50$	-0.185	[-0.628, +0.258]
tie-breakers up when behind, $w_{\text{behind}} = 1.00$	-0.890	[-1.384, -0.396]
tie-breakers down when behind, $w_{\text{behind}} = -0.50$	-0.075	[-0.552, +0.402]

Table 19: Two ideas for the information the engine does not carry between turns, each ablatable to zero and each refuted with a dose-response. All cells are duplicate deal-pairs against the champion. The gated and sign-reversed rows are the controls: both are flat, so the loss is the term firing where no duel exists rather than the duel being mispriced.

Style that responds to the match. Every objective term in v0.4 is a function of one candidate ask and the posterior. None knows the score, and none knows that the last four moves were the same two players passing one card back and forth. Two ideas were built to close that gap, each ablatable to zero, and both are refuted, not as nulls, but with a dose-response, which is the more informative outcome because the gradient confirms the mechanism rather than merely the sign. All cells are 200 duplicate deal-pairs against the champion.

Breaking the duel penalises taking back a card this seat just lost to that same opponent. Section 3 notes that the state graph is cyclic precisely because two opponents can trade a card back and forth; the question here was whether a player should decline to. They should not, and three times the penalty costs three times the sets. This is the scheduling argument of Proposition 3 arriving from a third direction: a card you watched an opponent take is a card you *know* they hold, so taking it back is a certain ask and a certain ask keeps the turn. Deferring it risks the turn to preserve an option that was not going anywhere, and the information concealed by not asking is worth less than the possession thrown away. It is worth recording that this is a manoeuvre strong players describe using, which the engine could not previously express at all — the measurement is the point, not the outcome.

Score-adaptive tie-breakers scale the weights by how far behind the team is, on the reasoning that a losing side should buy information and a winning side should buy certainty. Scaling them *up* when behind loses monotonically. Scaling them *down* is indistinguishable from the champion. The asymmetry is the finding: the loss surface is steep on the side of trusting the tie-breakers more and flat on the side of trusting them less, which is what one would expect if v0.3’s hand-tuned constants already sit at the point where the tie-breakers have given everything they have, and is a third piece of evidence, after the inverted-U sweeps and the substitution result of Section 8.6.1, that they are tie-breakers rather than an objective.

Neither idea is retried in another form. What both share is that they spend $P(\text{success})$ to buy something else, and this paper’s most reproducible finding is that nothing else is worth buying with it.

Belief-space possession-chain search. Not abandoned — this one resolved, and it is listed here because it spent more pairs than anything else in the study and because three of its four rounds belong in a section about what does not work. Three runs of the identical configuration returned +0.570 (200 pairs), -0.044 and +0.386 (500 pairs each); the first and third exclude zero and the second excludes the third’s estimate. Two 700-pair cells pre-registered as decisive pooled to [-0.019, +0.324] and settled nothing, because their size had been chosen against an estimate the selected cell inflated. A fourth round, pre-registered and sized against the unbiased figure, finally answered it: +0.104, 95% CI [+0.020, +0.189] over 6000 fresh pairs, with the six blocks agreeing at $\tau = 0$. Section 9.3.1 has the arithmetic. The lesson kept here is the shape of the estimate over rounds — +0.570, then +0.153, then +0.104, and that no amount of cleverness in the analysis substituted for pairs.

Inherited from v0.3, and not revisited. Determinized search (PIMC and Information-Set MCTS) lost to its own prior; a value network trained on belief features and applied inside determinized worlds

lost 34–6 to a train/inference distribution mismatch; quiescence extension made search worse; and an expected-value claim model predicted an optimal threshold near 0.70 that direct measurement falsified. None of these were re-attempted here, because v0.3’s diagnosis — that the binding constraint is the objective, not the depth of search — was the premise of this version and is not contradicted by anything we measured.

12.3 KRAKEN v1.1: seven directions, seven closures, nothing shipped

Version 1.0 established where the remaining errors are: allocation rather than ownership, and compelled rather than voluntary declarations. Version 1.1 took seven pre-registered runs at that residue. All seven closed. **None shipped**, and V06_DEPLOYED is byte-for-byte what it was — every parameter added in this round is bit-identical to the champion at its default, which tests4/ asserts rather than assumes.

A round that ships nothing is worth reporting only if the closures are informative, so this section is organised around what each one rules out. Two are stronger than a null: one is a *theorem*, and one is a correction to intervals published earlier in this paper.

12.3.1 The discipline: a futility screen before a duel

Each direction carried a pre-registered screen with thresholds fixed before any number was read — typically “does this knob change enough decisions to be worth n pairs?” **Four of the seven never reached a duel**, because the screen stopped them first. At the sizes those duels were registered for, that discipline saved roughly 9,000 duplicate-deal pairs.

A fifth direction also closed without a duel, and by a different route: the covariate of row 2 was refuted on its own pre-registered decision rule for the *belief*, not stopped by a futility screen, so it saved no duel pairs — none had been sized for it. Its lesson is separate and is the one worth carrying: it is worth +3,143 held-out nats at predicting *which half-suit a teammate asks in*, and makes the posterior over *where the cards are* worse at every weight tested. A better model of the ask is not a better model of the deal. The dual-population rule added later cuts both ways: P45’s one arm lost $2.8\times$ more in self-play than against SESTINA (§11.11.8), so the population added as the check on self-play optimism was the optimistic one.

This matters for reading the table below. A direction marked “no duel” is not untested; it is a direction whose own pre-registered condition for deserving a duel came back false.

#	direction	outcome
1	one γ per side of the table	refuted on a co-primary; withdrawn, no duel
2	the choice model’s covariate	unlocated_now beats depth by +3,143 nats on the <i>ask</i> and is refuted in the belief : worse at every weight, monotone in it
3	a pre-play convention in the rules	belief says yes three times, the duel says no twice: at the mis-priced gate (+1.750) and, re-priced to a free message that changes only tie-breaks (P45), $-0.31 / -0.86$; the price is not the message but the deterministic tie-break that decorrelates three seats
4	decompose the teammate ceiling	46% is interaction : $D + K = 1.84$ against $T = 3.41$
5	locate: price an ask by the declaration it enables	null , $+0.047 [-0.075, +0.168]$, diagnosed
6	declarability inside the possession search	closed by proof , no duel
7	claim_gamma and reach	both screens fired on each; no duel

Table 20: The seven v1.1 directions. Four closed without a duel because a futility screen with thresholds fixed in advance stopped them. Nothing in this table shipped.

12.3.2 Direction 6, closed by proof: an invariant of the possession search

The largest single class of remaining error is *allocation* — our team holds all six cards of a half-suit and names the wrong split. The natural fix is to make the belief-space search resolve which teammate holds which card, by scoring declarability at its leaves rather than only material.

That search cannot resolve it, and the reason is structural rather than empirical. Let $M[c, p]$ be the posterior mass on seat p holding card c . The search’s only update, `apply_success`, scales the *target*’s column by a constant and renormalises each row. Both operations preserve ratios among non-target entries, and in Literature the target of an ask is always an *opponent*. Therefore, for any two teammates t_1, t_2 and any card c the chain does not itself take:

$$\frac{M[c, t_1]}{M[c, t_2]} \text{ is invariant under the entire search tree}$$

— at every depth, under every beam width, and *under every leaf evaluation, including a perfect one*. The quantity a declaration must get right is exactly the quantity the search cannot move.

We assert the invariant numerically to 5×10^{-16} and price it over 20,472 (decision, live half-suit) pairs with no ground truth, by measuring what share of the allocation uncertainty lies in ratios the tree cannot reach. Over all pairs that share is 40.7%; restricted to the band where the allocation errors actually occur — $\mathbb{P}(\text{we own all six}) \geq 0.9$ — it is **99.96%** over 91 pairs. Deepening or widening this search, or giving it a perfect leaf, cannot buy the thing it would need to buy.

12.3.3 Direction 4: the teammate ceiling is 46% interaction

§8.7 priced perfect card-reading at +6.61 sets per game and split it: teammates’ cards are worth +3.41 [+3.16, +3.66], opponents’ +1.31 [+1.01, +1.61]. Direction 4 asks what the teammate half is *for*, by handing a seat the true deal in one channel at a time over 600 games:

arm	what the seat is given	ceiling
D	the truth only when it <i>declares</i>	+1.077 [+0.798, +1.355]
K	the truth only when it <i>asks</i>	+0.760 [+0.473, +1.047]
T	both	+3.410 [+3.162, +3.658]

$D + K = 1.84$ against $T = 3.41$. **The two channels together are worth 1.57 sets per game more than the sum of their parts** — 46% of the teammate ceiling is in neither channel alone. Knowing the deal only when you declare, or only when you ask, each buys about a set; knowing it in both buys three and a half.

That is the finding that makes directions 5, 6 and 7 the right things to have tried and the wrong things to have expected to work. Each attacks one channel. The interaction term is by construction unreachable from either.

12.3.4 How much is actually left, measured with power

An earlier draft of this paper carried an assertion of ours (“*what is left in this engine is not a defect to be fixed*”), that nothing measured supported. It is corrected in place, and these runs replace it.

channel	verdict
ask, one ply, objective alone	no headroom: -0.0188 [$-0.0951, +0.0575$]
ask, one ply, deployed champion	weak positive: $+0.1641$ [$+0.0046, +0.3236$], and not caused by the ask choice
declaration	correct in both directions
single-seat deviation	-0.608 [$-1.122, -0.095$]; uninformative by construction

Both ask intervals are clustered by deal (§12.3.7), which is why they are wider than the figures an earlier draft carried.

The declaration is taken at the right moment. `scripts4/declare_regret.py` fills a hole `ask_regret` names in its own source (if not `isinstance(chosen, Ask): continue`). With claim actions in the set, over 109 positions where a declaration was available, the load-bearing figure is a within-position contrast: where the policy asked, **best claim minus best ask** = -1.1849 ± 0.1051 over 98 positions, and the best claim beat the best ask in **3 of 98**.

This refutes the story that motivated four earlier directions. The teammate oracle declares at move 39.2 against our 70.6 because it is *certain*, not because early declaration is better, and that certainty cannot be bought: the public record holds 92.1% [91.2, 92.9] of which partner holds an unlocated card, and quadrupling the posterior draws moves it 0.2 points.

12.3.5 The one-step regret belongs to the turf, and the turf is the positions

The champion shows a small positive one-step regret while the ask objective alone does not. Read side by side those two runs appear to say the lookahead makes ask selection worse. They cannot support that: they differ in the positions harvested, in the continuation that scores them, and in whether the acting and evaluating posteriors agree.

`scripts4/actor_compare.py` holds all three fixed by construction — one set of positions, one set of worlds, one set of rollouts, *both* actors scored on them. Run as a 2×2 over the two remaining confounds, every interval clustered by deal:

harvest	continuation	pos	deals	champion – objective
objective	objective	129	5	$-0.0006 [-0.0409, +0.0396]$
objective	champion	110	4	$+0.0417 [-0.0439, +0.1273]$
champion	objective	125	6	$-0.0087 [-0.0843, +0.0669]$
champion	champion	103	5	$+0.0202 [-0.0465, +0.0870]$

The actors do not separate in any cell, the tightest at ± 0.041 , while picking a different ask in 34–36% of positions. So the lookahead is not a tie-break, and whatever it earns — it won $+0.104 [+0.020, +0.189]$ over 6,000 pairs — it earns *beyond* one ply. That is the first direct evidence of it.

The continuation, measured paired on identical positions, is null: $+0.0176 [-0.1010, +0.1362]$ over 213 positions in 9 deals. What is left is the position distribution, at $+0.0803 [-0.0786, +0.2393]$. So the answer is an **elimination**, not the ask choice, not the continuation, and the size of what remains is a lead rather than a finding. We priced settling it at 15–30 CPU-hours for an interval that would still touch zero, and did not spend it.

12.3.6 The endgame ask correction does not survive the rule change

§9.3.3 flipped the misdeclaration rule to the standard opponent-award baseline. The endgame ask correction `endgame_d_info = +2.0` was fitted against *void-era* one-ply targets (the fit landed twelve hours before the flip), and was withdrawn from the deployed configuration pending a refit. The refit does not bring it back, and it says something sharper than “null”.

rows	rule	grid’s info	held-out gain over champion
388	void	+2.00	$+0.0092 [-0.0282, +0.0467]$
457	award, same 79 deals	–1.00	$+0.0142 [-0.0118, +0.0403]$
764	award, all 156 deals	+0.10	$-0.0011 [-0.0033, +0.0011]$

Held to the same deals, so that only the rule moves, the grid’s pick **reverses sign**. The scale family still picks $k = 1$ under both rules, so that half of the original diagnosis stands: de-weighting the success probability does not help.

The part that is not about this knob: `info = +2.0` was nominated for a 4,000-pair duel on a held-out gain of $+0.0092$ whose interval covers zero and is four times the estimate. The duel itself was pre-registered and run as written, and its $+0.0835 [+0.0338, +0.1332]$ remains a true fact about the void rule.

What was wrong is upstream: the screen that chose where to spend the pairs could not tell its winner from zero, and reported a point estimate with no interval. **Standing rule adopted:** an offline screen nominates an arm for a duel only if its held-out gain has an interval excluding zero, clustered by whatever unit shares a deal.

12.3.7 A correction to intervals published earlier in this paper

ask_regret’s harvest walks games in order and emits *every* qualifying ply, returning as soon as it has enough. So a run reported as “162 positions” is 162 consecutive plies drawn from **eight deals**, sampled twenty to forty deep. Positions inside one deal share the hands, the history and every earlier decision. Every interval this instrument published divided by 162 rather than by 8.

Fixing the count alone would have replaced one too-narrow interval with another: cluster-robust variance is asymptotic in the *number of clusters*, and these runs have three to ten deals, so the standard error must be paired with a t critical value at $k - 1$ degrees of freedom rather than 1.96. At six clusters that is a 31% difference; at four, 62%.

run	as published	clustered by deal	deals
ask objective alone	−0.0188 [−0.0808, +0.0431]	−0.0188 [−0.0951, +0.0575]	8
deployed champion	+0.1641 [+0.0797, +0.2484]	+0.1641 [+0.0046, +0.3236]	8
champion, narrow run	−0.0893 [−0.2535, +0.0749]	−0.0893 [−0.7109, +0.5323]	3

No conclusion reverses, but one comes close: the champion’s positive one-step regret still excludes zero, by +0.0046 against a published lower bound of +0.0797, with an interval $1.89\times$ wider. One conclusion elsewhere *is* withdrawn — a regret figure over 11 positions from 4 deals, which no longer excludes zero and is reported as silent rather than supporting.

The audit extends to every other instrument that harvests positions this way. The four slopes in §9.4.1 are restated with deal-clustered standard errors; none reverses. **Duplicate-deal duels are unaffected:** the duel harness pairs on the deal and treats the pair as the independent unit, and already used a t critical value at the pair count. Every headline duel figure in this paper (including the +2.3466 against v0.7, since withdrawn for a bridge defect and not a clustering one, §11.10, and the +0.104 lookahead result) stands exactly as written with respect to clustering.

12.3.8 What v1.1 leaves

The interaction term of direction 4 is real, it is 46% of the largest information channel in the game, and directions 5, 6 and 7 establish that it is not reachable from the ask basis, not reachable from the possession search, and not reachable by the two auxiliary knobs. The theorem of §12.3.2 is the strongest of these because it holds against a perfect leaf evaluation rather than against the one we have.

What that leaves is a channel this engine does not have: something that couples the ask and the declaration rather than improving either. Direction 3 built the one such mechanism the rules permit (a pre-play convention), and the duel refused it twice: at a gate later found mis-priced, and (P45, §11.11.8) at a free-message gate that costs nothing in the objective’s own units and still loses −0.31 against SESTINA and −0.86 in self-play, a cost the belief-space instrument could not see because the price is the deterministic tie-break and not the message. Both halves of the channel are now measured in play: receiving is worth nothing (the free read, −0.002 [−0.127, +0.123] at 3,000 pairs) and sending costs at every gate tried. We do not have a candidate that survives both the theorem and those duels, and we say so rather than nominating one.

12.4 Limitations

The opponent model is a model. Everything else in the inference stack is deduction, validated against enumeration. Equation (6) is a hypothesis about how people choose, fitted with one parameter. Against opponents whose ask policy is unrelated to their depth in a half-suit it could in principle hurt, and against

an opponent who deliberately inverts the heuristic it is exploitable in the ordinary sense that any model is. We measure both cases rather than assert robustness, and γ is a single number that can be set to zero.

Depth is evaluated on the initial deal. The choice model is really about what a player held *when they asked*. Cards move, so the two diverge late in a game. The exact quantity would need a replay of the public history per candidate world, turning an $O(1)$ update into $O(|h|)$.

The uniform-posterior result is about *this* game and *these* opponents. We show that exact inference under an uninformative- choice assumption is worse than a biased sampler that accidentally encodes a choice model. That is a statement about which approximation is closer to the truth, not a general licence to prefer biased inference.

Exact solving is still bounded, and bounded in a way this section previously misstated. The tables solve the *perfect-information* game. Calling them absolute ground truth without that qualifier is an overclaim, and two measurements say how large a one.

First, the coverage. “Reaches positions with one, and now two, unresolved half-suits” describes the solver’s reach over positions handed to it, not the game’s. A tablebase value is legitimate only where the acting seat’s belief pins every live card, which `tablebase4.pinned_state` enforces. Over 2582 decisions from 25 games that gate opens on 2.59% of decisions: never at three or more live half-suits (0 of 188), 12.1% at two, 52.7% at one. It first opens at step 100.2 of 103.5, with 8.1 of 54 cards left, and in 5 of 25 games it never opens at all.

Second, the gap. Applied by determinizing over a belief, the tables overstate the mover’s team by 5.294 ± 0.159 sets on the positions play actually reaches, and by 8.18 with all nine half-suits still live, where the closed form $V = \text{sign} \times (2f - m)$ of §5 predicts the team on move takes 8.36 of the nine sets where play delivers +0.27. (That 8.36 is an average over the $m = 9$ positions play reaches, which is not the same as a fresh deal: footholds only shrink, so f has already fallen a little by then. At the deal itself (9) gives +8.794 exactly.) The closed form is not evidence that the Fish endgame is shallow; its own proof turns on “every ask can be made to hit, and a hit retains the turn”, which is what perfect information buys and nothing else does. Measured independently, a team given the true deal beats the champion by 17.483 sets per duplicate deal-pair against a maximum of 18.

So the honest statement is not that midgame positions remain out of reach pending a larger solver. It is that exact ground truth covers 2.6% of the decisions a player faces, all in the last 3% of the game, and that extending the perfect-information solver to three half-suits in full (the natural next computational step) would move that coverage by approximately zero, because the positions it would newly cover essentially never arise with every live card pinned. More ground truth requires solving the imperfect-information game, and the smallest Fish position carrying any hidden information already needs three players holding live cards, hence two on one side: a team-coordination problem with no two-player zero-sum reduction beneath it.

No equilibrium claim, and the one exception. We compute no equilibrium. We do now compute an exploitability *lower* bound, in Section 11.7: a single seat deviating optimally at one endgame decision gains +0.1111 per game with an interval clear of zero, so the all-champion profile is not an equilibrium. That is a bound against a deterministic realisation of the champion rather than against the champion as it plays, and it is a lower bound only — it says the profile is exploitable, not by how much. The appropriate solution concept for a game with pre-play agreement and no in-play communication is TMECor (Celli and Gatti, 2018), and nothing here approximates it. Every strength statement is relative to a named opponent or absolute against an exactly solved subgame.

The cross-engine margin is measured through a bridge that changes their agent. The largest limitation this paper acquired late, and the one a reader should weigh hardest. §11.10 reports the same two engines measured in *their* arbiter, where our withdrawn +2.3466 becomes a loss, and their whole released ladder reverses order with it. Three reconciliations are excluded by controls — the dialect and the opponent, by experiments inside their engine involving nothing of ours, and our bot package, by its

beating their v0.3 inside their own arbiter — and the fourth is the answer: our shim is stateless and their agent is not, for every version from v0.4 on.

So the honest reading of every cross-engine *absolute* in this paper is that it is a margin against their engine *as our bridge runs it*. That is a real opponent and an independently authored one, and it is not the opponent their arbiter runs. The paired contrasts are unaffected — an absolute measured through a bridge is a statement about the bridge as well as about the engines, whereas a contrast whose arms share the bridge is not (Appendix C) — and this paper is now an instance of the hazard it wrote that sentence about, which is the most useful thing it can say about it.

A rules discrepancy, stated. The printed rules allow a set to be declared at any point, including during an opponent’s turn. v0.3’s game loop only ever asked the player on move, so the rule was accepted by the engine but unreachable by any policy, which matters most for a seat that has run out of cards and therefore never gets a turn of its own while a teammate still has some. v0.4 adds off-turn declaration, offered identically to both generations so that the comparison stays a comparison of one thing. The primary evaluation is nevertheless run under v0.3’s own configuration, so that the numbers are comparable with the previous version, with the variant reported separately.

13 Conclusions

We set out to remove a defect the previous version had documented in its own code: a belief sampler that was not uniform over the deals consistent with the public record. We removed it, exactly. Hidden-state inference in Literature reduces to counting assignments in a degree-constrained bipartite system, and although that is $\#P$ -complete in general, the instances the game produces are tiny in the dimension that matters, roughly four distinct candidate sets even when twenty-six cards are unlocated, so a dynamic program returns the exact partition function, exact marginals, exact uniform sampling and exact joint probabilities. The disjunctive constraints that the program cannot absorb are handled by importance sampling with closed-form proposal densities, verified against exhaustive enumeration for support, density and marginals, with a control demonstrating that dropping the weights visibly breaks the estimate.

And it changed nothing. Against the previous champion, at matched seeds over 250 duplicate deal-pairs, provably correct inference moved the set differential by -0.008 per pair.

The instructive part is why, and that we could find out. Because the hidden hands exist in simulation, a posterior can be scored on the question it is actually trying to answer, and that measurement says the exactly-uniform posterior is *worse* at predicting reality than the biased sampler it replaced. Proposition 1 is a correct theorem about a false hypothesis: it assumes players’ choices carry no information beyond legality. The old sampler’s bias, an artefact of how it seeded the disjunctive constraints, happened to encode a crude model of how players choose. Removing the bias removed the model with it.

Replacing the accident with one line of modelling (a player asks in a half-suit in proportion to how many cards of it they hold) recovers the loss and goes past it, on all three instruments we have: posterior likelihood against the true hands, agreement with an exact solver, and duplicate-deal play. The last of these puts it at $+1.9$ sets per deal-pair against the previous champion, which is larger than the entire objective improvement that separated that champion from its own baseline.

Seven lessons we would carry to another game.

Correct is not the same as good. An inference procedure is correct relative to a model. When a cheaper, wronger procedure beats an exact one, the exact procedure’s *assumptions* are the thing that is failing, and the productive response is to find which assumption and relax it, not to add compute to the correct-but-misspecified thing.

Measure the intermediate quantity, not only the outcome. Every belief metric in the previous version was indirect, and a change that was provably correct and empirically useless would have been very hard to diagnose from win rates alone. Scoring the posterior directly against ground truth turned a puzzling null result into an identified cause in one experiment. Where a simulator gives access to the hidden state, refusing to use it for *evaluation* is a false economy; the discipline that matters is not letting it reach the acting policy.

Know what your experiment could not have found. The measured per-pair standard deviation of the set differential is around 3.8 sets, which means the previous version’s 600-pair cells could not resolve effects below about 0.44 sets per pair. Several of its reported nulls were therefore uninformative rather than negative. We report the minimum detectable effect alongside every interval for that reason.

The sharper form of this lesson took most of the study to find, and it is that the noise figure is not a constant of the problem. Under common random numbers a paired trial on which the two arms behave identically contributes exactly zero, so the standard deviation scales with $\sqrt{s}$, where s is the share of trials on which the two arms diverge at all, and across our 46 cells the remaining factor is nearly constant at 3.86 ± 0.29 , though Section 7.3 shows it drifts enough at low s to matter where the rule is used. The consequence is that sizing every experiment against a single number taken from two maximally different policies systematically over-sizes the comparisons that matter most, the near-identical ones. A candidate perturbing a tenth of trials needs $9.5\times$ fewer of them than the constant implies. The error is conservative, so no interval we published is too narrow; what it cost is experiments never run and nulls recorded from cells that could not have found anything, and unlike a bad interval, those leave no trace to audit later.

Replicate the winner, not the losers. A screening study over twenty-five cells will hand you an apparent winner under a complete null, and it will arrive with a plausible mechanism attached — ours came with a story about two independently-motivated changes stacking, which is exactly what the previous version had genuinely found for a different pair of terms. The cheap defence is not a multiplicity correction, it is one more run on fresh seeds. Ours cost forty minutes and saved a wrong champion.

And it is not a lesson one learns once. The same error appears five times in this study, each occurrence a level above where the last was caught: a selected screening cell, then a section built on three points, then the power calculation for the experiment correcting the first, then the estimator built to measure the ask objective, and finally, in a shape we did not anticipate, an impossibility claim that closed a research direction. None was caught by being more careful. Each was caught by measuring the quantity that had been reasoned about, which is the next lesson, and the reason it is stated separately.

A diagnostic you can optimise is not the thing you want. The sampler reports an effective sample size, and it is a real quantity honestly computed. Tuning the proposal to raise it worked, 83.7 effective draws became 105.8, and made the posterior $3.4\times$ worse, because effective sample size measures how *flat* the importance weights are and a proposal covering one corner of the target very well has beautifully flat weights over the corner it covers. Every intermediate metric in a pipeline is a proxy for something further downstream, and the moment it becomes an optimisation target it stops being a measurement of what it proxied. Section 8.3 keeps both halves of that result, and the parameter ships at zero.

An impossibility claim is a claim. We wrote, in this paper and in the code, that the belief tracker *cannot* attach to a determinized mid-game position, and closed the objective-learning line on it. The claim was specific, mechanistic, well-argued and supported by a number, and every one of those properties made it less likely to be checked. It was false: the tracker is anchored on the initial deal but does not have to be handed one, and given the current hand and the public history it back-computes a consistent deal in a few lines. The measurement that had been used to support the diagnosis, repeated with the strong policy actually running, moves from $+0.101$ to $+0.681 \pm 0.142$.

What made it survive is that a negative claim generates no experiment. A positive result invites replication; “this cannot be done” invites moving on, and the sentence that closes a direction is exactly the sentence nobody tests.

Having been caught once, we swept the engine for the rest of them. `fish4/NEGATIVE_CLAIMS.md` registers every load-bearing impossibility claim in the source with a verdict: one refuted (the above), one true in the limit and false where the engine operates (a proposal twist “cannot change what the policy computes” — correct asymptotically, and at the shipped 160 draws it made the marginals $3.4\times$ worse), and six that are either rules theorems or were measured rather than assumed. The rider that sweep produced is worth more than the register: an asymptotic “cannot” and a finite-sample “cannot” are different claims, and only one of them is about the engine that ships, so a measured negative should always carry the range it was measured over. The asymmetry is worth naming because the cost is asymmetric too: an inflated positive wastes the run that eventually corrects it, while an unexamined negative costs

everything downstream of the direction it closed, and leaves no anomaly to notice. Our stated defence (measure the quantity you reasoned about) applies verbatim, but we had been applying it only to numbers we wanted to be true.

A coefficient that decays may be a covariate that decays. The choice model’s fitted exponent falls from 2.00 at the opening to -0.02 by the endgame, which reads as players changing how they choose. Most of it is regression dilution: the model conditions on depth at the *initial deal*, cards move, and the covariate’s disagreement with the hand it describes grows monotonically through a game. Noise in a covariate attenuates its own coefficient, so noise that grows with time manufactures a trend out of drift. Refitting on the quantity the model is actually about leaves a decay that is real and half the size. The general form: before explaining a trend in a fitted parameter, check whether the thing it is fitted to is degrading at the same rate.

13.1 What the three cycles after v0.4 established

Everything above is the v0.4 conclusion, and it stood unchanged while three more cycles ran. What they added, in the order the paper reports them:

A rule was wrong for five versions, and fixing it re-priced the record. Every engine in this lineage nulled a misdeclared half-suit when the standard dialect awards it to the opposition. Under the corrected rule a saved misdeclaration is worth two points of differential rather than one, which doubles the value of every declaration-side mechanism and inverted at least one earlier verdict (§3). Results are labelled by era throughout because the two are not comparable.

The residual is allocation, not card-reading. 95.3% of the champion’s remaining declaration errors are wrong splits inside its own team rather than mistakes about where opponents’ cards are, and an oracle handed the teammates’ hands is worth $+3.41$ sets a game against $+1.31$ for the opponents’ — a ceiling, obtained by cheating, and never a strength figure (§8.7). That is the finding that reframed the programme: the inference machinery this paper opens with was aimed at the smaller of the two.

Exploitability is measured, not bounded by a search. Exact best responses over information sets put the deployed policy’s endgame exploitability at $+0.4405$ over 305 one-half-suit positions (§11), against a deterministic realisation of a policy that mixes at most of these decisions — a real limitation, and still an exact computation where the alternative is a search’s lower bound.

The engine was measured against one from outside the project, and the first number was wrong. The published $+2.3466$ over 10,000 games against a frozen build of an independently developed C++ engine (Nguyen, 2026b) measured our stateless bridge, which rebuilt their belief before every decision they made; it is withdrawn. Corrected through a persistent bridge KRAKEN loses by -0.6960 , by -0.6200 measured inside their arbiter, and by -0.5250 [$-0.6886, -0.3614$] through the repaired bridge at BRIDGE_REV 3 (§11.9, §11.10). The two systems play different rule dialects on the one clause their author’s own sweep shows matters (out-of-turn declaration), so even the corrected figure is a measurement under this paper’s rules and not a ranking of two projects.

And v1.1 shipped nothing. Seven pre-registered directions at the residual returned no adoption, and three registrations of the v1.2 programme (ten arms, each required to clear $+0.15$ against SESTINA and in self-play) returned none either; two of those arms raised the ask hit rate the cross-engine table said we lack, and both lost by a third of a set doing it. The pre-registered thresholds and outcomes for all forty-six registrations this project has written are in Appendix A; the forty-sixth (P46) is registered and running, and no result of it appears here. Two mechanisms measured there are real and positive and still do not ship, because the bar was fixed in advance and was not amended afterwards. That is the discipline this paper is actually arguing for, and it costs more than it looks: exactly one of the forty-six registrations produced a change to the deployed engine.

13.2 What we would do next

- **The declaration latency, which is where the cross-engine deficit turned out to live.** §11.11 is the one item here that arrived with its target already measured rather than argued: we sit on a half-suit our own team has completed for 16.80 plies against SESTINA’s 7.84, every ask made into

one meanwhile is certain to fail, and the excess is the right size to be the whole loss. It is also the item this study has now attacked twice by registration and closed at both ends: the ask end (P44 D1, P45 F1: both raised the hit rate and lost about a third of a set against SESTINA) and the declaration end (P44 D2 stopped by futility; the ownership sweep flat from 0.77 to 0.20). What a registration here would have to fix in advance is the split itself, which is what the latency is *for*: a completed half-suit is not declared because the team cannot place the split (the ownership sweep confirms it: relaxing the ownership test adds 0.103 good declarations a game, precision 0.92, and no more). An arm that shortens the wait is not starved by wrong declarations but by too few positions where the split is placeable at all: 0.1676 of the 0.1759 wrong declarations a game are allocation class, the exactness band 0.77–0.97 is empty, and the locating convention clears its information gate at +0.0408 and loses in play at every gate tried. The tail shape is the reason to think it is tractable anyway — a third of completions carry seven eighths of the cost, so the question is not “declare sooner” but “what is different about that third”.

- **A proper likelihood, not a proxy — two of three steps taken.** This item previously read: the choice model uses depth on the initial deal, and the right object is the acting policy’s own probability of the observed action given the hand, evaluated at the time of the ask. Two thirds of that turned out to be cheap.

At the time of the ask costs a lookup table, not a fixpoint. Cards move only on successful asks and successful asks are public, so at-ask-time depth is initial-deal depth plus a delta the log fixes identically in every world (Section 8.5). It is worth 4,654 nats on 17,005 real decisions.

The policy’s own probability was measured rather than assumed, by fitting the champion’s actual half-suit choice as a conditional logit, which is what Table 5 is. The exponent is 2.195 ± 0.033 at the ask, not the 1 the shipped model assumes.

What remains is the part that is genuinely not free: using that exponent requires the denominator the engine drops. Measured across worlds, the coefficient of variation of that denominator is a V centred on $\alpha = 1$ — 0.0000 there, 0.0352 at 1.207, 0.2255 at 2.195, and 0.0997 at the shipped $\gamma = 0.35$. So the engine currently pays more for the omission than the fitted exponent would, and reaching the exponent the data want means computing the denominator per world: tracking depth in *every* live half-suit per player during the draw rather than only the asked ones. That is a constant-factor cost in the sampler.

And the free version of it has now been played, which changes what this item is worth. At $\alpha = 1$ the denominator is exactly constant, so the better covariate can be used there at no cost at all. Six pre-registered blocks put that configuration at +0.1015, 95% [+0.0063, +0.1967] — real, below the +0.15 this study fixed in advance as the smallest effect worth adopting, and heterogeneous enough ($\hat{\tau} = 0.150$, the size of the effect itself) that a random-effects reading includes zero. The engine keeps $\gamma = 0.35$.

That is the useful bound on the expensive version. The cheap form of this idea, at the one exponent where the model is correctly specified, buys about a tenth of a set. Paying a constant factor in the sampler to reach $\alpha = 2.195$ has to beat that by enough to be worth the sampler time it takes away from draws, and draws are the largest measured effect in this study. We would now want a sizing argument for that trade before writing the code, not after.

A fixpoint is still not required, and that is the substantive update: the opponent being modelled is a policy whose choice rule can be fitted directly from its own play.

- **A coordinated best-responder.** The robustness tests here show the model is not harmed by an opponent built to violate its assumption, which is weaker than showing it cannot be exploited. Training a genuine best response, and reporting the value it extracts, is the only route to an absolute claim about how good the champion is.
- **Cross-play.** Once anything in this engine is *learned* from self-play rather than configured, cross-play between independently seeded runs becomes necessary rather than optional, for the reason Hanabi established: an agent that only wins alongside copies of itself has not solved the game.
- **Belief-space search, past the one point we tested.** Section 6.3 built the design this section

previously proposed (a lookahead over the posterior rather than over sampled determinizations), and Section 9.3.1 now settles it at +0.104 per deal-pair. Settling it took pairs rather than a cleverer analysis, which is the part worth carrying forward: the harness’s intervals are measured to be honest, so the 1200 pairs that started this were never going to separate a tenth of a set from zero, and no amount of pooling was going to fix that. The question answered is narrow: possession-chain value, a single sweep of quota coupling, beam 4, depth 3, entering additively. Proposition 3 says the version without the coupling could never have worked, which removes a whole branch of the space rather than one point in it. What is still open is a leaf evaluator other than banked cards, a converged rather than single-sweep coupling, and a horizon that does not stop at the end of one possession, and, now that the one tested point is known to be worth something, whether any of those is worth more.

- **What a fourth registration measures, and what is closed without one.** P46 (prereg/kraken_v12_belief_grid.md) is registered and running as this is written, and no result of it appears in this paper. It measures the one thing three programmes never measured: the belief about SESTINA’s cards. Every P43 sentence about the opponent model rests on a two-knob difference, because `gamma_team` defaults to `opponent_gamma`, so switching the model off on their three seats also switched it off on our two teammates. Stage 0 scores twenty cells of the model exponent, on their seats and on ours, against the dealt truth on frozen champion decisions at two budgets (720 and 2,880 draws), with the Kish ESS recorded per row and intervals clustered by game; it ships nothing. Stage 1 plays the single-knob decomposition of P43’s model-off arm (off on their seats only; off on ours only) at 300 deals on both parities in both populations, plus one upward cell licensed only by a rule fixed before Stage 0 was run. The registered prediction is that nothing clears either half of the bar, for a fourth time, and that the value is the table. Four doors are closed without a registration, and the record behind each is written into the registration so that the omission is a decision and not an oversight. The draws budget: the self-play half of the dual bar is already measured — 480 to 1,440 draws is +0.0945 [−0.002, +0.191] at 6,000 pre-registered pairs, not demonstrated, and the located deficit is flat in draws on the champion’s own decisions — so no draws arm can ship whatever SESTINA returns. `n_worlds`: inert on the act path by construction — `Posterior.worlds()` has no caller in the agent, the ask features, the claim logic or the lookahead, and the knob reaches play only through the SIS-failure fallback, so the games reproduced byte for byte at 64 and 128 worlds are a no-op and not a null. A joint-scored split for voluntary claims: closed by proof — the shortlist’s first entry is the per-card marginal argmax and is always joint-scored, and $P_{\text{joint}}(A) \leq \min_c M[c, A_c]$ holds exactly because both come from the same weighted batch, so any split with joint above 0.5 has every marginal above 0.5, is the shortlist head, and is already scored. Swapping or recalibrating the ownership estimator: closed by the measurement of §11.11.7. If P46 returns what it predicts, the residual half set is not reachable by any single knob this engine exposes, and the next attempt would have to change what the engine computes rather than what it is told.

A The pre-registered thresholds, verbatim

This paper uses the words “pre-registered”, “fixed in advance” and “registered before” about seventy times and, until this appendix, reproduced no registered threshold anywhere. The forty-six documents in `prereg/` are the record those phrases refer to. Each was committed to the repository before the material it governs was played: before the arm existed, before the runner existed, or — where a screening measurement had already been taken — with that measurement stated inside the document so it could not be re-described afterwards. The table below reproduces, for every one of them, the decision rule as written, whether the document stated a power limit in advance, what the run returned, and where in this paper the verdict appears.

The registrations that returned nulls, refutations, withdrawals and non-results are listed on exactly the same footing as the ones that returned effects, in the same table, in filename order — with the four written after this appendix was first compiled (P39–P42) appended at the end, in filename order among

themselves, so that the identifiers already cited in the text keep their referents. Twelve of the forty-six have *no outcome recorded in the registration document itself*; ten of those twelve were nonetheless run and are reported in this paper or in RESEARCH_FRONTIER.md, and two have not been run at all. A forty-sixth (P46) is registered and running at the time of writing, and is listed with its prediction and no verdict. Both facts are marked in the table rather than smoothed over, because a registration that is written and never resolved is the file-drawer risk this appendix exists to expose. Anything that could not be verified against a committed results file is marked **unverified**; it is not presented as though it had been.

Throughout: “pairs” are duplicate-deal paired games with the pair as the independent unit; “ship bar” is this project’s standing +0.15 sets/game threshold; a quoted string is verbatim from the registration. Two typographic liberties, and no others: a Markdown code span such as ‘ $D + K$ ’ is set as $D + K$, and where a quote joins one of a document’s own section headings to the text beneath it the heading is set in bold and closed with a full stop.

#	Registration	Registered condition	Power limit stated	Outcome	Where reported
P1	<code>channel_vs_precision.md</code> <i>Does the channel’s value shrink as the belief improves?</i>	“ REFUTED if the interval on D lies entirely below zero.” $D = d_{1440} - d_{180}$, clustered by game. 40 games, stride 4, seed base 880,000 , sender gate 0.05, $\beta = 0.8$.	Yes: “40 games, stride 4, and no extension”; anchor tolerance 0.010.	REFUTED . “ $D = -0.0372 [-0.0481, -0.0263]$ ”. <code>results/channel_precision.json</code>	Not in the paper; alluded to in §12.2
P2	<code>choice_basis.md</code> <i>Is there a richer teammate choice model worth building?</i>	“ $D \geq 1,000 \rightarrow$ licensed to build ”; $D < 200$ withdraw; $200 \leq D < 1,000$ fit only. $D =$ held-out nats, M5 over M2, 17,005 choices, 5 folds at the game level. No seed base fixed.	None .	REAL . “Primary outcome: $D = M5 -$ row 2 M2 = +3,143 nats, against a bar of 1,000.” <code>results/choice_basis.json</code>	§12.3, Table 20
P3	<code>claim_gamma.md</code> <i>A separate action-model γ for the declaration?</i>	“ If declarations per game do not rise by at least 5% ... no duel is run. ” Duel, never reached: 3,000 pairs, ship bar +0.15, $\gamma = 0.7$. Screen: 480 games self-play.	Duel half-width “ ~ 0.11 ” named; no unresolvable-effect clause.	SCREEN FIRED , NO DUEL . “Declarations per game FELL by 1.9% against a rule requiring a 5% rise.” <code>results/path_ledger_self_claim_gamma0.7.json</code>	Table 20 row 7
P4	<code>concentration_v2.md</code> <i>Does a corrected concentration term cut allocation errors?</i>	“Point estimate $\geq +0.15$ and interval lower bound > 0 ... Here the document governs. ” Primary: paired set margin vs v0.7, “fresh seed block” (unnumbered).	Yes: “a standard error near 0.044 and a 95% interval of about ± 0.086 ... Analysed once, at the end. No interim looks.”	WITHDRAWN on condition 1. “Allocation errors were required to fall. They rose , monotonically in the weight.” <code>results/concentration_confirm.json</code>	§9.1
P5	<code>convention.md</code> <i>Can a pre-play naming agreement beat the depth heuristic?</i>	“An arm licenses a duel only if, on the teammate pool: 1. its paired NLL interval lies entirely below zero , and 2. its paired top-1 interval is not entirely below zero . Both, or nothing.” Plus $V1 \geq 25\%$ carry, $V2 \geq 20\%$, $V3 \geq 50\%$, V4 bit-identity. 40 games per sender gate, $\sim 1,040$ decisions, $n_{\text{draws}} = 720$, seed base 560,000.	None .	REAL . “The channel is real and the depth code book clears both gates at every sender gate tested.” V2 failed as written and passed only as amended — disclosed in the document. <code>results/convention_posterior.json</code>	Table 20 row 3; §12.4

continued overleaf

The pre-registered thresholds, continued

#	Registration	Registered condition	Power limit stated	Outcome	Where reported
P6	convention_aimed.md <i>Does the aimed code book replicate on fresh seeds?</i>	“replicates only if ... 1. the paired NLL interval lies entirely below zero, and 2. the paired top-1 interval lies entirely above zero, and 3. both point estimates are at least half the exploratory magnitudes — NLL at most -0.0356 , top-1 at least $+0.0176$.” 40 games, stride 4, $n_{\text{draws}} = 720$, fresh seed base 880,000.	None.	REAL, then re-priced. “Replicated on all three conditions” (2026-08-29); the 2026-08-31 reproduction records that the replication run committed no results file , so $-0.0535/+0.0392$ are unverified . Re-run: results/convention_replication.json	Not in the paper
P7	convention_duel.md <i>Does the aimed convention win games?</i>	“ Ship bar: +0.15 sets/game ... lower bound of the interval > 0 AND point estimate $\geq +0.15 \rightarrow$ ships.” Primary: paired sets/game, $Y - X$. n_pairs: 3,000.	Yes: “a 95% half-width near 0.093 sets, which resolves the +0.15 bar with room to spare”; extend if paired SD > 4.0 .	REFUTED (defect). “The convention as registered LOSES , heavily, and the reason is a mis-specified gate.” Stopped at 40 pairs; the registered 3,000 were not run . results/convention_duels.json	Table 20 row 3
P8	convention_freeread.md <i>Is the agreement already on the wire worth reading?</i>	“The ship bar is +0.05 sets/game , not the $+0.15$... because those two spend a resource and this spends nothing.” “n_pairs: 1,200 to start. If the interval straddles the bar, extend to 3,000 rather than reading a straddle as a result.”	Partial: the extension rule only; no numeric half-width.	NULL. “3,000 duplicate-deal pairs. Y minus X: $-0.002 [-0.127, +0.123]$.” results/convention_duels.json, results/convention_freeread_belief.json	Not in the paper
P9	convention_locate.md <i>Should a code book carry a location instead of a count?</i>	“supersedes the depth book only if ... 1. its best arm satisfies both gates of prereg/convention.md ... and 2. its best arm’s paired top-1 point estimate is at least as good as the best passing depth arm’s” ($+0.0351$, later re-based to $+0.0260$). Gates 0.02 and 0.05; $\beta \in \{0.25, 0.5, 0.8, 1.2, 2.0\}$.	Partial: “reported as underpowered rather than as a result” clause; no numeric limit.	PASSED, then RE-VERSED. 2026-08-29 “clears its gate”; 2026-08-31 “ condition 2 fails and the locating book does not supersede ”. 2026-08-29 figures unverified (no results file committed). results/convention_replication.json	Not in the paper
P10	convention_mixture.md <i>Is the flat decode weight mis-specified?</i>	“supersedes the flat weight only if ... its best arm’s paired NLL point estimate is at least as negative as the best flat arm’s.” Falsifiable side-prediction: “the spread of dtop-1 must be less than half the spread over the four flat arms at $\beta \leq 1.2$ (which is 0.0202).” $q \in \{0.4, \dots, 0.8\}$, same transcripts as P5.	None.	REFUTED on its own withdrawal condition. “The mixture is beaten everywhere ... Both conditions fail.” Side-prediction “2 of 3”. results/convention_posterior.json	Not in the paper
P11	deadline_signalling.md <i>Was the signalling gate set below the measured free-turn threshold?</i>	“ Ship bar. Point estimate $\geq +0.15$ with the interval clear of zero.” Primary: paired set margin vs v0.7, $C - A$. 1,000 games per arm ; “fresh seed block”; withdrawal 1: B must reproduce $+0.068$ within its interval.	Yes: “paired SD of ≈ 1.15 ... at about ± 0.071 .”	REAL, SUB-BAR. “C is $+0.1220$... It does not ship , and the bar is not amended after seeing the number.” results/signal_gate_confirm.json	§10.1

continued overleaf

#	Registration	Registered condition	Power limit stated	Outcome	Where reported
P12	<code>declarability_leaf.md</code> <i>Score declarability inside the search, not beside it?</i>	“The confirmation runs at the SMALLEST w in $\{0.25, 0.5, 1.0, 2.0\}$ whose bite falls in [15%, 60%] ... If no w lands in that window, no duel is run. ” Amended ceiling $w = 6$ fixed by argument. Duel, never reached: 3,000 pairs, ship bar +0.15.	Yes: “3,000 pairs at fresh seeds. Half-width ~ 0.11 , which resolves the bar.”	CLOSED BY PROOF, NO DUEL. “Bite does not reach 15% even at $w = 6$... The falsifier did not fire.” <code>results/declare_bite.json, results/allocation_locus.json</code>	Table 20 row 6; §12.3.2
P13	<code>declaration_timing.md</code> <i>Is the teammate ceiling card-reading or declaration-timing?</i>	“ CONFIRMED if the interval on $D - A$ lies entirely above zero and $S_D \geq 0.50$ and ($S_K < 0.25$ or the interval on $K - A$ contains zero). REFUTED if $S_K \geq 0.50$ and $S_D < 0.25$. SPLIT otherwise.” 600 games , SEED0 = 5,500,000. V1–V4 validity conditions.	None.	SPLIT. “ $S_D = 0.316$, $S_K = 0.223$... Neither. SPLIT ... the hypothesis ... is not supported.” V1 and V2 both voided as mis-specified, disclosed. <code>results/declaration_timing.json</code>	(Direction 4)
P14	<code>defer_gate_at_power.md</code> <i>Is the deferred gate real on the current engine, at power?</i>	“ $D = \text{margin}(C_{\text{defer}}) - \text{margin}(A_{\text{shipped}})$... REAL if the interval is clear of zero and positive. REFUTED if clear of zero and negative. NULL AT POWER if it covers zero.” Seed base 10,900,000 , barred from seven earlier bases. 2,000 deals $\times$ 2 parities , clustered on the deal.	Yes: “At 2,000 deals expect about 0.038, which separates +0.058 from zero.”	REAL. “+0.0455 [+0.0134, +0.0776], 2,000 deal clusters.” Below the +0.15 bar; does not ship. <code>results/signal_vs_defer.json</code>	Not in the paper
P15	<code>endgame_ask_on_combined.md</code> <i>Does the endgame ask correction still pay on top of what ships?</i>	“1. CI entirely above 0 — it stacks. Ship into V04_COMBINED and into <code>api/engine.py</code> ’s <code>WEB_SPEC</code> ... 2. CI straddles 0 with a positive point estimate — cannot resolve. Do not ship ... 3. Point estimate at or below 0 — ... Do not ship.” 2,000 pairs, base seeds 991000–992875 (16 blocks of 125 after Amendment 1).	Partial: “a run that can see +0.09 can see whether it is still there”; no numeric half-width.	NOT RECORDED in the registration. The run exists; the paper reports +0.1220 [+0.0711, +0.1729] over 2,000 pairs. <code>results/v04_duels.jsonl</code>	§11.9
P16	<code>endgame_ask_weights.md</code> <i>Does the endgame ask correction pay in play?</i>	“1. CI entirely above 0 — the correction pays in play. Ship it ... 2. CI entirely below 0 — it costs ... 3. CI straddles 0 — no detectable effect at this resolution.” 2,000 pairs , base seed 771000, agent seed 7710; replication at 881000–881750.	Yes, unusually explicit: “A test with ± 0.047 resolution against an effect that may be a fifth of that is underpowered by construction ... No outcome licenses running more pairs until the interval clears zero.”	REAL, then WITHDRAWN. “+0.0835 ... [+0.0338, +0.1332]”; replication +0.0980. Amendment 2: “The correction is withdrawn permanently ” under the award rule. <code>results/v04_duels.jsonl</code> , and the <code>ii_ask_fit</code> refits	§11.9; §12.3
P17	<code>endgame_m3.md</code> <i>Does the correction pay at three live half-suits?</i>	“1. CI entirely above 0 — it extends. Ship <code>endgame_m = 3</code> ... 2. CI straddles 0 with a positive point estimate — unresolved. Do not ship. 3. Point estimate at or below 0 — it does not extend.” 8 blocks of 250 pairs, base seeds 771000+, agent seeds 6610–6617; weight not refitted .	None (its “what this cannot tell us” is about the weight, not power).	SHIPPED, then WITHDRAWN. “Outcome 1 fired by this document’s rule ... It is now withdrawn : ... $m = 3$ is <i>worse</i> than no correction against v0.3 by -0.438 $[-0.829, -0.047]$.” <code>results/xplay_sweep.json</code>	§11.9

continued overleaf

#	Registration	Registered condition	Power limit stated	Outcome	Where reported
P18	endgame_m_ladder.md <i>How far up does the ask correction go?</i>	“CI entirely above 0 — ship $k+1$ and take the next step. Anything else — stop. Do not ship, do not try $k+2$, do not re-run with more pairs.” 8 blocks of 250 pairs a rung; endgame_d_info = +2.0 “never refitted” . Amendment 2 replaces the ladder with a foreign-opponent sweep.	None.	ALL RUNGS ABOVE $m = 2$ WITHDRAWN. “every rung of this ladder was measured against the sibling and therefore measured ... the poisoning of its inference rather than better play.” results/v04_duels.jsonl, results/xplay_sweep.json	§11.9; Table 6
P19	exact_endgame_in_play.md <i>Is exact $m = 1$ endgame play worth anything in play?</i>	No decision rule or threshold is registered. Five graded predictions instead, e.g. “one comes out positive, and at least +0.05” and “all is the worst of the three”. No sample size, no seed base fixed.	Yes, informally: “ ± 0.083 at 40 pairs ... about 300 pairs resolves ± 0.03 .”	NOT RECORDED in the registration. The run exists; the paper reports pooled +0.0765 [+0.0422, +0.1108] over 1,200 pairs. results/exact_endgame_one.json	§11.8
P20	forced_exhaustive.md <i>Should the last forced declaration enumerate the whole team space?</i>	“THE DECISION RULE IS NOT THIS PROJECT’S USUAL ONE ... the ship bar of +0.15 sets/game ... is the wrong instrument here and is deliberately not used. ... Primary outcome, fixed now. Forced-declaration accuracy at one live half-suit ... Ships if accuracy rises with the interval clear of zero.” 2,400 self-play games; 1,000 vs v0.7. Three guards, three withdrawal conditions.	Yes: “roughly 12,000 games to resolve from zero. Every duel this project could afford would return a null.”	REAL — SHIPPED. “Every condition from the pre-registration fixed in advance is met ... This ships. ” 0.2984 $\rightarrow$ 0.4070 self-play; 0.3218 $\rightarrow$ 0.4253 vs v0.7. results/forced_exhaustive_self.json, results/forced_exhaustive_v07.json	§8.7.2
P21	foreign_opponent_m2.md <i>Does the $m = 2$ correction survive a foreign opponent?</i>	“1. CI entirely below 0 — the correction’s in-play gain was also a sibling artefact. Withdraw it ... 2. CI straddles 0 — no reversal detected. The correction stays ... 3. CI entirely above 0 — it transfers.” 250 deals $\times$ 2 rotations = 500 paired games per arm, seeds 331000–331249, agent seed base 3310. “No second run to chase a straddle.”	Yes, explicitly: “this run is underpowered to confirm an effect of sibling size ... What it is a reversal of deception scale. ”	NOT RECORDED in the registration (“Amendments: None yet”). The run exists; the paper reports -0.0300 [$-0.1081, +0.0481$], and -0.0440 [$-0.1344, +0.0464$] rule-matched. results/foreign_m2_check.json	§11.9
P22	gamma_policy_specific.md <i>What does the wrong-signed opponent model cost?</i>	“Minimum interesting effect 0.15 sets/game ... G2 runs ONLY if G1’s arm C clears the bar ... C must clear the bar against v0.3 as well, on the SAME fixed setting fitted against v0.7, with no refitting.” Arms $\gamma \in \{+0.35, 0.0, -1.00\}$; “400 deals $\times$ 2 rotations per arm ”, “fresh seed block 1,800,000+”; clustered over deals.	None (an expected outcome is stated, not a power bound).	NOT RECORDED in the registration. The run exists and both arms lose; G2 correctly does not run. Deviation: run at 800 deals $\times$ 2 = 1,600 games, twice the registered size, with no amendment. results/g1_gamma_cost.json	§8.5
P23	gamma_split.md <i>Does each side of the table deserve its own γ?</i>	“A cell licenses a play experiment only if, on the teammate pool: 1. its paired NLL interval entirely below zero ... and 2. its paired top-1 interval is not entirely below zero . Both conditions, or nothing. ... If no cell satisfies both, the direction is withdrawn and no play experiment is run. ” 60 games, stride 4, $n_{\text{draws}} = 720$; unit is the decision.	Partial: only the withdrawal clause “reported as underpowered on the smaller pool” if the two pools differ by more than $5\times$.	REFUTED. “Condition 1 passes and condition 2 fails ... Withdrawn. No play experiment is run. ” results/gamma_split.json. Deviation: the file carries an unregistered $\gamma_{\text{team}} = 2.0$ cell.	§12.2; Table 20 and row 1; Figure 8

continued overleaf

#	Registration	Registered condition	Power limit stated	Outcome	Where reported
P24	ii_bounds.md <i>Is “+0.3250 is a lower bound on the whole $m = 2$ layer” right?</i>	“1. $\bar{U} \leq H$ — the claim is refuted ... 2. $\bar{L} > H$ — the claim holds by attainment ... 3. $\bar{L} \leq H < \bar{U}$ — not settled .” $H = +0.3250$. 60-game run; budgets 400,000 nodes / 90 s upper, 120 s lower; support > 400 excluded.	Yes, as a budget clause: “a wide interval is never read as a statement about the game when it is a statement about the clock.”	NOT SETTLED (outcome 3), recorded in the amendments rather than under an outcome heading. “The pre-registered decision rule is untouched and its verdict is unchanged: outcome 3, not settled .” results/ii_bound_analysis.json	§11.4
P25	ii_memo_rerun.md <i>What should the corrected solver do to the $m = 1$ numbers?</i>	Five graded predictions, not a threshold: “1. The pinned control stays 344/344. 2. mean_champion lands within 0.02 of +0.0670. 3. n_solved falls below 308. 4. mean_optimum and mean_gain come out HIGHER ... 5. Zero negative gains.” 200 games; budget amended 60 s $\rightarrow$ 300,000 nodes before any run reached game 20.	None (it names an “interesting way to be wrong” at $ \Delta < 0.01$, which is not a power bound).	GRADED 4/5 . “Four of five, and the one that missed is the one whose mechanism was argued rather than provable.” results/ii_memo_effect.json	Not in the paper
P26	information_ceiling_split.md <i>Is the information we lack our teammates’ or theirs?</i>	“ Ship bar. None. Nothing here can ship, ever. ” Primary: “paired difference of set margins against v0.7: T minus A, and O minus A, reported separately and never summed ”. 600 games per arm , both parities, award rule pinned, “fresh seed block”.	Yes: “600 games puts the standard error near 0.11 sets/game ... sized for the <i>smallest</i> arm rather than the largest.”	REAL; PRE-DICTION RE-FUTED . T +3.4100 [+3.1625, +3.6575], O +1.3067 [+1.0070, +1.6063]. “T beats O by a factor of 2.6 ” against a registered prediction that O would win. results/ceiling_split.json	§8.7.2 ; Figure 5
P27	locate_term.md <i>Can an ask be priced by the declaration it enables?</i>	“ Primary: mean paired difference in sets/game, $Y - X$, 95% interval, on the confirmation run only . Ship bar: +0.15 sets/game. ” Screen $w \in \{0.15, 0.3, 0.6\}$ at 400 pairs; confirmation 3,000 pairs, seed base 1,250,000 . Withdrawal: the screen must be “monotone or single-peaked in w ”.	Yes, in the amendment: “half-width ~ 0.106 , which resolves the bar” — and the document records that the registered shape condition “was untestable at this sample size, and that is my error”.	NULL . “Y — X = +0.047 [−0.075, +0.168] ... The term does not ship. ” Backed only by a record inside results/v04_duels.jsonl; no dedicated results file.	Table 20 row 5
P28	precision_generality.md <i>Is “the effect grows with draws” the channel or the instrument?</i>	“ INSTRUMENT PROPERTY if D’s interval lies entirely above zero ... SPECIFIC TO THE CHANNEL if D’s interval lies entirely below zero, or covers zero ... VOID if the 480 anchor misses by more than 0.010.” Cells $n_{\text{draws}} \in \{180, 360, 480, 720, 1440\}$; arms $w_{\text{unlocated}} \in \{0.0, -4.0\}$; seed base 720,000 ; 40 games.	Yes: “40 games, stride 4, no extension ... $k = 40$ clusters against a floor of 8.”	INSTRUMENT PROPERTY . “D = +0.0094 [+0.0016, +0.0172] ... Entirely above zero.” Anchor reproduced to 0.0000. results/precision_generality.json	Not in the paper; alluded to in §12.2
P29	reach_term.md <i>Should an ask be charged for the entry point it spends?</i>	“ gate + forced declarations per game must fall ... no duel is run. Total declarations per game must not fall by more than 10% ... no duel is run at this weight. ” Weight chosen by the same [15%, 60%] bite window $\rightarrow w = 0.80$. Screen 480 games; duel, never reached, 3,000 pairs at ship bar +0.15.	Yes: “3,000 pairs at fresh seeds, half-width ~ 0.11 .”	BOTH SCREEN RULES FIRED, NO DUEL . “They were 0.427. They are 0.800 — nearly doubled ... They fell 16.3% .” results/path_ledger_self_w_reach0.8.json	Table 20 row 7

continued overleaf

#	Registration	Registered condition	Power limit stated	Outcome	Where reported
P30	rules_award_baseline.md <i>Does everything survive the mis-declaration rule becoming standard?</i>	Six designs, each with outcomes fixed first. R1: 500 pairs, seed block 400000+, “CI clear of zero, positive: ...the new estimate REPLACES +1.85”. R2: seeds 332000+. R3: 200 pairs each, “registered expectation is <i>no detectable difference</i> 0.90..0.999”. R4: 500 pairs. R5: 500 pairs, seeds 440000+. R6: screen 333000+, confirm 334000+, guard vs v0.3, “the screen picks, the confirm proves”.	None for any of R1–R6.	NOT RECORDED in the registration. All six ran; results are reported across the paper (R1 +2.412 [+2.088,+2.736]). results/award_headline.json, results/r6_screen.json and four others	§3; §8.6.1; §9.3.3; §10.1; §11.9
P31	signal_budget.md <i>Is the signalling exchange rate flat?</i>	“ $D = \text{margin}(C_budget6) - \text{margin}(B_uncapped)$ ” — “ Not ‘the best cap against the incumbent’.” Verdicts: “clear of zero, positive = REAL ...negative = REFUTED ...covers zero, half-width ≤ 0.05 = NULL AT POWER ...covers zero, half-width > 0.05 = UNDERPOWERED , and it retires nothing.” Seed base 11,700,000 (nine bases barred), 4,000 games an arm, 16,000 total.	Yes, the sharpest in the set: “ It cannot resolve an improvement of +0.03 , which is a plausible size for this intervention.”	NOT RUN . No results file exists at seed 11,700,000; RESEARCH_FRONTIER.md records it as “registered, implemented and pending a run”.	Not in the paper
P32	signal_no_repeat.md <i>Does suppressing repeated signalling asks help?</i>	“ $D = \text{margin}(C_norepeat) - \text{margin}(B_incumbent)$, paired per game ... SHIP-CANDIDATE if the interval excludes zero and D is positive ... REFUTED if the interval excludes zero and D is negative ... INCONCLUSIVE if it covers zero.” Seed base 9,700,000 , re-run at 9,900,000 after the replication gate was found mis-specified; 4,000 games per arm, clustered on the deal.	None stated in advance.	REFUTED . “ $D = -0.0715$ [−0.1068, −0.0362].” Citation error : the document cites results/signal_no_repeat.json , which now holds the later seed-10,100,000 run (−0.1075); the quoted figure is in results/signal_no_repeat_9900000.json .	§10.1
P33	signal_value_after_exhaustive.md <i>Does signalling’s +0.1220 survive the engine change?</i>	“ SURVIVES if the interval is clear of zero AND contains +0.1220. SHRUNK if the interval is clear of zero and lies entirely below +0.1220 ... GONE if the interval covers zero.” $D = \text{margin}(B_incumbent) - \text{margin}(A_shipped)$. Seed base 10,100,000 ; 4,000 games per arm; clustered on the deal.	Yes, explicitly: “This run is powered to answer GONE, and is not powered to distinguish SURVIVES from SHRUNK when both are positive.”	SURVIVES . “+0.1435 [+0.0971,+0.1899], 2,000 deal clusters.” Still below the +0.15 bar, “which is not amended now either”. results/signal_no_repeat.json	Not in the paper
P34	signal_vs_defer_additivity.md <i>Are signalling and the deferred gate one effect or two?</i>	“ $I = (D_both - B_signal) - (C_defer - A_shipped)$... ONE EFFECT if I ’s interval excludes zero and I is negative ... TWO EFFECTS if I covers zero AND d is clear of both b and c ... INCONCLUSIVE otherwise.” Seed base 10,500,000 ; 4,000 games per arm, four arms on the same deals.	Yes: “expect roughly 0.05–0.07 ... It cannot finely estimate a partial overlap.”	NOT RUN — degenerate design . “ B_signal and D_both produce identical margins in all 400 games ... So the run is not executed.” The 400-game probe is results/signal_vs_defer_probe200.json .	Not in the paper

continued overleaf

#	Registration	Registered condition	Power limit stated	Outcome	Where reported
P35	stuck_claim_gate.md <i>Should the 0.5 doomed-ask gate read p_team?</i>	“Point estimate $\geq +0.15$ with the interval clear of zero: ships ... Point estimate in $[+0.05, +0.15]$ with the interval clear of zero: ships only if the ledger confirms the mechanism ... and the v0.3 guard does not sit below zero. Otherwise: does not ship.” 1,000 pairs per rung vs dylan_v07; seed block 2,400,000; clustered over deals.	Yes: “ ± 0.071 — enough to resolve the predicted effect from zero, and honest about the fact that this is a 0.16-events-per-game intervention.”	NULL IN PLAY, MECHANISM CONFIRMED (withdrawal condition 1). “Neither interval is clear of zero ... Nothing ships. ” The registration names no results file ; verified against results/stuck_gate_confirm.json.	§8.7.2
P36	tempo_regime.md <i>Should the objective stop charging for a turn that is free?</i>	“ Ship bar. Point estimate $\geq +0.15$ with the interval clear of zero.” 1,000 games per arm. Addendum, fixed before any game of it: 8,000 games per arm , TEMPO_SEED0 = 7,700,000, “analysed once , at the end. No interim looks, no stopping early”, and “ Bar, disambiguated in advance. Ship only if both readings agree ”.	Yes: “this puts the standard error near 0.040, which is the point at which the two readings of the bar coincide.”	UNRESOLVED, then WITHDRAWN. 1,000 games $+0.2280$ $[+0.0076, +0.4484]$; 8,000 games -0.0163 $[-0.0973, +0.0648]$. “ B does not replicate ... condition 3 ... Triggered. Withdrawn. ” results/tempo_rep8k_confirm.json	Not in the paper
P37	unlocated_belief.md <i>Does unlocated_now make the BELIEF better, not just the fit?</i>	“ SHIP-CANDIDATE requires BOTH , each with a 95% interval excluding zero: team NLL improves, AND team top-1 improves. REFUTED if team NLL improves and team top-1 does not ... REFUTED if the best cell is $w_{\text{unlocated}} = 0.0$, the incumbent.” Grid $\{0.0, -0.5, -1.0, -2.0, -4.0\}$; 40 games at stride 4 ; clustered by game, floor 8. Clamp $u = \max(\text{unlocated_now}, 1)$ fixed in advance.	Yes: “If 40 games turn out to be too few ... that is a finding about power ... it does not license a longer run against the same hypothesis.” The document also records that n_{games} was <i>omitted</i> from its first version and added before any interval was read.	NOT RECORDED in the registration. The run exists and refutes the covariate at every weight; the verdict lives only in RESEARCH_FRONTIER.md and in this paper. results/unlocated_belief.json (no verdict field)	§12.3 , Table 20 row 2
P38	where_the_margin_lives.md <i>Which channel of the margin identity does signalling move?</i>	Four numbered predictions with a named falsifier, not a ship rule: “2. their_wrong_per_game(B_signal) — their_wrong_per_game(A_shipped) is about $+0.13$ a game, and clear of zero.” “If the opponent’s error rate does not move under signalling ... the finding is withdrawn and the defect is the result.” Seed 11,300,000 , 2,000 deals, three arms.	None.	NOT RECORDED / NOT RUN. No results file exists at seed 11,300,000; RESEARCH_FRONTIER.md records the run as still in flight when the document was committed.	Not in the paper
P39	signal_dose Law.md <i>Is the transfer across opponents multiplicative or logistic?</i>	“ MULTIPLICATIVE — the interval covers $+0.0426$ and excludes $+0.0143$ ”; LOGISTIC the reverse; NEITHER if it excludes both, named in advance as “the outcome that would matter most”. Both predictions fitted before the run to dylan_v07’s measured rise; heuristic at 2,500 paired deals, seed base 14,900,000 .	Yes: UNDERPOWERED if the half-width exceeds 0.0142.	LOGISTIC. $+0.0178$ $[+0.0048, +0.0308]$, covering the logistic prediction and excluding the multiplicative one; realised half-width 0.0130. One registered input was later found to be a screen figure good to ± 4 points; the verdict survives recomputation at the measured baseline. results/signal_dose_law.json	§10.1

continued overleaf

#	Registration	Registered condition	Power limit stated	Outcome	Where reported
P40	signal_dose_linearity.md <i>Is the log-odds shift linear in the dose, or is there a dead zone below three signals a game?</i>	“LINEAR — the interval covers the linear prediction and excludes zero”; THRESHOLD the reverse; NEITHER if it excludes both, named as “the outcome that would hurt most”. A direct replication at power of the 1.477-dose arm, 5,000 paired deals, seed base 15,700,000 .	Yes: UNDER-POWERED if the half-width exceeds 0.0095.	NEITHER. +0.0088 [−0.0007, +0.0169] at a realised dose of 1.381: excludes zero, so there is no dead zone, and excludes the linear prediction of +0.0201, so the shift is not proportional to the dose either. results/signal_dose_linearity.json	§10.1
P41	signal_generality.md <i>Is the signalling effect a property of the convention or an exploit of one opponent’s declaration policy?</i>	“GENERAL — the margin effect clears zero and the declaration error rate rises outside its Wilson interval”; DYLAN-SPECIFIC if both cover zero at a half-width at or below 0.0734; CONTROL VIOLATED overrides everything if self shows an effect. 800 paired deals per opponent, seed 12,100,000 .	Yes: half-width 0.0734, with UNDERPOWERED named.	ANSWERS NEITHER — the design is dose-confounded. The control holds (+0.0200 [−0.0232, +0.0632]) and ev_claim does not move (−0.0525 [−0.1415, +0.0365]), but the two arms fired at 8.94 and 2.17 signals a game, so opponent and dose are not separated and the registered verdicts are not read. Superseded by P42. results/signal_generality_ev_claim_12100000.json, results/signal_generality_self_12100000.json	§10.1
P42	signal_matched_dose.md <i>The same question with the dose held equal by construction.</i>	“GENERAL — ev_claim ’s opponent channel clears zero at matched dose”; DYLAN-SPECIFIC if it covers zero at a half-width at or below 0.0209; INFEASIBLE if the calibration gate finds no common dose. Calibration bank scores nothing; was played, and the scored bank 2,500 paired deals, seed base 14,300,000 , dose tolerance $\pm 15\%$.	Yes: UNDER-POWERED above a half-width of 0.0414; re-sized from the calibration bank before any scored deal was played, and the change recorded in the document.	GENERAL. Both opponents inside tolerance at 2.766 and 2.872; dylan_v07 +0.0454 [+0.0326, +0.0582] and ev_claim +0.0172 [+0.0048, +0.0296]. The effect transfers as a rate <i>factor</i> (+0.0160 predicted, inside) and not as a fixed rise in points (+0.0423, outside). results/matched_dose_scored.json, results/matched_dose_calibration.json	§10.1

continued overleaf

#	Registration	Registered condition	Power limit stated	Outcome	Where reported
P43	kraken_v12_vs_sestina. md <i>Can anything beat SESTINA v1.0 through the repaired bridge, and does it still hold in self-play?</i>	Ships only on +0.15 sets/game with the interval clear of zero in <i>both</i> populations: against SESTINA through BRIDGE_REV 3, and against the v1.1 champion in self-play. Three candidates fixed in advance (opponent-model γ , avoid_doomed_asks, depth_mode); screen 300 deals $\times 2$ parities at seed base 9,300,000 , confirm 600×2 at 9,400,000 , never pooled. A candidate clearing against SESTINA and failing self-play is reported as opponent-specific and not shipped.	Yes, and stated as the expected outcome: that none of the three clears the 0.6 gap and v1.2 does not beat SESTINA. Multiplicity across three candidates is named in the document rather than in a discussion section.	NOTHING SHIPS , as predicted. No arm clears either half of the bar, so none reaches confirm and the stacking arm is not licensed: $\gamma=0$ at -0.5867 / -1.0667 , $\gamma=\hat{\alpha}$ at -0.9033 / -2.0033 , at -0.0933 / -0.0900 , depth_mode at -0.0033 / -0.0867 (vs SESTINA / self-play). Two lessons outrank the nulls. R0's re-measured exponent survives the bridge repair (-1.1055 against -1.0041), and acting on it is the worst arm in the study — a correct measurement of an opponent does not license a change to us, and only the self-play half of the bar separates those. And removing the opponent model entirely costs -0.59 / -1.07 , so the mis-signed model is an <i>asset</i> : being wrong about an opponent's propensity beats having no opinion. results/v12_screen.json, results/v12_screen_c3.json	§11.10

continued overleaf

#	Registration	Registered condition	Power limit stated	Outcome	Where reported
P44	kraken_v12_declaration_latency.md <i>The deficit is a declaration latency, not an ask deficit — can either end of it be reached?</i>	Same dual-population bar as P43. Two arms, each replacing a CERTAINTY test with a graded one: <code>dead_ask_threshold</code> = 0.5 (D1) restricts the ask when the top-scoring one's half-suit is more likely than not already ours; <code>claim_owned_threshold</code> = 0.77 (D2) declares a half-suit we are certain we own at a lower exactness. D2's 0.77 is a break-even from this project's own turn price and declaration accuracy, stated with its approximation, not the best of a sweep. Futility screen fixed first: 2% of ask decisions (D1), 0.25 declarations/game (D2), at seed base 9,500,000 ; screen 300×2 at 9,600,000 . D3 (both) licensed only if both cleared independently.	Yes: that D1 clears futility and fails the duel, with the mechanism named — a failed ask publicly proves the asker holds another card of that half-suit, so D1 removes the engine's own way of resolving the allocation it is stuck on. D2 was predicted to be "the one with a case".	NOTHING SHIPS. §11.11 D2 fired on 0.235 declarations/game against its 0.250 bar and never reached a duel, which is itself the result: it gates on $p_{\text{team}} \geq 0.99$, and the belief puts 0.3120 on a genuinely dead half-suit being ours; a later sweep (<code>results/ownership_estimators.json</code>) shows relaxing that test to 0.20 adds only 0.103 declarations a game, so the gate is bound by the exactness band, not ownership. D1 reached the duel and returned a significant negative in both populations, -0.3233 $[-0.527, -0.119]$ and -0.2300 $[-0.448, -0.012]$ — while raising the ask hit rate it targets, 0.5245 against the champion's 0.5184 on the same deals. It did what it was built to do and cost a third of a set doing it, so the ask hit rate is not a quantity to maximise and the 1.95-point gap of §11.10 is not necessarily a gap to close. <code>results/p44_futility.json</code> , <code>results/p44_screen.json</code>	§11.11

continued overleaf

The pre-registered thresholds, continued

#	Registration	Registered condition	Power limit stated	Outcome	Where reported
P45	kraken_v12_free_message. md <i>The convention's gate was priced in the wrong currency — does the channel win once the price is right?</i>	Same dual-population bar as P43 and P44. One arm, F1, the free-message gate: convention_max_cost = 10^{-9} in the <i>re-priced</i> ob- jective units, so the agreed card is named only when it <i>ties</i> the card the objective already chose and the message costs nothing in the currency the old gate mis-read — a median 0.3596 objective units out of a range of 1.5. Book locate, $q = 0.5$, neither swept. The decoder half is not re-run: the free read settled it at 3,000 pairs, $-0.002 [-0.127, +0.123]$. Futility screen fixed first — the arm must raise the carry rate by 2 points over the champion on the same deals, seed base 10,150,000 ; screen 300×2 at 10,200,000 .	A null in both populations, near the free read's, because F1 is that same channel plus 4.8 points of carry. Stated as <i>not expected to clear</i> .	NOTHING SHIPS. The screen passed at +5.80% carry gain, arm carry 39.5% against the 40.1% the earlier document predicts for this gate. The duel then returned a signif- icant negative in <i>both</i> populations, -0.3067 $[-0.607, -0.006]$ against SESTINA and -0.8567 $[-1.069, -0.645]$ in self-play, so the con- firm stage was not paid for. The prediction was wrong about the size: sending changes the game in a way receiving does not, and the free read's null does not bound it. Self-play is where F1 loses <i>most</i> , the opposite of the opponent-specific shape the dual population was added to catch. A gate costing nothing in the objective's own units costs 0.86 sets a game — it replaces a <i>random</i> tie-break with a <i>deterministic</i> one, and the uniform draw it removes is what decorrelates three seats sharing an objective. And the ask hit rate rose again, 0.5287 against 0.5136, replicating P44's finding on an independent chan- nel. results/p45_ futility.json, results/p45_ screen.json	§11.11.8

continued overleaf

#	Registration	Registered condition	Power limit stated	Outcome	Where reported
P46	kraken_v12_belief_grid. md <i>Is the belief about SESTINA's cards improvable at any exponent, on either side, at any budget — and which half of P43's C1b loss is about them and which is about us?</i>	Written before any instrument position was scored or candidate game played. Stage 0 is an instrument that ships nothing: the champion's own games, frozen after <code>act()</code> at every fourth scored decision, re-weighted at 20 cells — $\gamma_{\text{opp}} \in \{-1.1055, -0.5, 0.0, 0.35, 0.7, 1.0\}$ $\times$ both a worse $\gamma_{\text{team}} \in \{0.0, 0.35, 0.7\}$, plus <code>opp_lambda</code> $\in \{0.3, 0.9\}$ at the incumbent $(0.35, 0.35)$ — at two budgets , 720 and 2,880 draws, block A cross-engine against SESTINA and block B self-play, at 720, ESS <code>rNLL/Brier/top-1</code> per truth pool paired against the incumbent cell, intervals clustered by game, Kish ESS per row. Decisions where the incumbent takes the exact DP are excluded from every paired contrast. Licensing rule , dual like the ship bar: a cell licenses a play arm only if its opp-pool NLL interval lies entirely below zero at both budgets in both blocks and no top-1 interval lies entirely below zero; only $(0.7, 0.7)$ can be promoted. Calibration clause : if the rule licenses $(0.0, 0.0)$ or $(-1.1055, 0.35)$ — the belief twins of P43's dueled C1b and C1c — or fails to fail the latter, the rule is recorded as uncalibrated and the grid closes nothing. Stage 1, the same dual-population design as P43–P45: E0 <code>opponent_gamma</code> = 0.0 (C1b replicated within deal), E1 off on SESTINA's seats only, E2 off on our teammates only, and G <code>opponent_gamma</code> = 0.7 played only if Stage 0 licenses it. Seeds 10,600,000 (instrument), 10,700,000 (screen, 300×2 per arm), 10,800,000 (confirm, 600×2 , never pooled).	Yes: nothing ships, for the fourth time. Stage 0: no cell with γ_{opp} outside $\{0.0, 0.35\}$ licenses; the fitted exponent $\times$ both a worse belief about their cards and a thinner sample (opp-pool NLL $+0.015$ $[+0.005, +0.025]$ at 720, ESS <code>rNLL/Brier/top-1</code> per truth pool $(0.7, 0.7)$ fails the licensing rule, so G is <i>not</i> played. Stage 1: C1b decomposes, $E1 - E0 = +0.30$ $[+0.02, +0.58]$, and the sentence “being wrong about an opponent's propensity beats having no opinion” is retained only if $E1$'s vs-SESTINA interval lies entirely below -0.15 .	REGISTERED; RUNNING. Stage 0 and Stage 1 results are recorded in the document and in <code>results/p46_belief_grid.json</code> , <code>results/p46_screen.json</code> when complete.	§13.2

How the forty-six resolved. Seven returned an effect that still stands on the rule as registered (P2, P5, P11, P14, P20, P26, P33). Five returned an effect and then lost it to later evidence, in every case by a rule the same document had already written down (P6, P9, P16, P17, P18). Eight were **refuted** on a registered condition (P1, P4, P7, P10, P23, P28, P32, P36). Five returned an explicit null, an explicit SPLIT, or an explicit “not settled” (P8, P13, P24, P27, P35), and three (P43, P44, P45) returned no arm clearing a dual-population ship bar, two of them with significant negatives in both populations. Four were closed by a screen or by proof before the registered duel was paid for, saving roughly 9,000 duplicate-deal pairs (P3, P12, P29, P34) — and a fifth arm was, P44's D2, inside a registration already counted above. One was a set of graded predictions rather than a threshold and scored four of five (P25). Twelve have **no outcome recorded in the registration at all** (P15, P19, P21, P22, P30, P31, P37, P38, P39, P40, P41, P42) — ten of them run and written up elsewhere, two of them never run. And one (P46) is registered and running, with its prediction on record and no verdict yet.

The last four of those are the ones added after this appendix was first compiled, and they resolved on their own registered verdicts rather than on the ship bar, because none of them was a candidate for shipping: P39 returned LOGISTIC, separating two laws of transfer and refuting the multiplicative one; P42 returned

GENERAL, the first result in this line to hold at a dose matched by construction across two engines; P40 returned NEITHER, which its own document had named in advance as the outcome that would hurt most, killing both the dead-zone reading and the proportional one at once; and P41 answered neither of its verdicts, because its two arms fired at 8.94 and 2.17 signals a game and the design could not separate opponent from dose — a confound that was in the registration rather than in the run, and the reason P42 exists.

Exactly one of the forty-six — P20 — shipped, and it is the one whose registration argued in advance that this project’s standing $+0.15$ ship bar was the wrong instrument for it.

A.1 The six conditions the headline claims rest on

The rules below are quoted in full because they are load-bearing: each one is the reason a claim elsewhere in this paper is stated the way it is, and in every case the rule was written before the number that would be judged against it existed.

P20, prereg/forced_exhaustive.md — the only change in this round that shipped, and the only one that refused the standing bar in advance.

“THE DECISION RULE IS NOT THIS PROJECT’S USUAL ONE, and here is why, in advance. The predicted margin effect is $+0.028$ sets/game ... Against a paired-difference SD near 1.15 that needs roughly 12,000 games to resolve from zero. Every duel this project could afford would return a null, and that null would read as ‘no effect’ when the truth is ‘an effect two orders of magnitude below what this instrument can see’. So the ship bar of $+0.15$ sets/game ... is **the wrong instrument here and is deliberately not used.** ... **Primary outcome, fixed now.** Forced-declaration accuracy at one live half-suit, arm B against arm A on identical deals, fresh seeds, both parities. Ships if accuracy rises with the interval clear of zero. **Secondary outcome, fixed now.** The paired margin against v0.7 with its interval, reported whatever it says, and **explicitly not a ship criterion.**”

The accuracy claim of §8.7.2 is exactly this primary. The registration also honoured its own secondary the hard way: the first self-play run gave every seat the arm, producing “ $+0.0000$, 95% $[-0.0142, +0.0142]$... an arithmetic identity wearing a confidence interval”, and nothing shipped until the registered v0.7 arm was actually run.

P23, prereg/gamma_split.md — the two-condition belief rule that every later belief measurement in this paper inherited.

“A cell licenses a play experiment only if, on the teammate pool: 1. its paired NLL interval lies **entirely below zero** (a real improvement in the proper score), **and** 2. its paired top-1 interval is **not entirely below zero** (it does not significantly worsen how often the true holder is named). **Both conditions, or nothing.** ... A model can improve NLL purely by spreading probability mass while naming the holder correctly less often. For a policy whose dominant error is *naming a split*, that is a worse belief, not a better one — the allocation decision reads the argmax, not the entropy. **If no cell satisfies both, the direction is withdrawn and no play experiment is run.**”

Condition 2 is what refuted the split (Figure 8), and it was then imposed unchanged on the convention (P5), on the aimed and locating code books (P6, P9), and on *unlocated_now* in the belief (P37). The three “refuted in the belief” entries of Table 20 are all this rule firing.

P12, prereg/declarability_leaf.md — the futility screen that saved four duels, with its window fixed before any bite was measured.

“The confirmation runs at the SMALLEST w in $\{0.25, 0.5, 1.0, 2.0\}$ whose bite falls in $[15\%, 60\%]$. Smallest, not largest: every point of bite is objective displaced ... **15% floor** is $\sim 4 \times \text{locate’s } 3.9\%$. A term that re-ranks 4% of asks produced a 95% interval of ± 0.121

centred on 0.047; nothing that small can be the 1.57. **60% ceiling**: above that the term is not tie-breaking the objective, it is overriding it ... **If no w lands in that window, no duel is run.** ... Bite is a mechanism quantity measured on the *champion's own* transcripts. It cannot be traded for a margin, which is what makes selecting on it legitimate where selecting on a screen's point estimate would not be."

§12.3 reports that four of seven directions never reached a duel. This is the screen that stopped two of them, and the invariant of §12.3.2 was found while diagnosing why it fired.

P13, prereg/declaration_timing.md — the three-way rule behind the “46% is interaction” result.

“CONFIRMED if the interval on $D - A$ lies entirely above zero **and** $S_D \geq 0.50$ **and** ($S_K < 0.25$ **or** the interval on $K - A$ contains zero) ... **REFUTED** if $S_K \geq 0.50$ **and** $S_D < 0.25$... **SPLIT** otherwise — reported as a decomposition with both shares, licensing neither conclusion. This is the honest outcome for a genuinely mixed effect and is not a failure of the design. ... $D + K$ is **not** required to equal T and any report saying so is wrong ... T is carried precisely so that $D + K = T$ is visibly a question rather than an assumption.”

The run returned $S_D = 0.316$ and $S_K = 0.223$: SPLIT, and the hypothesis the document was written to test is *not supported*. The finding this paper reports — $D + K = 1.84$ against $T = 3.41$ — is the question the last clause insisted on carrying, not the one the registration set out to answer.

P11, prereg/deadline_signalling.md — the ship bar that was missed by 0.028 and not amended.

“Ship bar. Point estimate $\geq +0.15$ with the interval clear of zero. The usual bar, used here without amendment ... **Expected outcome, written down in advance.** C beats B ... and neither clears $+0.15$. The honest prediction is another decomposable null ... **If C clears $+0.15$ I should check the seating and the ledger before believing it,** because that would be three times what the previous measurement of the same mechanism returned.”

C returned $+0.1220$ [$+0.0291, +0.2149$] — clear of zero, short of the bar. The same threshold then declined $+0.1435$ (P33), $+0.0580$ (P35) and $+0.0455$ (P14), each time without amendment, and the registrations record the refusal explicitly: “it is not amended now because it came within 0.0065 of being met”.

P26, prereg/information_ceiling_split.md — the only registration in the set whose bar is that nothing can ever ship.

“A SAFETY RULE THAT GOVERNS THIS ENTIRE FILE. Every arm here except the baseline **cheats** ... Its margins are **bounds on what better inference could buy**, and they are not strength measurements. They must never appear in a strength ladder, never be quoted beside an honest engine's margin as though the two were comparable, and never enter the head-to-head table. Any figure this produces carries the word *ceiling* wherever it is written down. ... **Ship bar. None. Nothing here can ship, ever.** There is no bar because there is no decision: an arm that cheats is not a candidate.”

Every appearance of $+3.41$, $+1.31$ and $+6.61$ in this paper is governed by that paragraph. The same document also recorded, in advance, the prediction that the *opponent* arm would win, and then recorded that it lost by a factor of 2.6 — which is why the teammate ceiling, not the opponent one, is the number the rest of the paper chases.

B Corrections this revision makes to its own record

Every entry below was in a previous draft of this paper and is wrong. They are collected because a correction that appears only where the error was is a correction most readers never see, and because the

count itself is information: a paper that reports no corrections to itself has either not been checked or is not saying.

Nine were found by an adversarial audit of this document against the results files, the registrations and the source, run in September 2026. The audit is not part of the paper’s method and its findings are not a result; it is recorded here because the alternative is to present the corrected text as though it had always read that way.

C1. Three claims about the other project, withdrawn. The paper described that engine’s policy as “frozen upstream after self-play tuning”, quoted its author as reporting that its advantage “does not transfer”, stated that a “feint” deception feature of theirs cost -2.2 on replication, and said their attribution study credited one coordinate with essentially the whole gain of a cycle. [Nguyen \(2026b\)](#) gives that policy’s fit objective as minimax regret over an opponent panel and argues in §12.1 that the advantage is *not* an artefact of self-play coordination; contains no such quotation; lists *feint* as a member of its opponent *panel*, which it beats by $+6.07$ percentage points; and puts the coordinate’s leave-one-out contribution at $+1.69$ [$+0.80, +2.57$] of a $+4.73$ whole while stating that the stronger claim is not supported at its power. All four are withdrawn.

C2. The deployed configuration was described three ways, one false. The paper said the public table “has been playing V04_COMBINED all along”. It has not: the site runs V06_DEPLOYED, which differs in carrying the endgame correction at `endgame_m = 0` — that is, not at all — after this paper withdrew it. A reader could have finished the document believing a withdrawn arm was live. Both configurations are now given verbatim in Appendix [E](#).

C3. “Pre-registered and replicated” for a figure run once. The $+0.1220$ endgame correction was replicated against the champion at $+0.0907$ [$+0.0555, +0.1260$]; the $+0.1220$ itself, on the stronger base, is a single 2,000-pair run. The word “replicated” was attached to the wrong one.

C4. Two mixed engine pools, undisclosed. `scripts4/check_engine_provenance.py` finds two duel pools whose blocks did not execute the same program, both quoted in this paper as primary results, and the paper cited the checker’s own Q statistic as ordinary evidence that the blocks agreed. One of the offending builds is not in the repository’s history, so those cells cannot be re-run. Disclosed in §[9.4.3](#).

C5. A table backed by its own numbers. The four turn-price bands are not recoverable from any results file: `scripts4/turn_price.py` does not persist the per-decision values the bands are cut on, and `scripts4/tempo_regime.py` carries the four prices as a constant copied from this paper. The file that appears to back the table contains the table. Disclosed in §[10.5](#).

C6. A selective table. Table [3](#) showed ten of the fourteen configurations its campaign scored, and the omissions included the one weight at which the opponent model is worse than switching it off. Two further campaigns scored 21 more. The caption now says so.

C7. $\pm$ meant two things. The symbol denoted a one-standard-error figure in some places and a 95% half-width in others, in the same document, with no convention stated. §[7](#) now fixes it and the violating sites are converted. One overclaim went with it: two turn-price bands described as “free, twice over” carry 95% intervals of $[-0.374, +0.288]$ and $[-0.276, +0.284]$.

C8. Same model, two calibration verdicts. The half-suit value model was described as “reliable in every decile” in one section and as miscalibrated in five of ten in another. The second is correct.

C9. A wrong file, a wrong percentage, and a mechanism claim beyond its instrument. The no-repeat result cited a results file that a later registered run had overwritten; the repeats are 92.7% of fires, not “ninety-six percent”; and the paragraph concluded that signalling is “a stall that happens to communicate”, which reads one team’s declaration ledger as the whole scoreboard. The instrument counts what our three seats declare and discards what the opposing three do. The conclusion is withdrawn.

C10. Two nulls that were not nulls. The signalling protocol and the deferred gate were both filed as “a real reduction that buys no sets”. Both have since been re-run at power and both buy sets: +0.1180 and +0.1220 for signalling, +0.0455 [+0.0134, +0.0776] for the gate. Neither clears the +0.15 bar, so neither ships, and that was always the reason — but “buys no sets” was a claim the thousand-game runs could not support.

C11. Smaller corrections. The introduction said three iterations where the paper covers four. A margin-decomposition table was introduced as 600 games when every row is the 10,000-game block. An avoided misdeclaration was priced at +1.7898, the deferral gate’s exchange rate, where the signalling arms’ own is +1.2784 [+1.0642, +1.4926] and excludes it. The volume and conversion halves of the acquisition gap were reported as two contributions of equal size; a landed ask retains the turn, so the measured hit rates alone predict the observed ask-count edge and the volume term is what the conversion term implies. The conclusions covered only the first of four development cycles. The reproduction section said v0.3’s code was untouched, after fifteen commits including the award-rule correction, and gave no repository URL. An earlier draft asserted that a half-suit’s split is frozen once a team holds all six, which does not follow from the rule that makes it tempting, and an interpretation of the depth-decay term was stated behaviourally where the evidence supports only a dilution correction.

Corrections to individual figures are recorded inline, beside the number each one corrects, rather than gathered here: a reader holding an earlier version finds the correction where the disagreement is.

C Interoperating with a foreign engine

Every cross-engine number in this paper — the retracted +2.3466 headline and the −0.6960 that replaced it, the declaration decomposition, the transfer failures — exists only because two independently written engines can be made to play one game. That is an artifact in its own right, and it has a failure mode severe enough to have already destroyed one set of results (§11.9), so it is documented here rather than left in the code.

Three directions, and only the third settled anything. A bridge is asymmetric: whoever owns the game loop owns the rules, the RNG, and every tie-break. If only one side can host, every published number is measured through that side’s arbiter and the two claims “our engine is stronger” and “our arbiter favours us” are not separable. This project therefore ships every direction it can, which for most of its life was two — and neither of those two could put our champion in *their* arbiter, so the argument stood unmeasured until the third arrived.

	<i>their engine inside ours</i>	<i>our engine driven by a host</i>	<i>our engine as their bot package</i>
entry point	external_v07/ shim_decide.cpp	fish4/decide.py	fishlab/ bot.py
host owns	our arbiter	the caller’s	their arbiter
transport	one decision per process	one decision per line	one decision per line
cards identified by	integer id	card name	card name
state from caller	counts, scores	nothing derived	counts, scores
dialect spoken	ours	ours	theirs

The two asymmetries are the lesson. They are not stylistic. Their shim takes card *integers*, and the two engines number the deck over different permutations of the half-suits; a host passing its own integers gets a bot playing a scrambled hand — legally, badly, and with no error raised anywhere. It does not look like a broken integration. It looks like a weak bot. Names are unambiguous between any two implementations of a 54-card deck, so the reverse bridge takes names and refuses anything else.

The second asymmetry is the same argument applied to derived state. Their shim accepts hand counts and scores on every event line, so a caller that miscounts silently poisons the policy’s view of the table. Ours *derives* both by replaying the event log, which means the only thing a caller can send is the public

record: there is no field through which a mistake (or a hostile host) can reach the policy. An event stream that no deal could have produced is rejected by name, with the three ways it usually happens listed, rather than raising the contradiction from inside the belief propagator.

A bridge is a measurement instrument and is calibrated as one. Bridge revision 1 substituted a legal move whenever the foreign engine returned one our arbiter rejected. That is not a neutral repair — it plays for them, badly, and it was worth -0.08 sets a game *to them*. Revision 2 removes substitution entirely and every reported match records zero fallbacks; the 10,000-deal headline was replayed under both revisions to price the difference rather than assert it. The principle this project now works to: *an absolute measured through a bridge is a statement about the bridge as well as about the engines, whereas a paired contrast whose arms share the bridge is not*. Every absolute here is reported with its bridge revision, and that is why.

Verification. The reverse bridge is checked by playing complete games in the normal way and, at every decision, serialising the position onto the wire and requiring the bridge to answer what a freshly seeded agent answers — freshly seeded rather than the in-game agent, because a stateless bridge cannot share an RNG stream with a long-lived one. Over two games this compares 150+ decisions and exercises both the ask and the declaration path. The JSON-transport handout runs the same comparison at a larger scale as its self-test: 329 decisions, 0 mismatches, plus 520 off-turn polls.

A third direction, and the one that finally measured something. The two bridges above are both *ours*: we wrote the shim that runs their policy here, and we wrote the service that lets a host drive ours. Neither could seat our champion in a batch run of *their* arbiter, which is why the separability argument at the head of this appendix went eleven months without a number attached to it. That project has since published an uploadable bot-package format, so our champion now installs into their engine as an ordinary seat and is matched from their command line (`fishlab/`, built by `fishlab/build.py`; §11.10 reports what it found).

Two things about it are worth recording, because both are the same lesson this appendix already teaches from the other side. Their format offers a bridge to *our* JSON dialect, which would have been one line of manifest, and taking it is not safe for this bot: our `ClaimEvent` wants the true holders at resolution and their arbiter publishes none on a wrong declaration — correctly, since that is what a person at the table sees. So the package speaks their dialect natively. And the first version of it made our declaration policy *unreachable*: it answered their pre-move poll with certainties only at both turn states, and the ask path then discarded any declaration the policy returned. The tell was a figure that looked like a triumph — 100.00% declaration accuracy over 1,200 games — and was in fact the signature of a policy that had been prevented from deciding. A bot that is never wrong about anything is not evidence of strength; it is evidence that something upstream is answering for it.

The opponent moved under the pin, and that is a measurement too. Every game reported here was played against their tree at commit `d017fbc`. That project subsequently released the identical configuration under a new name, SESTINA v1.0 (Nguyen, 2026b), and put five commits under our pin — among them 124 new lines in the factory function this bridge calls to construct their agent. A margin against “their v0.7” is a claim about a policy, so a change beneath it threatens the claim; and their release note saying that nothing measured changed is their assertion, not our measurement.

Rerunning the head-to-head would have been the wrong instrument: it answers a noisy question expensively, and a null there is equally consistent with a real behavioural change too small for 10,000 games to resolve. The shim is a *pure function* of the script fed to it — it boots, replays the public record, answers one decision and exits — so the two builds can be compared exactly instead. We captured the `stdin` script the bridge sends for every decision of real games and replayed each one through a binary built from the old pin and one built from the release:

decisions compared	6,329
asks	5,817
declarations	472
mismatches	0

Their freeze artifact is byte-identical across the same range, so the configuration is unchanged as well as the play, and the pin moves with the published margin intact. The limit is worth stating precisely: this exercises the decision distribution *our games reach*, which is the distribution our claim is about, and says nothing about positions our policy never creates.

The same question, asked of the rest of their ladder, is §8.6.1’s provenance check, and it is answered by a different instrument for a mechanical reason — this shim calls a function their v0.4 tree did not have, so it cannot be built against those headers, and editing it to compile would mean verifying a shim other than the one our measurements run through.

D Reproducing this work

Availability. The engine, agents, experiment artifacts, the pre-registrations and the source of this document are in one repository: <https://github.com/kv1514/fish-research>12. v0.3’s engine lives in `fish/` and v0.4 onward in `fish4/`.

An earlier version of this paragraph said v0.3’s code was untouched, so that its published results remained reproducible alongside the newer ones. That is no longer true and the exception matters: the misdeclaration award rule was corrected in `fish/rules.py`, whose default is now "opponent", so a reader who runs a v0.3 entry point today gets award-rule behaviour and not the void-era behaviour v0.3’s numbers were measured under. Pass `wrong_distribution_outcome="null"` to reproduce a void-era figure. Every result in this paper is labelled by rule era for the same reason.

How the numbers in this document were checked. `scripts4/check_paper_numbers.py` re-reads 324 watched figures from `results/` and fails if any no longer appears in the text as its results file holds it; `scripts4/check_verdicts.py`, `scripts4/check_seeds.py`, `scripts4/check_prereg_backing.py` and `scripts4/check_tex.py` check that every registered verdict is reported, that no registration is scored on the deals that motivated it, that every pre-registration claim has backing, and that the document has no broken reference. All five run in continuous integration.

One does not. `scripts4/check_engine_provenance.py` exits non-zero on this repository, truthfully, for the reason given in §9.4.3, and is deliberately excluded from continuous integration because a check that is permanently red is one people learn to scroll past. It is the one guard a reader should run by hand.

Environment. Python ≥ 3.11 , then `pipinstall-rrequirements.txt-rrequirements-dev.txt`. Fitting the learned models additionally needs `requirements-learn.txt` (PyTorch); nothing else in the repository requires it, and the engine, the duels, the exact solver and every figure in the strength ladder run on numpy alone. Figures are regenerated by `scripts4/make_figures.py`, which needs `matplotlib`. The commands below are written with `py`, the Windows launcher; use `python3` elsewhere.

```
py -m pytest tests -q      # v0.3: rules, fuzz, leakage proofs, belief
                           # soundness, statistics, exact solver
py -m pytest tests4 -q     # v0.4: exact counting vs brute force, sampler
                           # support/density/marginals, batch equivalence,
                           # information-boundary invariance, exact solver v2

py scripts4/duel.py jobs/j6_ladder.json 4      # the strength ladder
py scripts4/duel.py jobs/j11_final.json 6      # headline + ablations
py scripts4/summarise_duels.py              # every duel, with its MDE

py scripts4/posterior_accuracy.py 14 3        # posteriors scored on the true hands
```

```

py scripts4/exact_bench4.py          # agreement with provably optimal play
py scripts4/fit_hsvalue.py 200 3    # fit the half-suit value model
py scripts4/analytics4.py 120 3     # strategy statistics
py scripts4/check_tex.py          # structural check on this document
py scripts4/check_paper_numbers.py  # every watched figure vs results/

py -m kraken.decide --self-test      # the drop-in handout, end to end
py -m pytest tests4/test_decide_bridge.py -q  # the reverse bridge

```

Two entry points exist for running this engine outside this repository, both documented in §C: `kraken/decide.py` speaks one JSON object per line and is the one to use from another language, and `fish4/decide.py` speaks the line protocol that mirrors the C++ shim. `kraken/README.md` is a complete integration guide with no research code in it.

Every duel appends to `results/v04-duels.jsonl` recording both random seed streams, the full agent specifications, the rule configuration and the result. `summarise_duels.py` prints, alongside every interval, the smallest effect that run could have detected, which is the column that stops a null being read as a negative.

E The champion, and the deployed configuration

Two different objects, named separately because the paper has confused them before. The first is the reference point every ablation is measured against; the second is what the public site runs.

The reference champion.

```

from fish4.registry4 import make_agent
agent = make_agent(("fishbot4", {"opponent_gamma": 0.35}))

```

The deployed configuration , `V06_DEPLOYED` in `fish4/registry4.py` and `WEB_SPEC` in `api/_engine.py`, which agree key for key:

```

{"opponent_gamma": 0.35, "n_draws": 480,
 "w_lookahead": 0.25, "lookahead_depth": 3, "lookahead_beam": 4,
 "endgame_m": 0, "claim_forced_exhaustive": 1}

```

`endgame_m` is 0: the endgame ask correction is *not* deployed, having been withdrawn under the award rule (§11.9). `V04_COMBINED`, which several measurements in this paper use as their arm, carries it at `endgame_m = 2` and is not what the site runs. Signalling and the pre-play convention are off (`signal_mode` defaults to "off"), so nothing in §12.3 changes the deployed engine.

A stochastic policy, re-seeded on every move. The deployed session is stateless by construction: the client carries the game and every request rebuilds it from a sealed token plus the public action log. Rebuilding the game rebuilds the six agents, and until this release it re-seeded each of them from the deal — so a paced client, which asks for one engine move per request, took every move from a freshly seeded agent. The same observation appears above in a narrower setting, where a stateless bridge cannot share an RNG stream with a long-lived one (§C); this is what it costs when the stateless thing is the engine itself.

It costs something because this policy is not deterministic. The posterior’s sampler draws from the agent’s stream, the tie-break among equal-scoring asks draws from it, and the exact endgame solver draws from it again once the public record pins every live card. Reset that stream before each move and the engine’s answer becomes a pure function of the position — so a position that recurs produces the move that produced it, and two sides both playing well can pass the last two half-suits back and forth without end.

The size of it is measurable on the deployed path, holding the deal fixed and changing only the seeding. Over 1,000 deals in each of two modes, so 2,000 games an arm: **56 never ended with the stream reset before every move, against 0 with it advancing**. Every one of the 56 had exactly seven of

nine half-suits resolved — two live half-suits, twelve cards, passed back and forth with neither side ever holding all six of either. The longest game that did finish ran 329 actions, against a median of about 105.

Nothing measured in this paper is affected. Every figure here is taken through `fish4/match.py` or `fish4/evalx/harness.py`, both of which play a whole game with one agent and therefore always had the advancing stream, and the deployed policy tuple is unchanged. What changed is the harness the site wraps around it. It is recorded here because *the policy is stochastic* is an assumption every evaluation in this paper makes silently, and one deployment quietly violated it.

Everything else is at its default, and the defaults are deliberately v0.3’s own ask objective ($w_{\text{suit}} = 0.06$, $w_{\text{turn}} = 0.60$, $w_{\text{scarce}} = 0.20$). The improvement reported here comes from inference, not from re-tuning the objective on top of it, and Section 9.1 is the record of trying to re-tune it and failing.

F Information-boundary tests

The property asserted is that every quantity v0.4 computes is bit-identical when the hidden cards of the other five seats are permuted into a different layout equally consistent with the public record. Candidate permutations are drawn from the acting seat’s own belief sampler — uniform permutation almost never lands on a consistent world, because the public record is very constraining, and each candidate is then validated by an *independent* replay of the whole public history, so the test does not assume the sampler is correct. Checked quantities: posterior marginals under four combinations of the opponent-model parameters, all eleven ask-objective terms, and the end-to-end policy action.

There is a second boundary, and it is worth recording that the first set of tests did not protect it. The permutation tests establish that the *policy* never reads a card it should not; they say nothing about what the *server* puts on the wire. An earlier revision of the multiplayer table included the deal seed in the per-seat view, as a debugging convenience. Every hand is a deterministic function of that integer, so one number defeated the entire property the permutation tests were built to guarantee, and it was readable without joining the table, since a table code alone fetches a view. The lesson generalises past this bug: a leak test that only exercises the component you were thinking about will keep passing while the leak moves to the component you were not. The transport now has its own test, which walks the serialised view recursively and fails on any key or value that could reconstruct the deal, for a seated player and for an anonymous visitor alike.

References

Dylan Nguyen. FishBot: an engine for six-player Literature, versions 0.2–0.7. Software repository and technical reports, <https://github.com/dylann4500/fishbot>, 2026. Version 0.7 was pinned at commit d017fbc for every game reported here, and the pin has since moved to f9b4743 (the SESTINA v1.0 release) on the evidence of §C: the two builds answer all 6,329 of our decisions identically. Nguyen (2026b) records a history rewrite during that project’s evaluation, so an earlier version of this entry warned that d017fbc might not resolve; the published history contains it, and contains the release commits of every rung of §8.6.1 as well. The frozen configuration string in `external_v07/v07_spec.txt` identifies the build we played.

Dylan Nguyen. Bridge, not edge: an independent re-measurement of KRAKEN v1.1 against SESTINA v1.0. FishLab technical report, 2 September 2026. Produced at f9b4743 of <https://github.com/dylann4500/fishbot>; artifacts, drivers and the repair patch under `research/kraken-check/`. Received 8 September 2026, after this paper had independently withdrawn the cross-engine headline on other evidence and before it had identified the defect correctly.

Dylan Nguyen. SESTINA v1.0: exact inference, robust evaluation, and adversarial testing in six-player Canadian Fish. Zenodo preprint, 31 August 2026. <https://doi.org/10.5281/zenodo.22196900>. Developed and evaluated as FishBot v0.7; the frozen v07:r12=25, . . . of its Appendix B is the configuration this paper plays against.

- Andrea Celli and Nicola Gatti. Computational results for extensive-form adversarial team games. In *AAAI Conference on Artificial Intelligence*, 2018.
- Jeffrey Richard Long, Nathan R. Sturtevant, Michael Buro, and Timothy Furtak. Understanding the success of perfect information Monte Carlo sampling in game playing. In *AAAI Conference on Artificial Intelligence*, pages 134–140, 2010.
- Nolan Bard, Jakob N. Foerster, Sarath Chandar, Neil Burch, Marc Lanctot, H. Francis Song, et al. The Hanabi challenge: a new frontier for AI research. *Artificial Intelligence*, 280:103216, 2020.
- Ivona Bezáková, Alistair Sinclair, Daniel Štefankovič, and Eric Vigoda. Negative examples for sequential importance sampling of binary contingency tables. *Algorithmica*, 64(4):606–620, 2012.
- Yuguo Chen, Persi Diaconis, Susan P. Holmes, and Jun S. Liu. Sequential Monte Carlo methods for statistical analysis of tables. *Journal of the American Statistical Association*, 100(469):109–120, 2005.
- Peter I. Cowling, Edward J. Powley, and Daniel Whitehouse. Information set Monte Carlo tree search. *IEEE Transactions on Computational Intelligence and AI in Games*, 4(2):120–143, 2012.
- Ian Frank and David Basin. Search in games with incomplete information: a case study using Bridge card play. *Artificial Intelligence*, 100(1–2):87–123, 1998.
- Mark Goadrich, Aiden Morenville, and Éric Piette. Valet: a standardized testbed of traditional imperfect-information card games. arXiv:2603.03252, 2026.
- Hengyuan Hu, Adam Lerer, Alex Peysakhovich, and Jakob N. Foerster. “Other-play” for zero-shot coordination. In *International Conference on Machine Learning*, 2020.
- Mark Jerrum, Alistair Sinclair, and Eric Vigoda. A polynomial-time approximation algorithm for the permanent of a matrix with nonnegative entries. *Journal of the ACM*, 51(4):671–697, 2004.
- Jean-Charles Régin. Generalized arc consistency for global cardinality constraint. In *AAAI Conference on Artificial Intelligence*, pages 209–215, 1996.
- Douglas Rebstock, Christopher Solinas, Michael Buro, and Nathan R. Sturtevant. Policy based inference in trick-taking card games. In *IEEE Conference on Games*, 2019.
- Christopher Solinas, Douglas Rebstock, and Michael Buro. Improving search with supervised learning in trick-based card games. In *AAAI Conference on Artificial Intelligence*, volume 33, pages 1158–1165, 2019.
- Christopher Solinas, Douglas Rebstock, Nathan R. Sturtevant, and Michael Buro. History filtering in imperfect information games: algorithms and complexity. arXiv:2311.14651, 2023.
- Bernhard von Stengel and Daphne Koller. Team-maxmin equilibria. *Games and Economic Behavior*, 21(1–2):309–321, 1997.
- Ian Waudby-Smith and Aaditya Ramdas. Estimating means of bounded random variables by betting. *Journal of the Royal Statistical Society Series B*, 86(1):1–27, 2024.
- Leslie G. Valiant. The complexity of computing the permanent. *Theoretical Computer Science*, 8(2):189–201, 1979.